

# **Efficient Methods and Hardware for Deep Learning**

**Song Han**  
**Stanford University**

May 25, 2017



# Intro



**Song Han**  
PhD Candidate  
**Stanford**



**Bill Dally**  
Chief Scientist  
**NVIDIA**  
Professor  
**Stanford**

# Deep Learning is Changing Our Lives

## Self-Driving Car



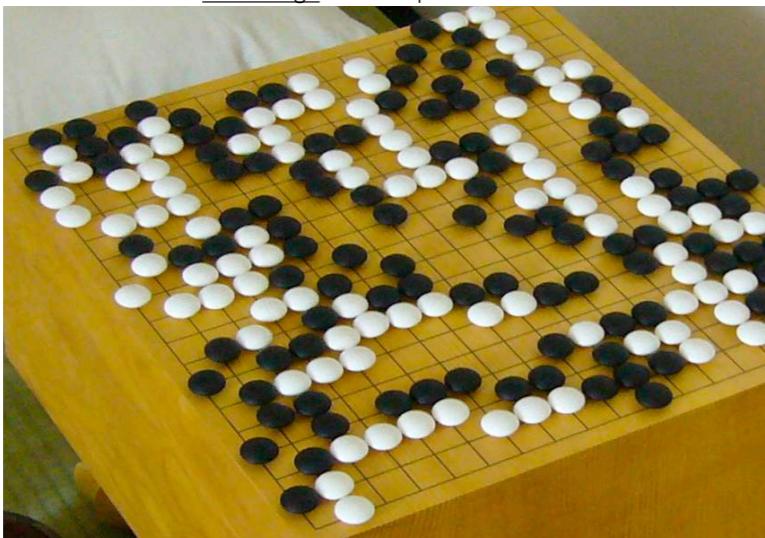
[This image is licensed under CC-BY 2.0](#)

## Machine Translation



[This image is in the public domain](#)

[This image is in the public domain](#)



[This image is licensed under CC-BY 2.0](#)



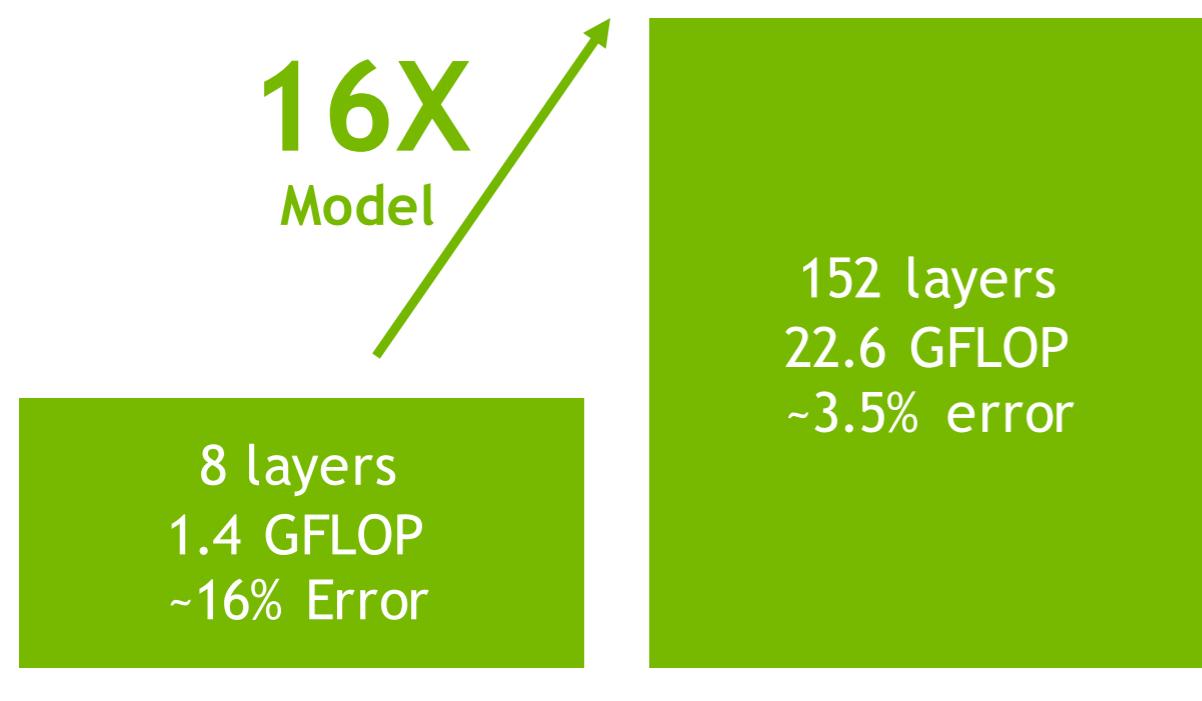
## Smart Robots

3

## AlphaGo

# Models are Getting Larger

## IMAGE RECOGNITION

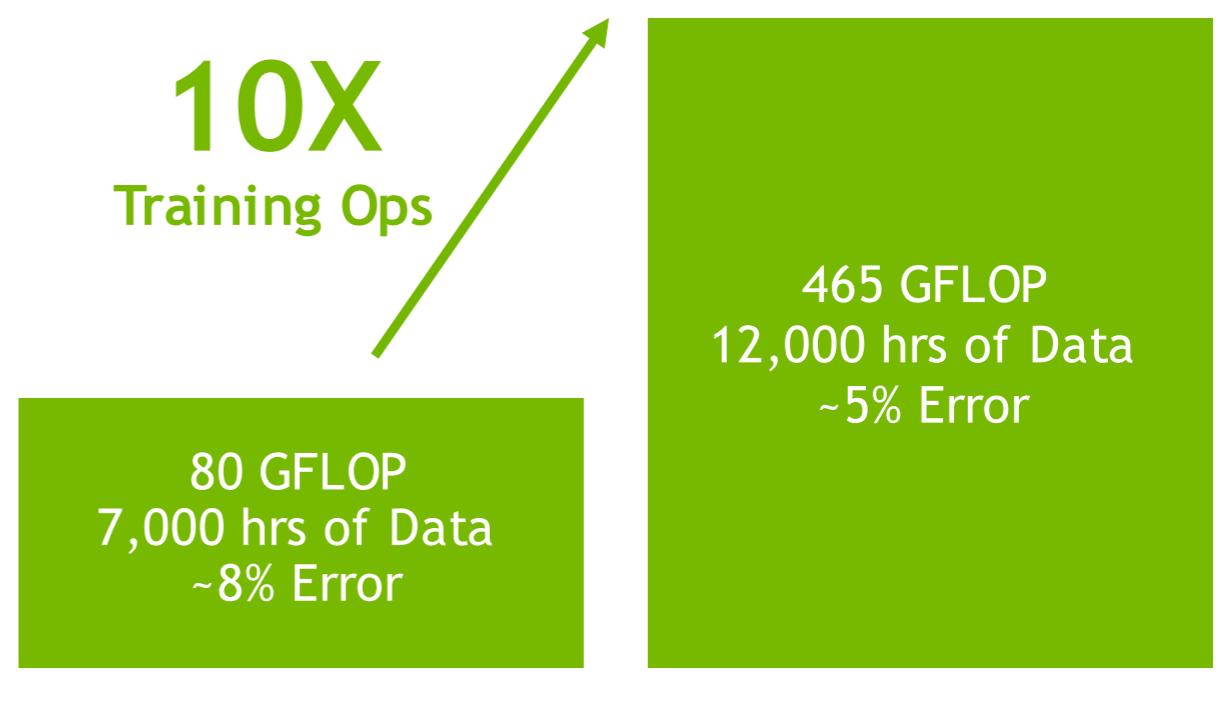


2012  
AlexNet

2015  
ResNet

Microsoft

## SPEECH RECOGNITION



2014  
Deep Speech 1

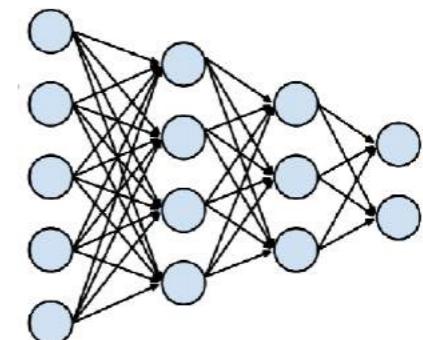
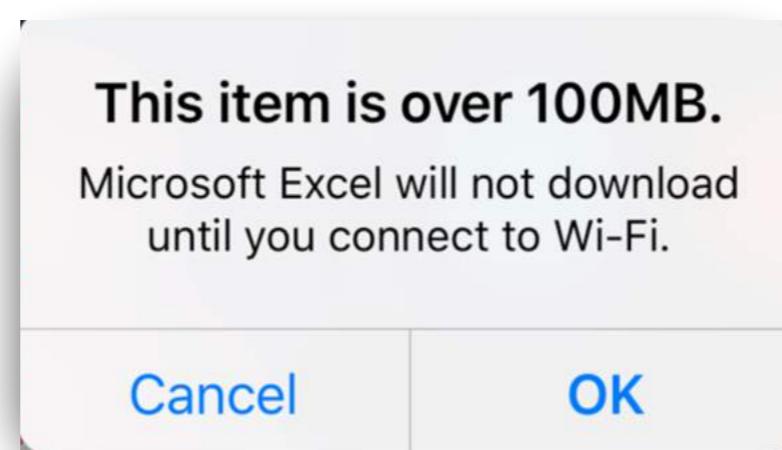
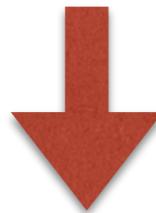
2015  
Deep Speech 2

Baidu

Dally, NIPS'2016 workshop on Efficient Methods for Deep Neural Networks

# The first Challenge: Model Size

Hard to distribute large models through over-the-air update



App icon is in the public domain  
Phone image is licensed under CC-BY 2.0

This image is licensed under CC-BY 2.0

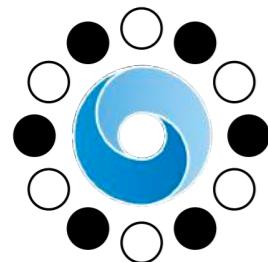
# The Second Challenge: Speed

	Error rate	Training time
ResNet18:	10.76%	2.5 days
ResNet50:	7.02%	5 days
ResNet101:	6.21%	1 week
ResNet152:	6.16%	1.5 weeks

Such long training time limits ML researcher's productivity

Training time benchmarked with fb.resnet.torch using four M40 GPUs

# The Third Challenge: Energy Efficiency

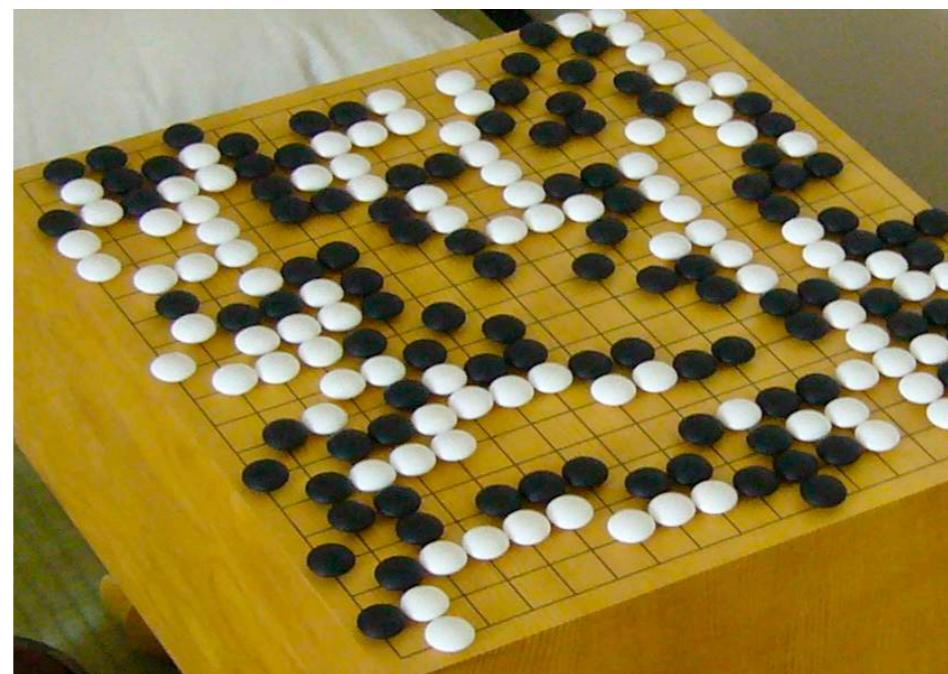


[This image is in the public domain](#)

AlphaGo: 1920 CPUs and 280 GPUs,  
**\$3000 electric bill** per game

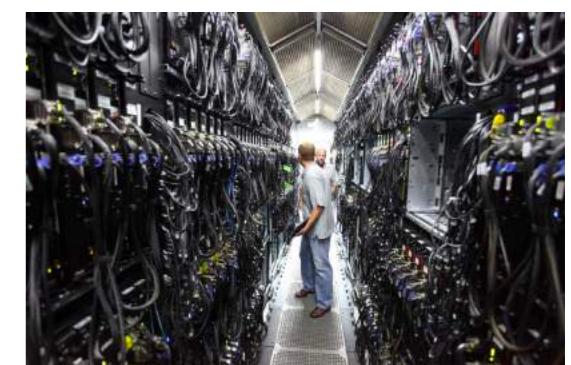


[This image is in the public domain](#)



[Phone image is licensed under CC-BY 2.0](#)

on mobile: drains battery  
on data-center: increases TCO



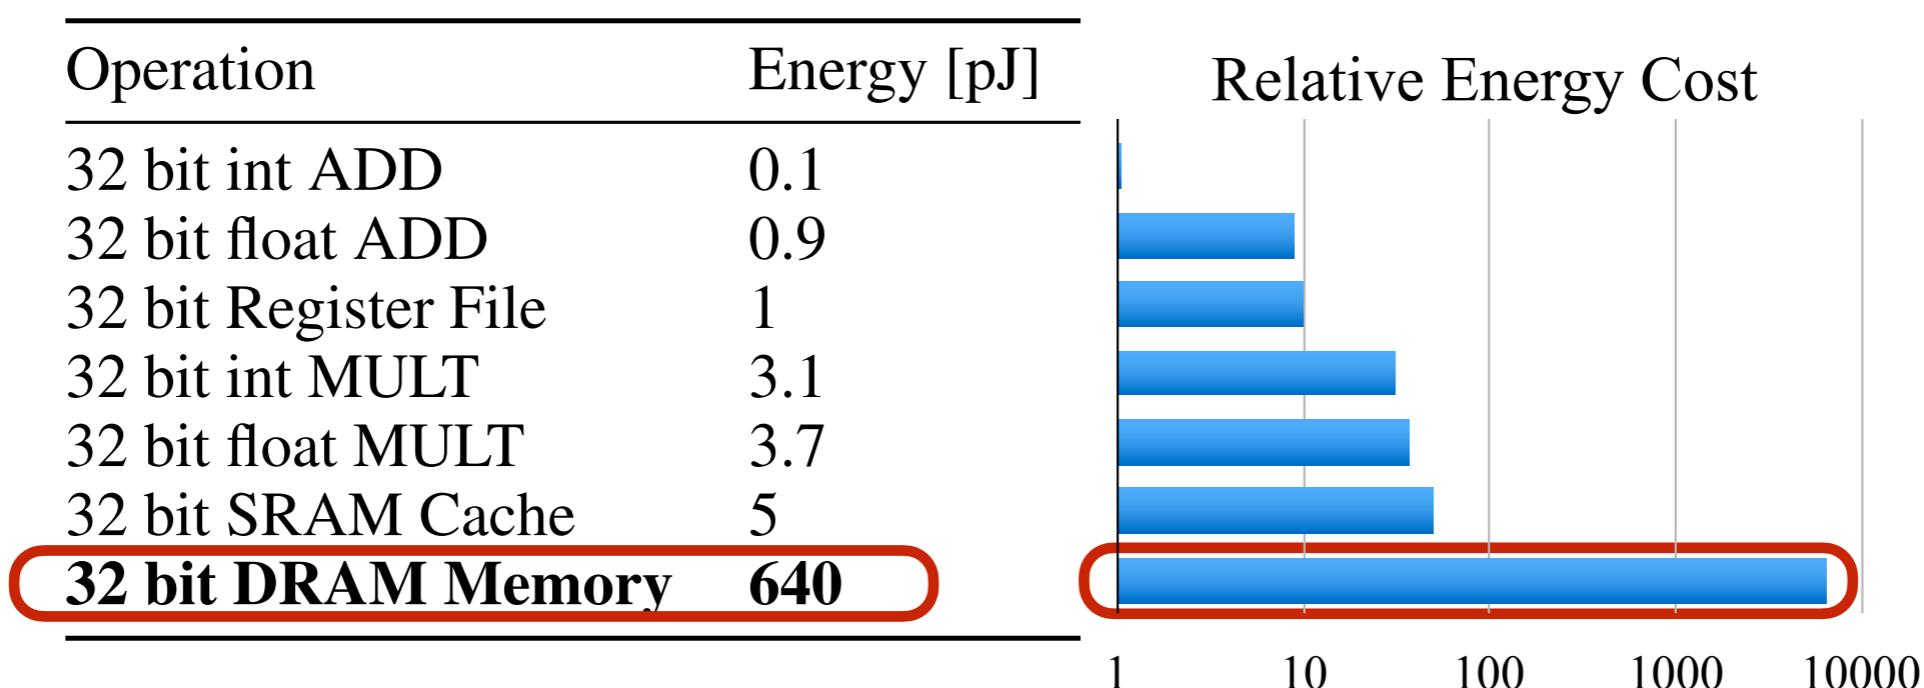
[This image is licensed under CC-BY 2.0](#)

# Where is the Energy Consumed?

**larger model => more memory reference => more energy**

# Where is the Energy Consumed?

larger model => more memory reference => more energy

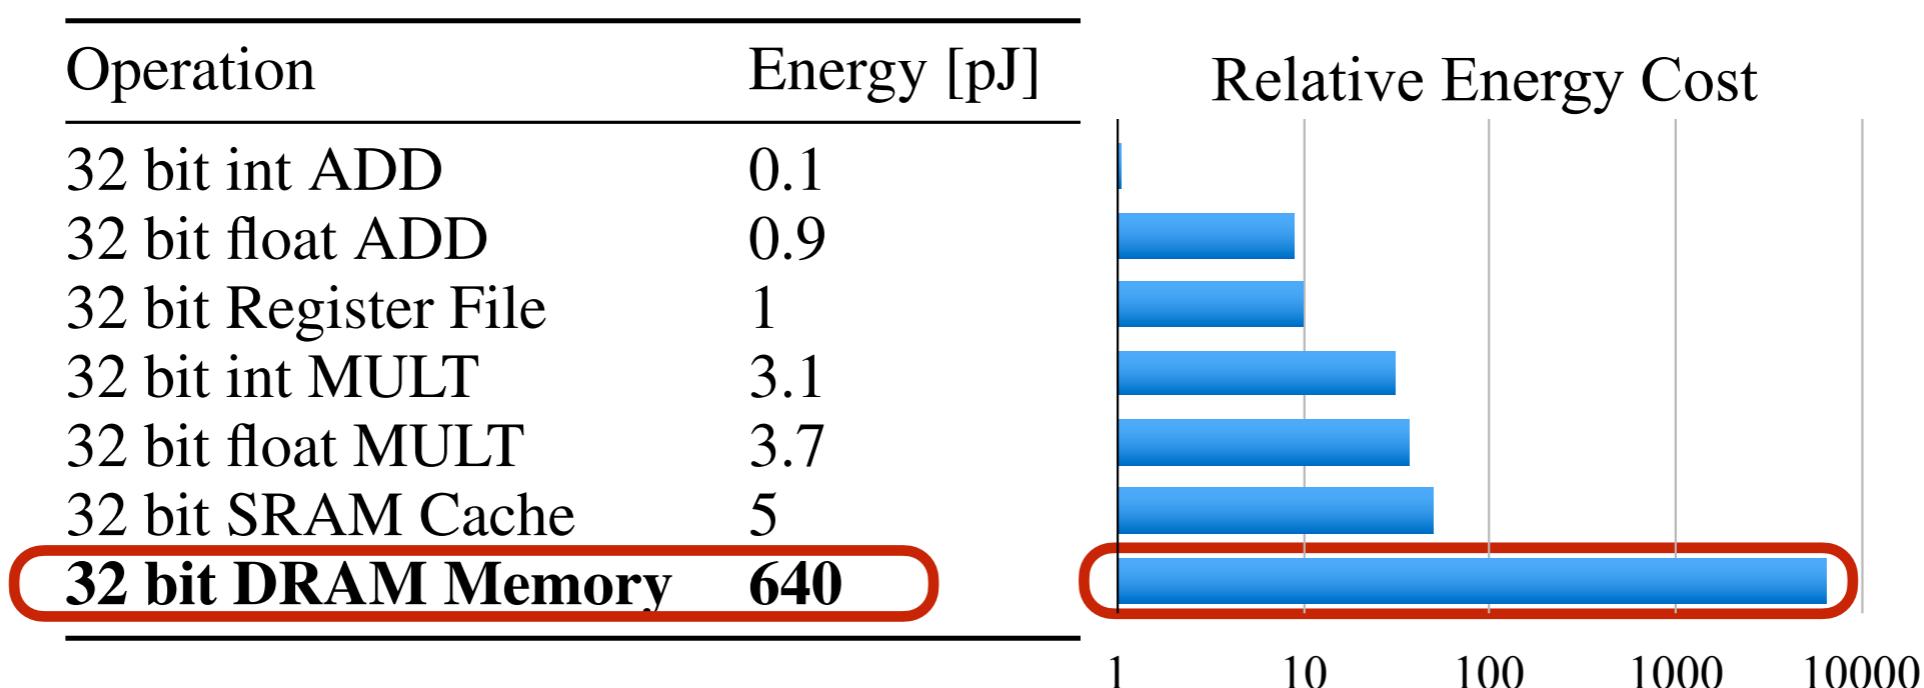


$$1 \text{ RAM module} = 1000 \times \text{Memory Accesses}$$

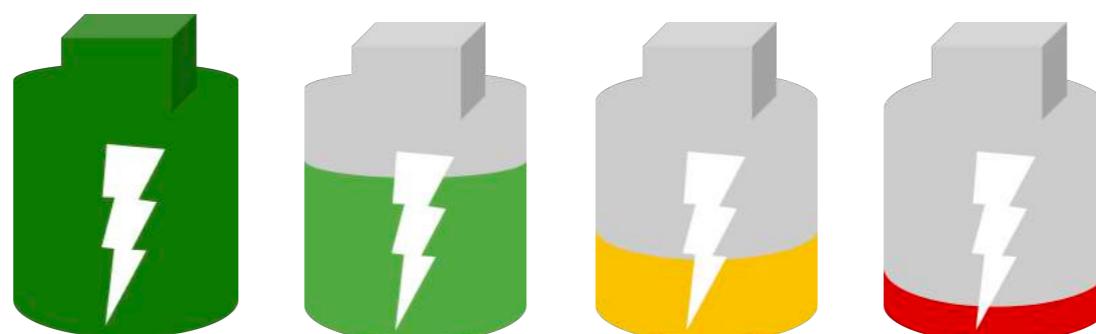
This image is in the public domain

# Where is the Energy Consumed?

larger model => more memory reference => more energy



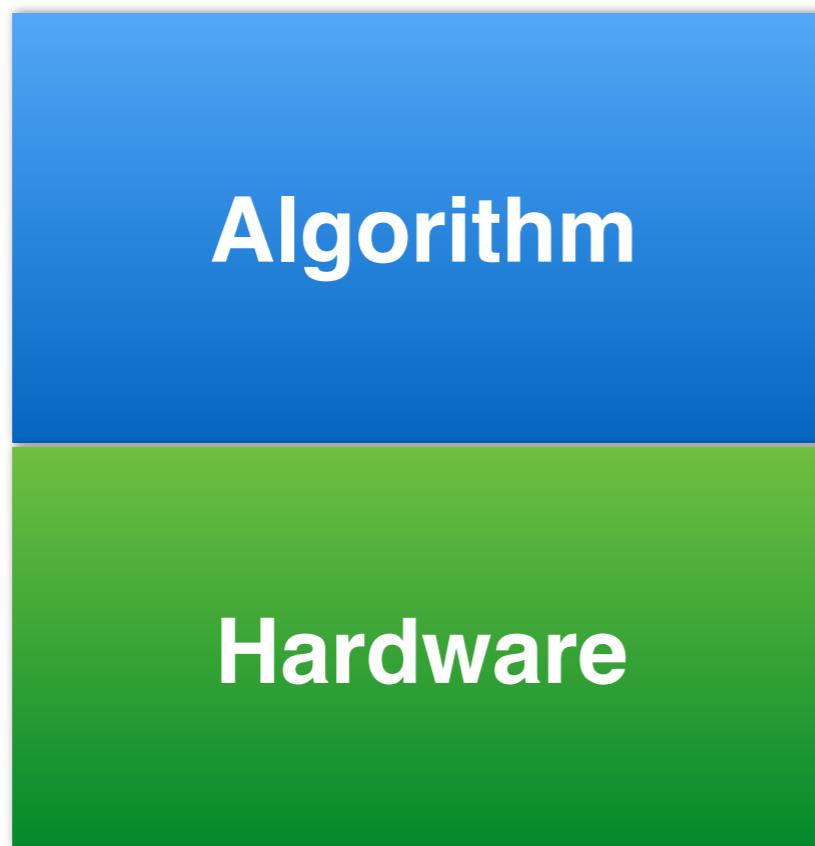
how to make deep learning more efficient?



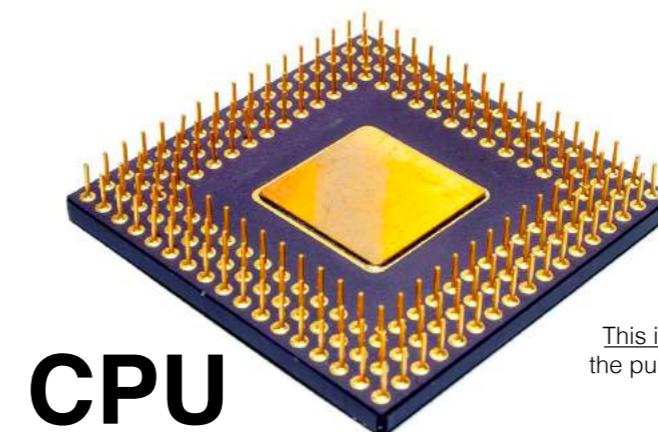
Battery images are in the public domain  
[Image 1](#), [Image 2](#), [Image 2](#), [Image 4](#)

# **Improve the Efficiency of Deep Learning by Algorithm-Hardware Co-Design**

# Application as a Black Box

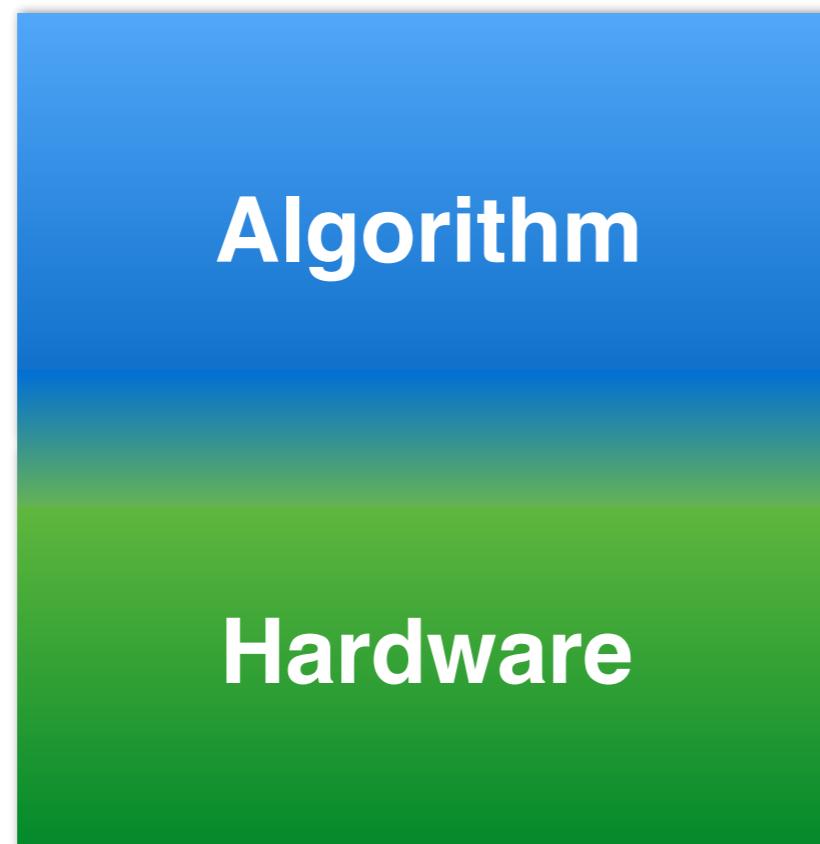


[This image](#) is in  
the public domain

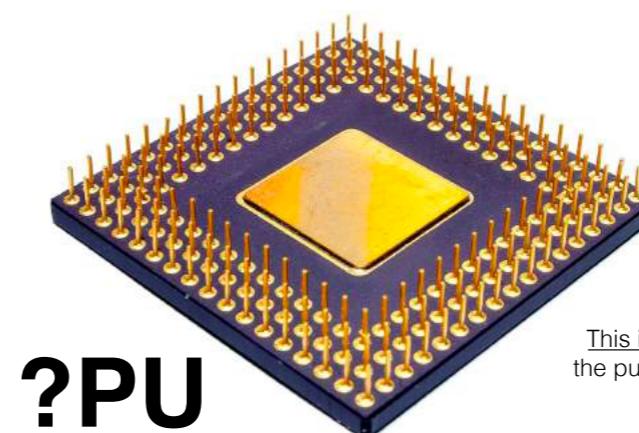


[This image](#) is in  
the public domain

# Open the Box before Hardware Design



This image is in  
the public domain



This image is in  
the public domain

Breaks the boundary between algorithm and hardware

# Agenda

Inference

Training

# Agenda

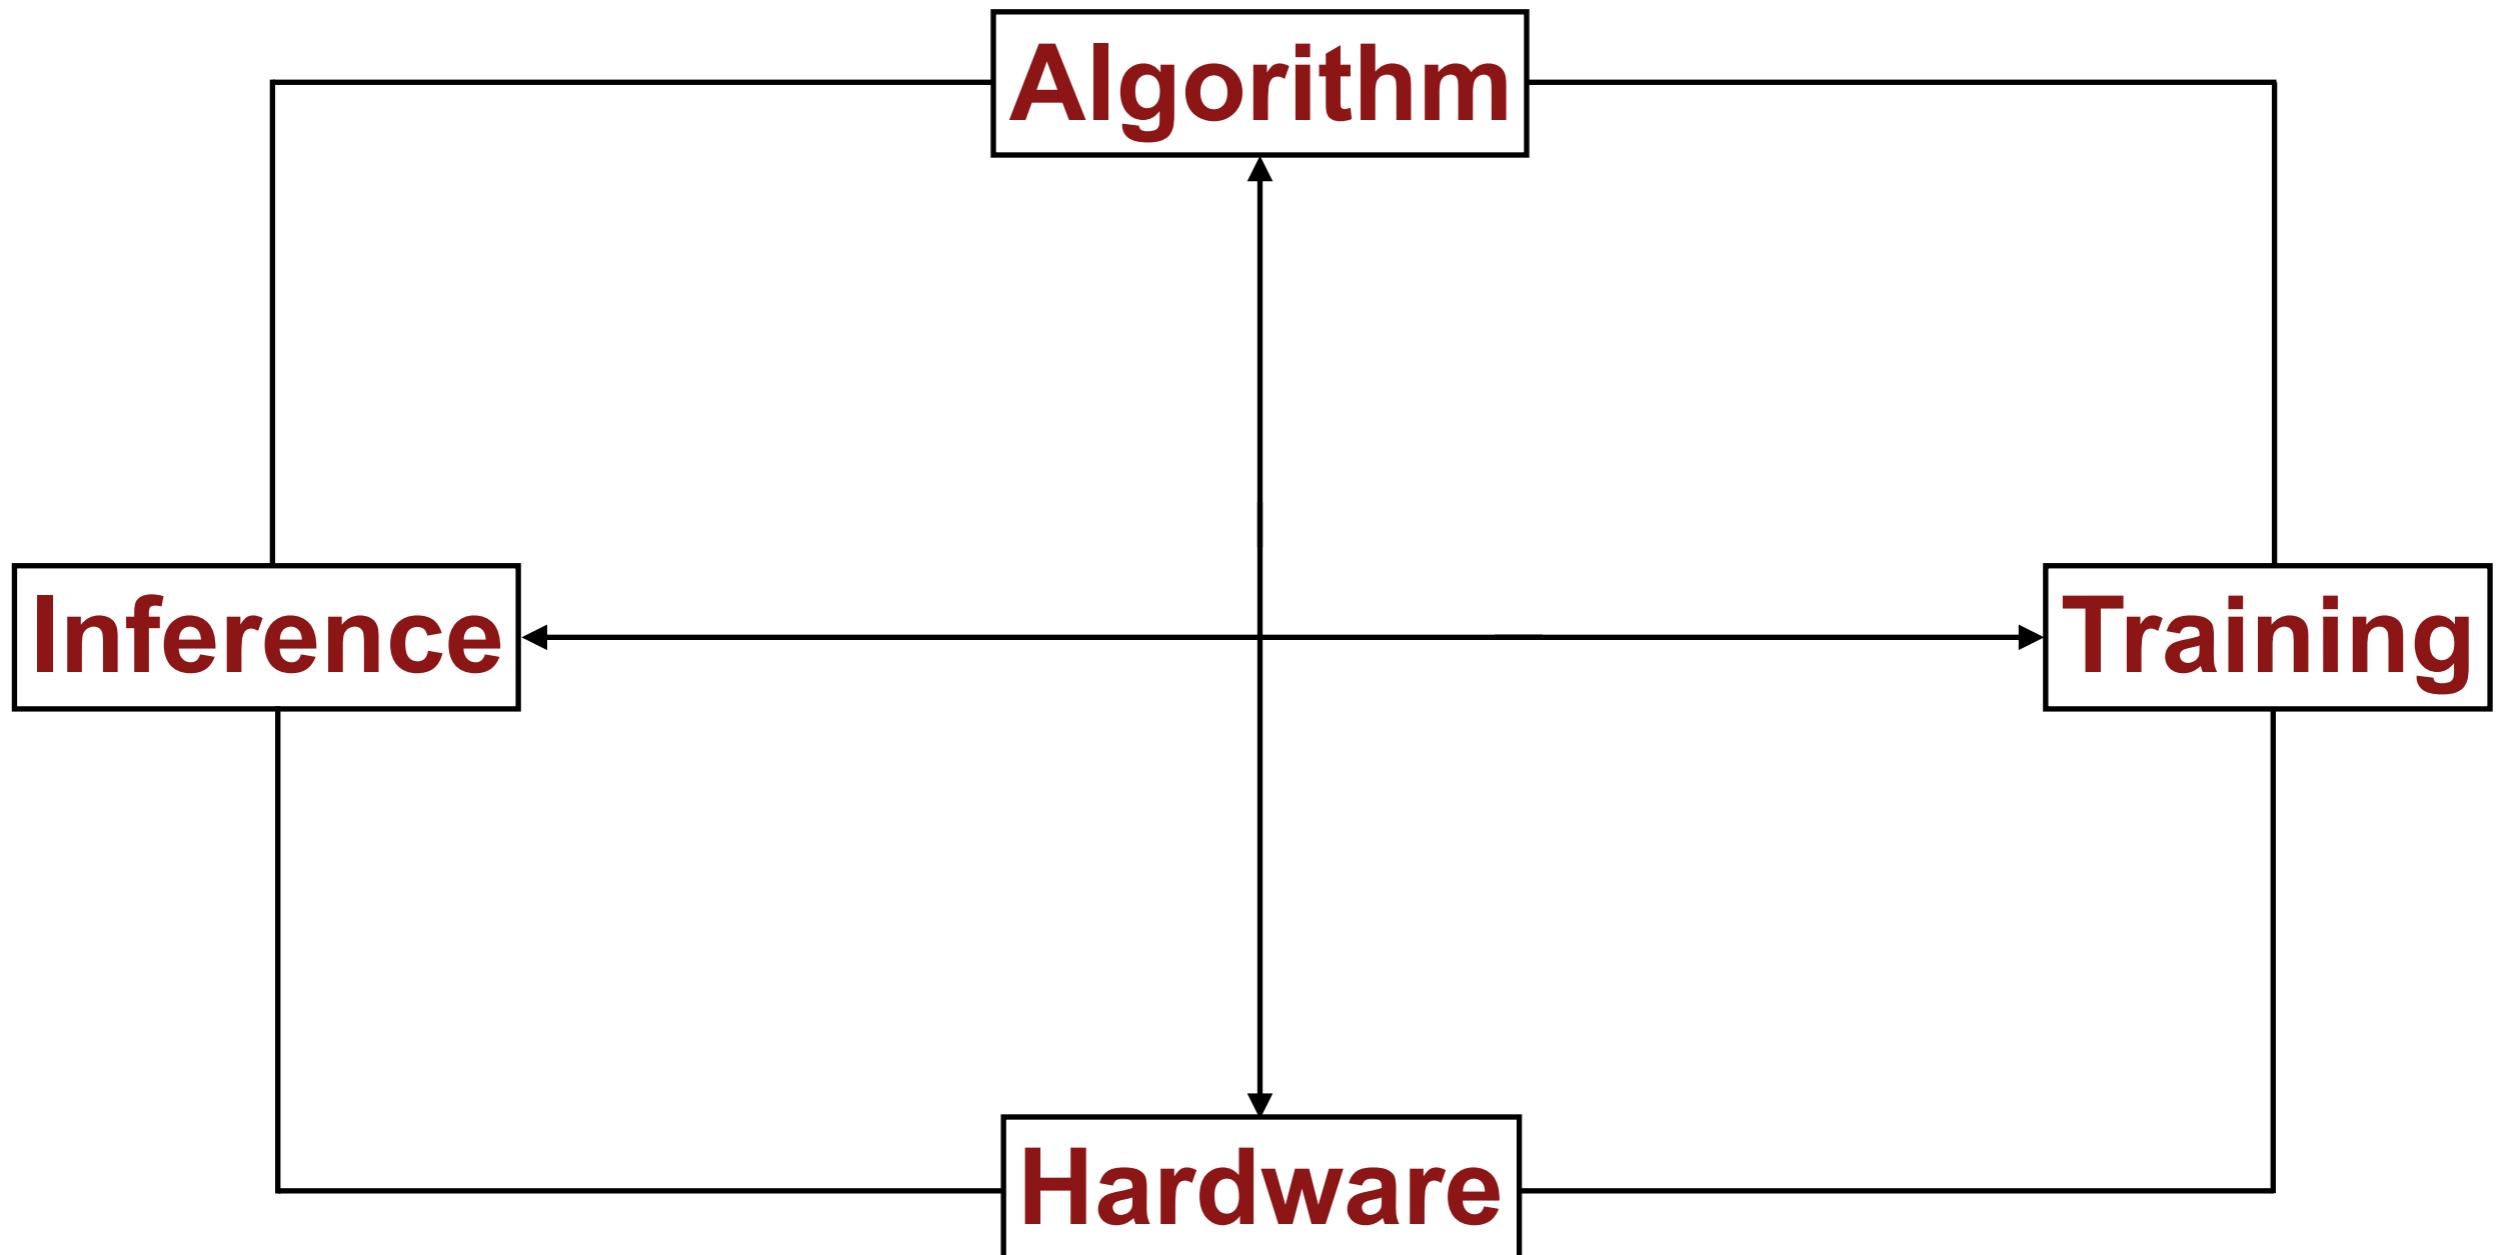
Algorithm

Inference

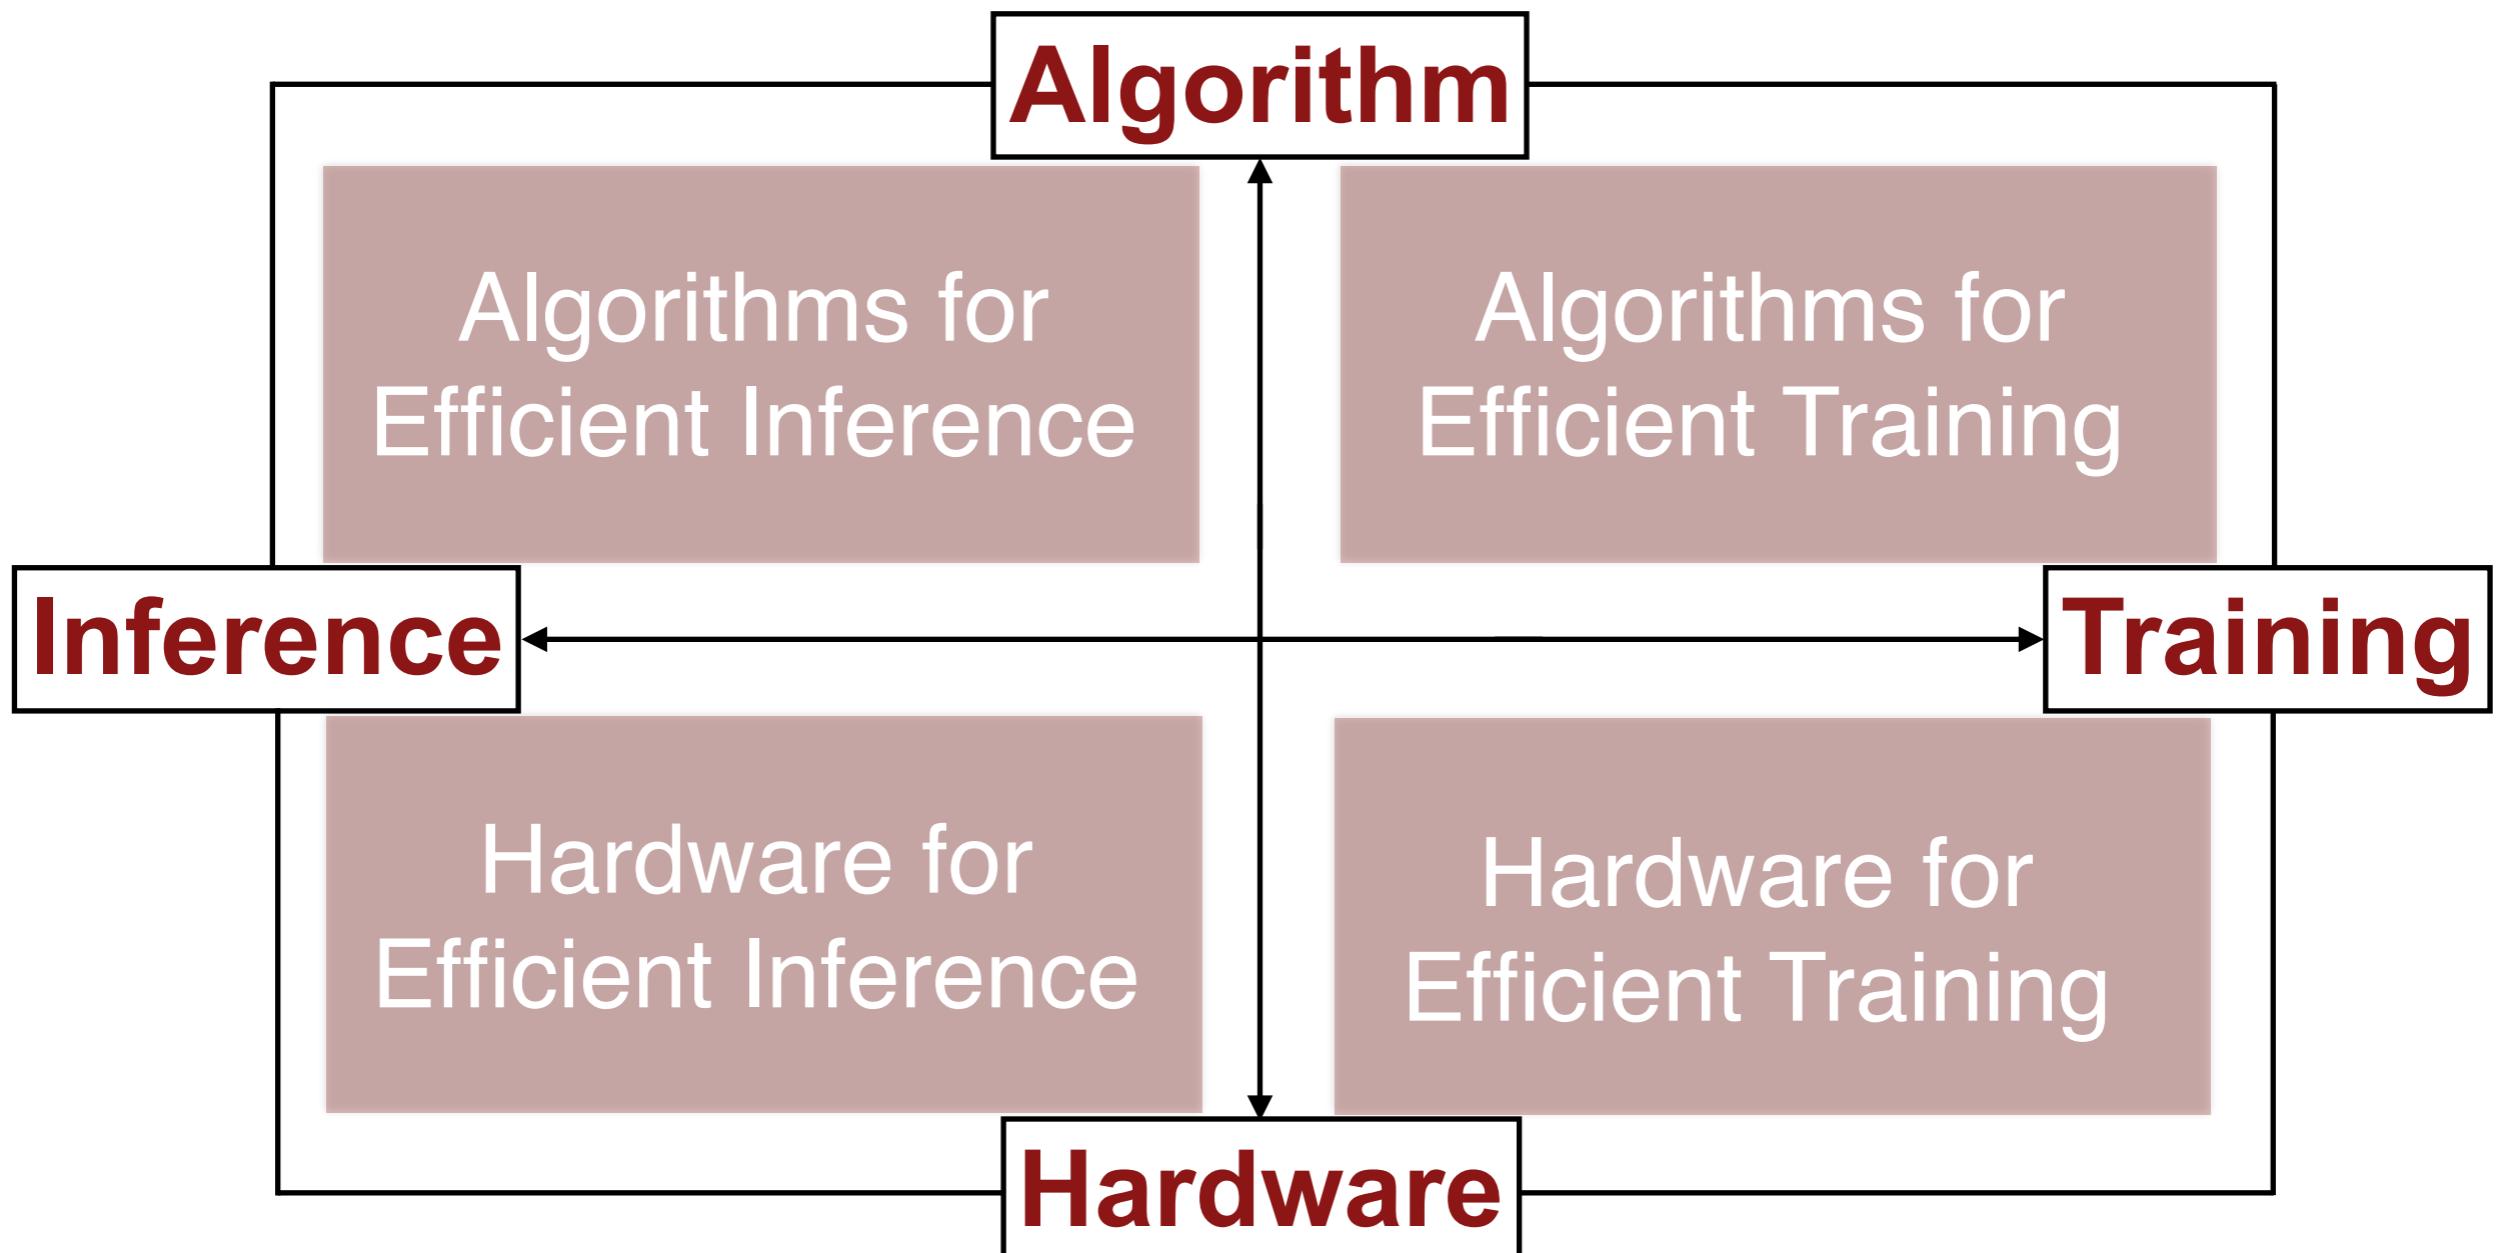
Training

Hardware

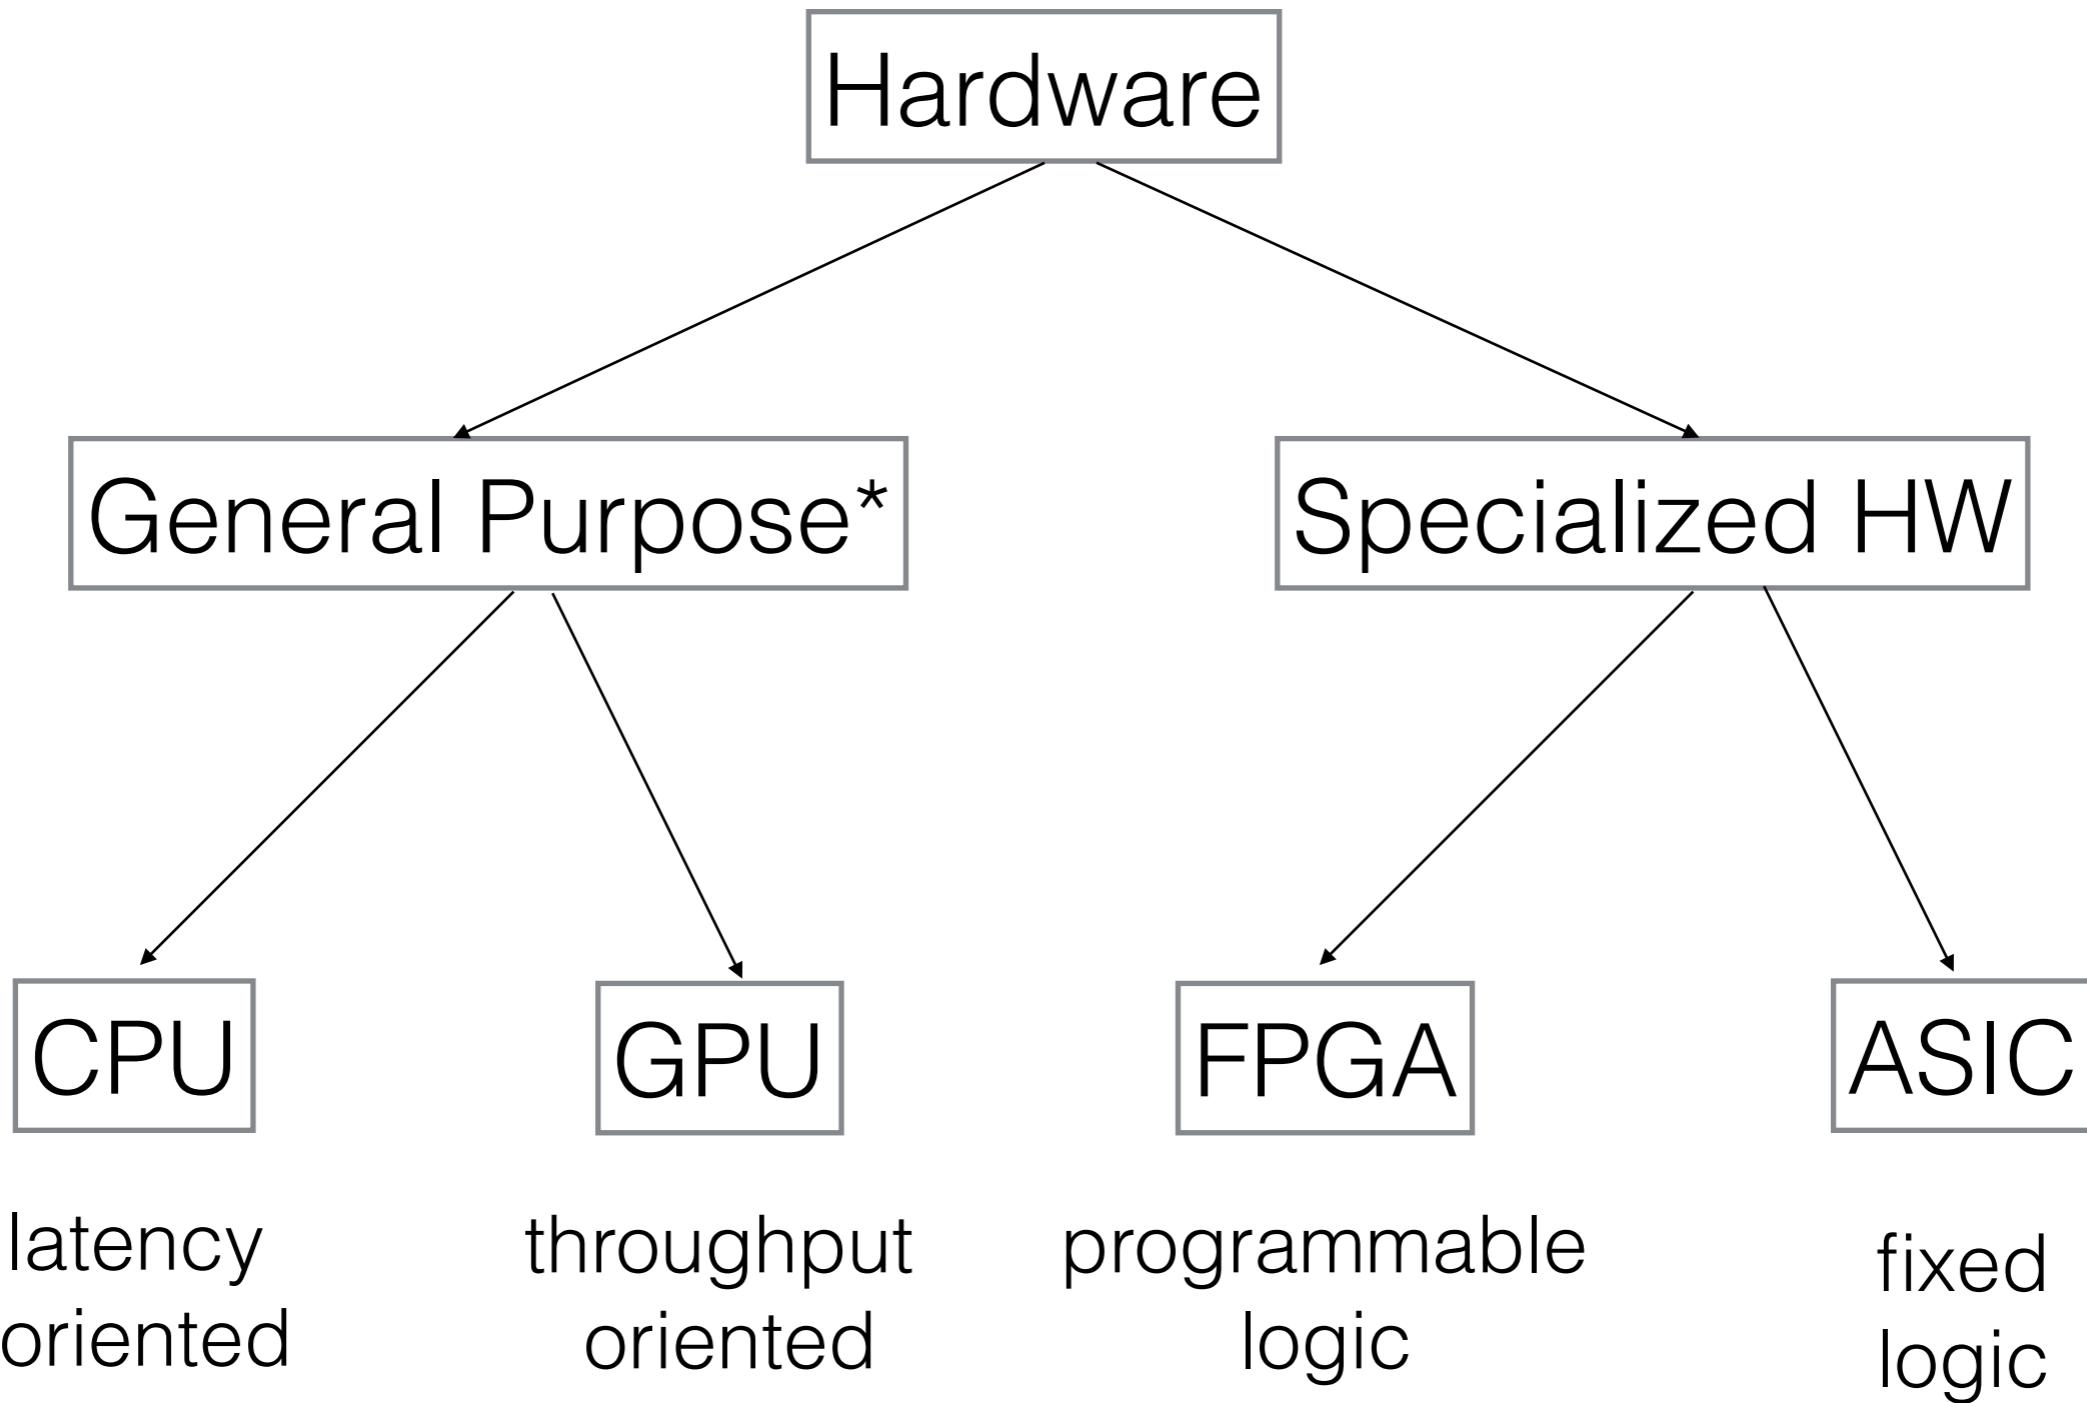
# Agenda



# Agenda



# Hardware 101: the Family



\* including GPGPU

# Hardware 101: Number Representation

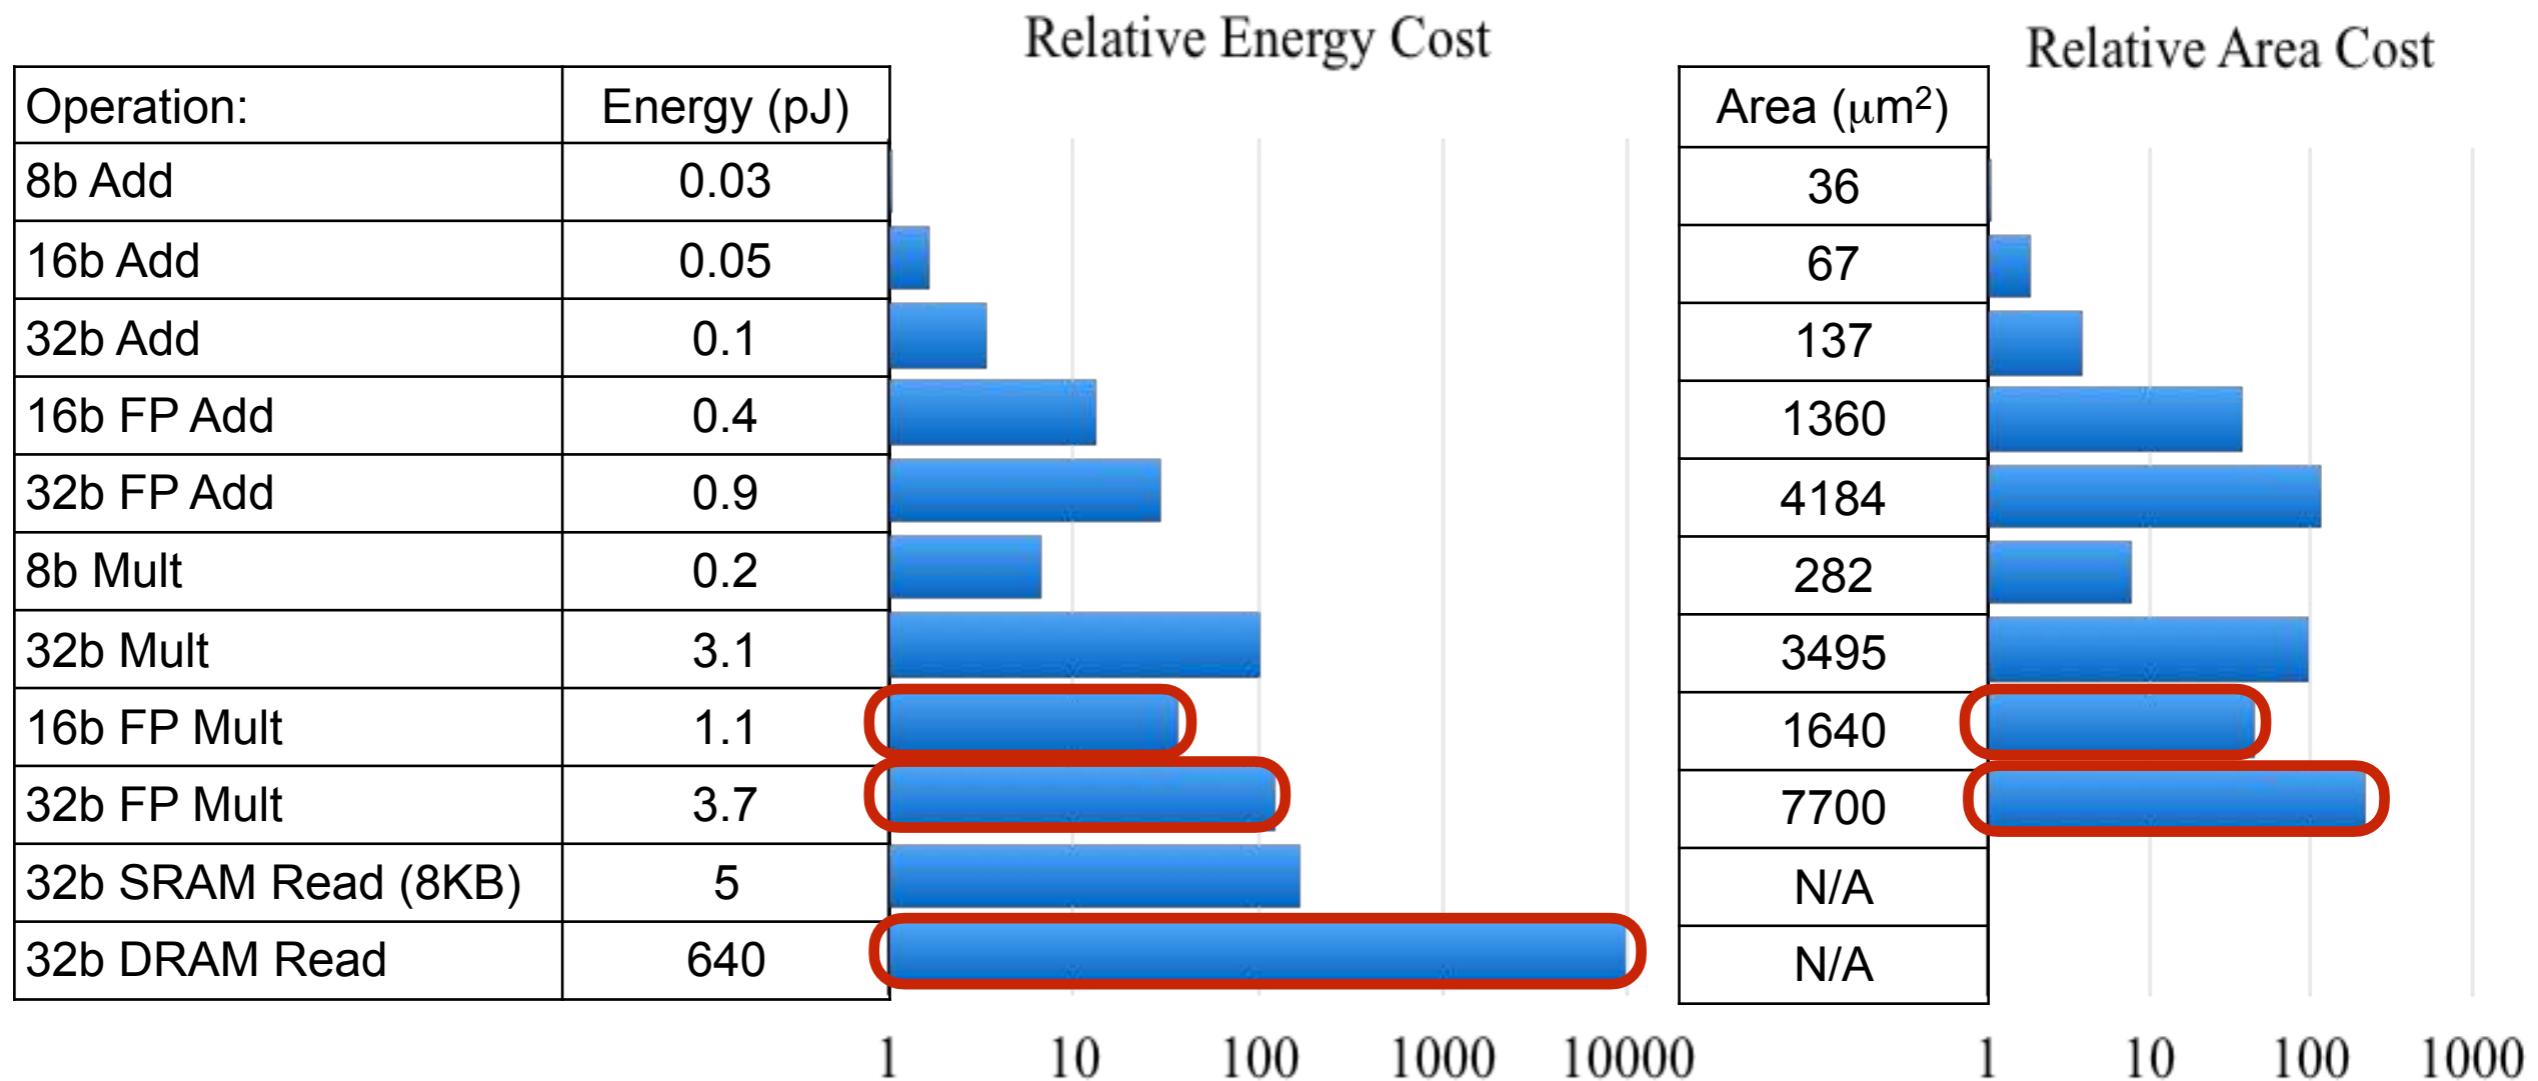
$$(-1)^S \times (1.M) \times 2^E$$

	S	E	M	Range	Accuracy
FP32	1	8	23	$10^{-38} - 10^{38}$	.000006%
FP16	1	5	10	$6 \times 10^{-5} - 6 \times 10^4$	.05%
Int32	1		31	$0 - 2 \times 10^9$	$\frac{1}{2}$
Int16	1		15	$0 - 6 \times 10^4$	$\frac{1}{2}$
Int8	1		7	$0 - 127$	$\frac{1}{2}$
Fixed point	S	I	F	-	-

radix point ↑

Dally, High Performance Hardware for Machine Learning, NIPS'2015

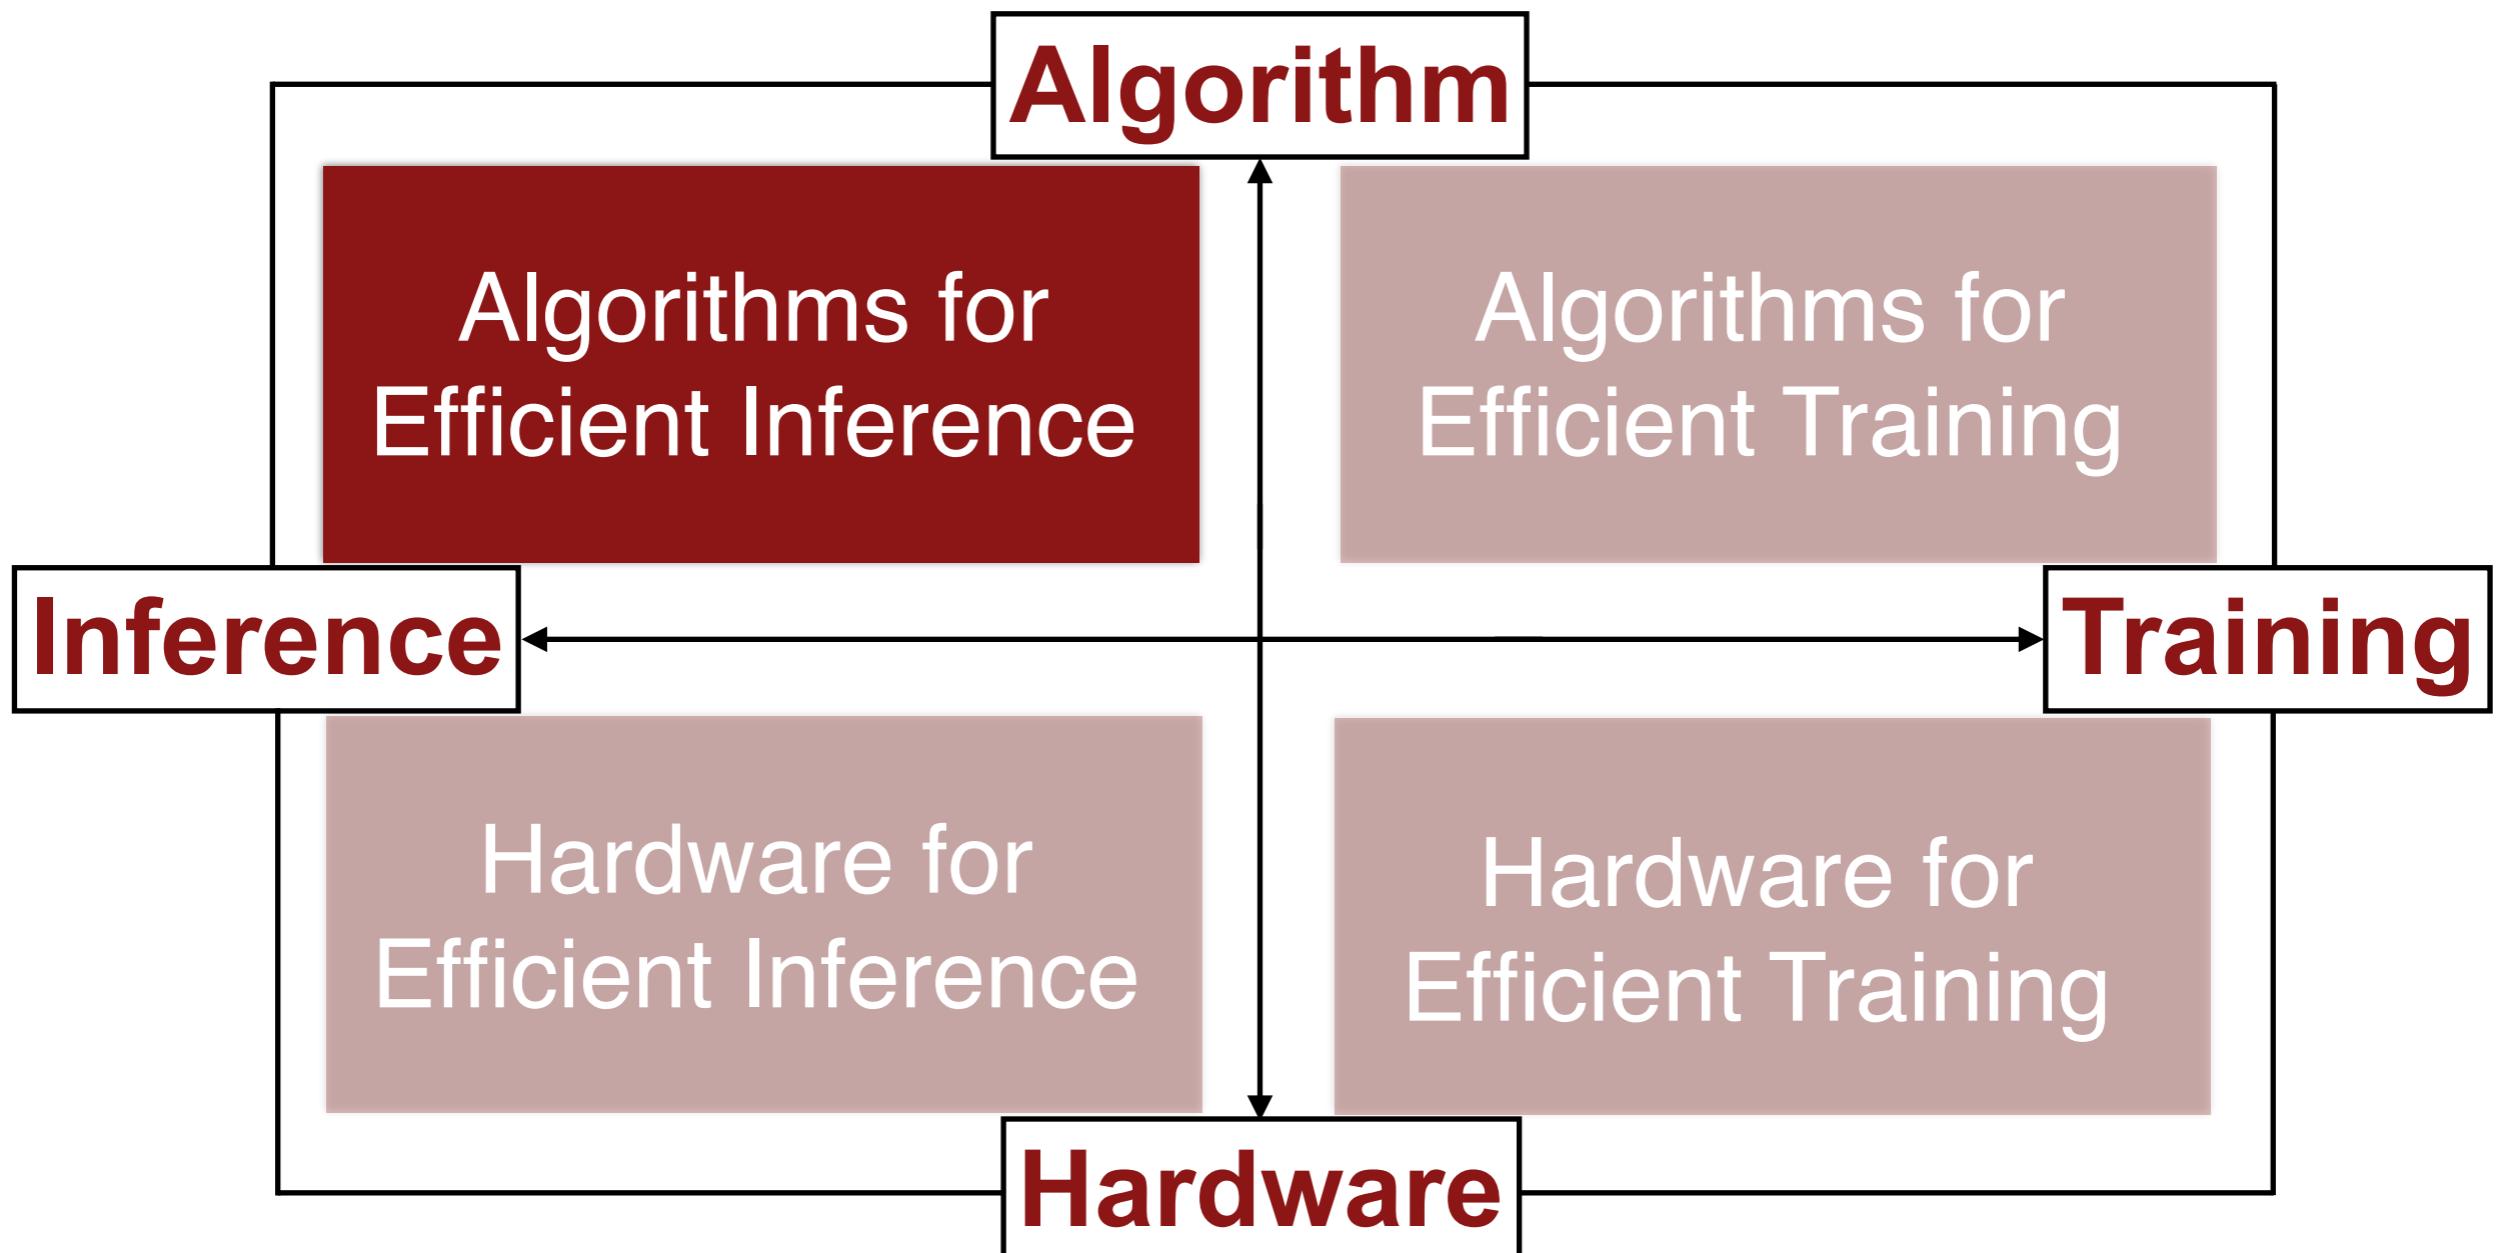
# Hardware 101: Number Representation



Energy numbers are from Mark Horowitz “Computing’s Energy Problem (and what we can do about it)”, ISSCC 2014

Area numbers are from synthesized result using Design Compiler under TSMC 45nm tech node. FP units used DesignWare Library.

# Agenda



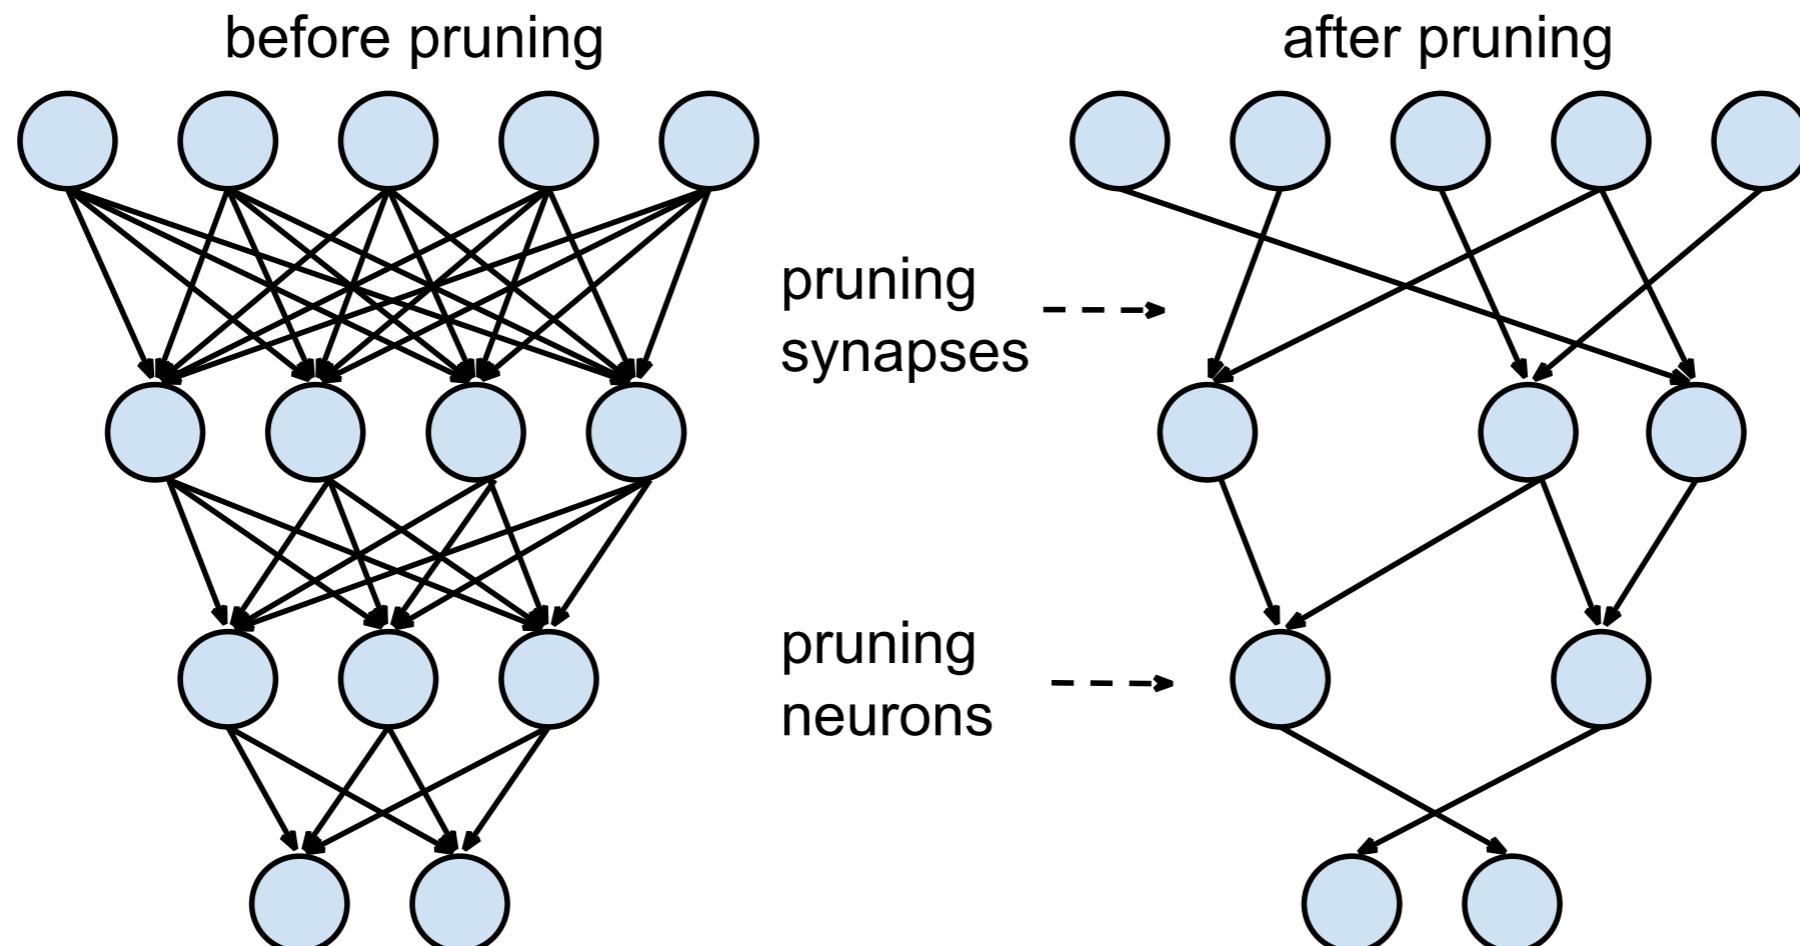
# Part 1: Algorithms for Efficient Inference

- 1. Pruning
- 2. Weight Sharing
- 3. Quantization
- 4. Low Rank Approximation
- 5. Binary / Ternary Net
- 6. Winograd Transformation

# Part 1: Algorithms for Efficient Inference

- 1. Pruning
- 2. Weight Sharing
- 3. Quantization
- 4. Low Rank Approximation
- 5. Binary / Ternary Net
- 6. Winograd Transformation

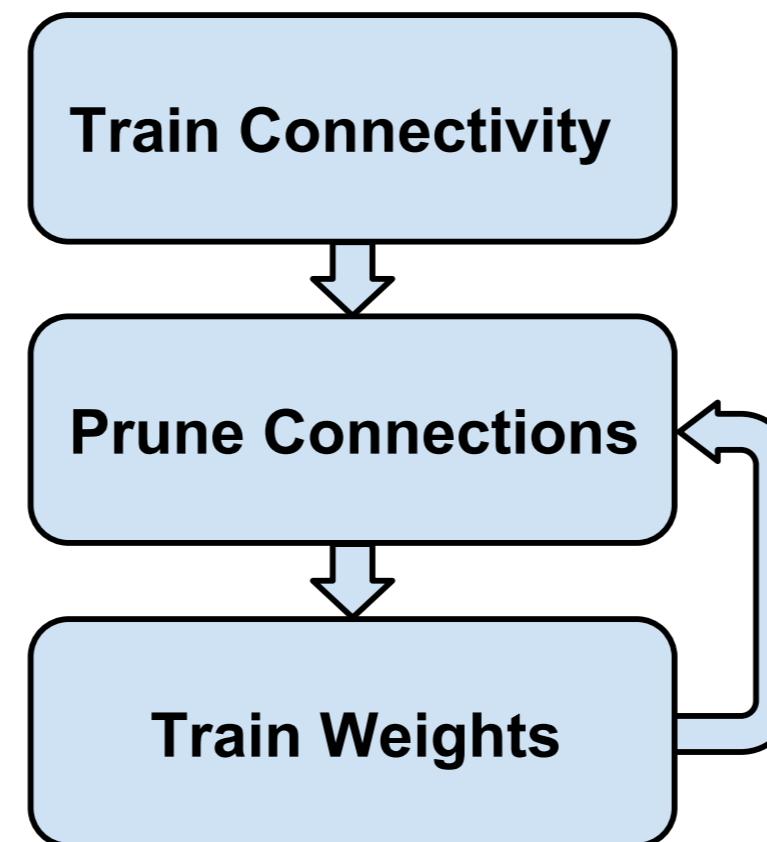
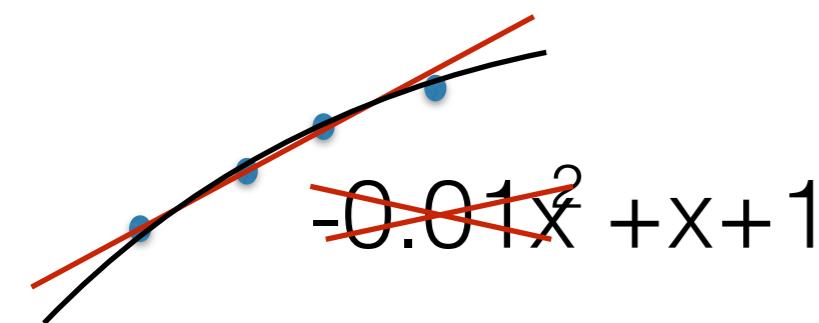
# Pruning Neural Networks



[Lecun et al. NIPS'89]

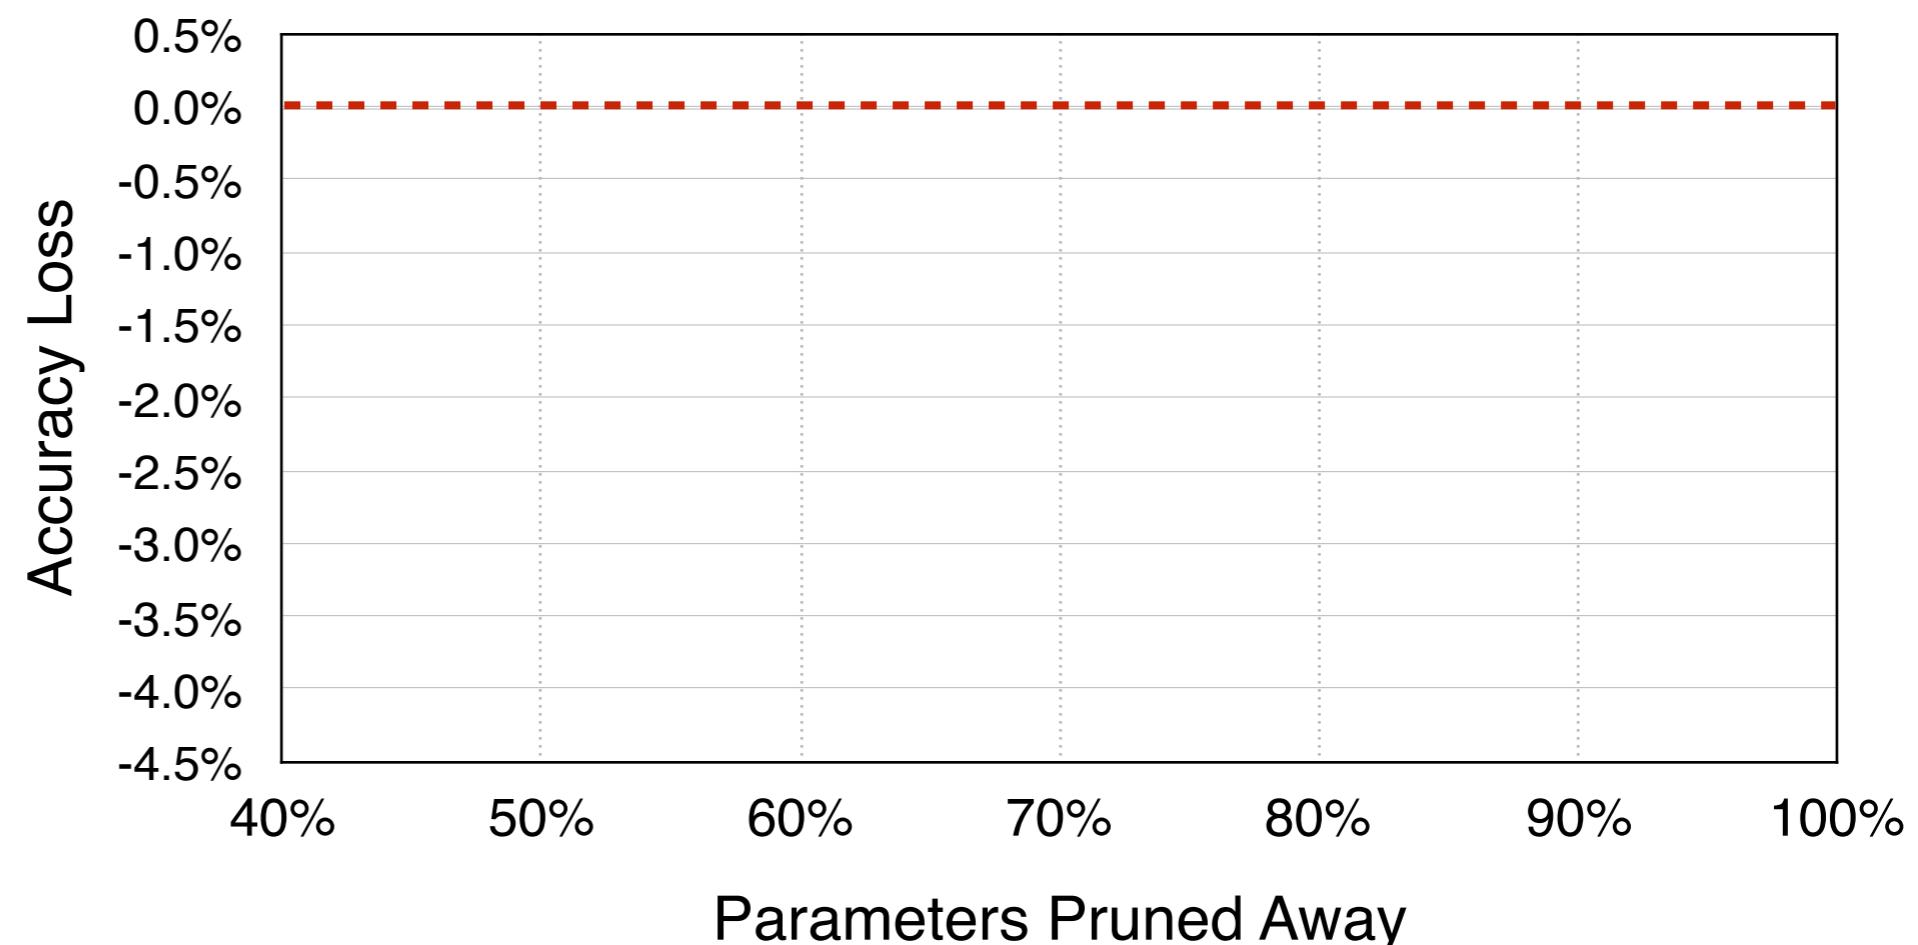
[Han et al. NIPS'15]

# Pruning Neural Networks

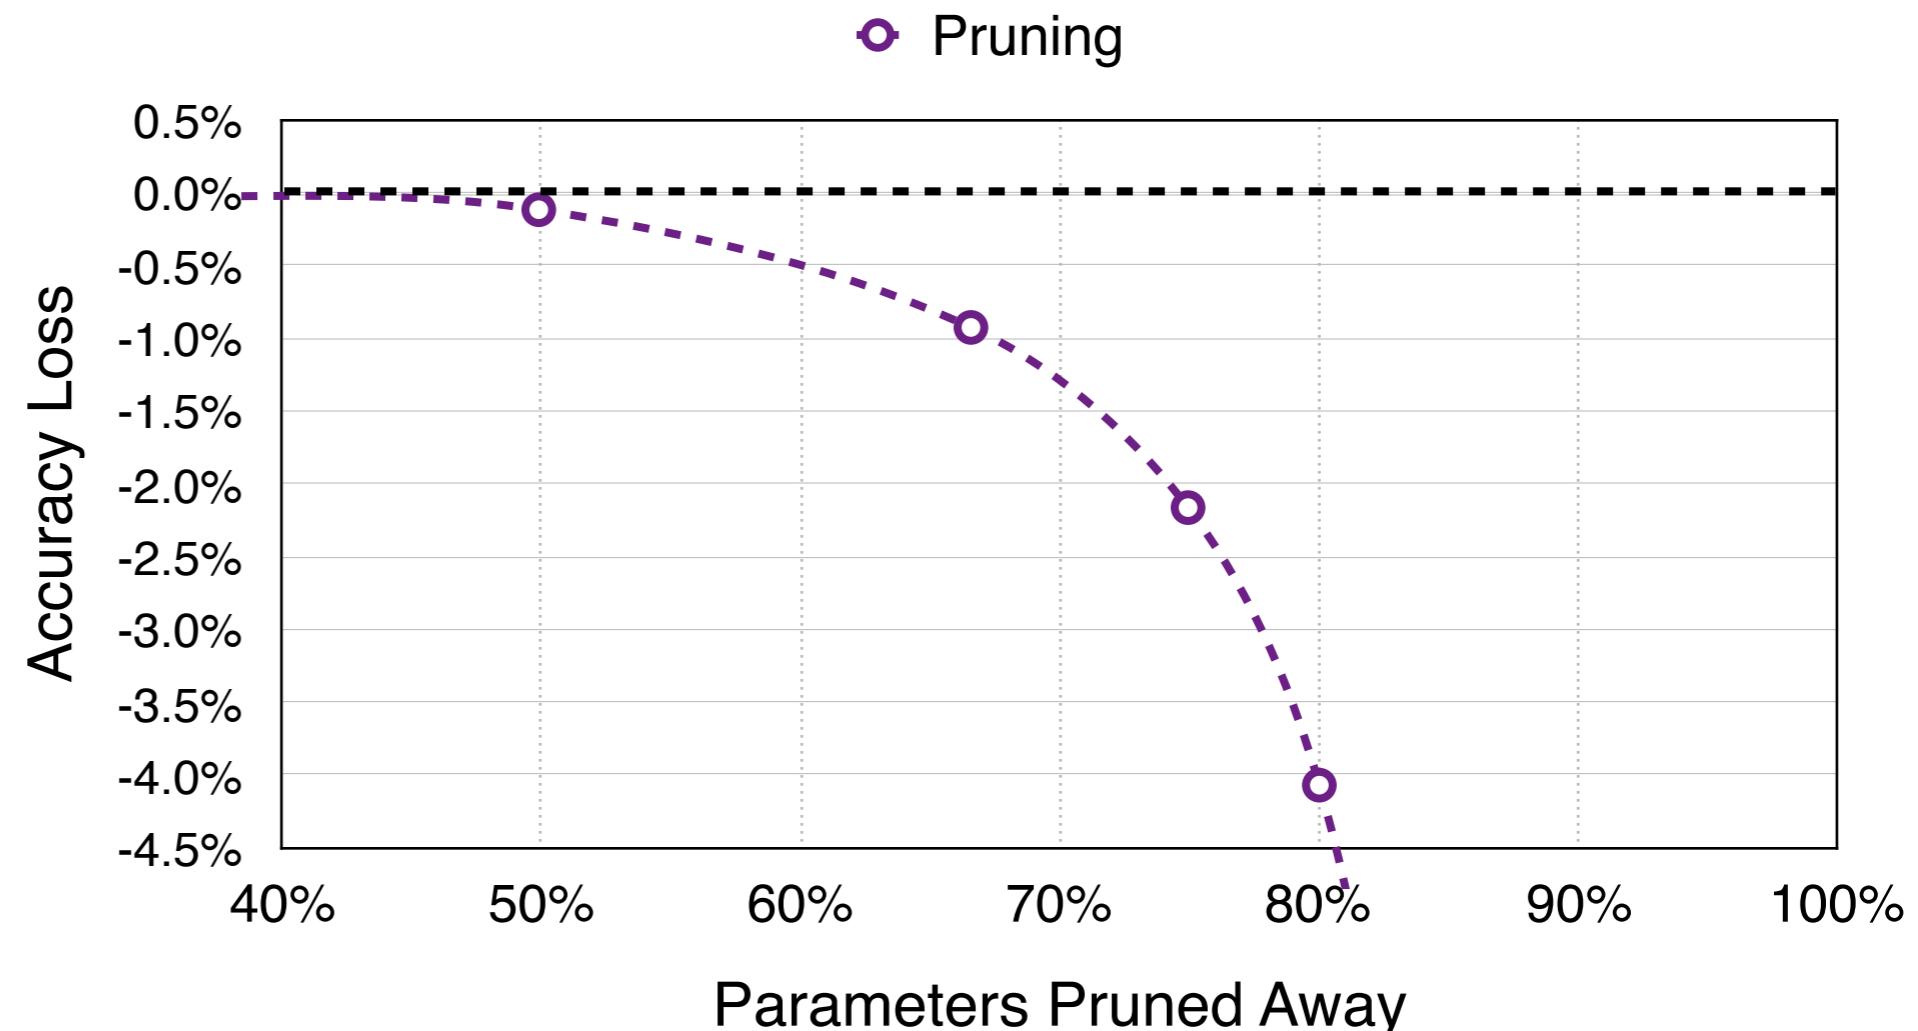
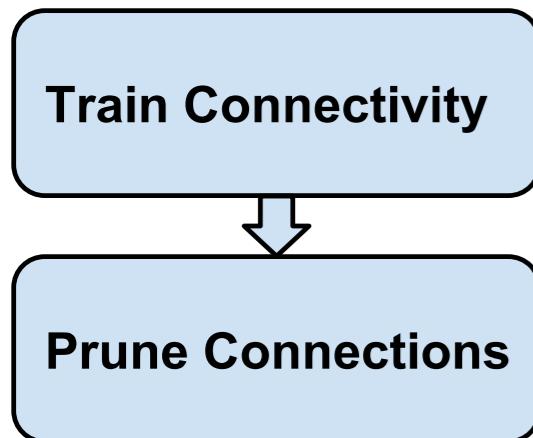


# Pruning Neural Networks

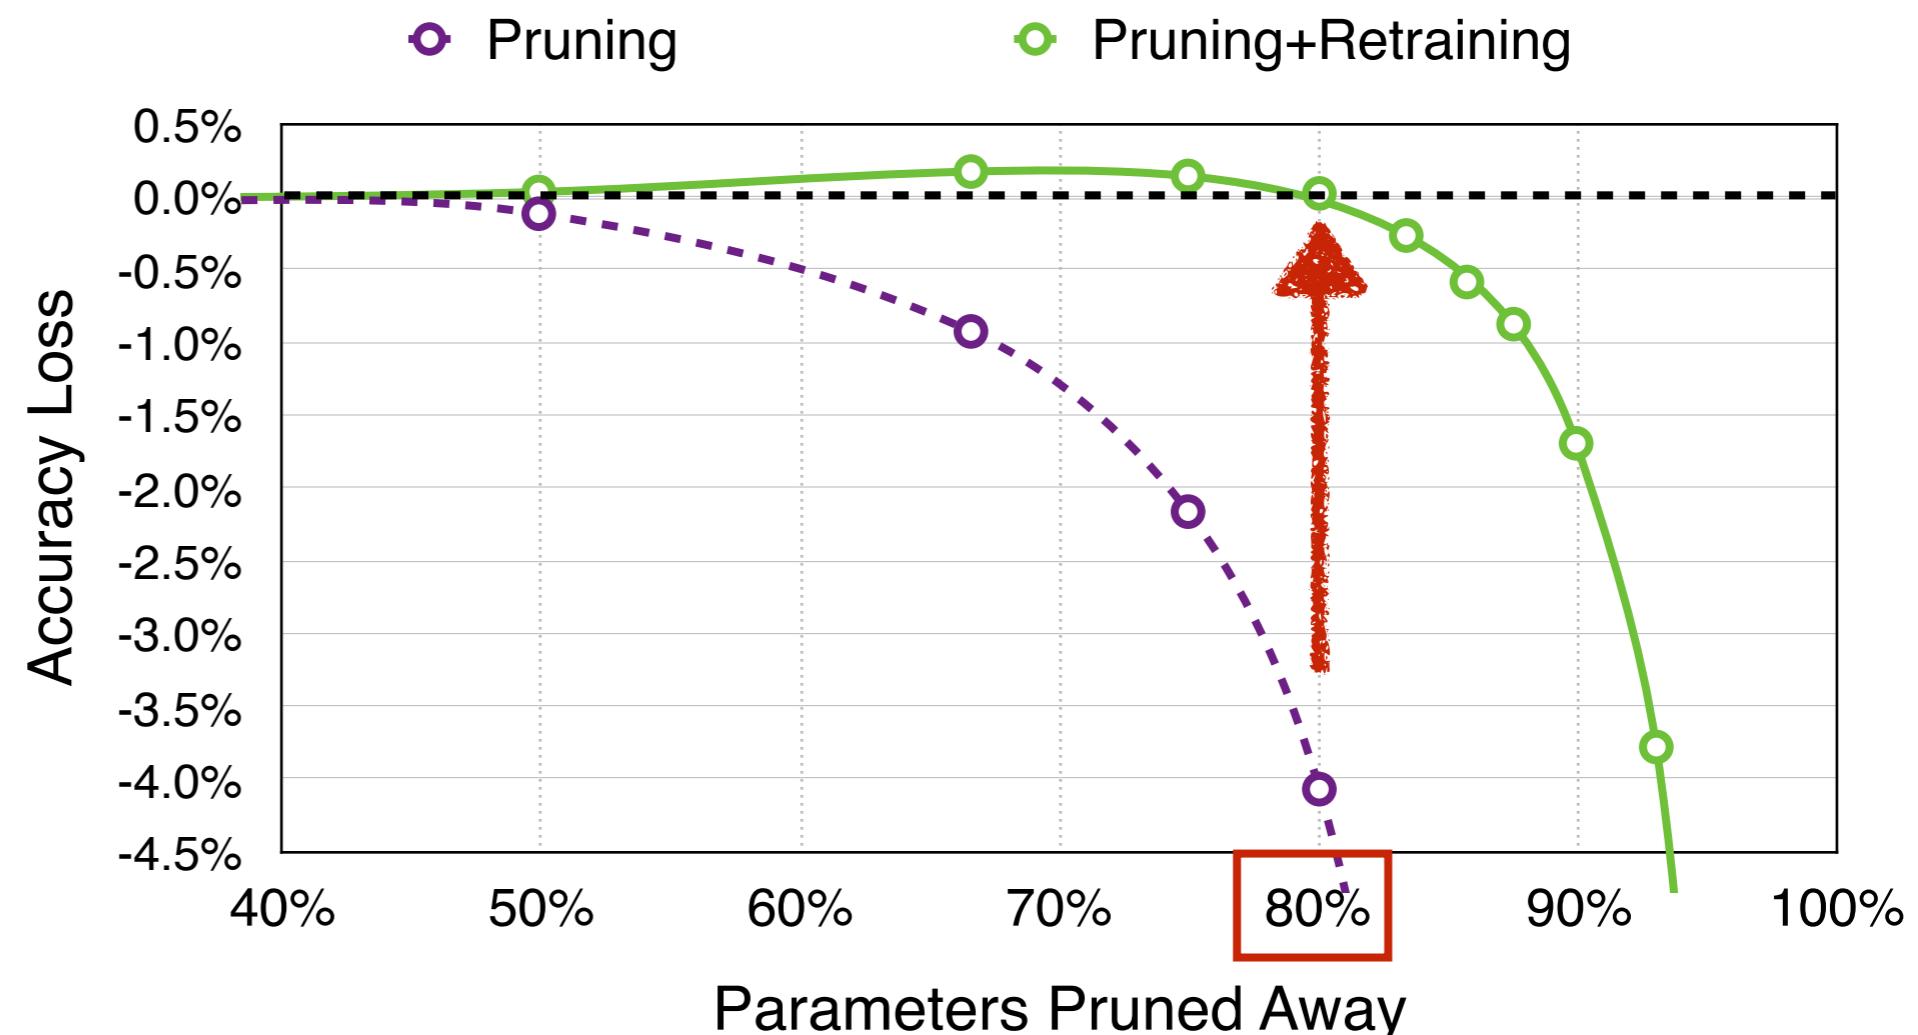
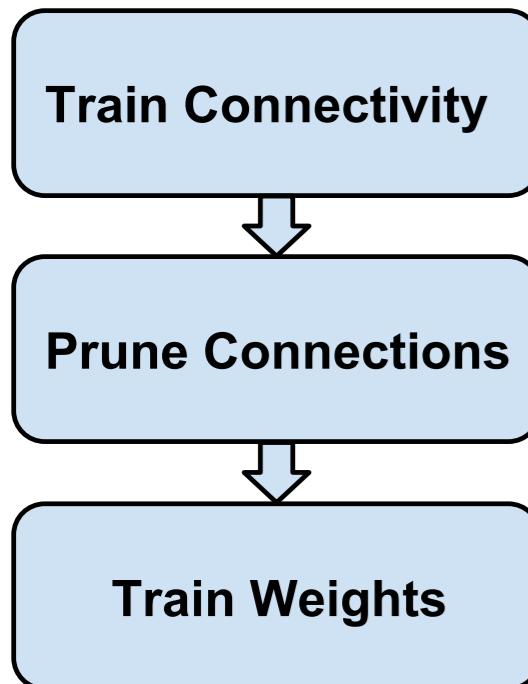
Train Connectivity



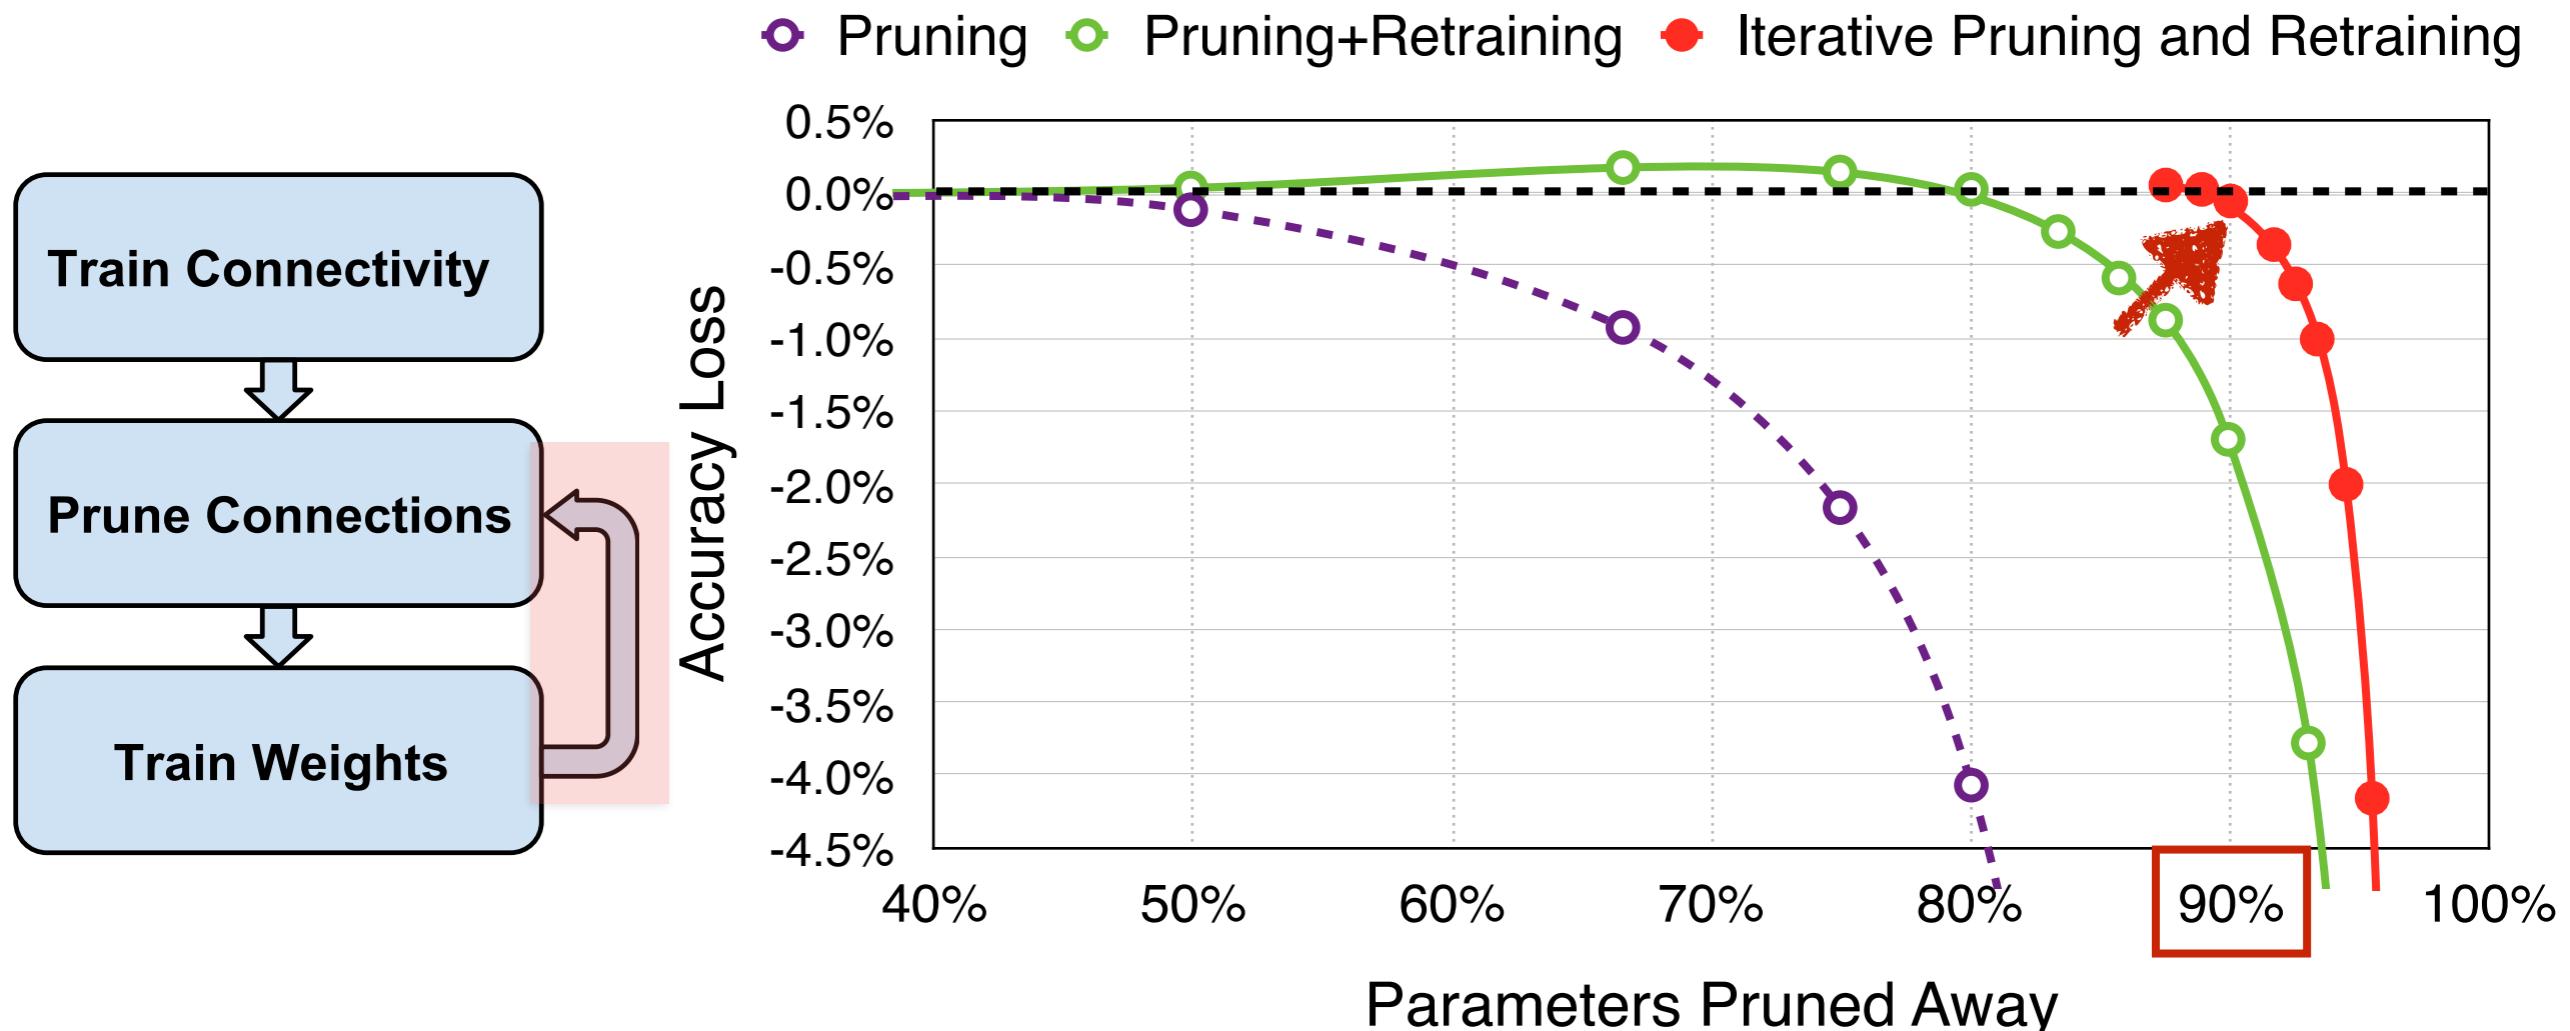
# Pruning Neural Networks



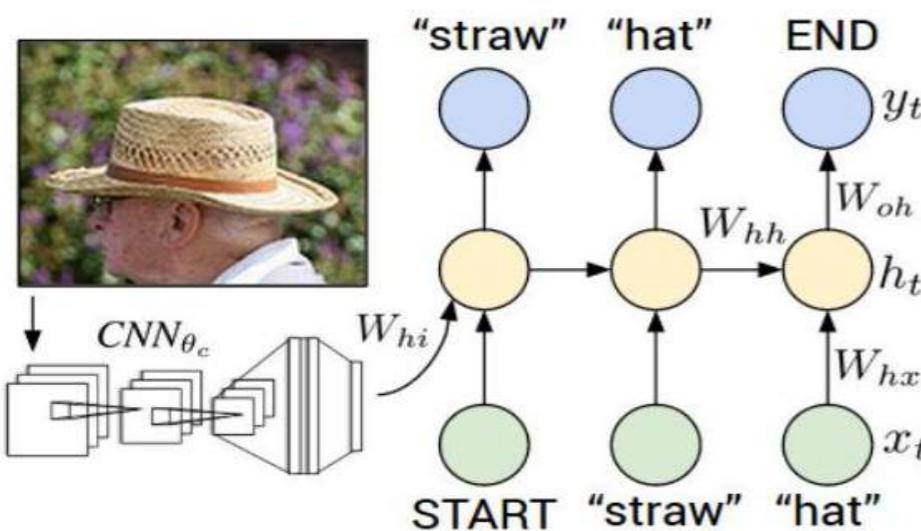
# Retrain to Recover Accuracy



# Iteratively Retrain to Recover Accuracy

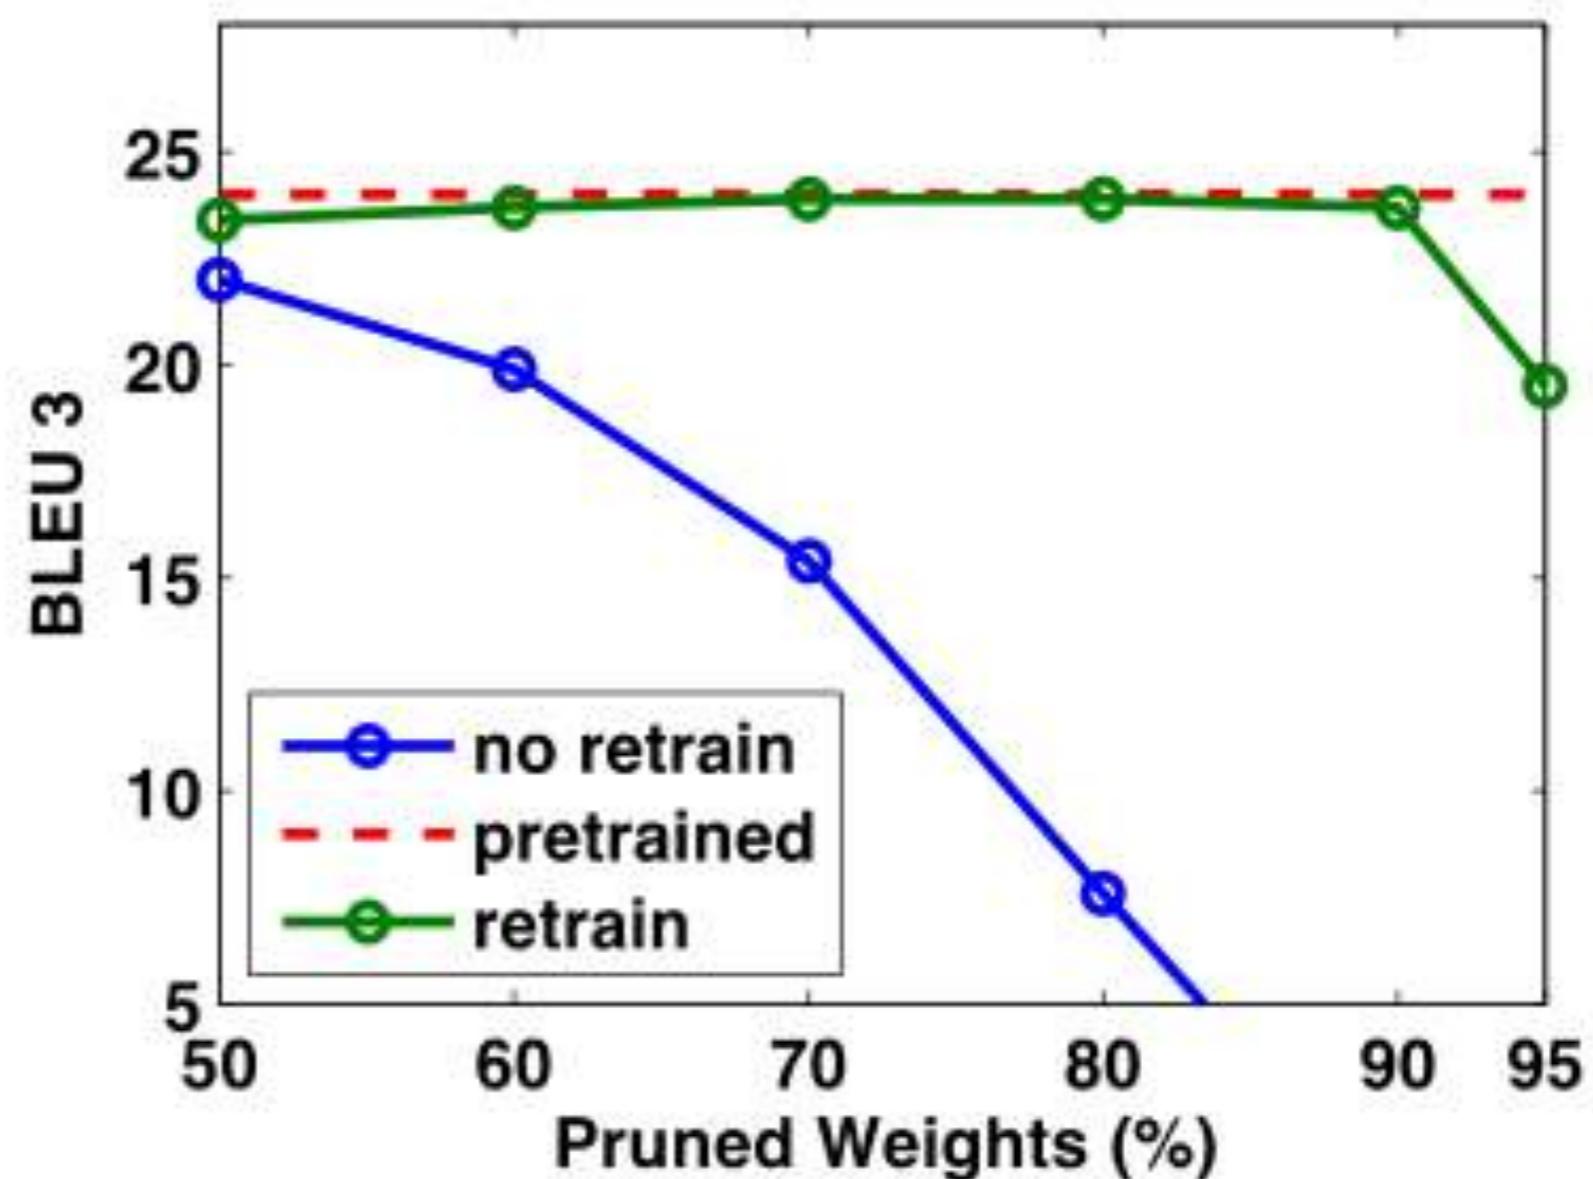


# Pruning RNN and LSTM



\*Karpathy et al, "Deep Visual-Semantic Alignments for Generating Image Descriptions", 2015.

Figure copyright IEEE, 2015; reproduced for educational purposes.



# Pruning RNN and LSTM

90%



- **Original:** a basketball player in a white uniform is playing with a **ball**
- **Pruned 90%:** a basketball player in a white uniform is playing with **a basketball**

90%



- **Original :** a brown dog is running through a grassy **field**
- **Pruned 90%:** a brown dog is running through a grassy **area**

90%



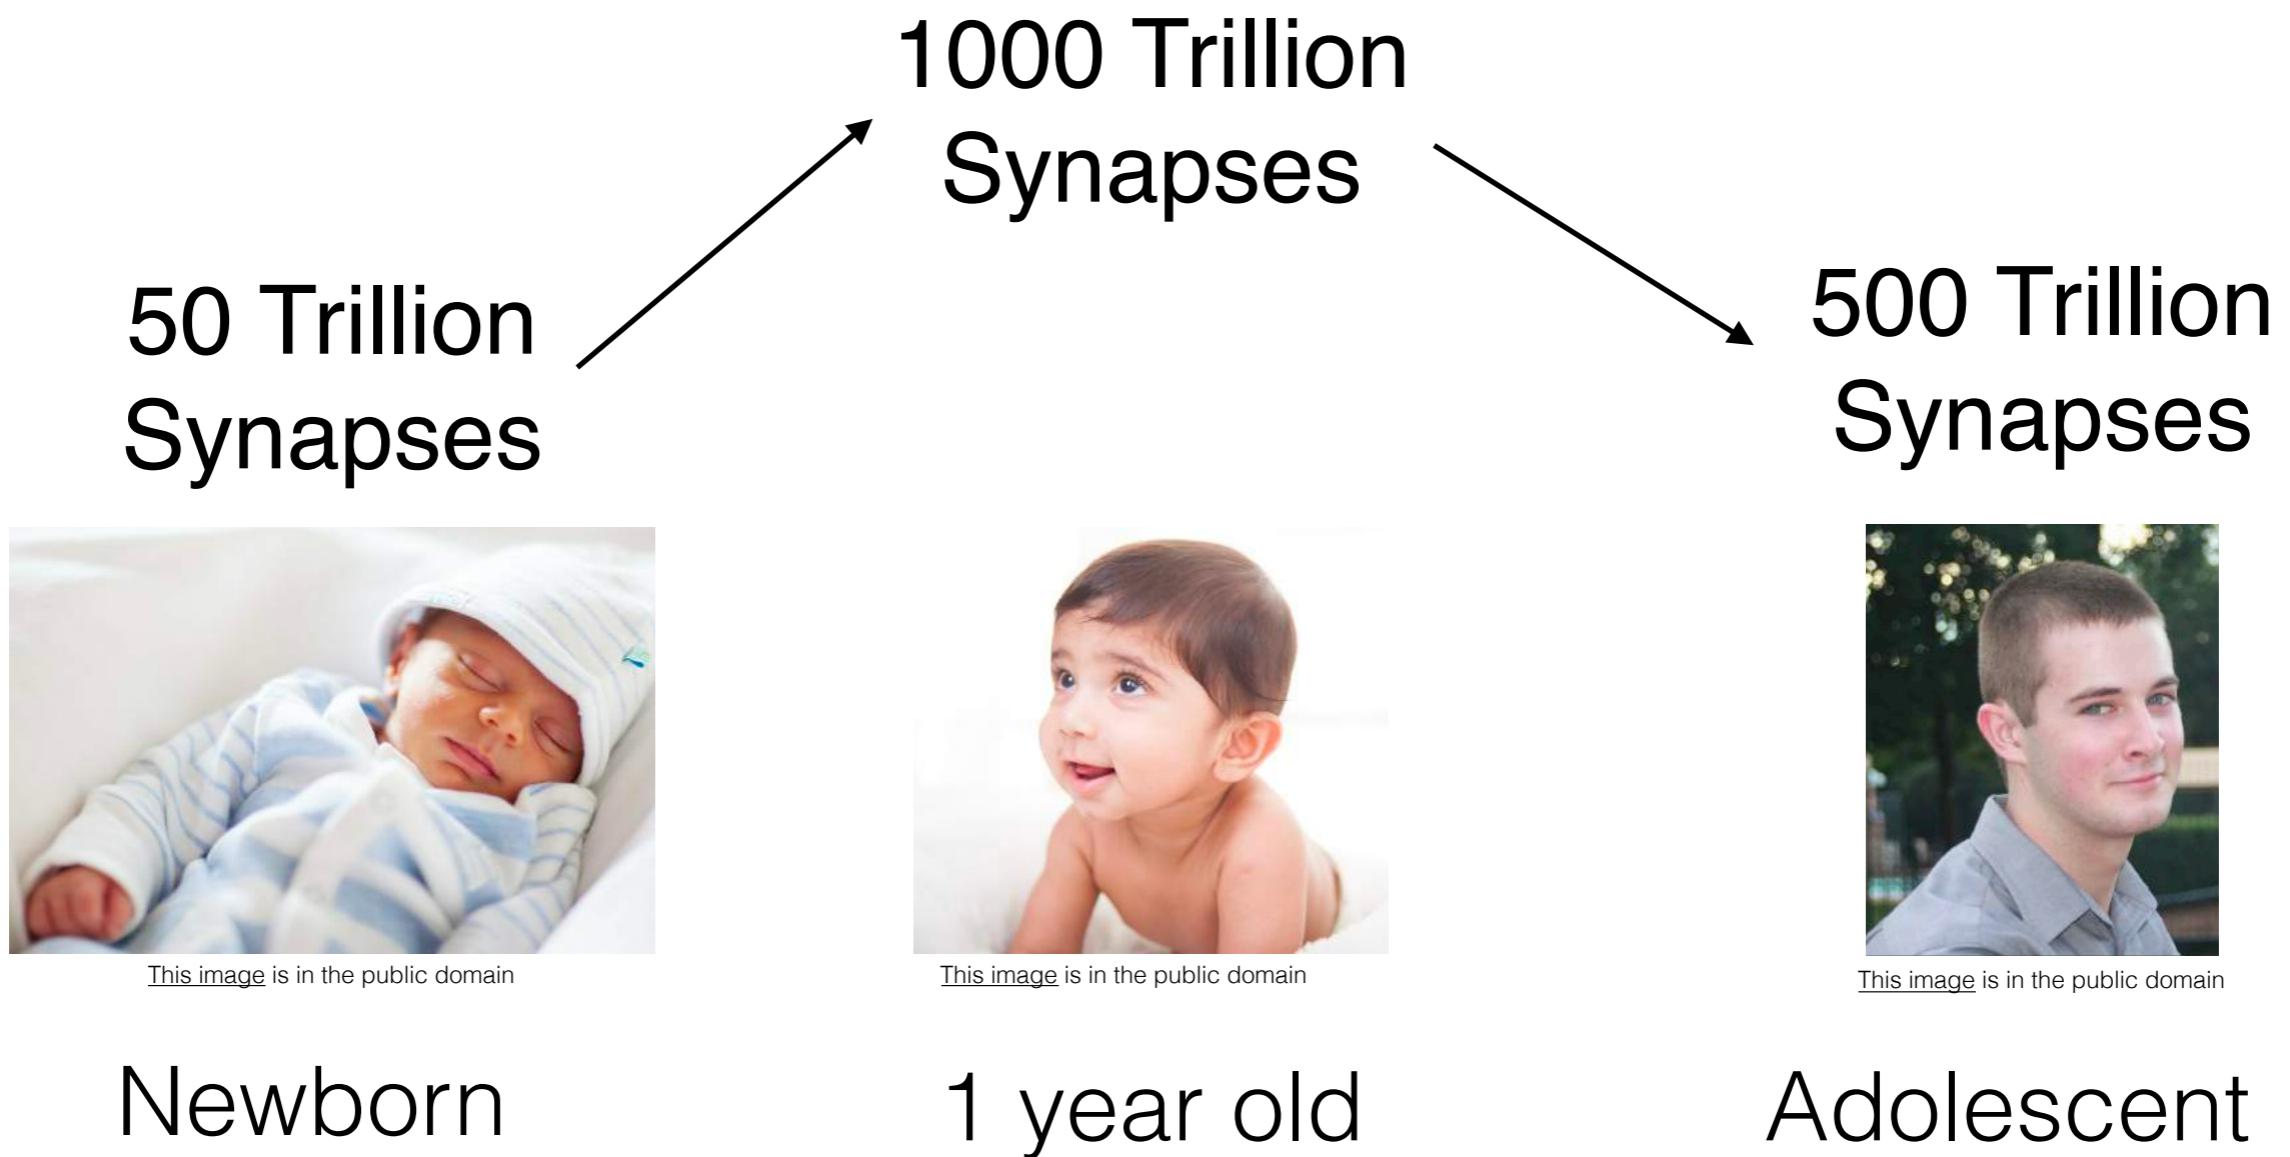
- **Original :** a man is riding a surfboard on a wave
- **Pruned 90%:** a man in a wetsuit is riding a wave **on a beach**

95%



- **Original :** a soccer player in red is running in the field
- **Pruned 95%:** a man in **a red shirt and black and white black shirt** is running through a field

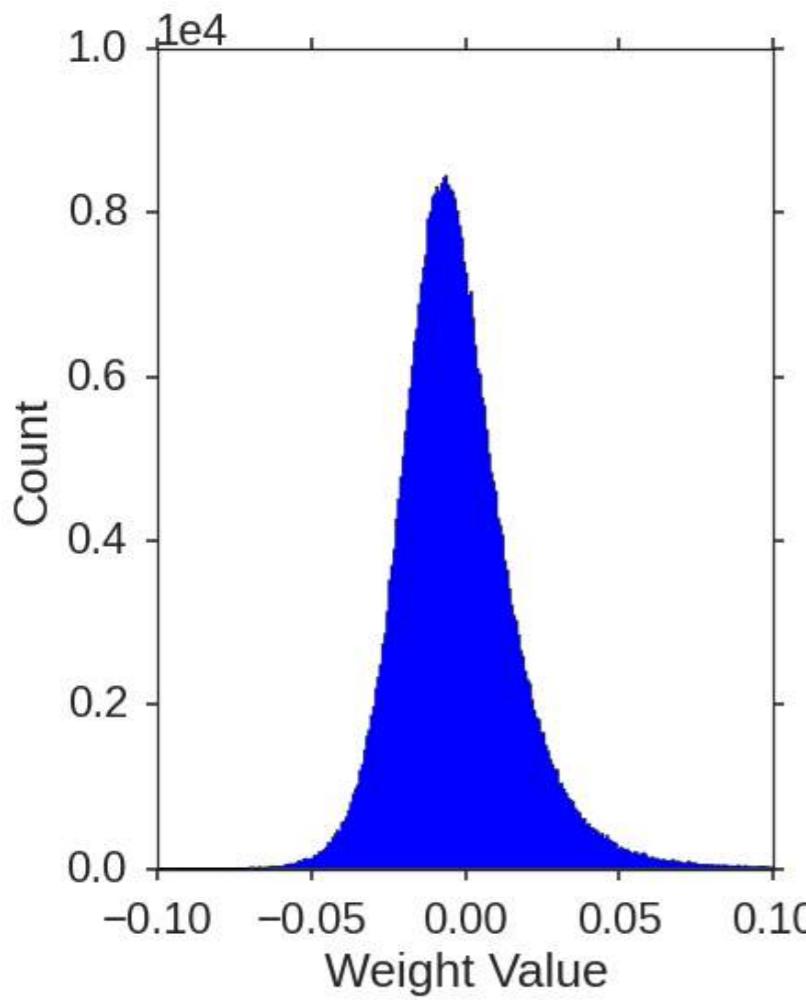
# Pruning Happens in Human Brain



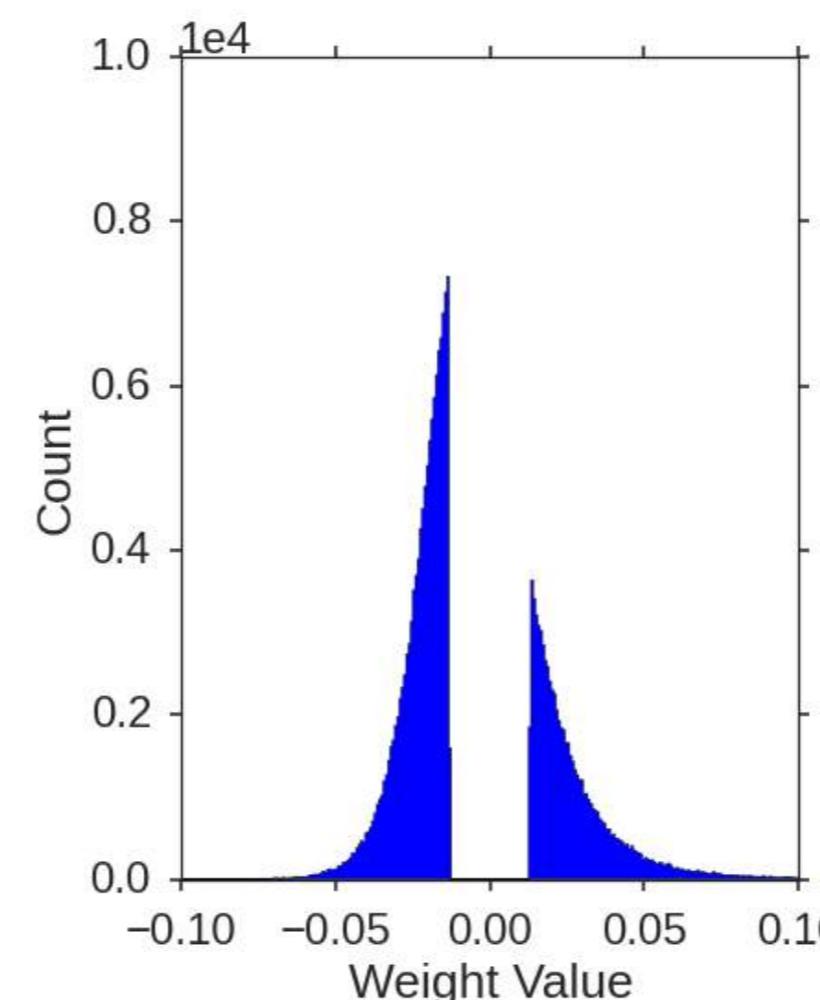
Christopher A Walsh. Peter Huttenlocher (1931-2013). Nature, 502(7470):172–172, 2013.

# Pruning Changes Weight Distribution

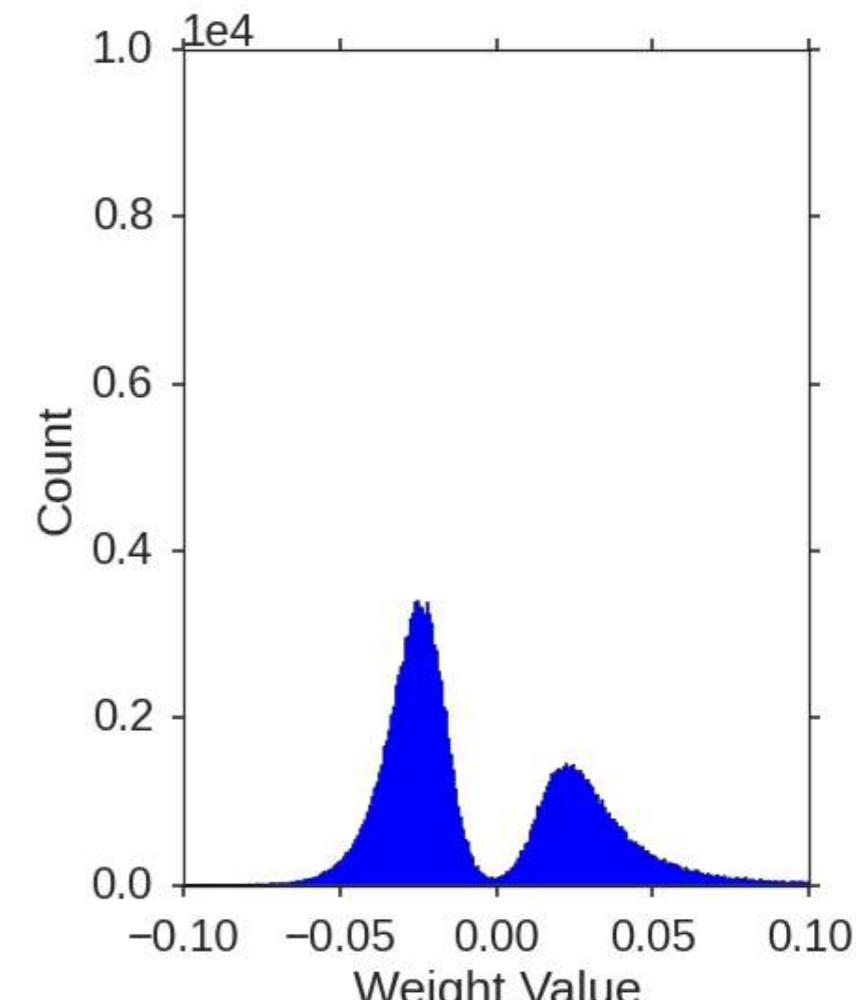
**Before Pruning**



**After Pruning**



**After Retraining**

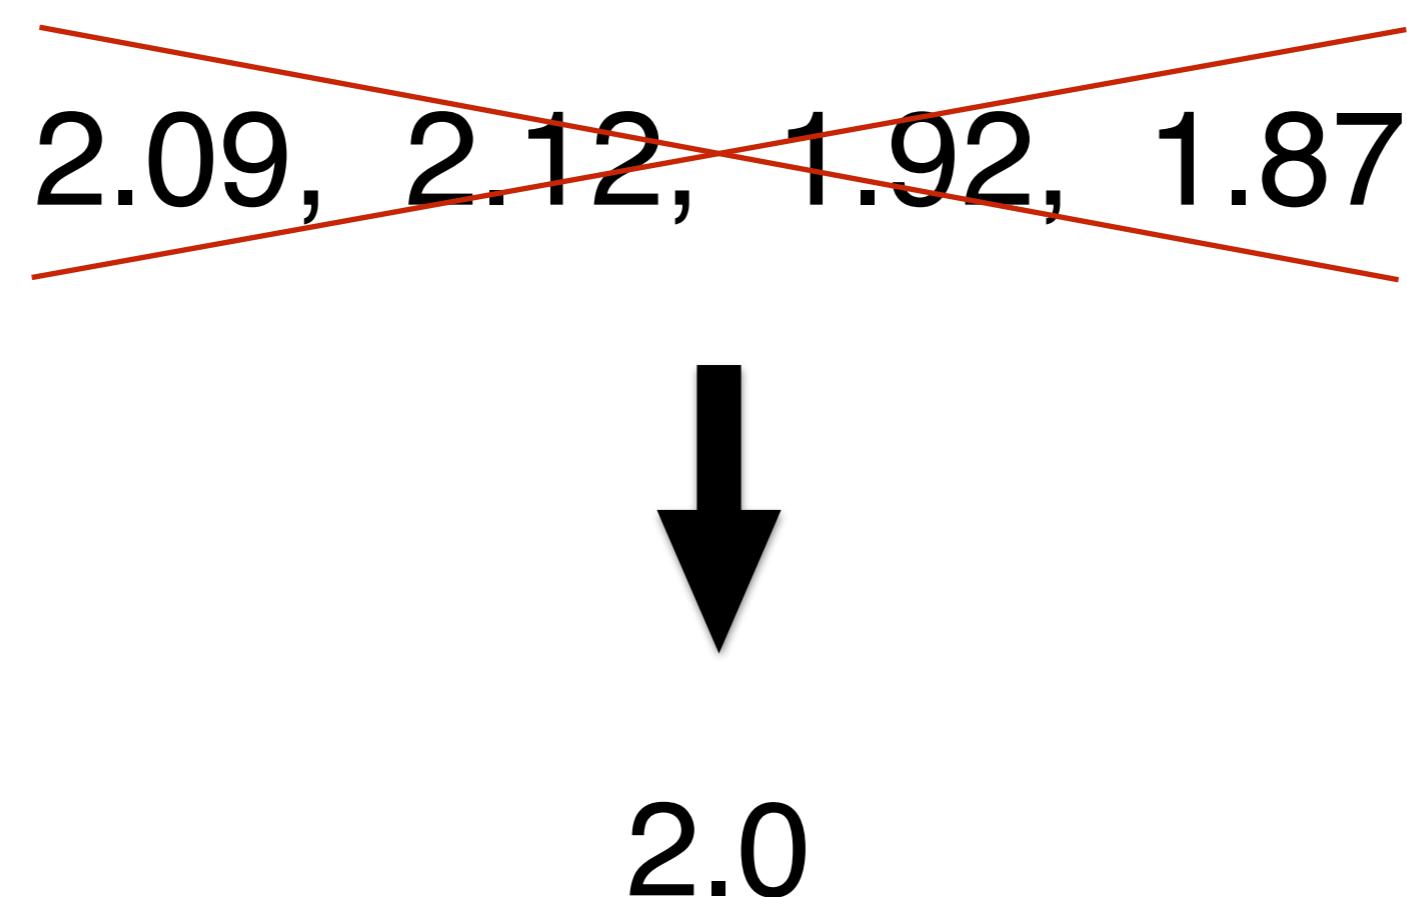


Conv5 layer of Alexnet. Representative for other network layers as well.

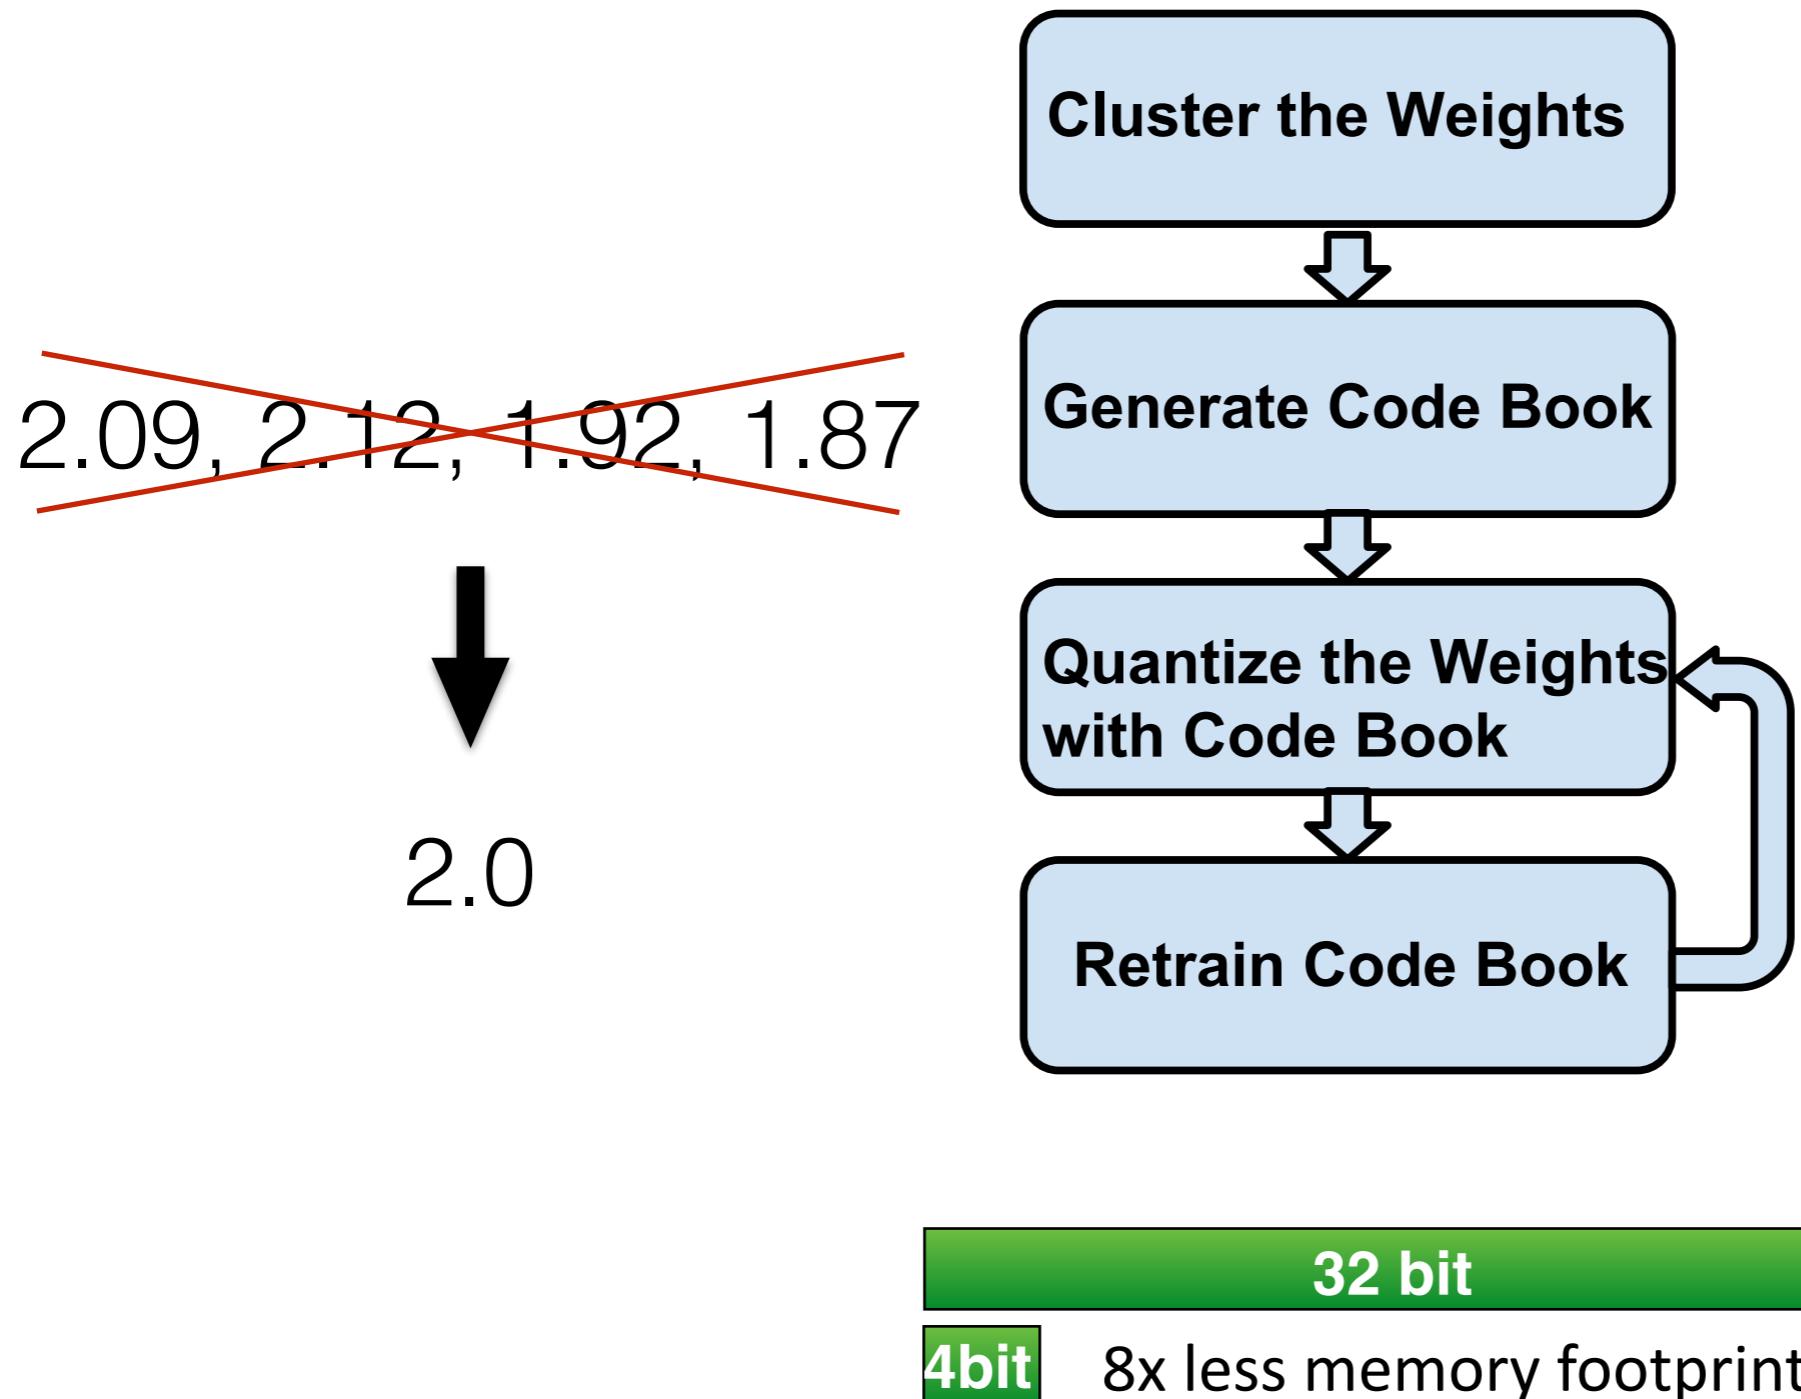
# Part 1: Algorithms for Efficient Inference

- 1. Pruning
- 2. Weight Sharing
- 3. Quantization
- 4. Low Rank Approximation
- 5. Binary / Ternary Net
- 6. Winograd Transformation

# Trained Quantization



# Trained Quantization

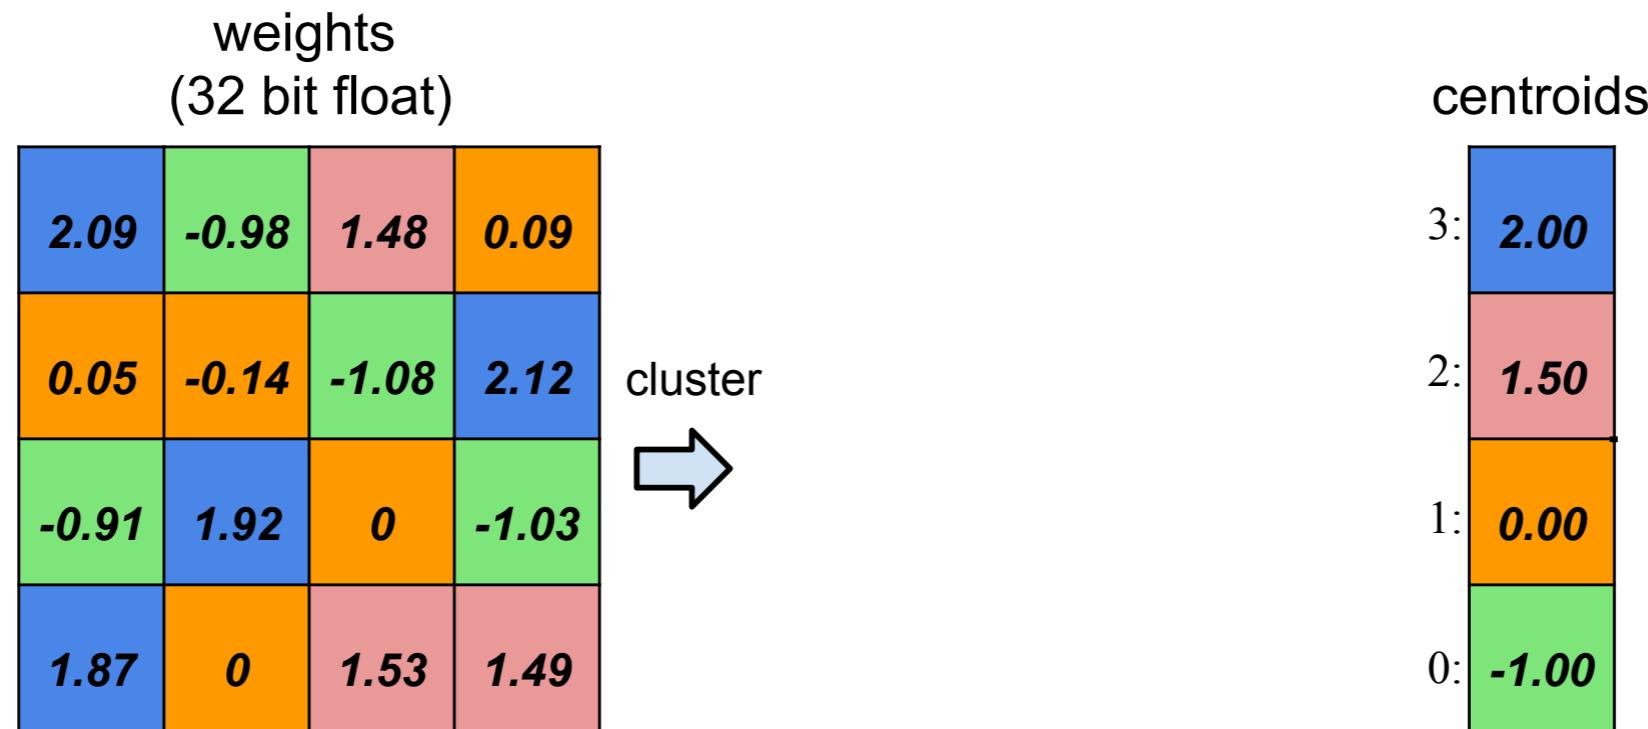


# Trained Quantization

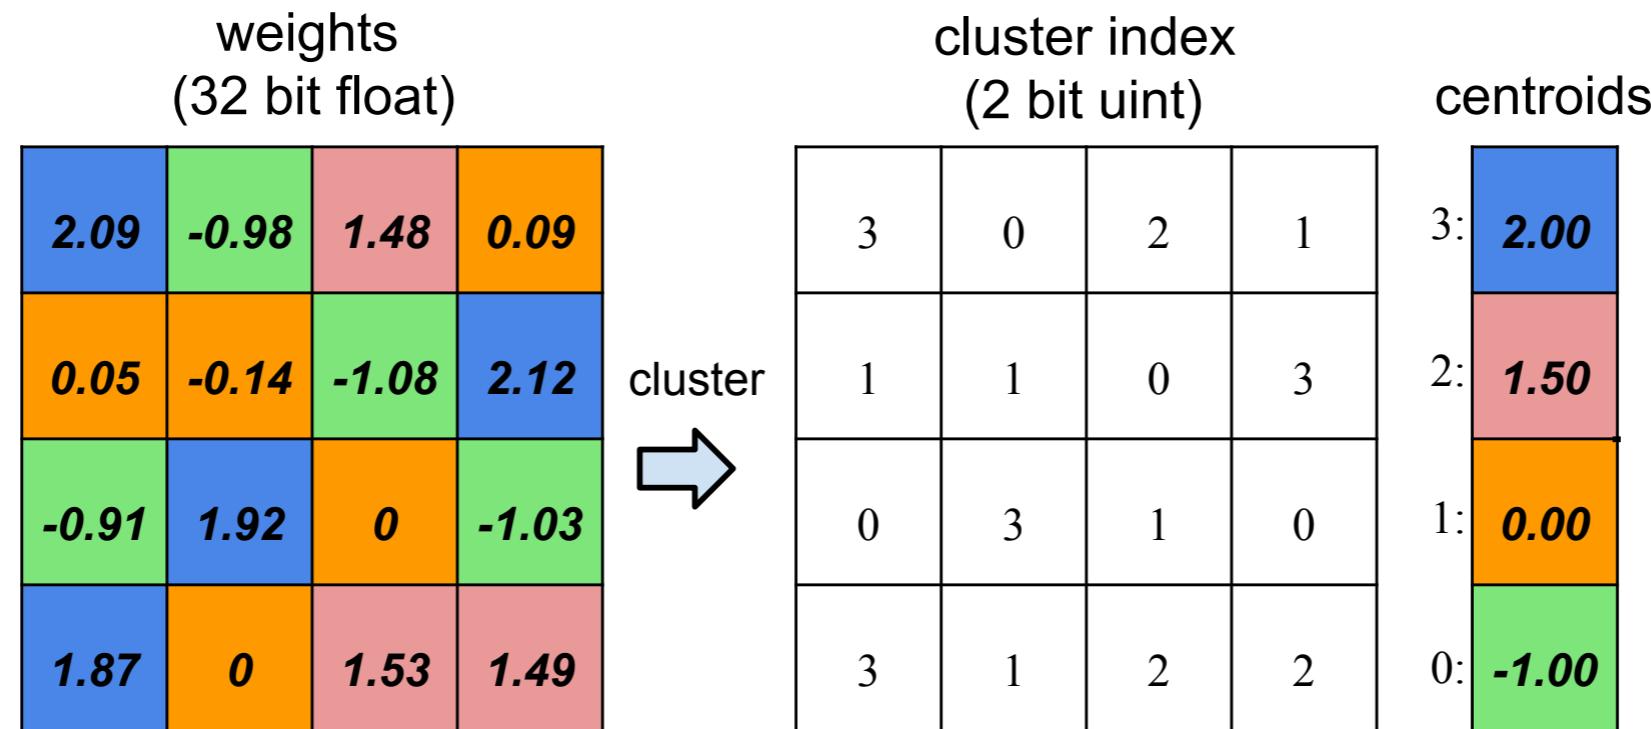
weights  
(32 bit float)

2.09	-0.98	1.48	0.09
0.05	-0.14	-1.08	2.12
-0.91	1.92	0	-1.03
1.87	0	1.53	1.49

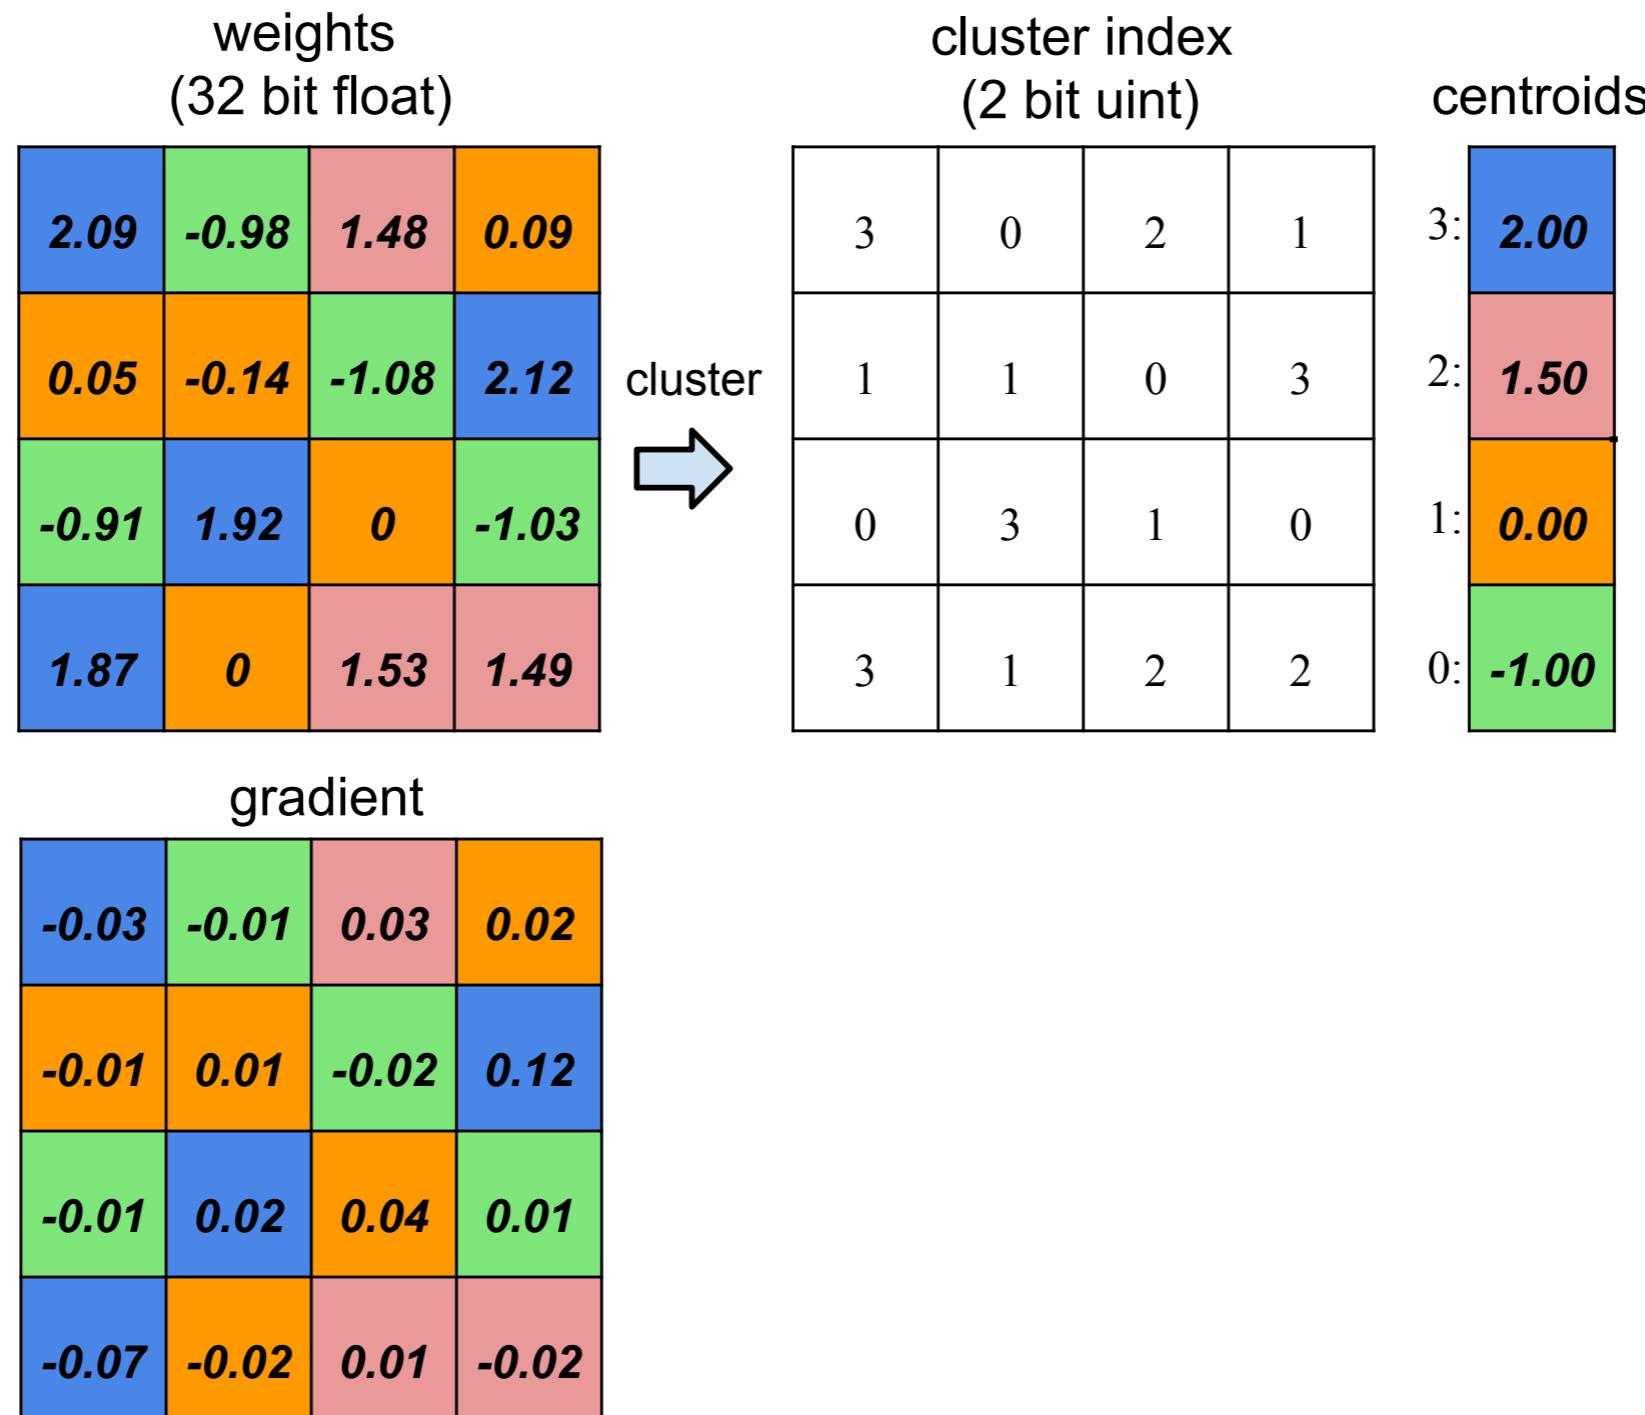
# Trained Quantization



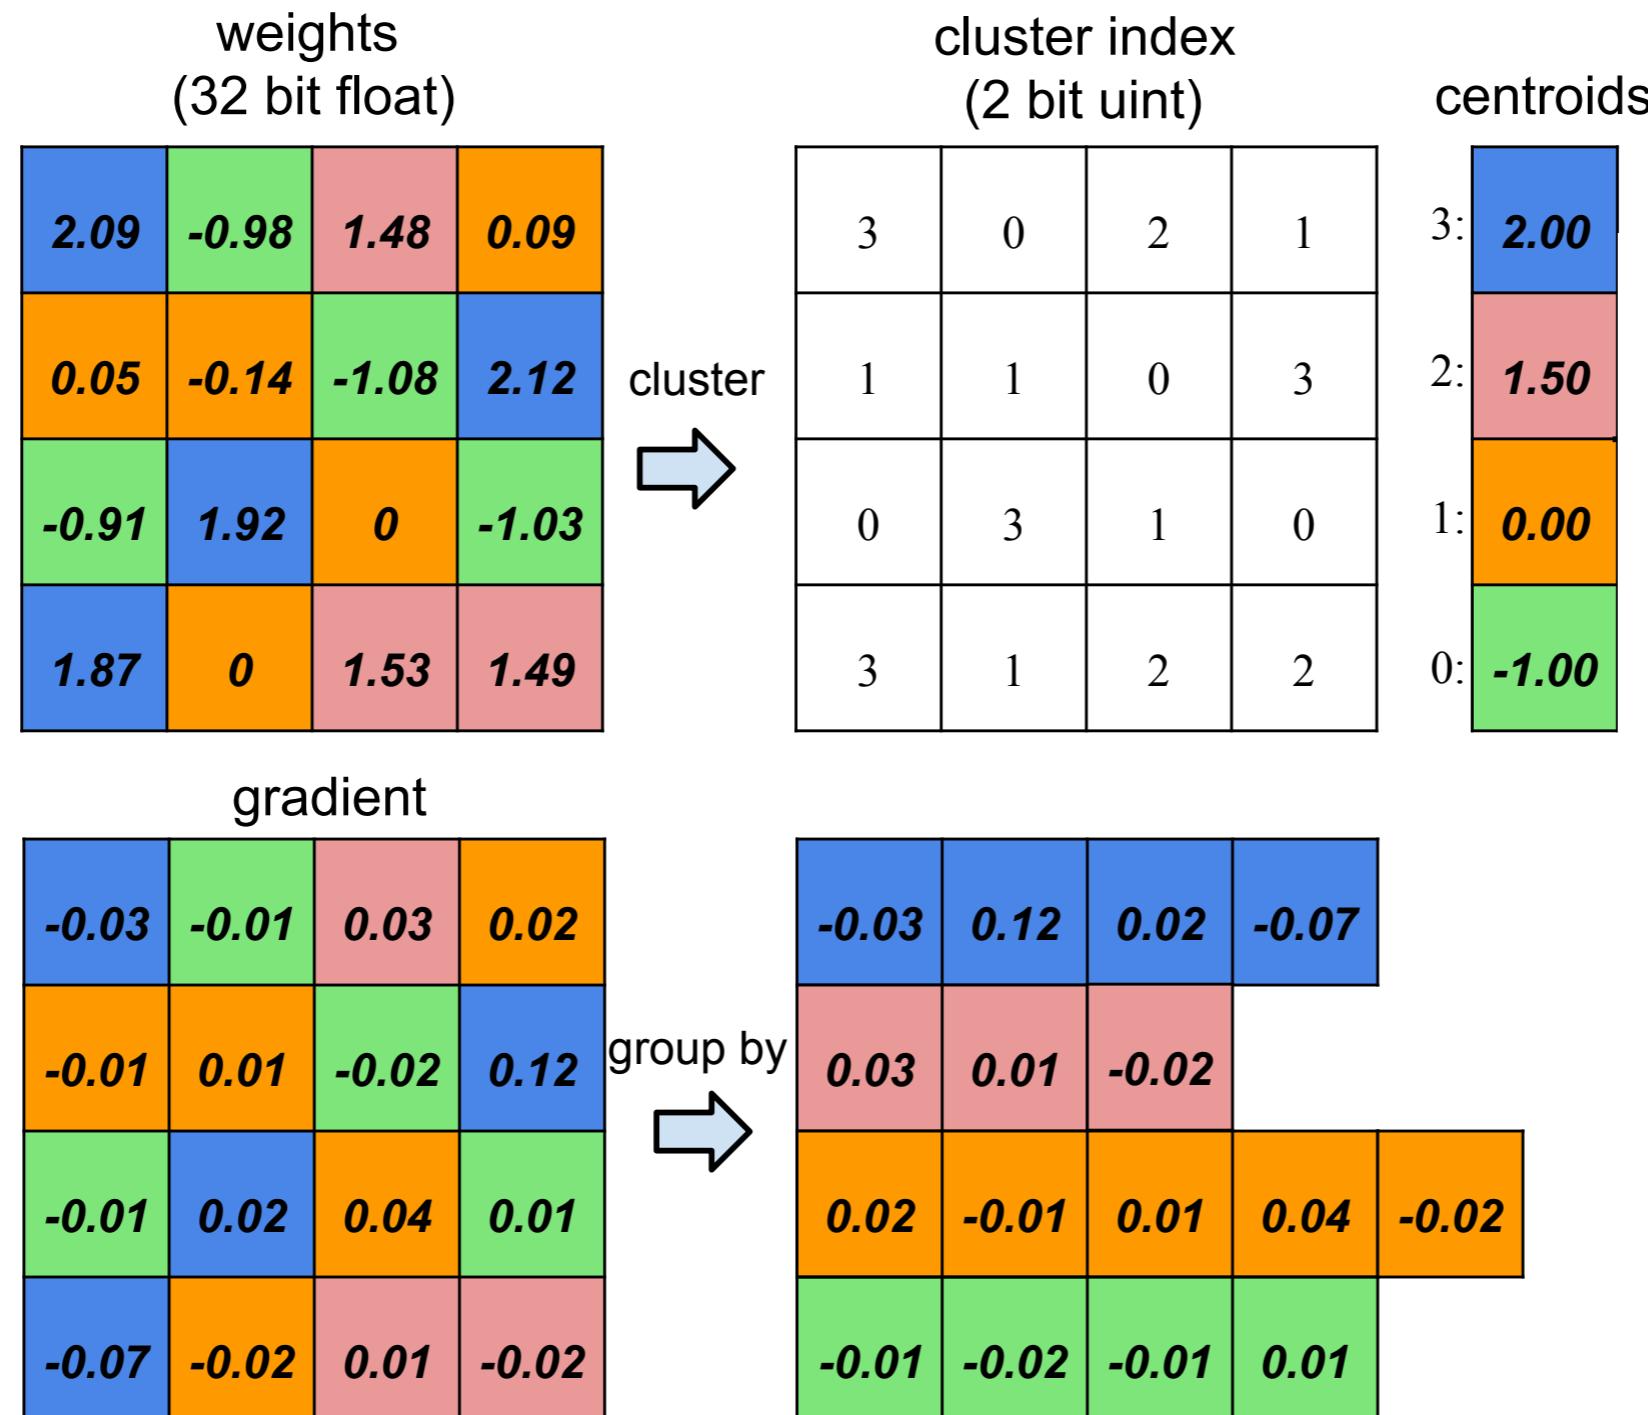
# Trained Quantization



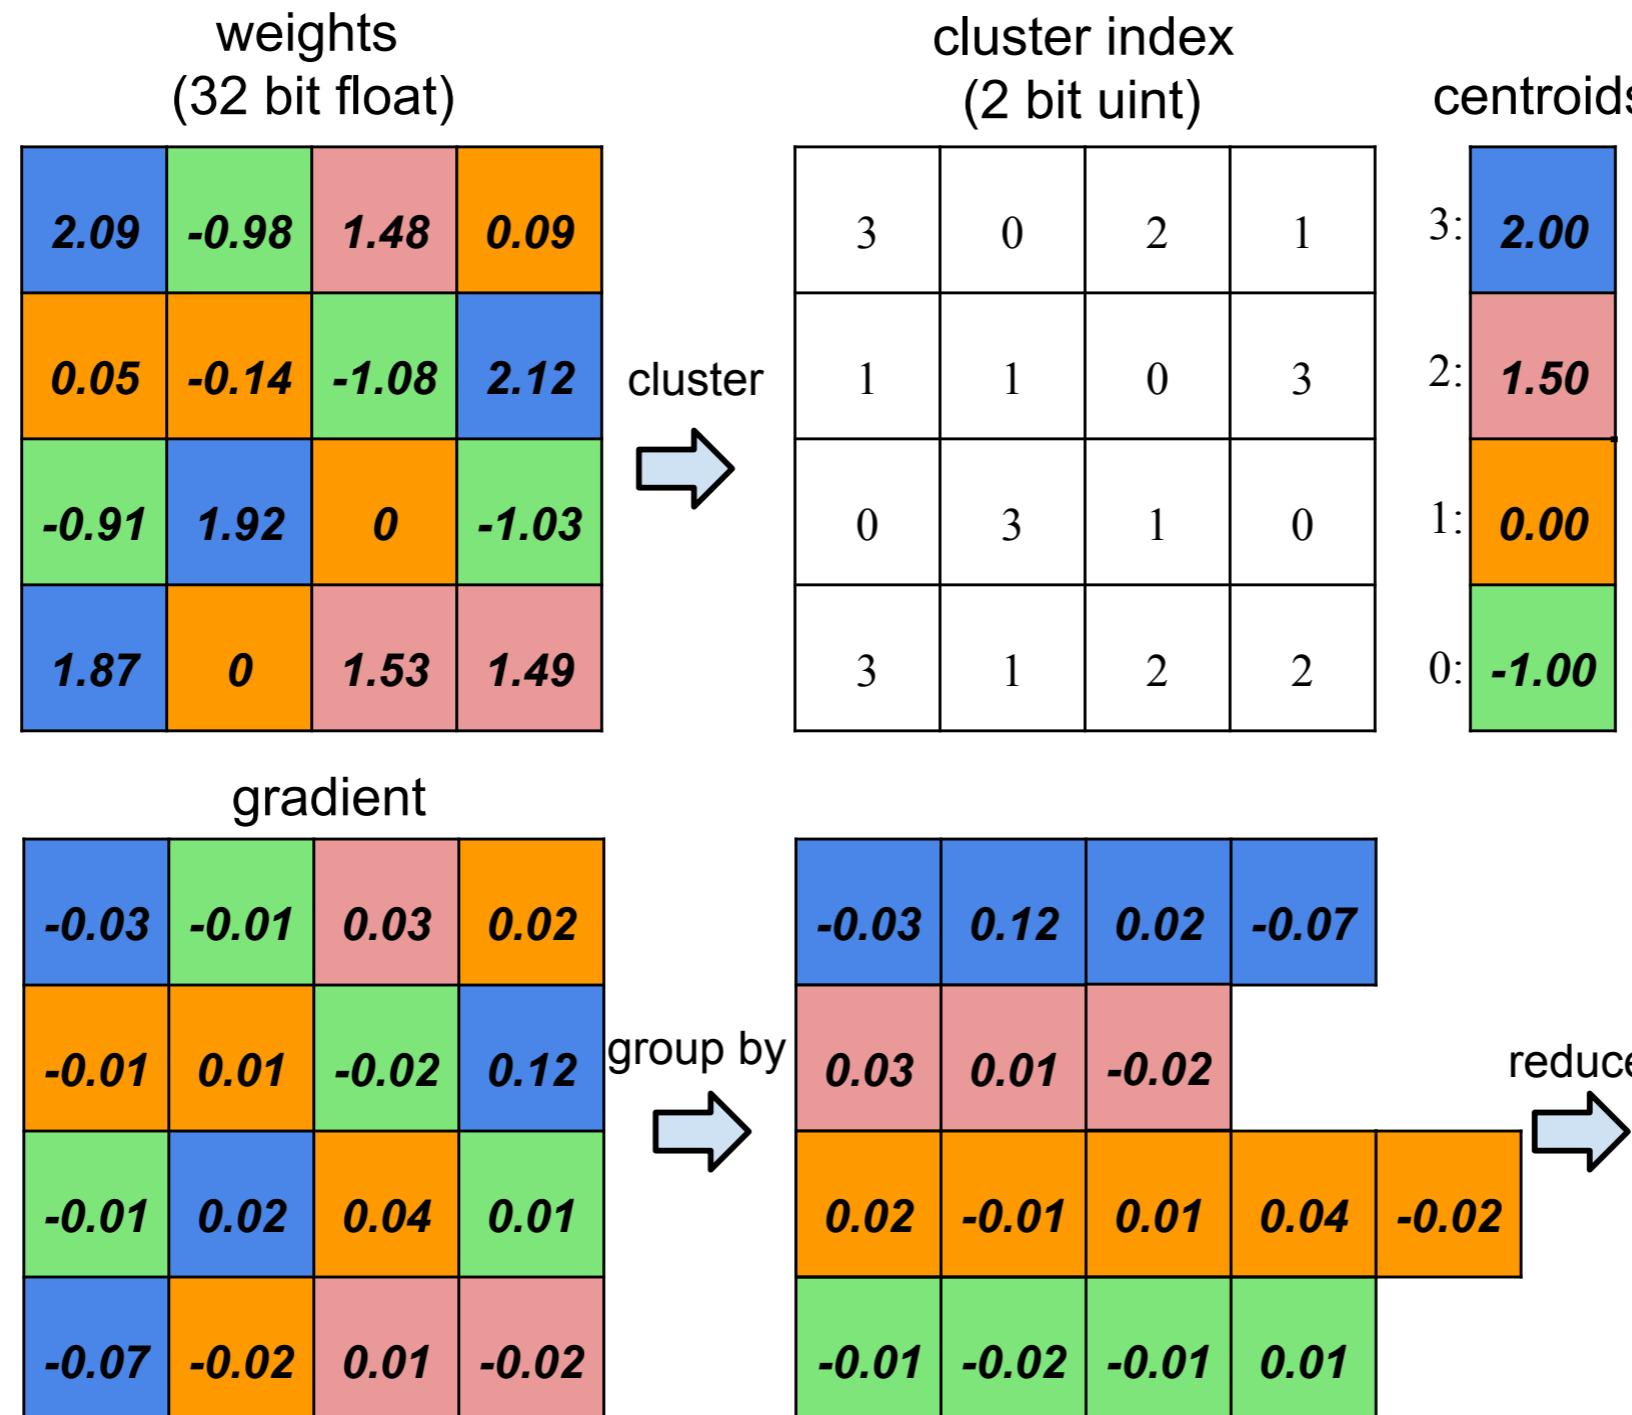
# Trained Quantization



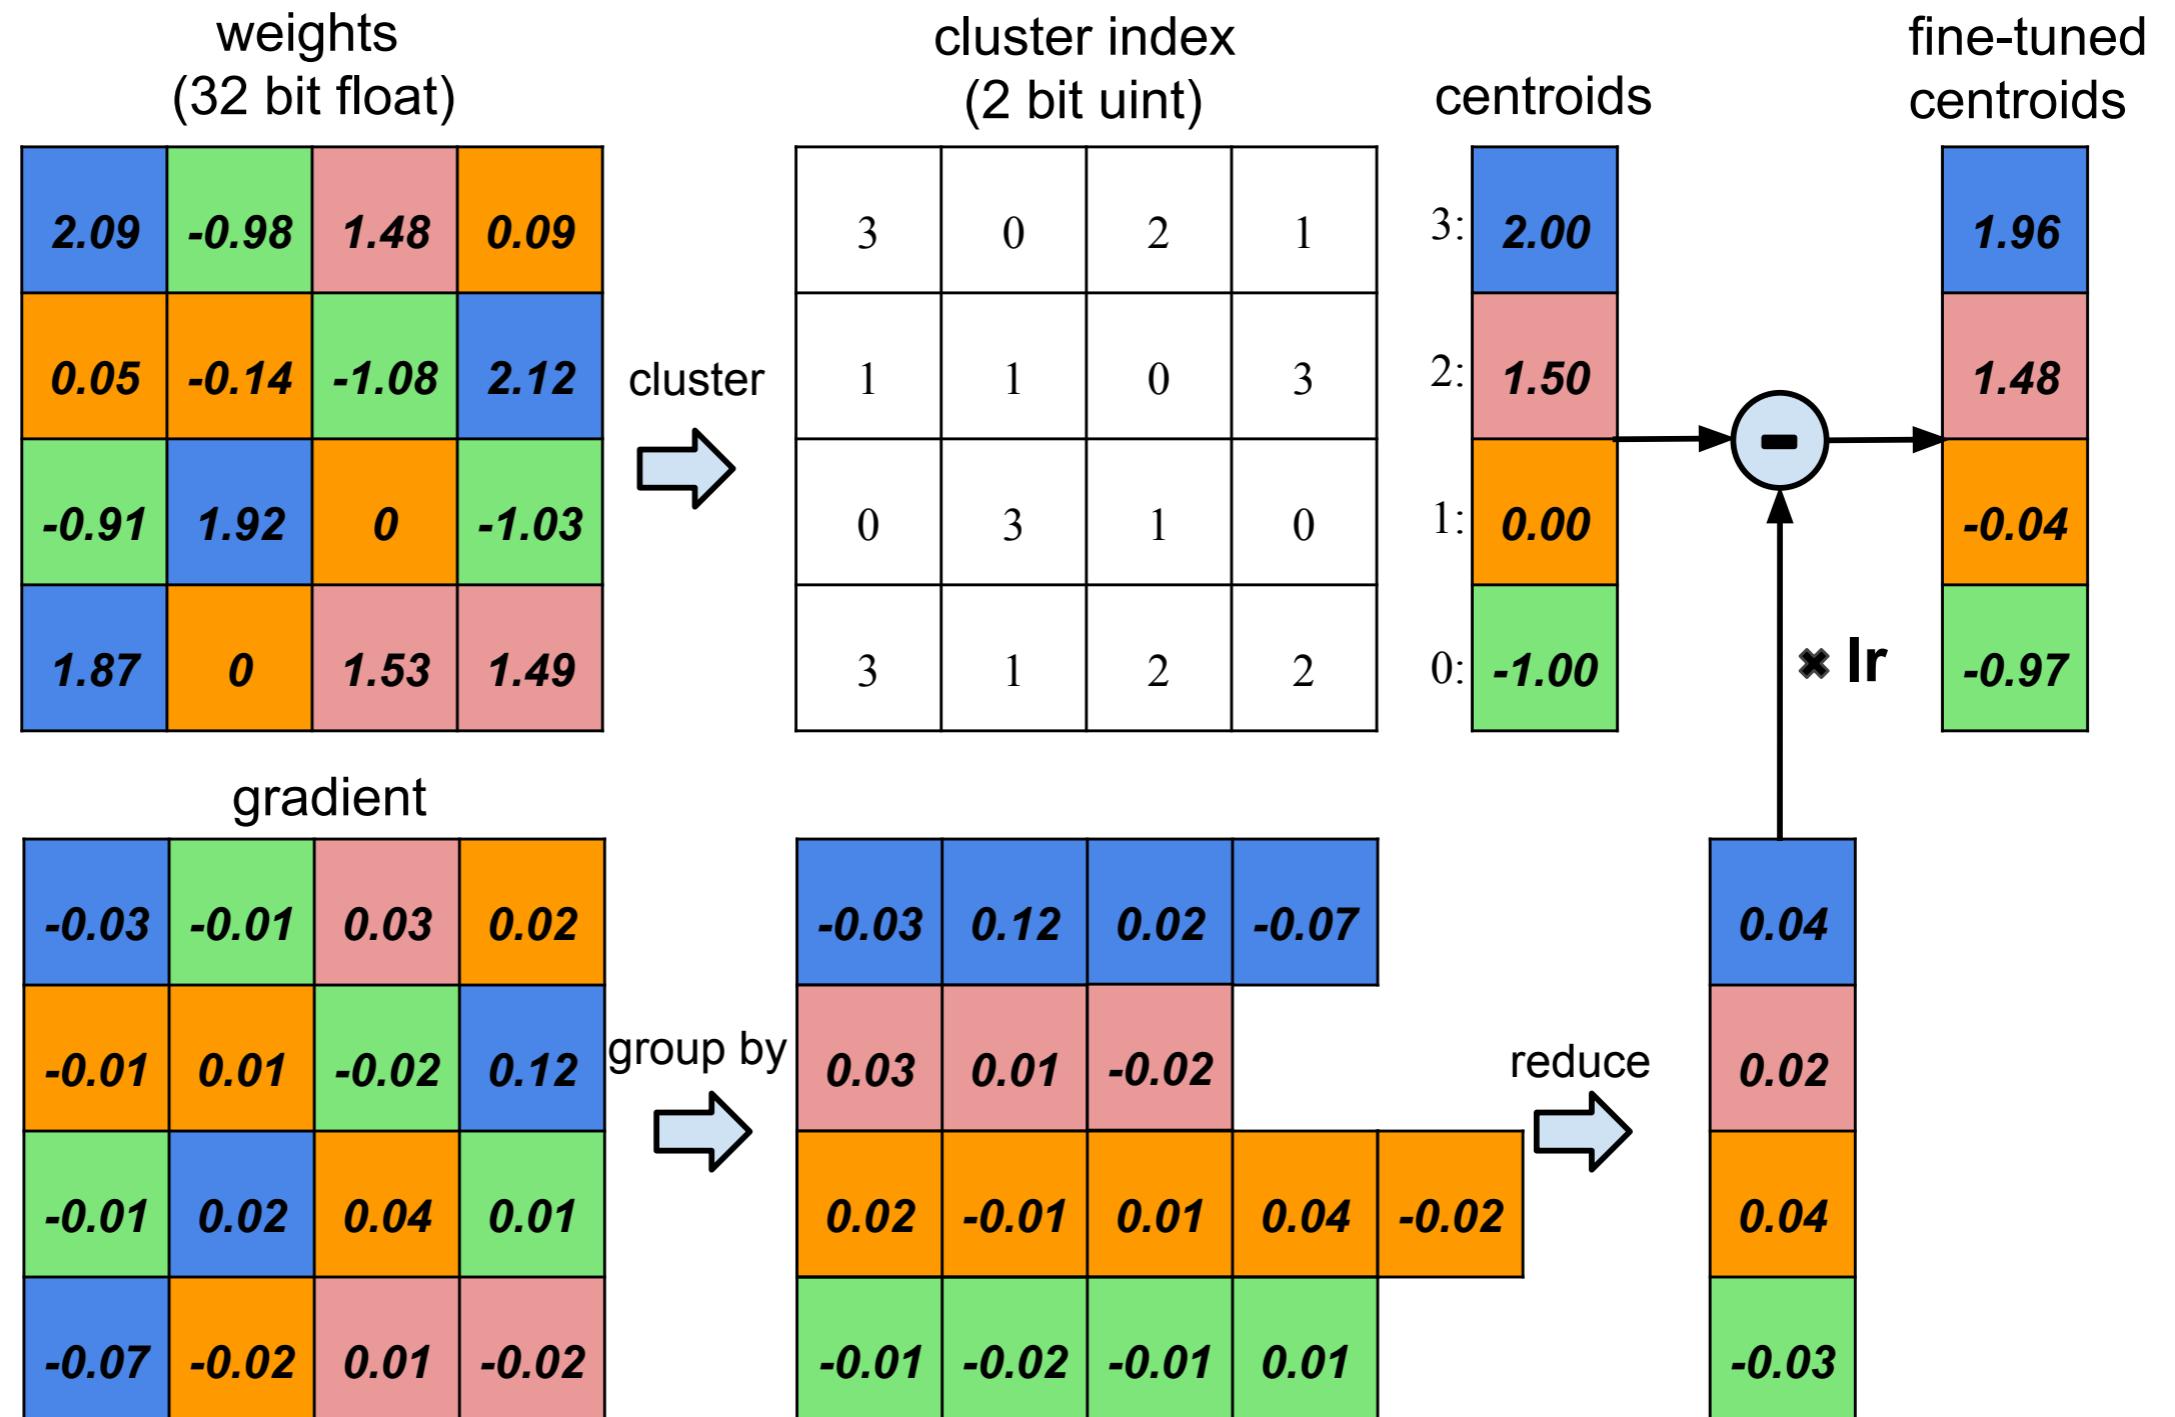
# Trained Quantization



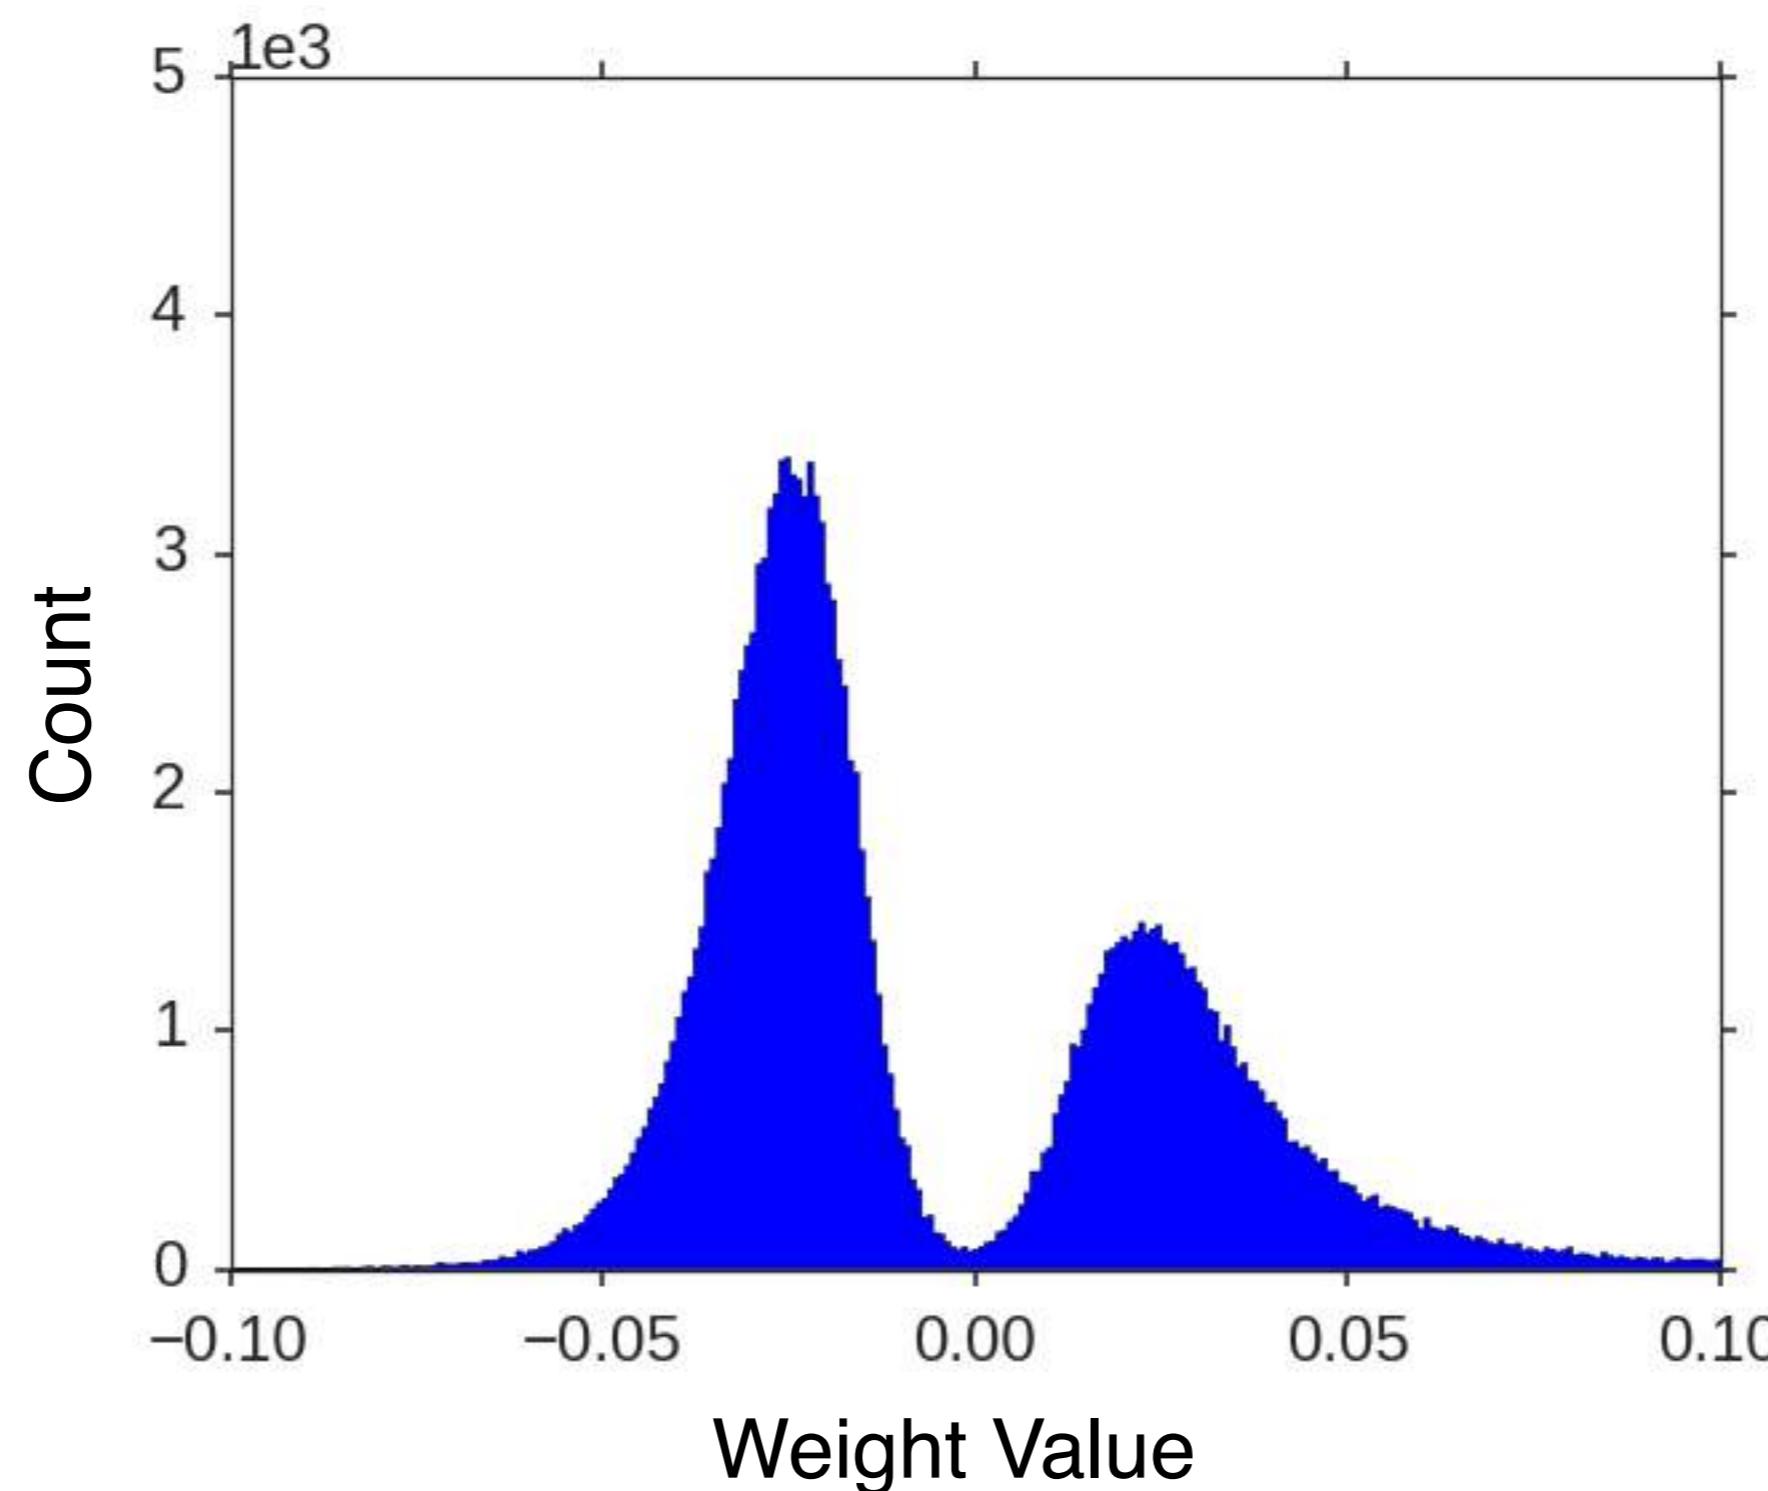
# Trained Quantization



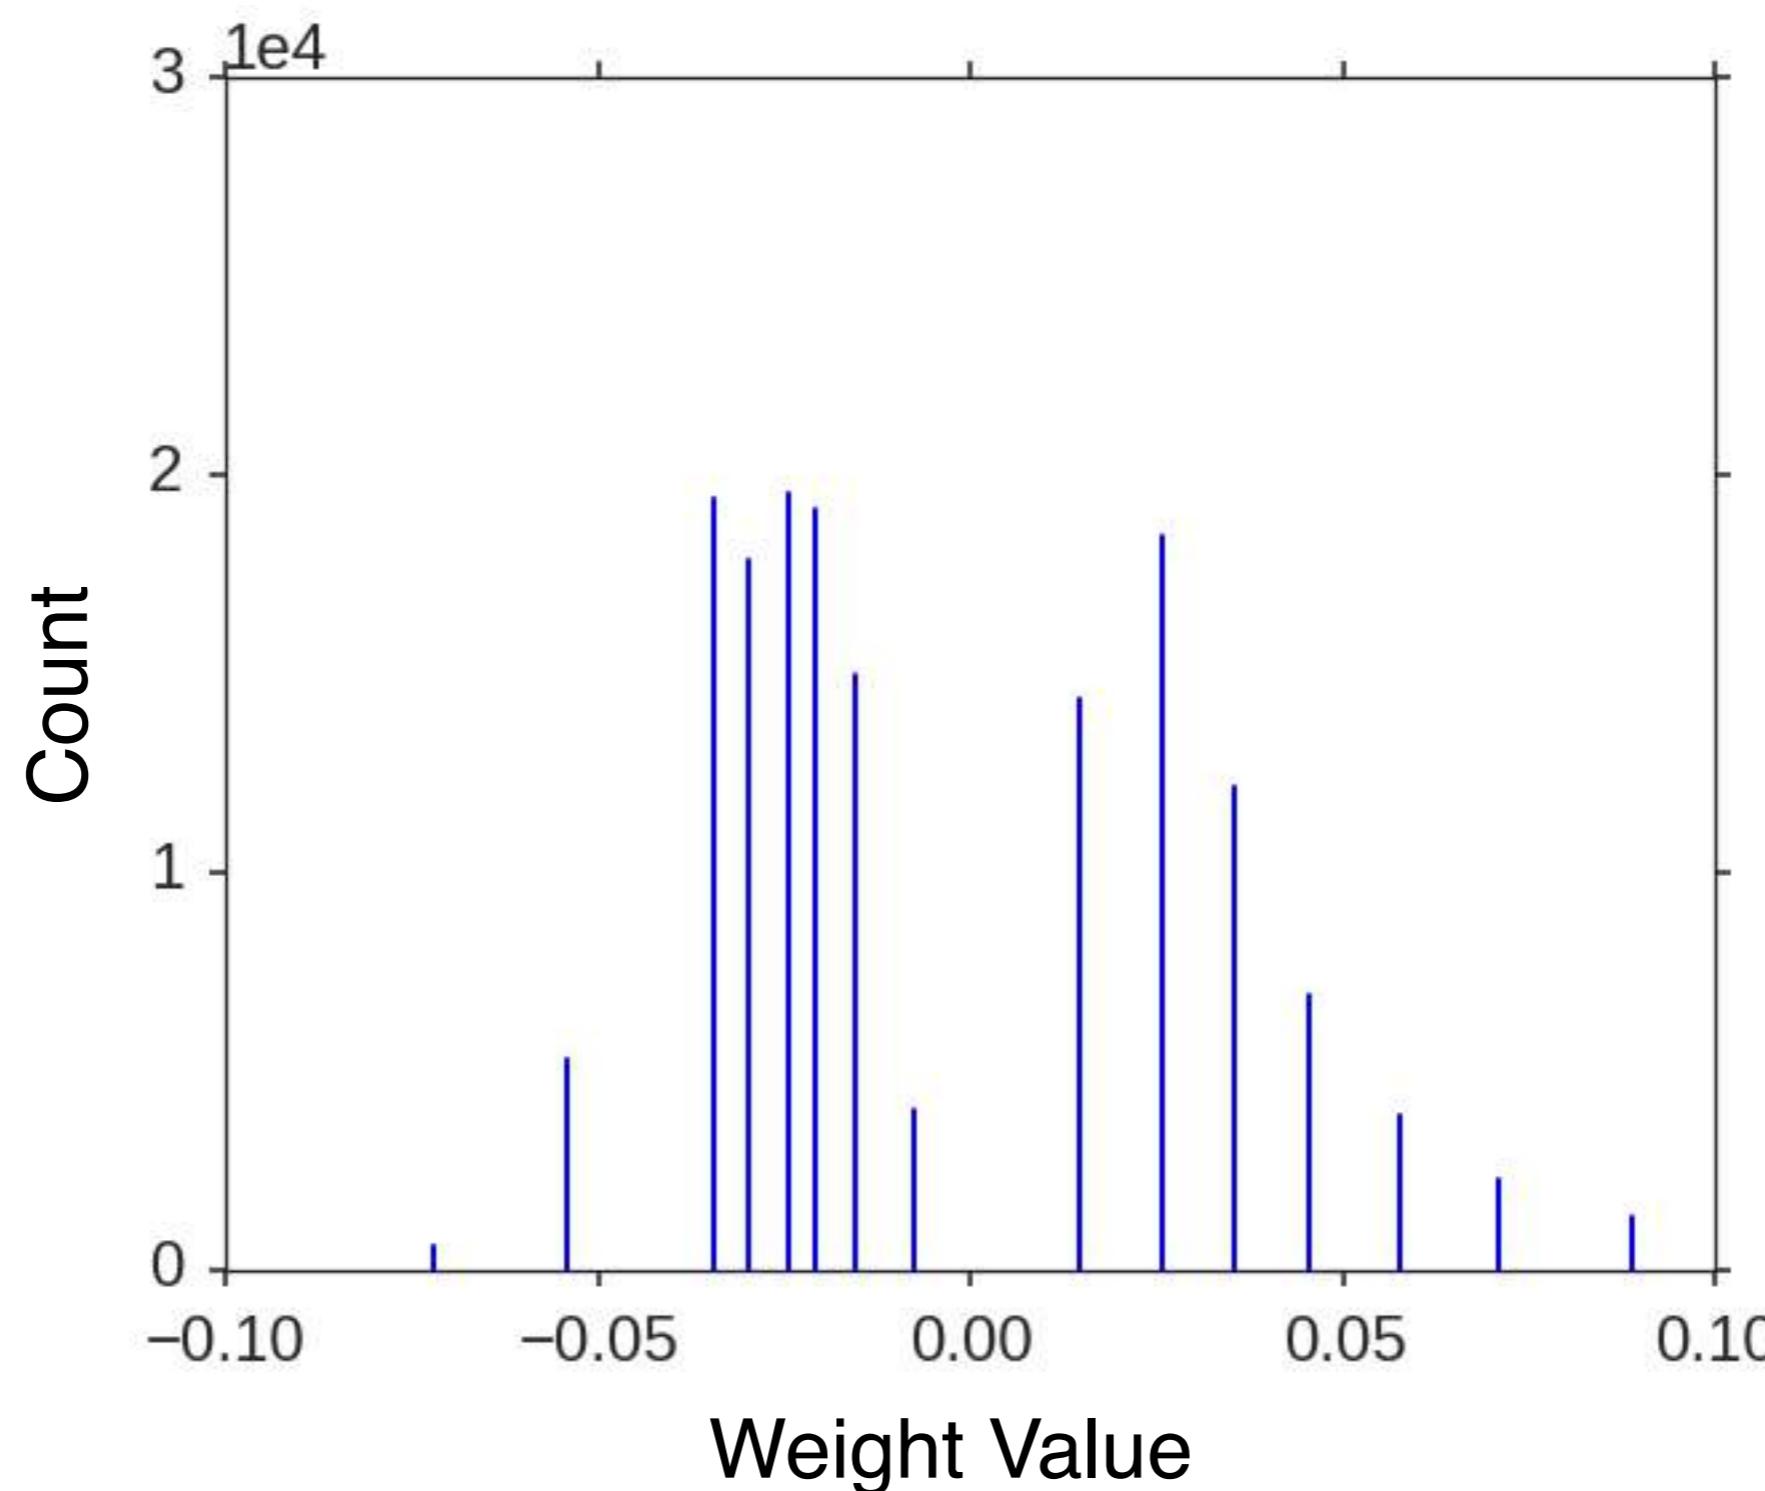
# Trained Quantization



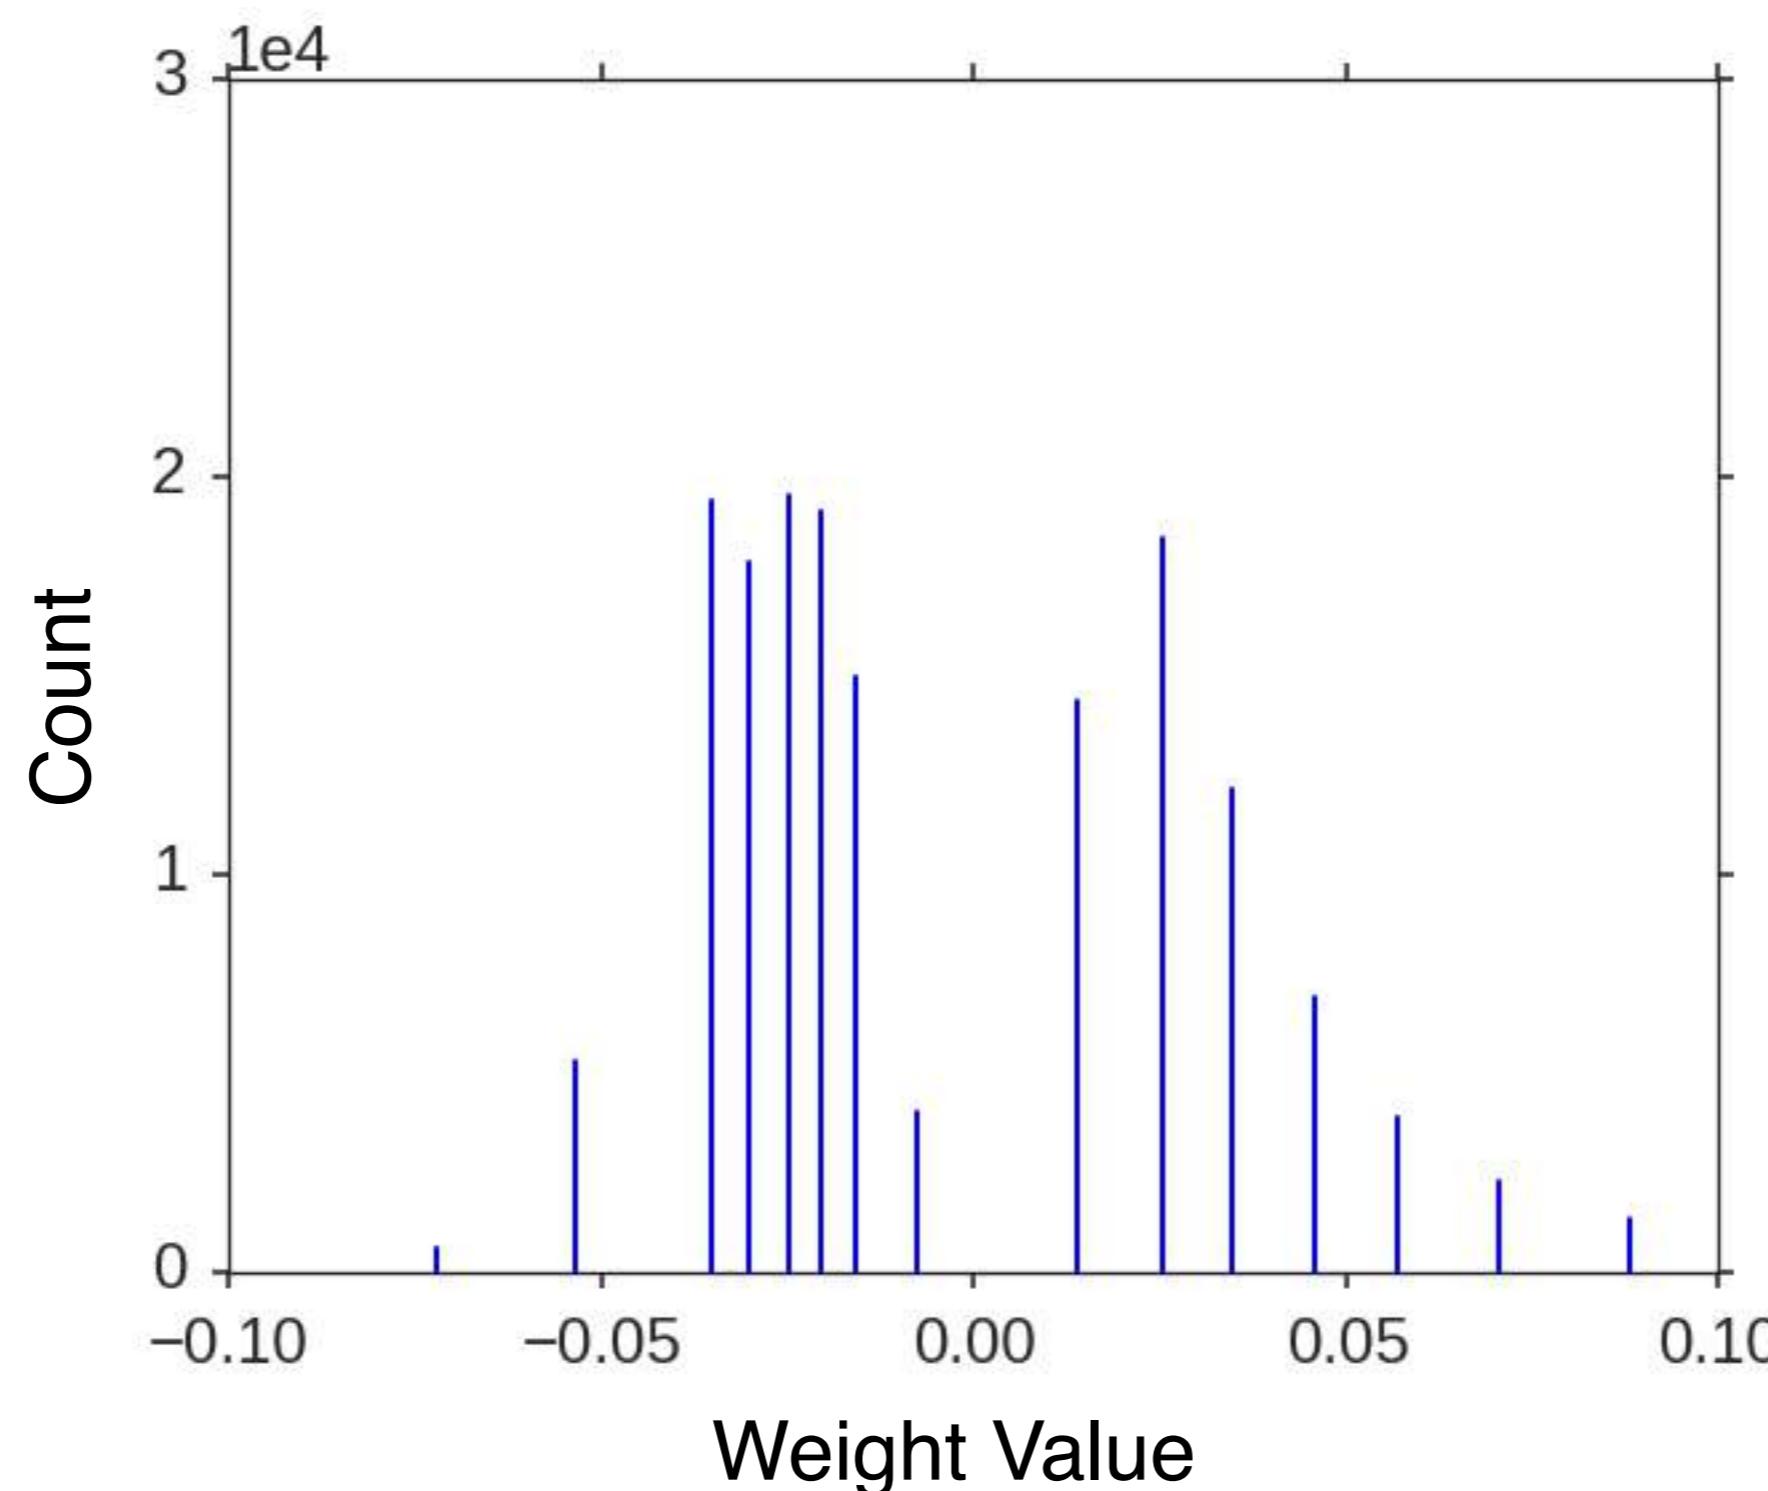
# Before Trained Quantization: Continuous Weight



# After Trained Quantization: Discrete Weight

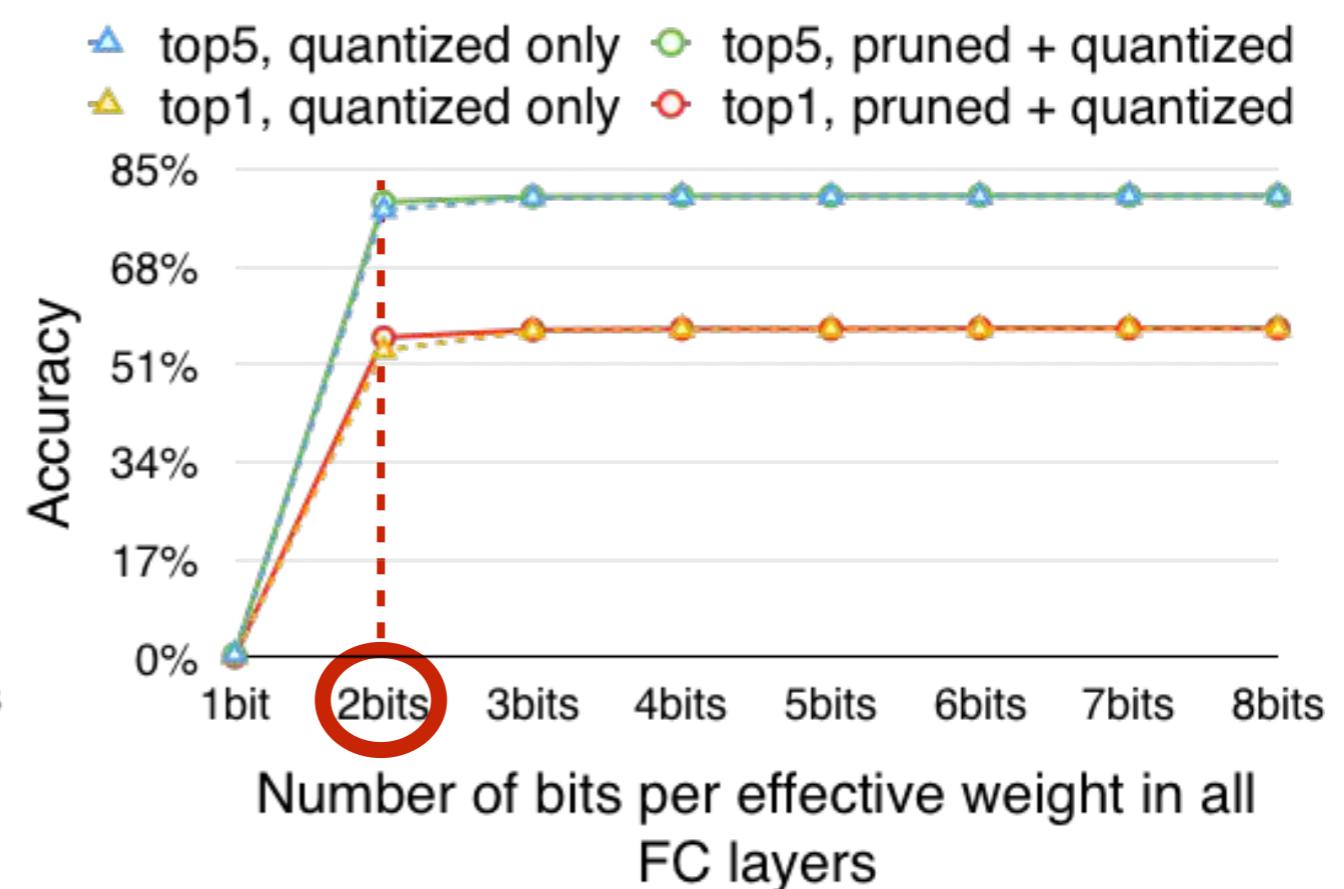
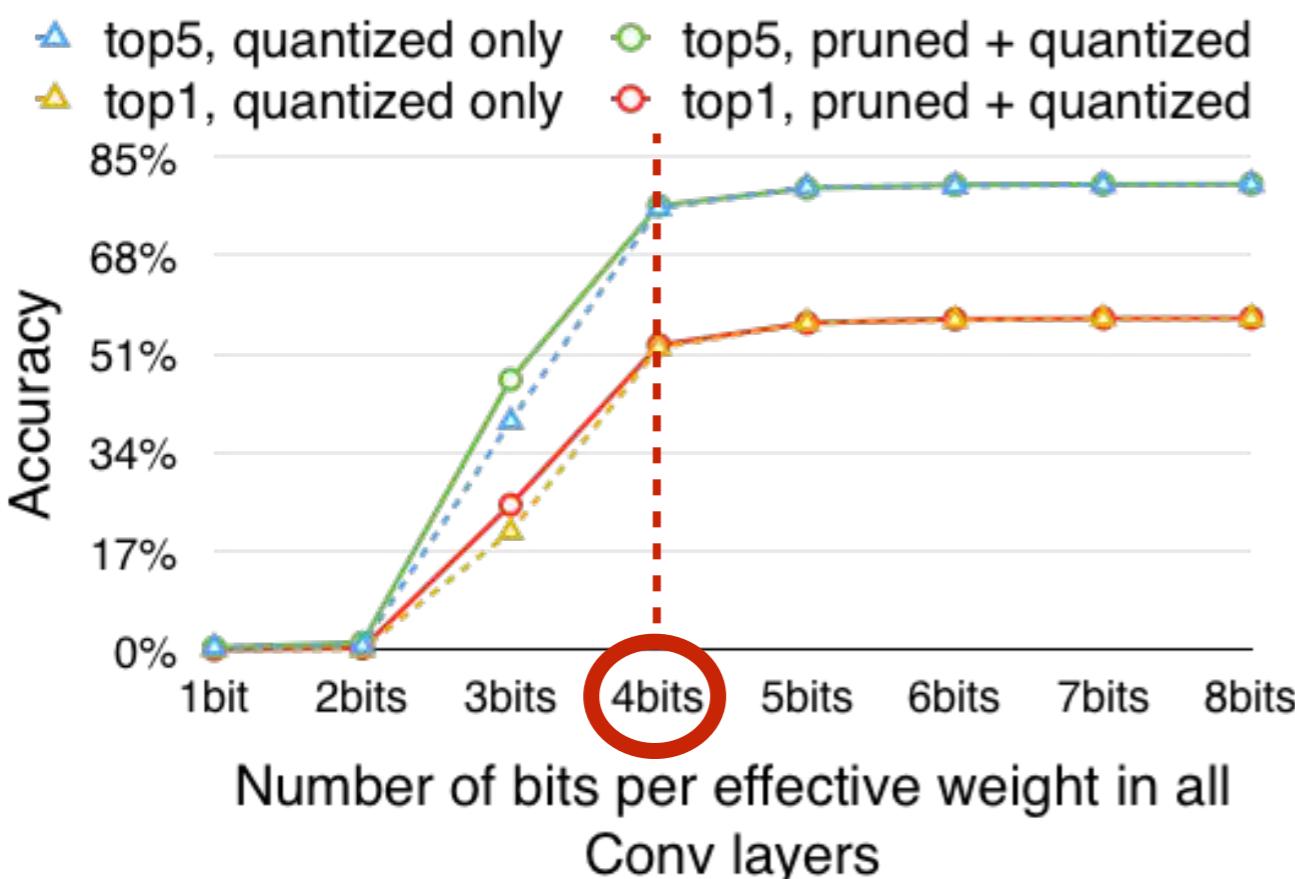


# After Trained Quantization: Discrete Weight after Training



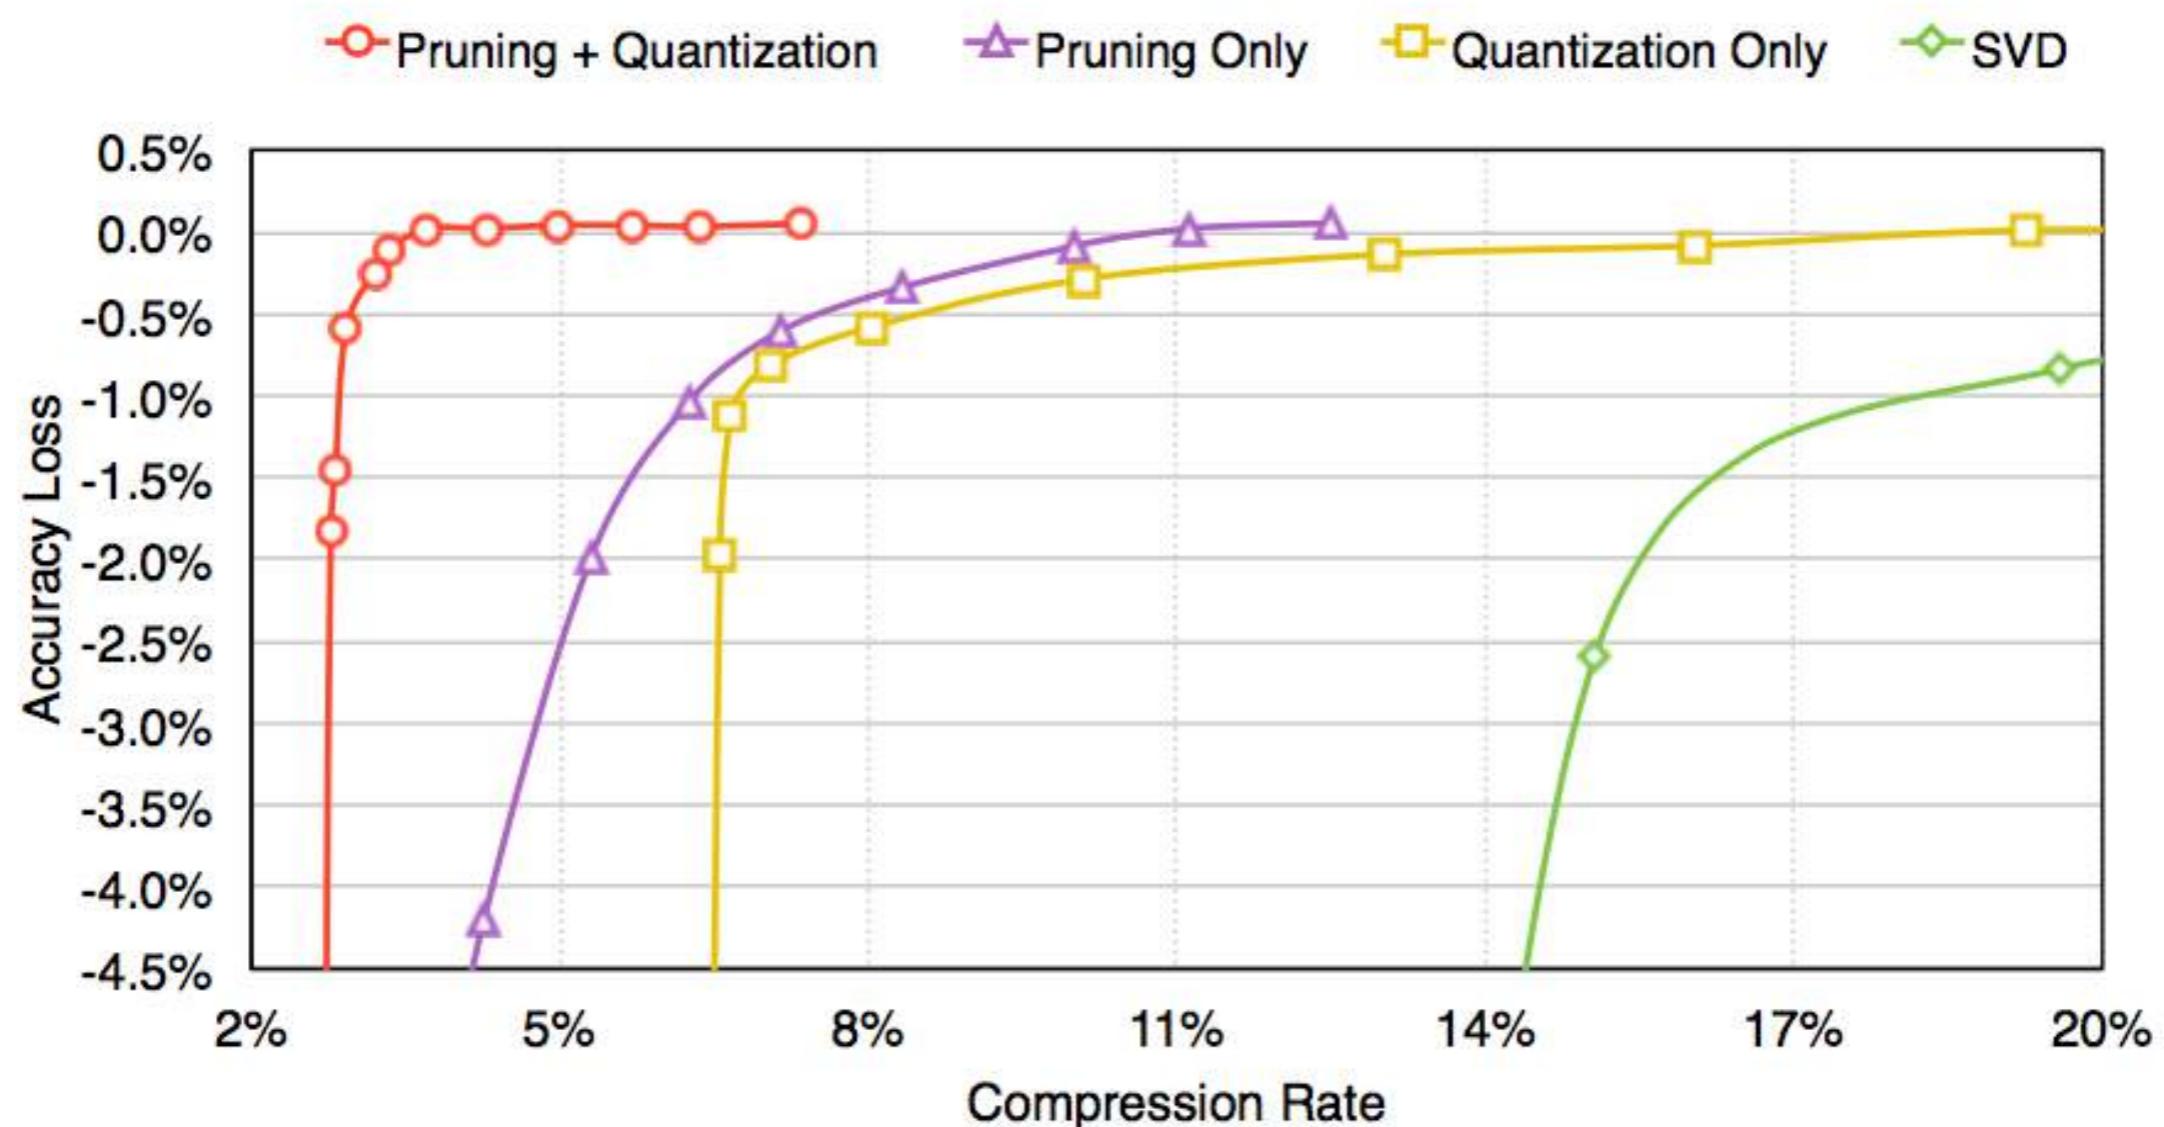
# How Many Bits do We Need?

# How Many Bits do We Need?



# Pruning + Trained Quantization Work Together

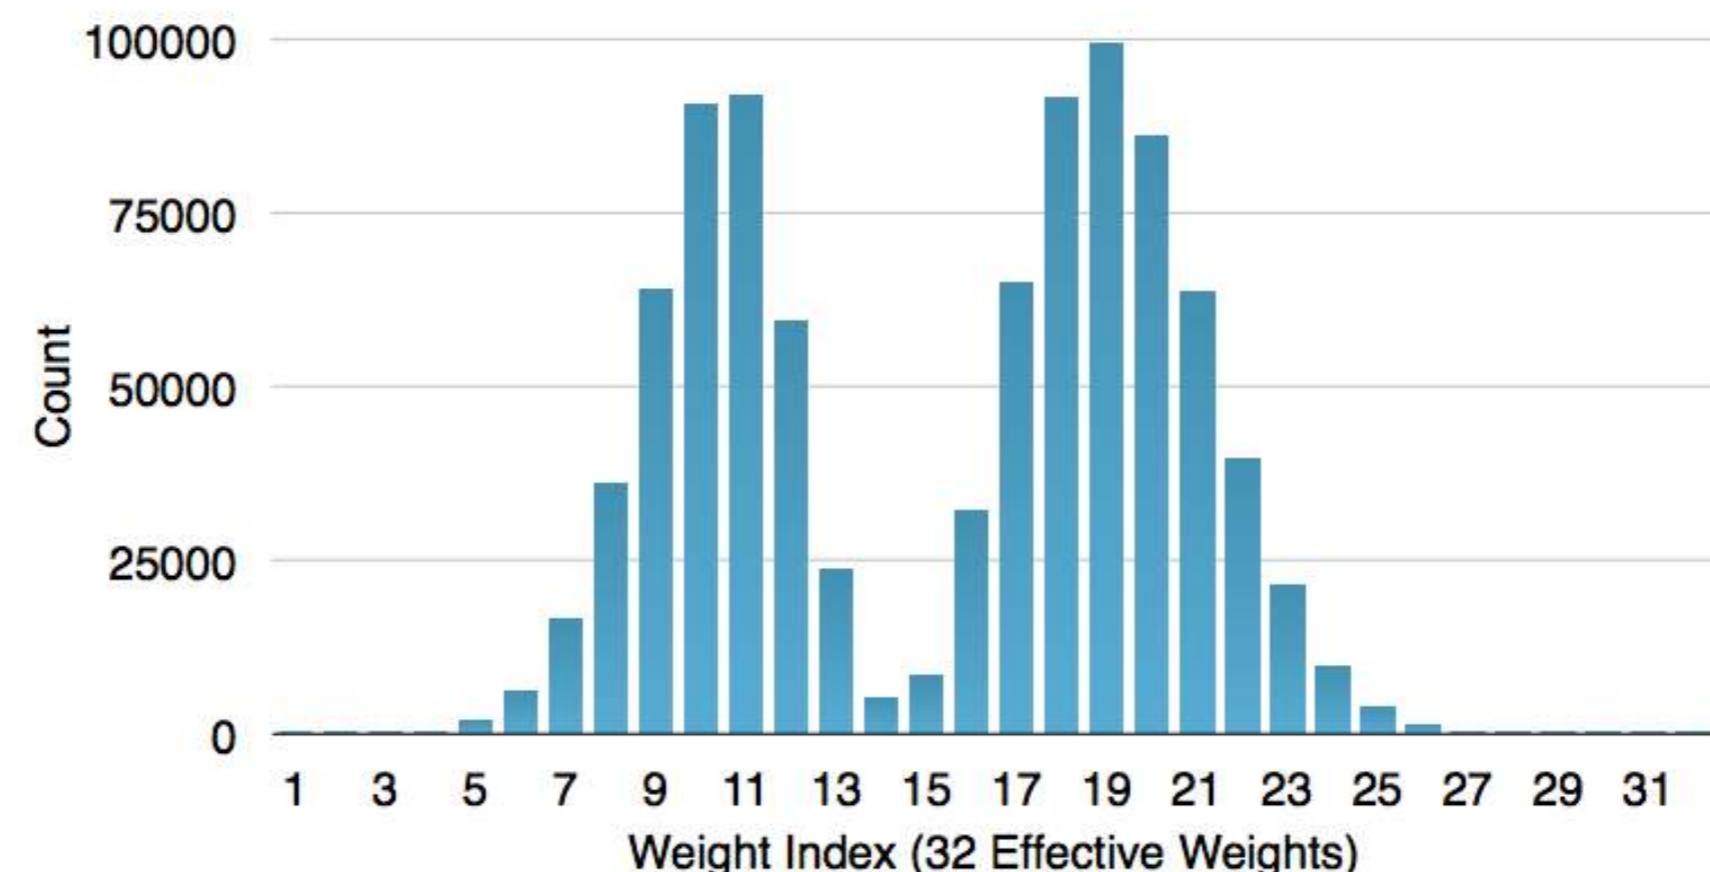
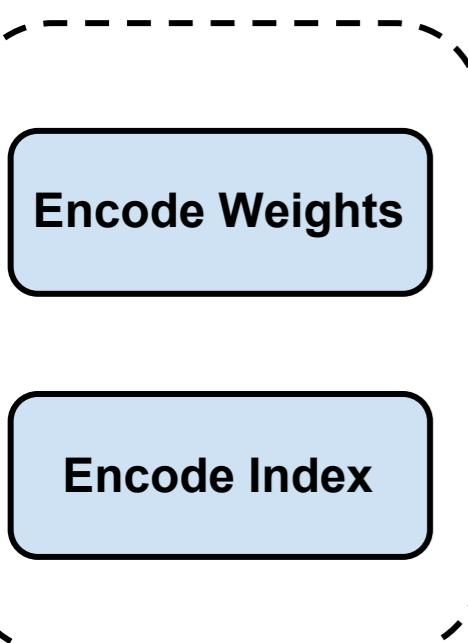
# Pruning + Trained Quantization Work Together



AlexNet on ImageNet

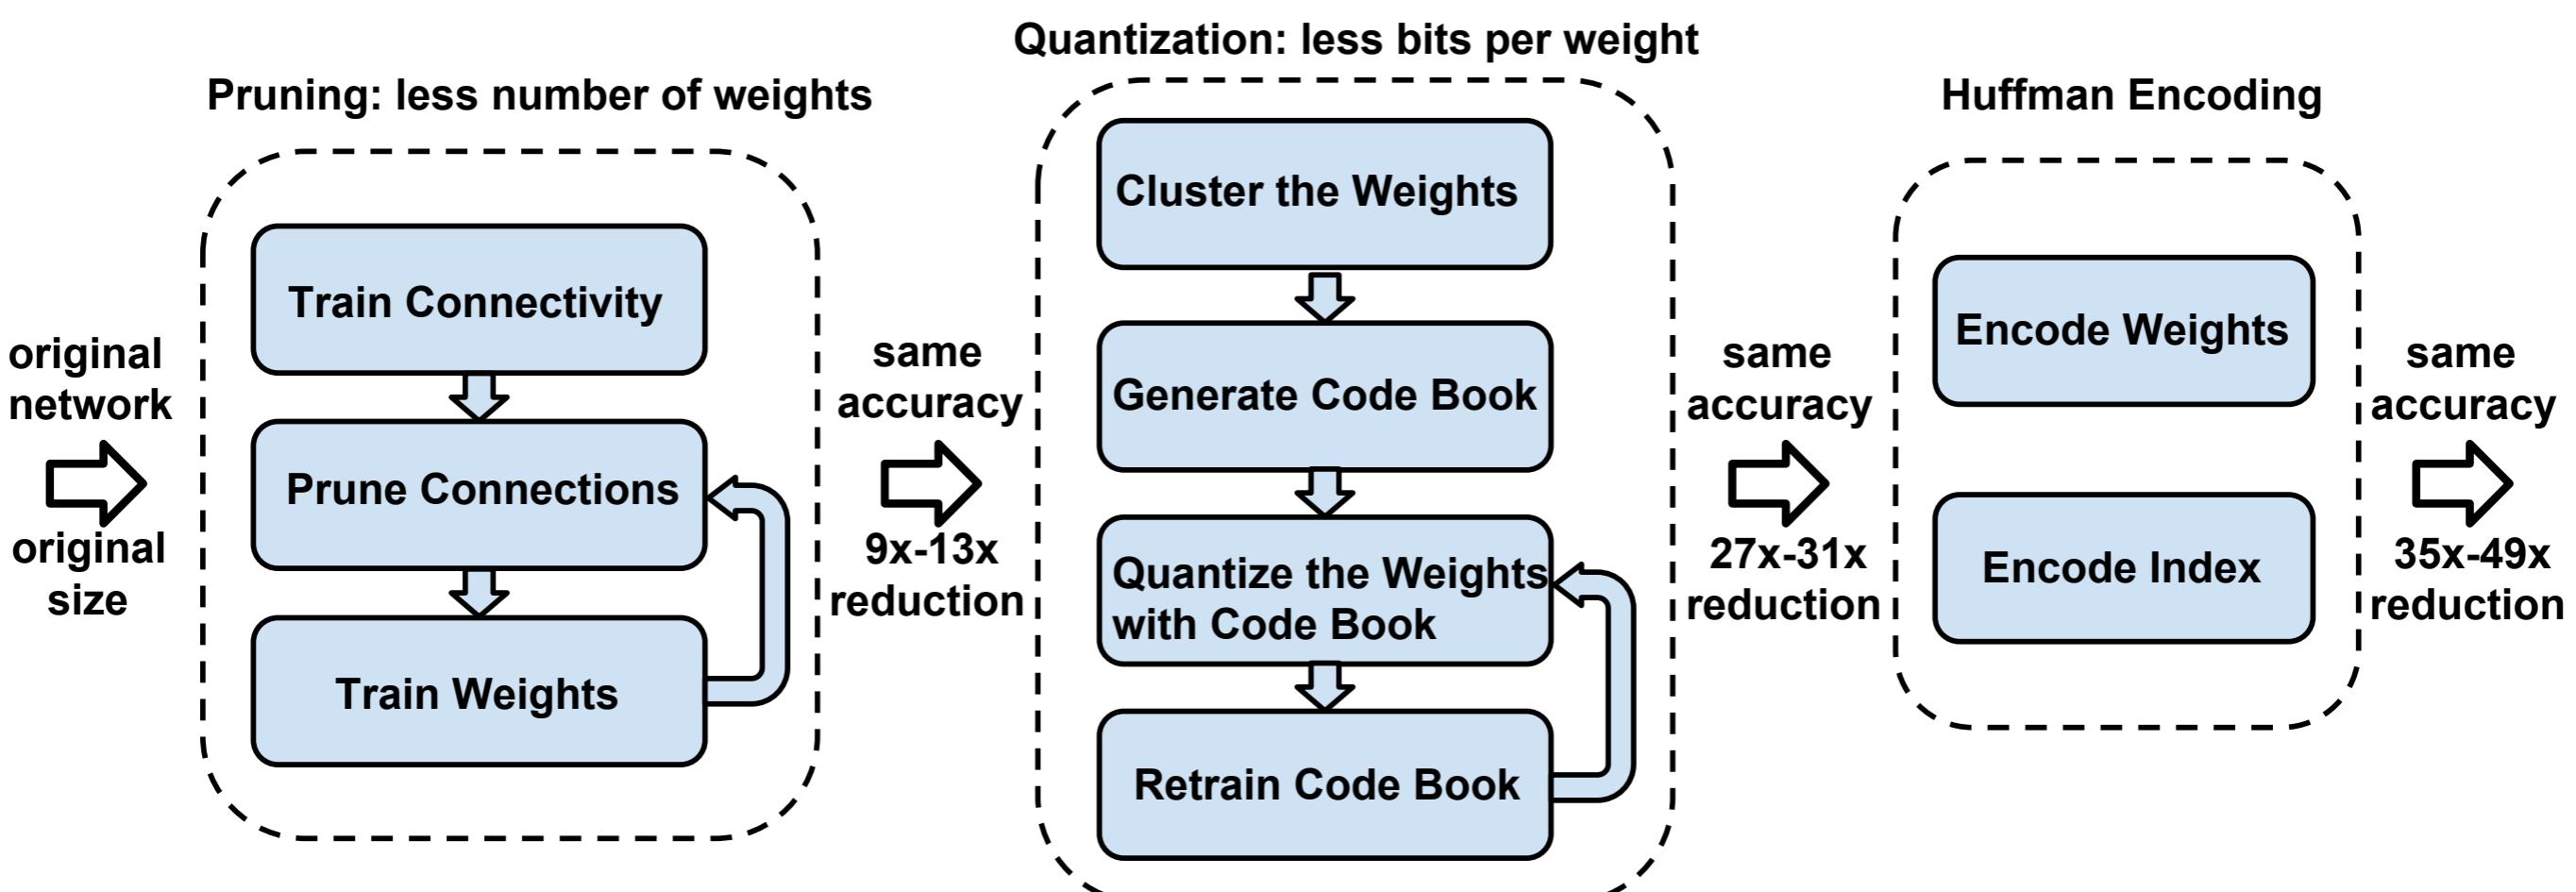
# Huffman Coding

## Huffman Encoding



- In-frequent weights: use more bits to represent
- Frequent weights: use less bits to represent

# Summary of Deep Compression

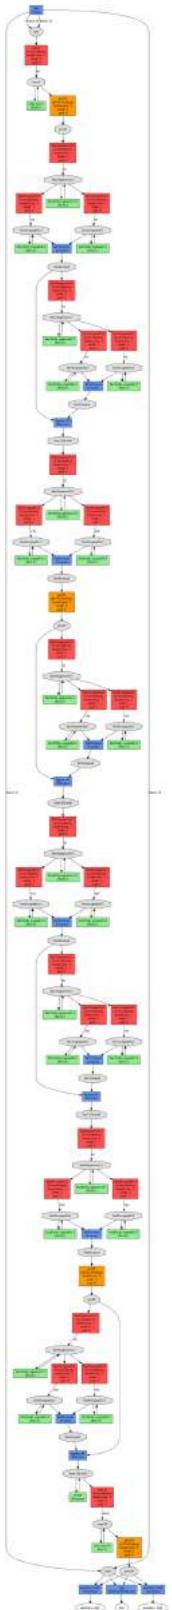
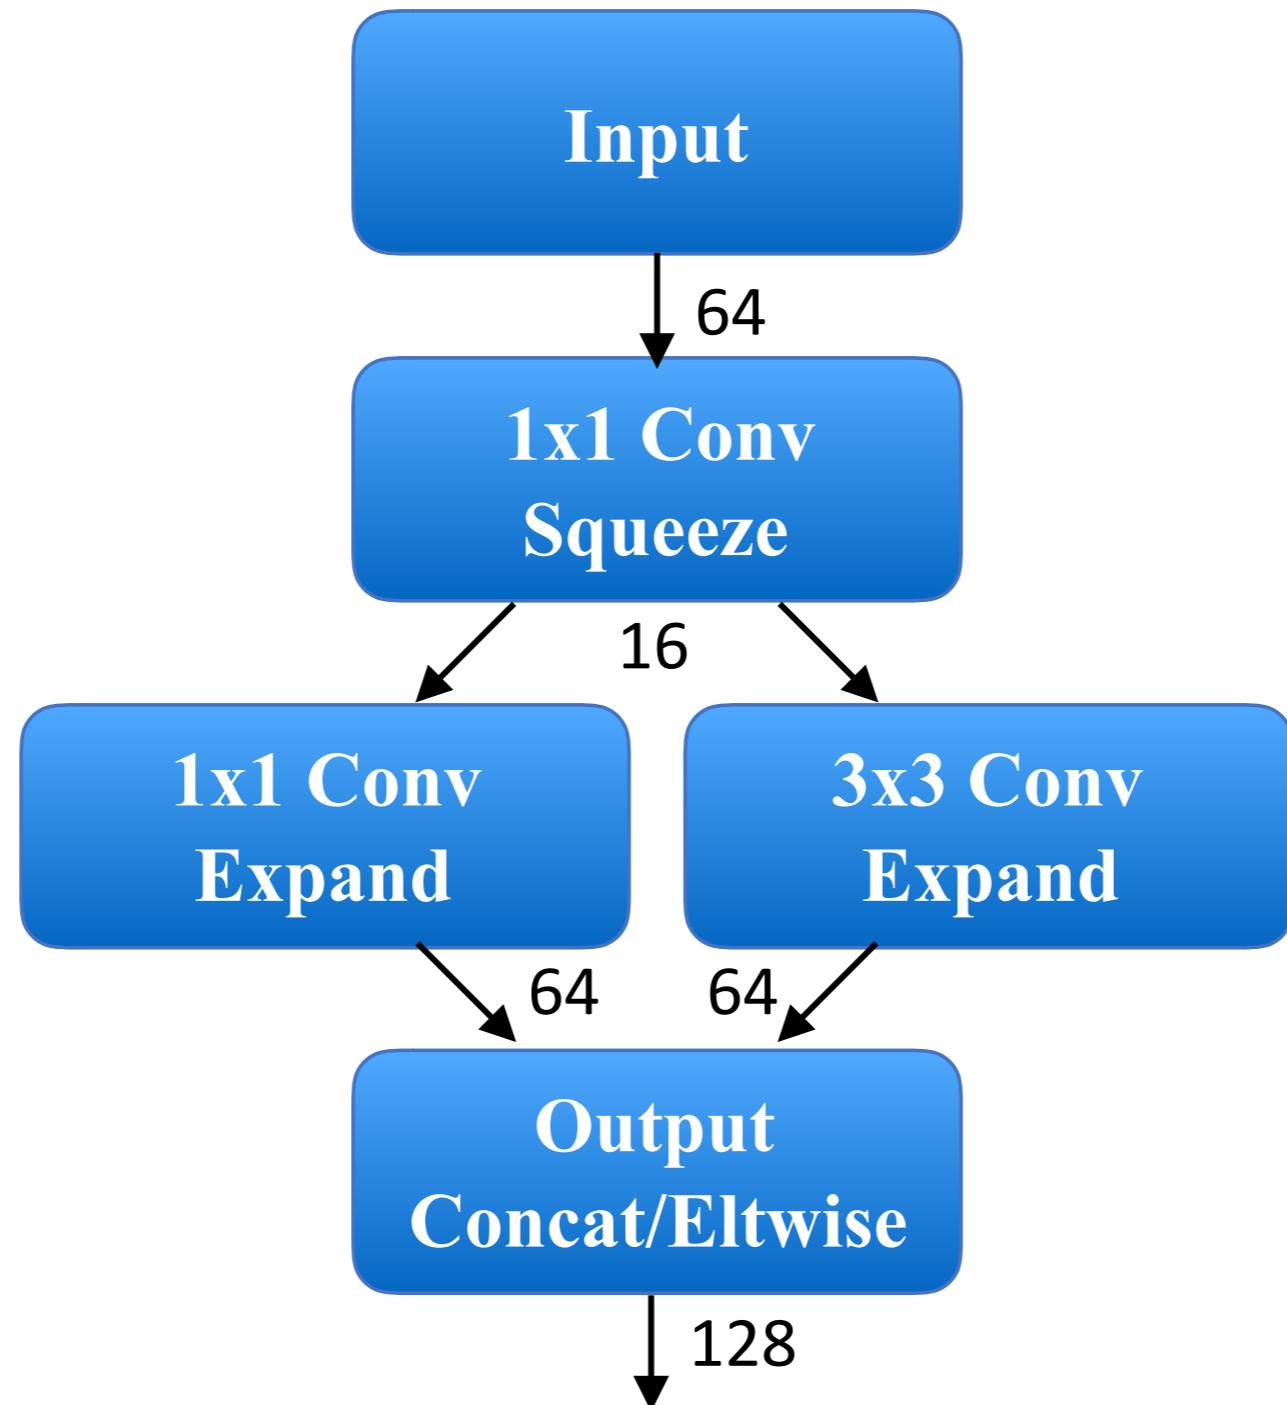


# Results: Compression Ratio

Network	Original Size	Compressed Size	Compression Ratio	Original Accuracy	Compressed Accuracy
LeNet-300	1070KB	→ 27KB	<b>40x</b>	98.36%	→ 98.42%
LeNet-5	1720KB	→ 44KB	<b>39x</b>	99.20%	→ 99.26%
AlexNet	240MB	→ 6.9MB	<b>35x</b>	80.27%	→ 80.30%
VGGNet	550MB	→ 11.3MB	<b>49x</b>	88.68%	→ 89.09%
GoogleNet	28MB	→ 2.8MB	<b>10x</b>	88.90%	→ 88.92%
ResNet-18	44.6MB	→ 4.0MB	<b>11x</b>	89.24%	→ 89.28%

Can we make compact models to begin with?

# SqueezeNet



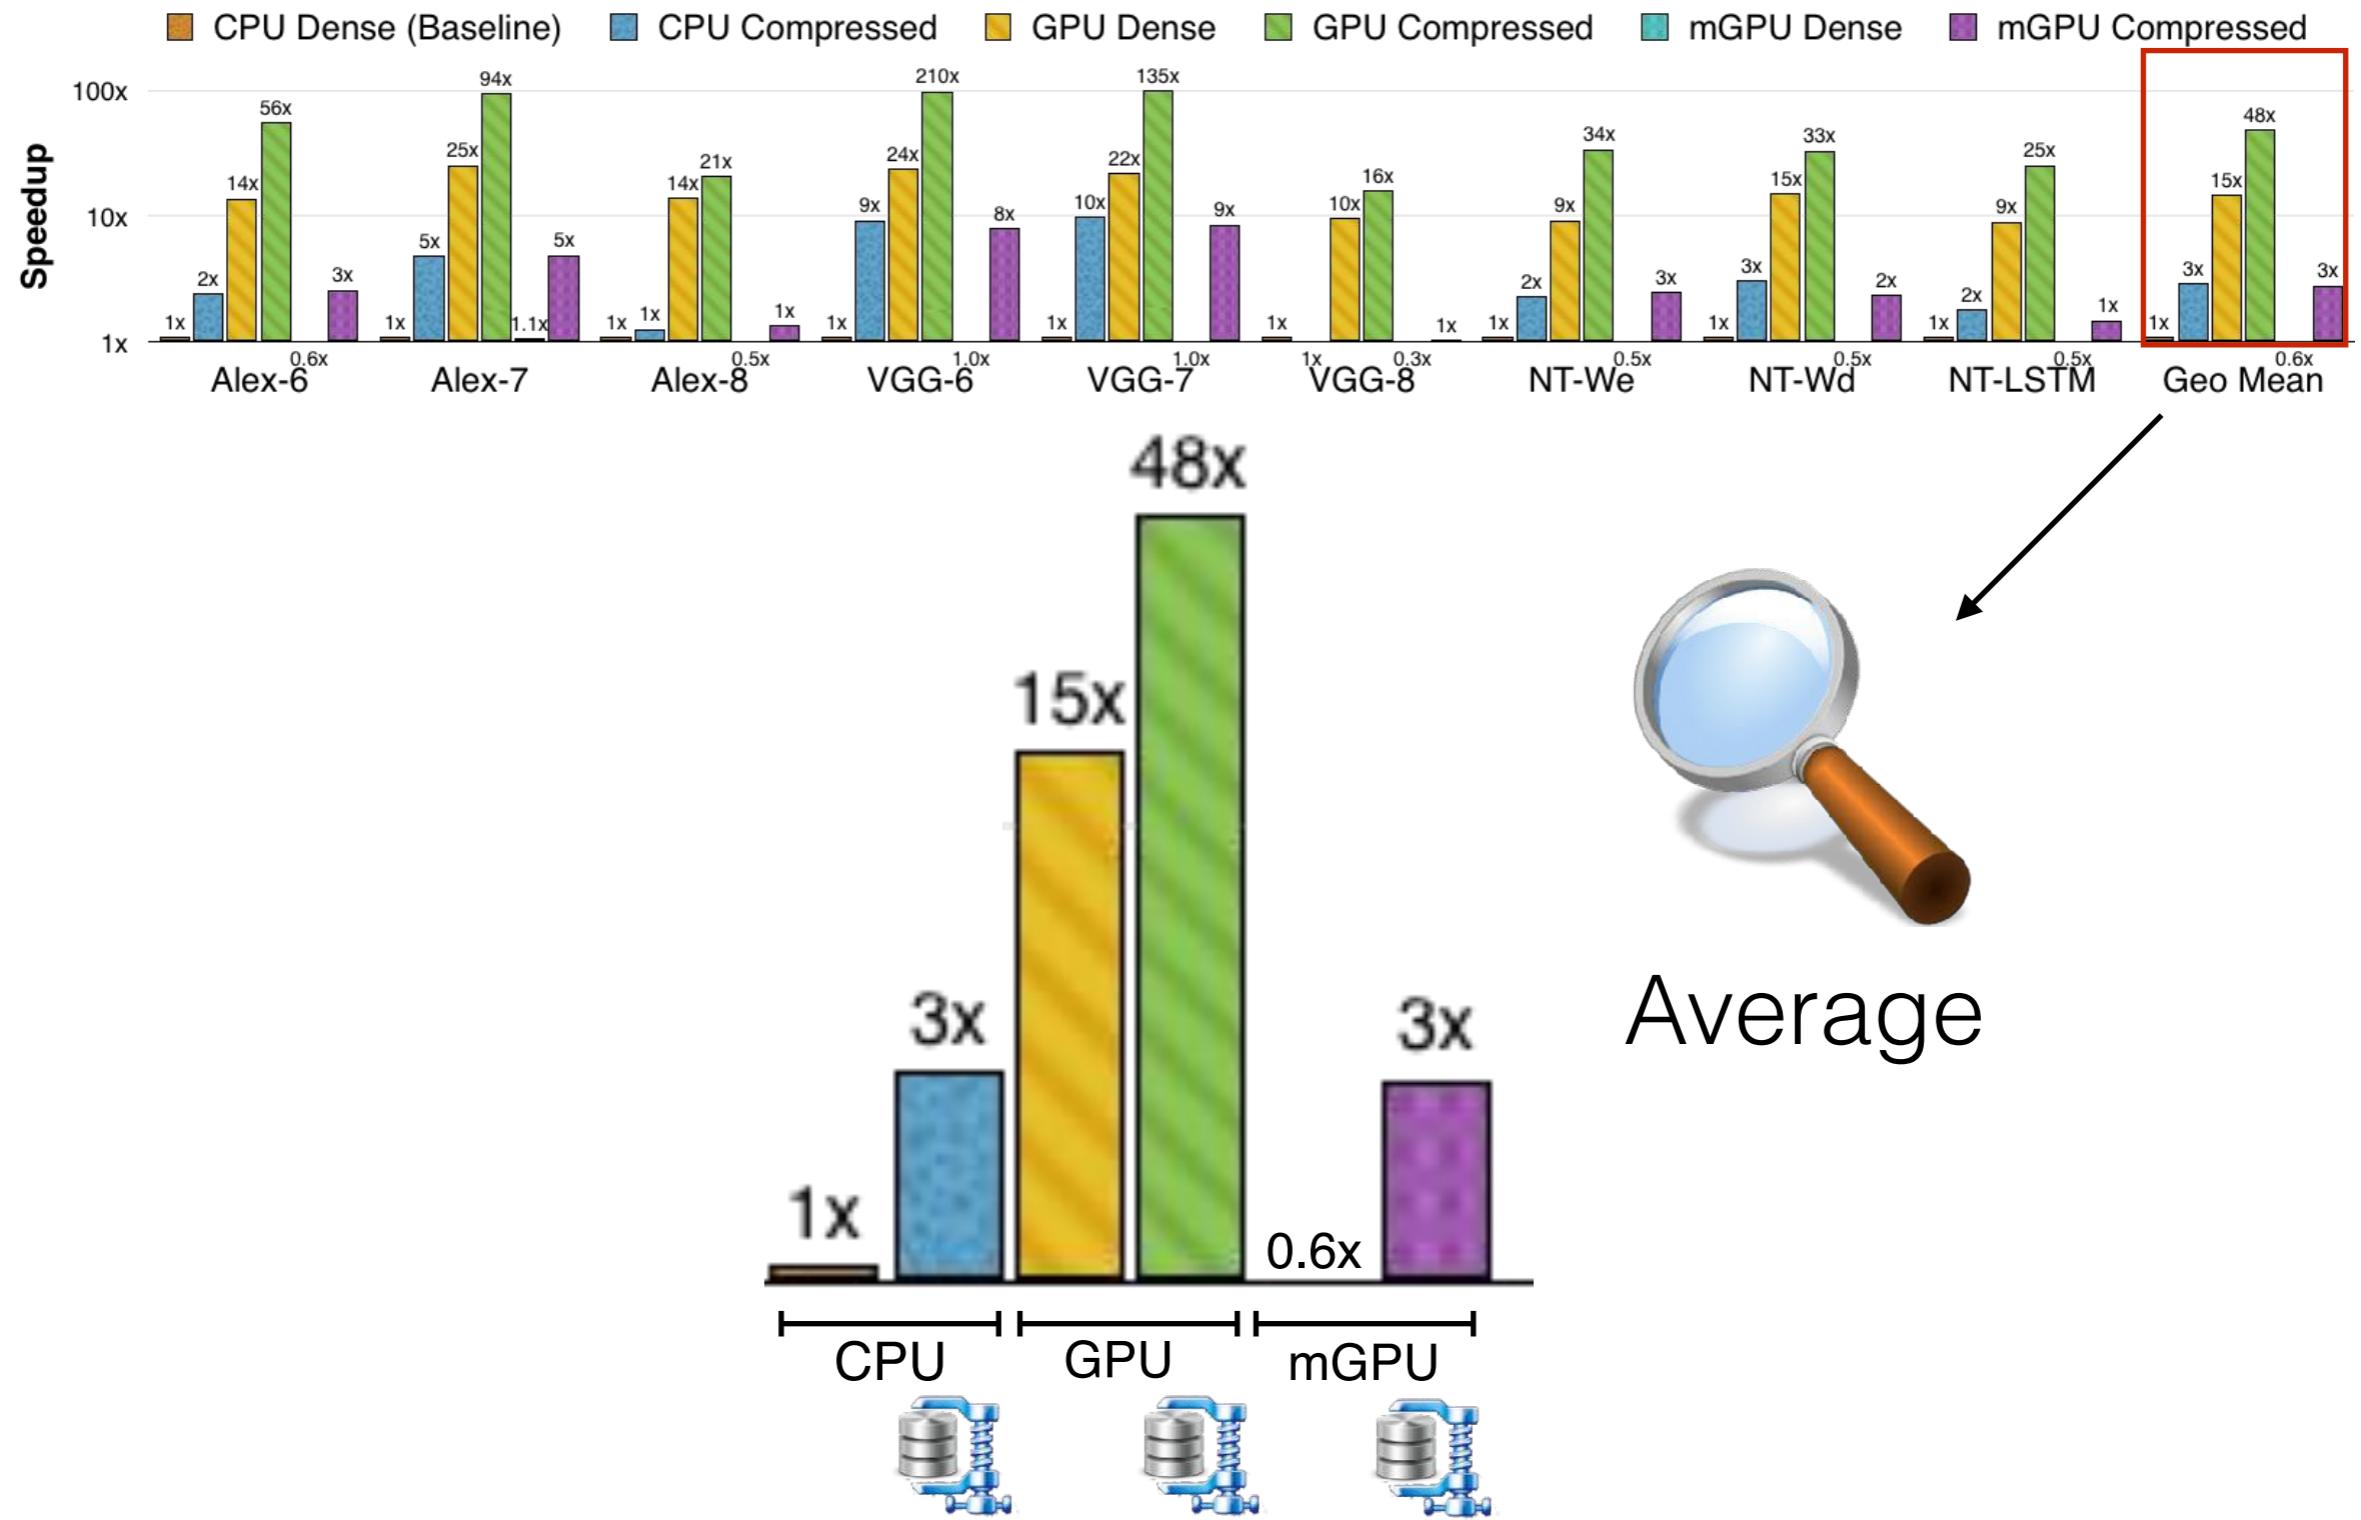
Iandola et al, "SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and <0.5MB model size", arXiv 2016

# Compressing SqueezeNet

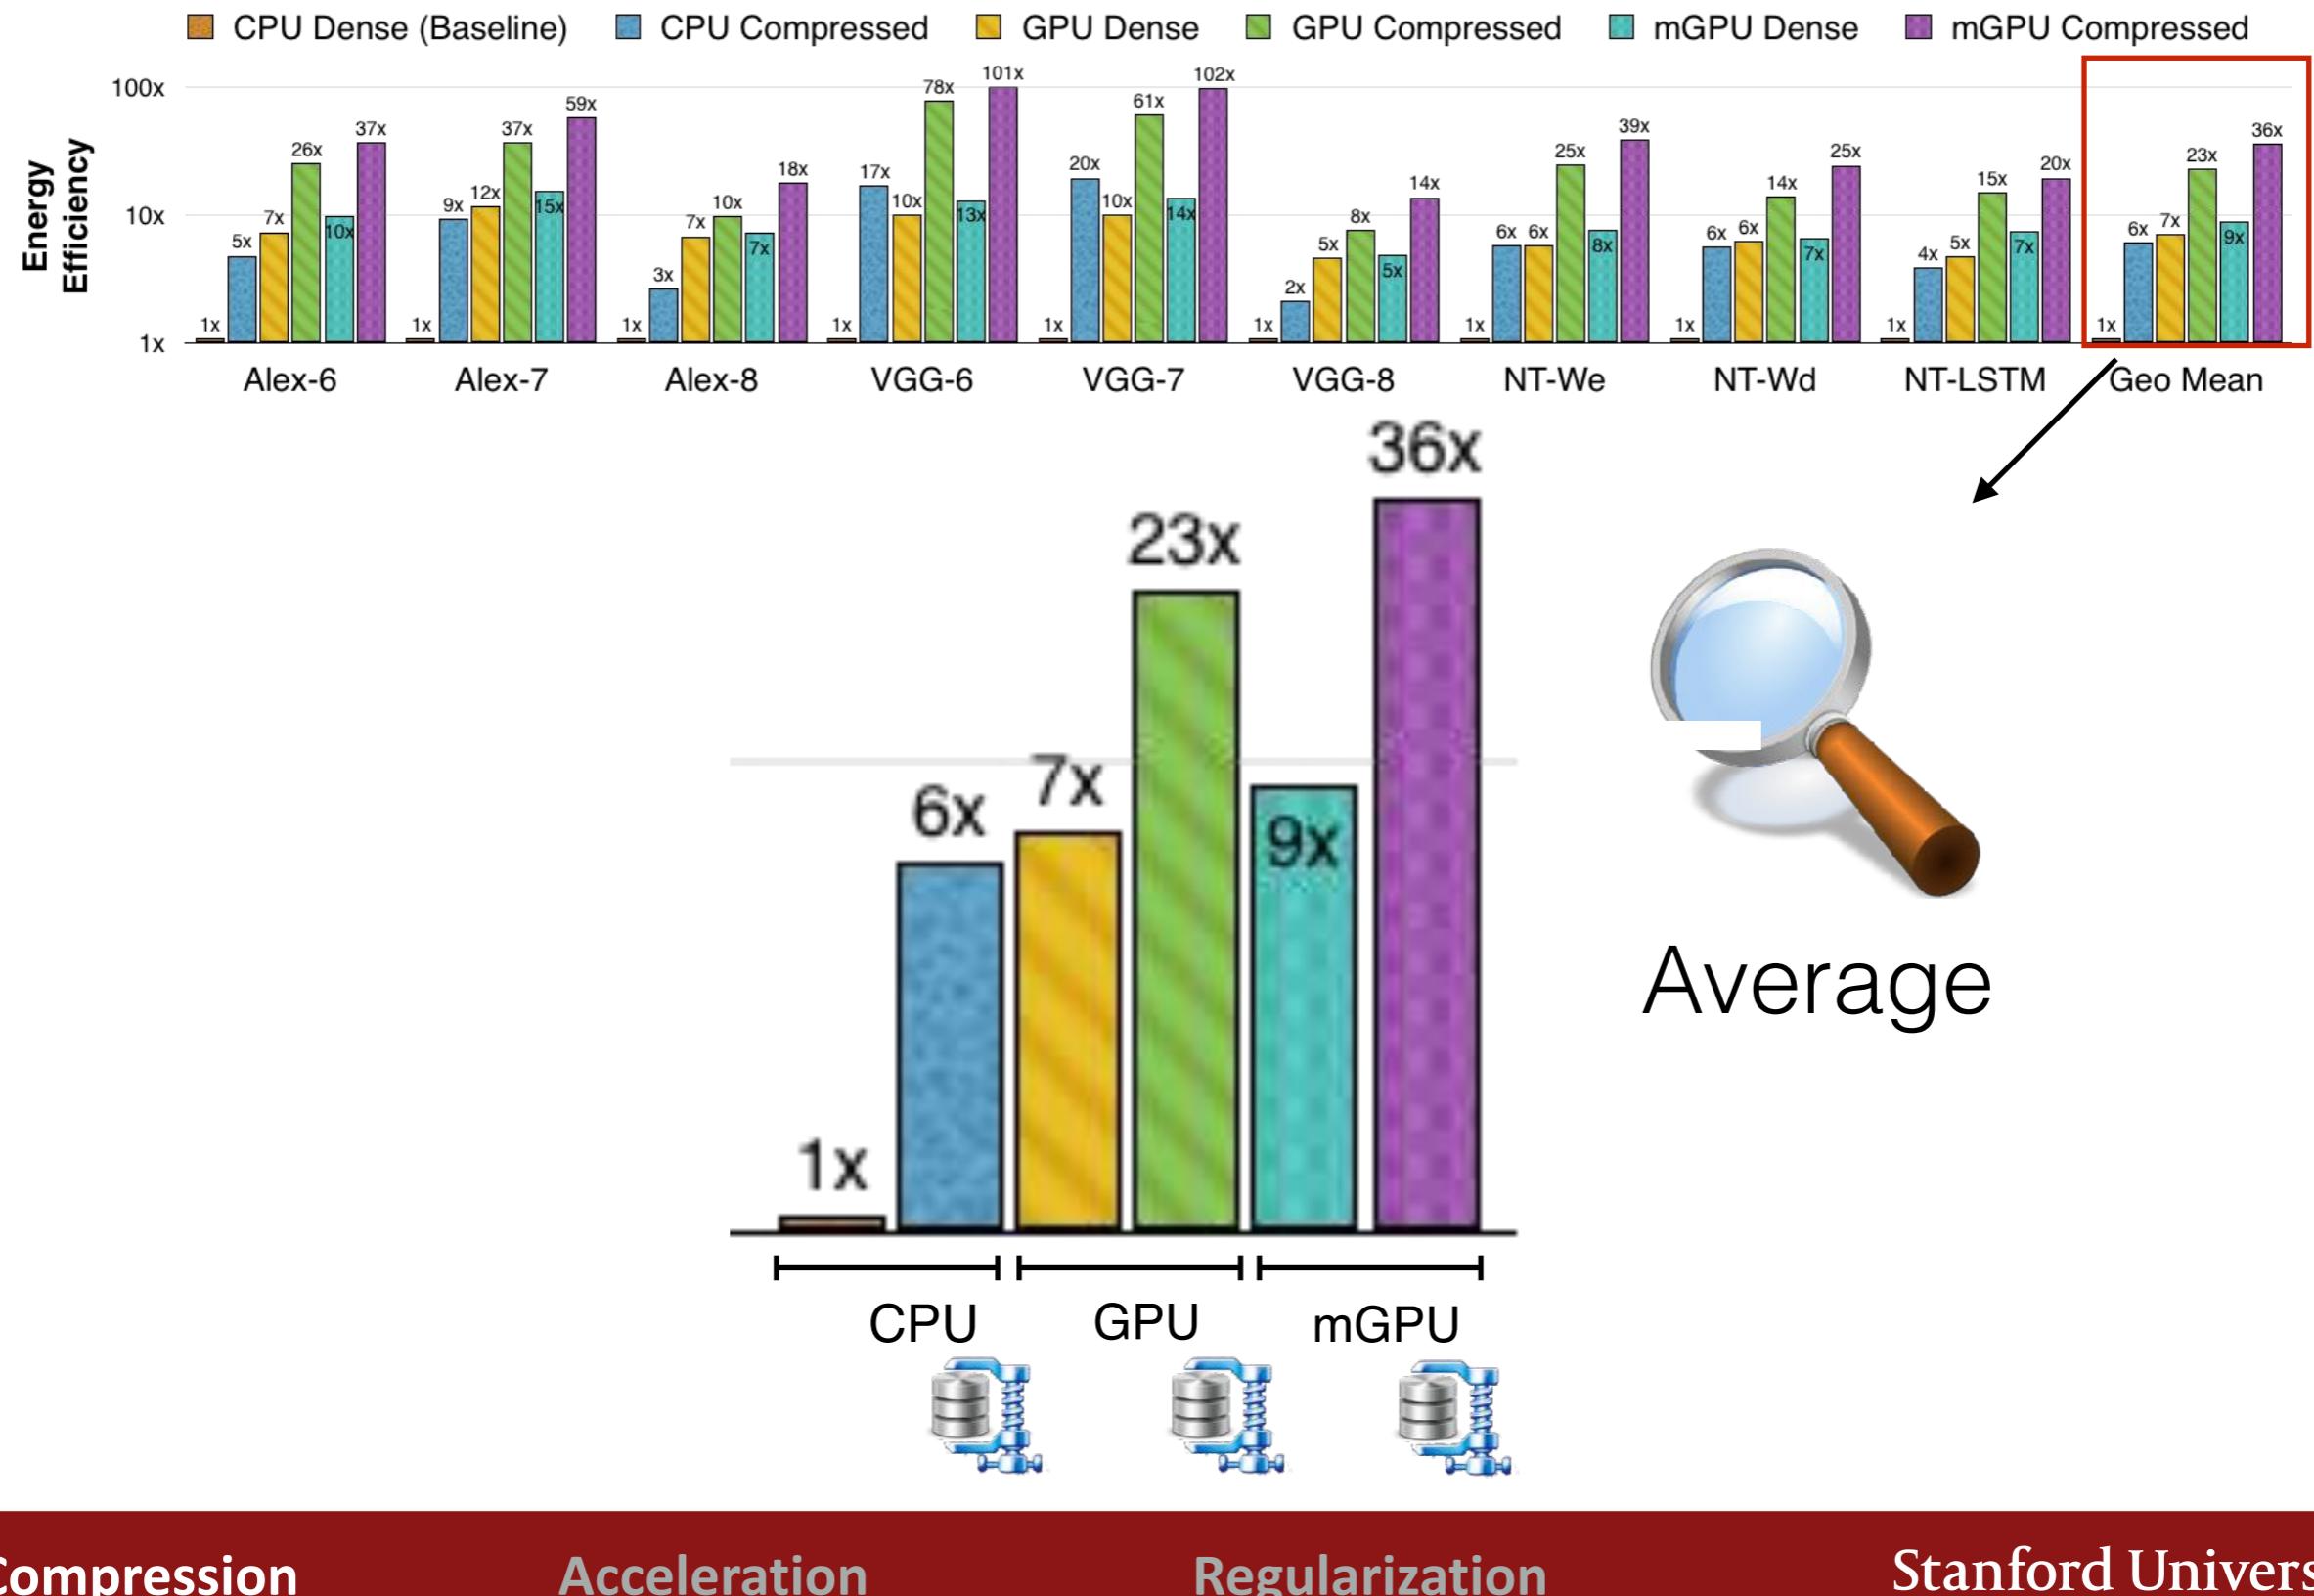
Network	Approach	Size	Ratio	Top-1 Accuracy	Top-5 Accuracy
AlexNet	-	240MB	<b>1x</b>	57.2%	80.3%
AlexNet	SVD	48MB	<b>5x</b>	56.0%	79.4%
AlexNet	Deep Compression	6.9MB	<b>35x</b>	57.2%	80.3%
SqueezeNet	-	4.8MB	<b>50x</b>	57.5%	80.3%
SqueezeNet	Deep Compression	0.47MB	<b>510x</b>	57.5%	80.3%

Iandola et al, "SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and <0.5MB model size", arXiv 2016

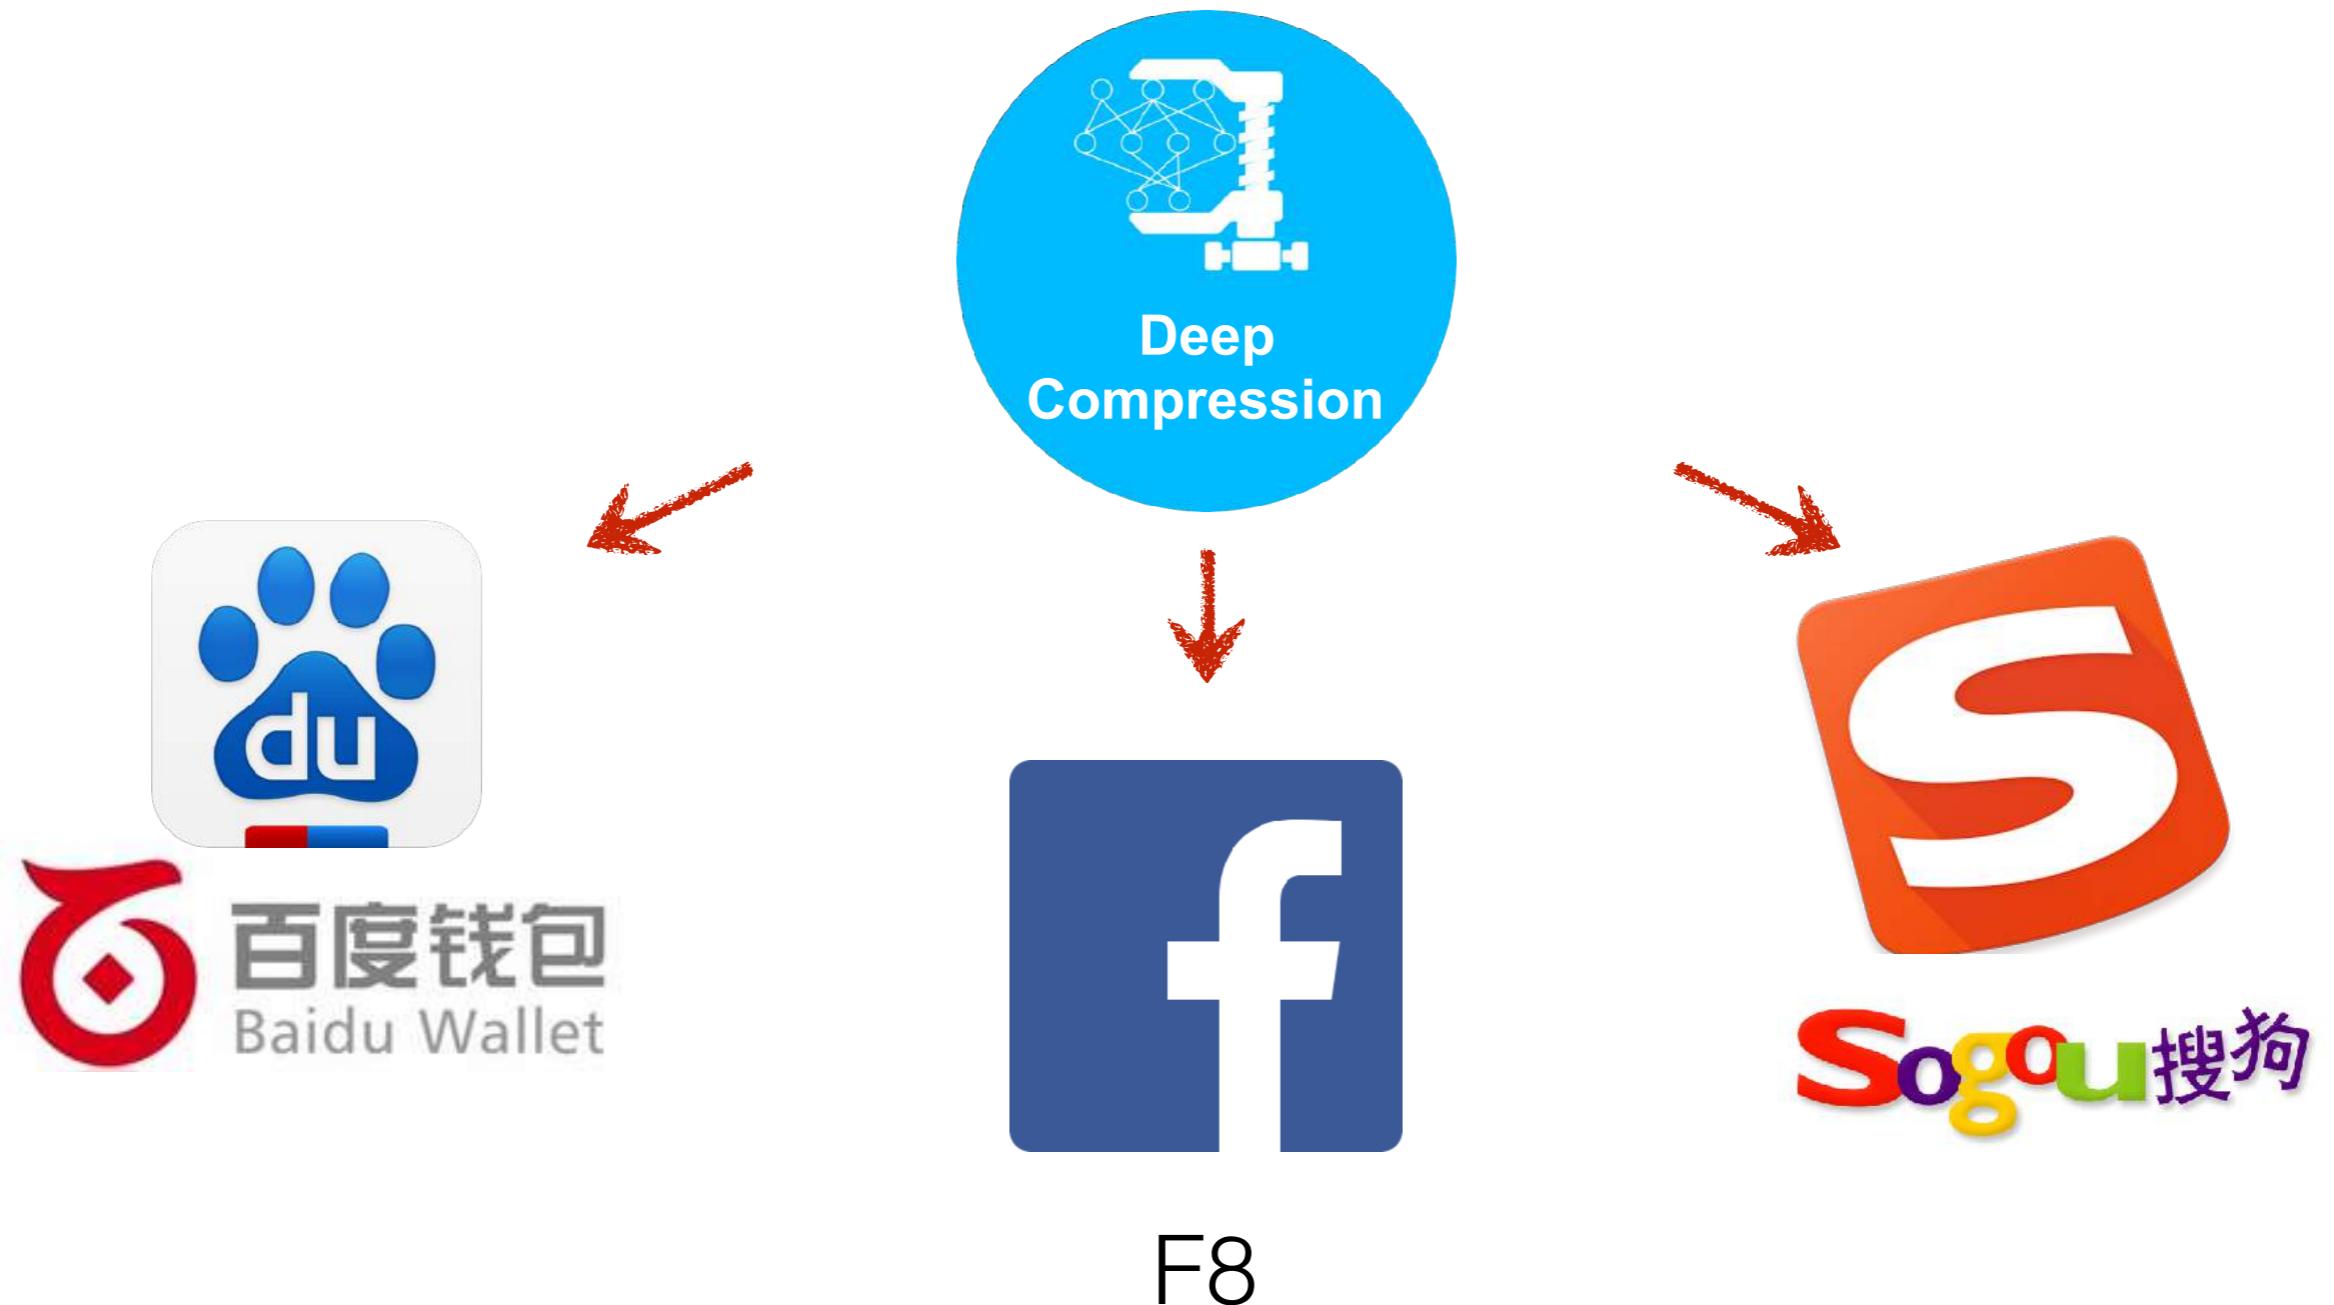
# Results: Speedup



# Results: Energy Efficiency



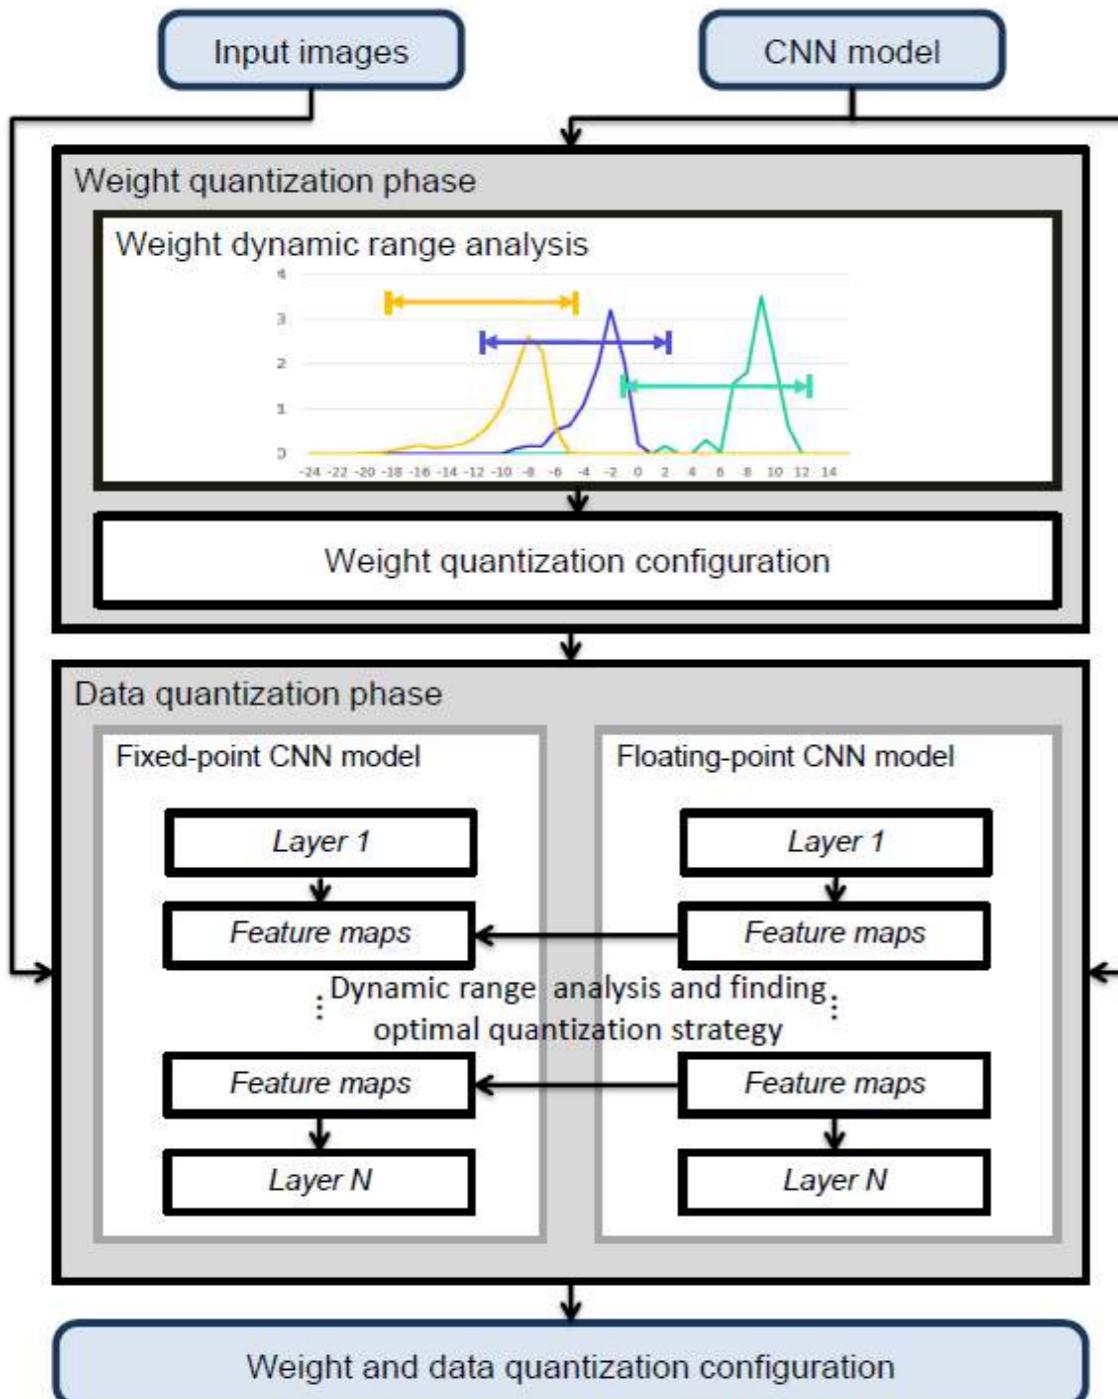
# Deep Compression Applied to Industry



# Part 1: Algorithms for Efficient Inference

- 1. Pruning
- 2. Weight Sharing
- 3. Quantization
- 4. Low Rank Approximation
- 5. Binary / Ternary Net
- 6. Winograd Transformation

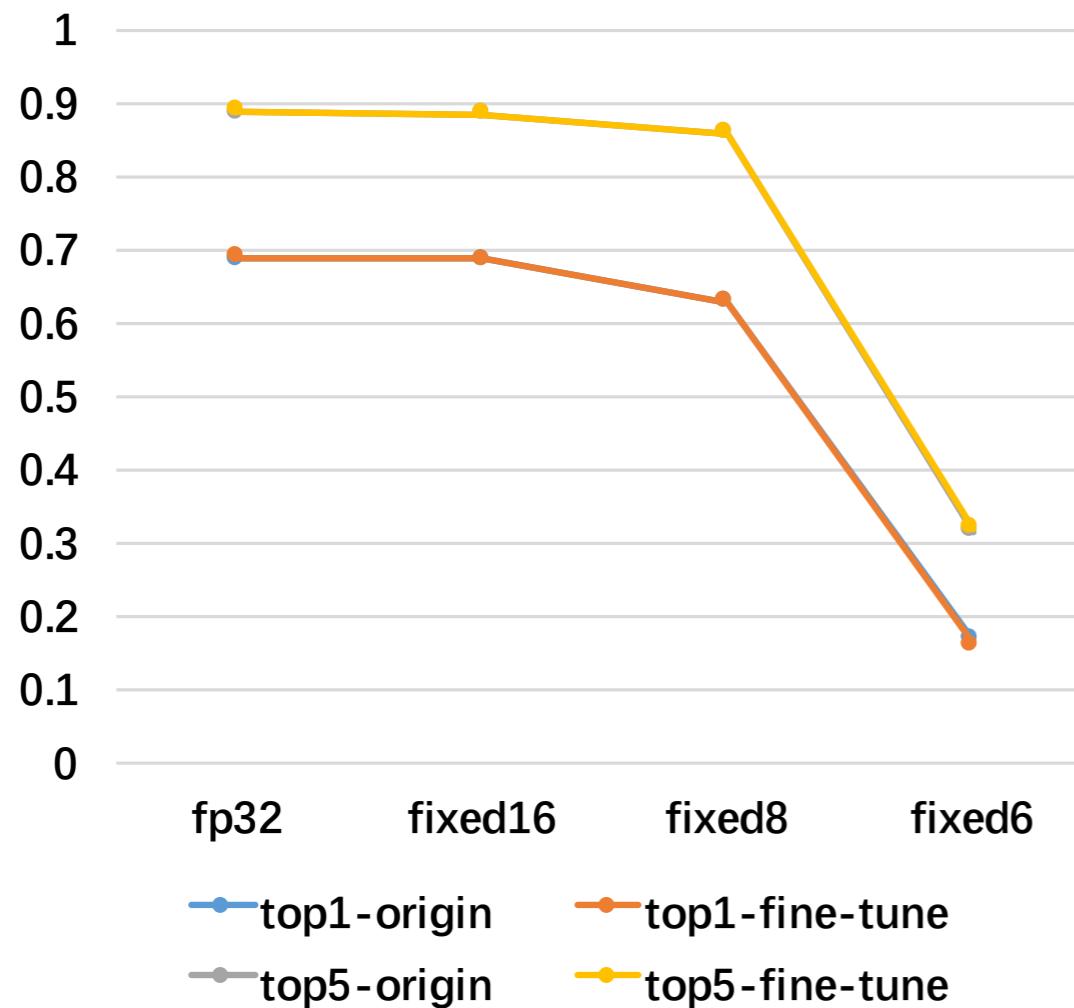
# Quantizing the Weight and Activation



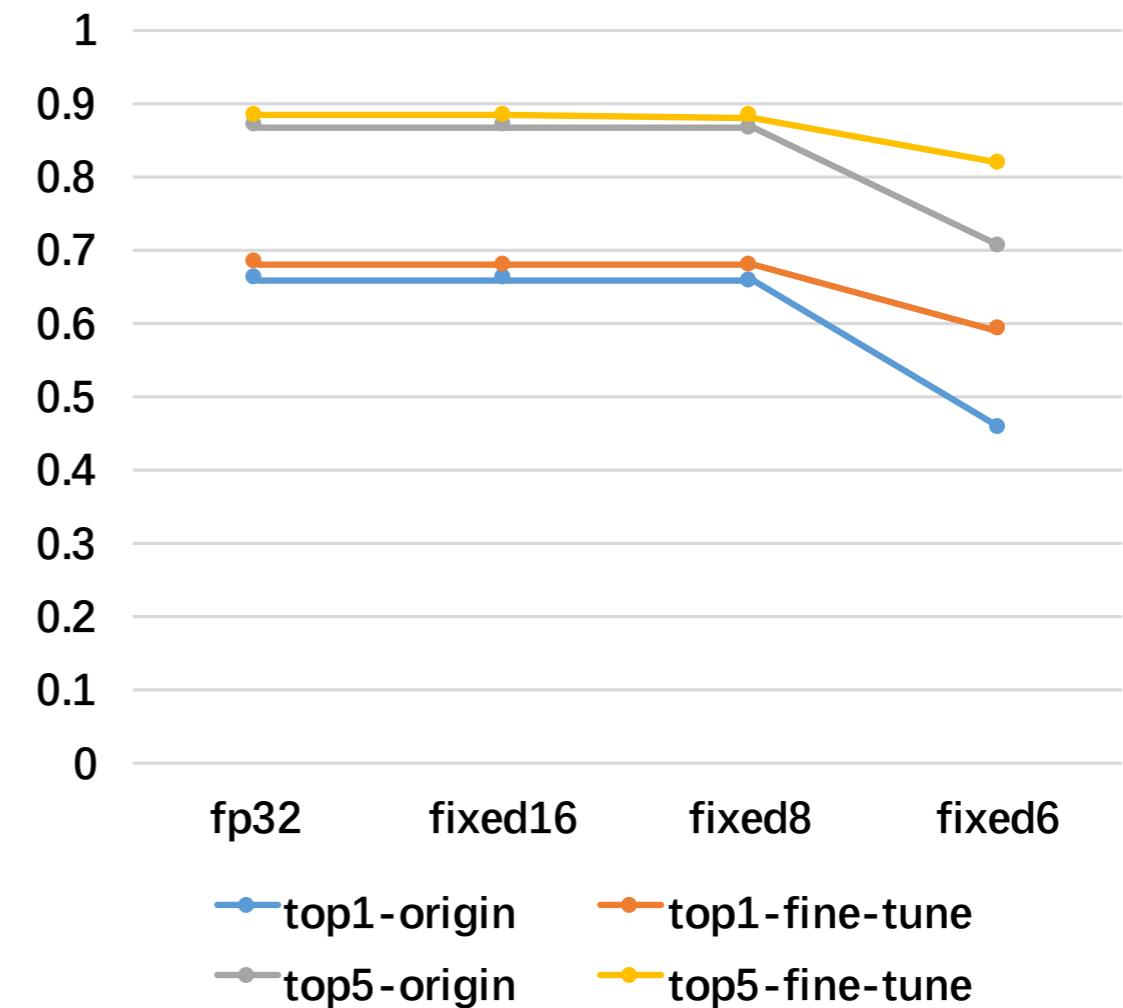
- Train with float
- Quantizing the weight and activation:
  - Gather the statistics for weight and activation
  - Choose proper radix point position
- Fine-tune in float format
- Convert to fixed-point format

# Quantization Result

GoogLeNet



VGG-16

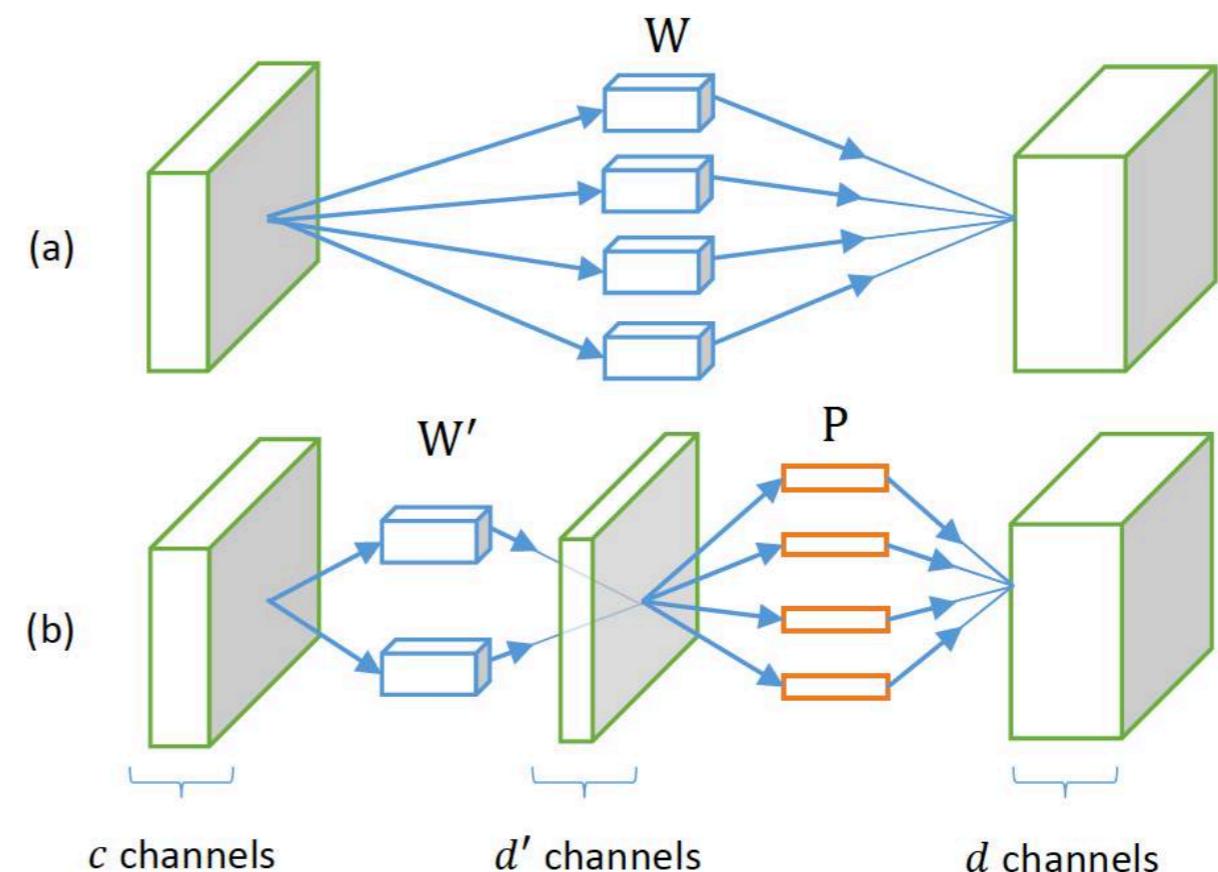


# Part 1: Algorithms for Efficient Inference

- 1. Pruning
- 2. Weight Sharing
- 3. Quantization
- 4. Low Rank Approximation
- 5. Binary / Ternary Net
- 6. Winograd Transformation

# Low Rank Approximation for Conv

- Layer responses lie in a low-rank subspace
- Decompose a convolutional layer with  $d$  filters with filter size  $k \times k \times c$  to
  - A layer with  $d'$  filters ( $k \times k \times c$ )
  - A layer with  $d$  filter ( $1 \times 1 \times d'$ )



# Low Rank Approximation for Conv

speedup	rank sel.	Conv1	Conv2	Conv3	Conv4	Conv5	Conv6	Conv7	err. ↑ %
2×	no	32	110	199	219	219	219	219	1.18
2×	yes	32	83	182	211	239	237	253	<b>0.93</b>
2.4×	no	32	96	174	191	191	191	191	1.77
2.4×	yes	32	74	162	187	207	205	219	<b>1.35</b>
3×	no	32	77	139	153	153	153	153	2.56
3×	yes	32	62	138	149	166	162	167	<b>2.34</b>
4×	no	32	57	104	115	115	115	115	4.32
4×	yes	32	50	112	114	122	117	119	<b>4.20</b>
5×	no	32	46	83	92	92	92	92	6.53
5×	yes	32	41	94	93	98	92	90	<b>6.47</b>

# Low Rank Approximation for FC

Build a mapping from row / column indices of matrix  $\mathbf{W} = [W(x, y)]$  to vectors  $\mathbf{i}$  and  $\mathbf{j}$ :  $x \leftrightarrow \mathbf{i} = (i_1, \dots, i_d)$  and  $y \leftrightarrow \mathbf{j} = (j_1, \dots, j_d)$ .

TT-format for matrix  $\mathbf{W}$ :

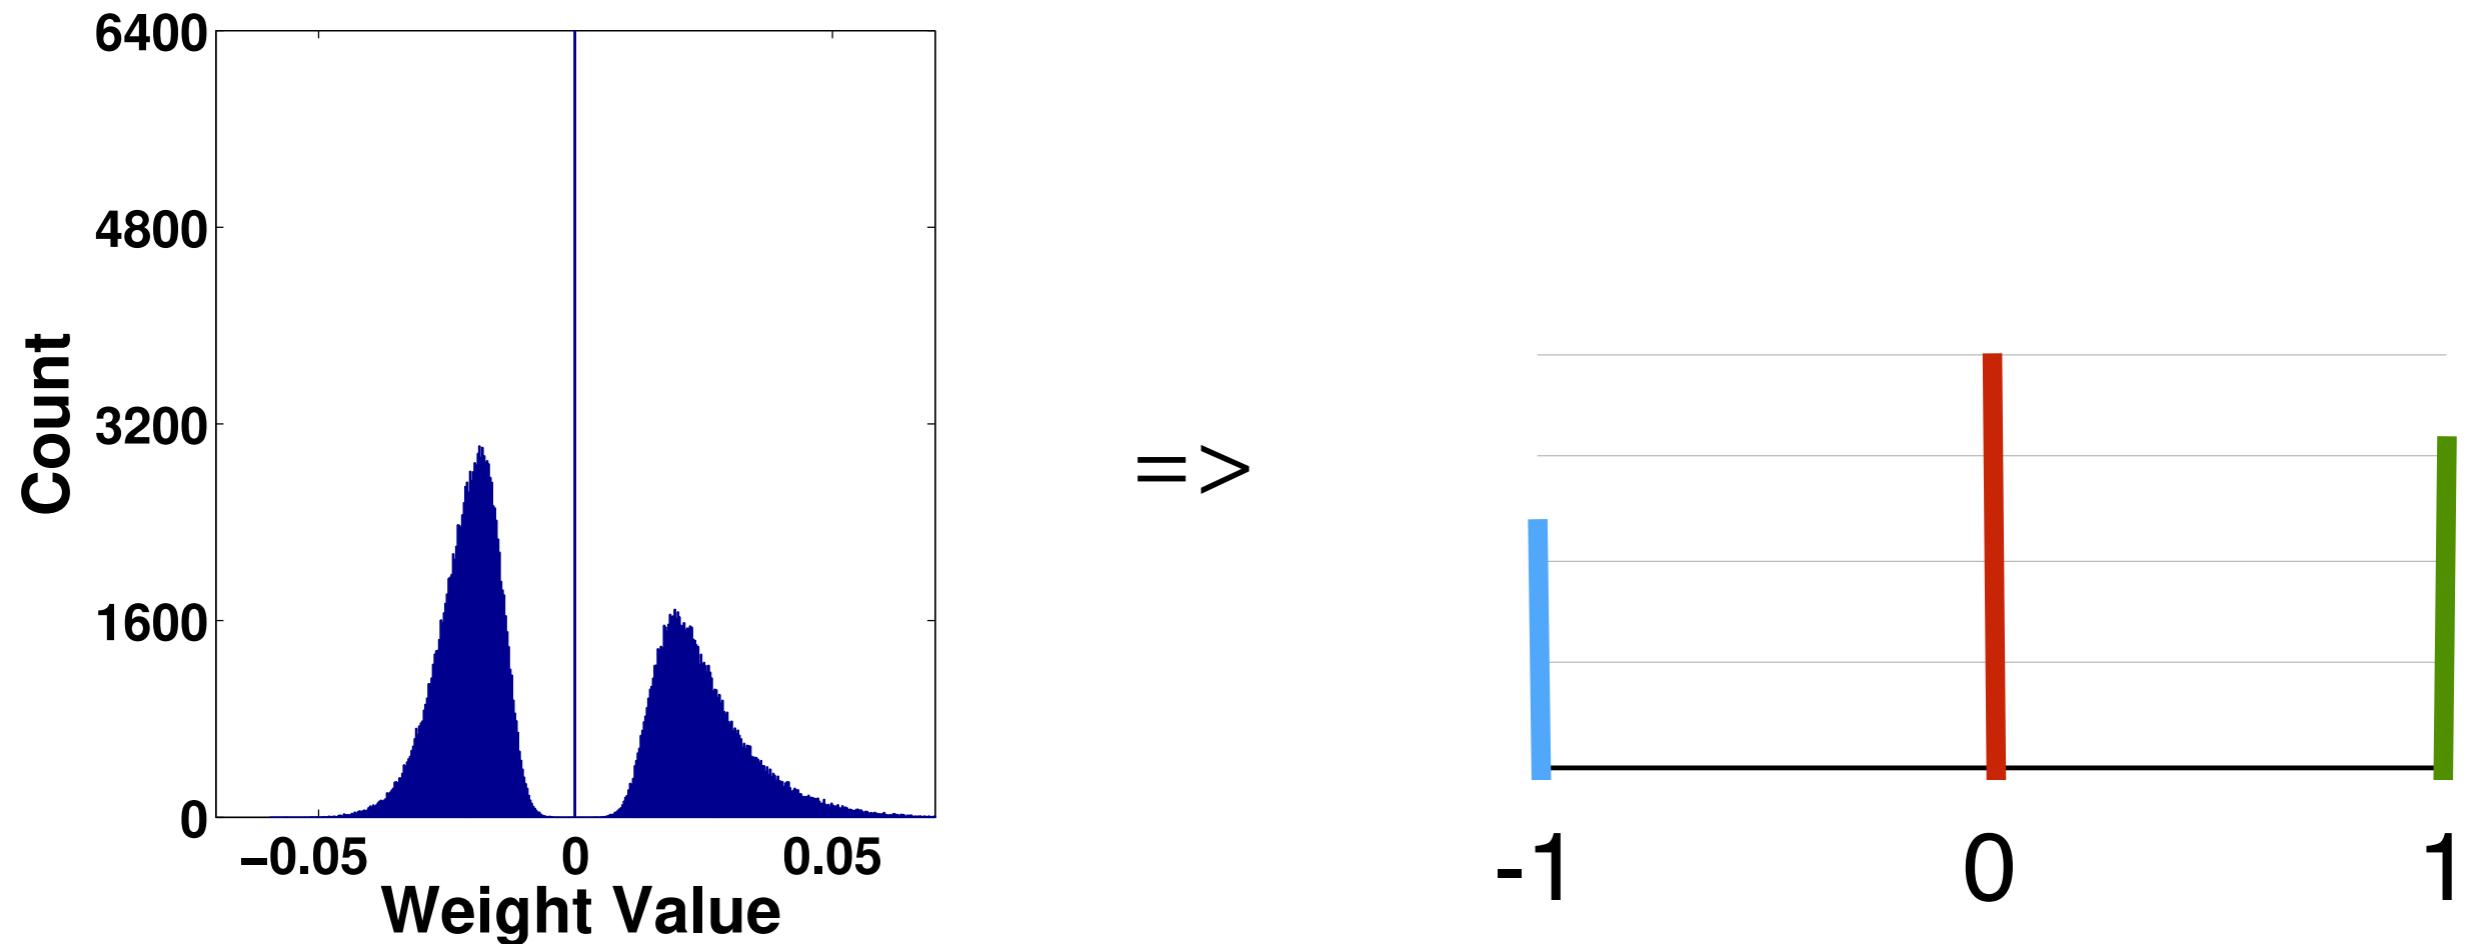
$$\mathbf{W}(i_1, \dots, i_d; j_1, \dots, j_d) = \mathbf{W}(x(\mathbf{i}), y(\mathbf{j})) = \underbrace{\mathbf{G}_1[i_1, j_1]}_{1 \times r} \underbrace{\mathbf{G}_2[i_2, j_2]}_{r \times r} \dots \underbrace{\mathbf{G}_d[i_d, j_d]}_{r \times 1}$$

Type	1 im. time (ms)	100 im. time (ms)
CPU fully-connected layer	16.1	97.2
CPU TT-layer	1.2	94.7
GPU fully-connected layer	2.7	33
GPU TT-layer	1.9	12.9

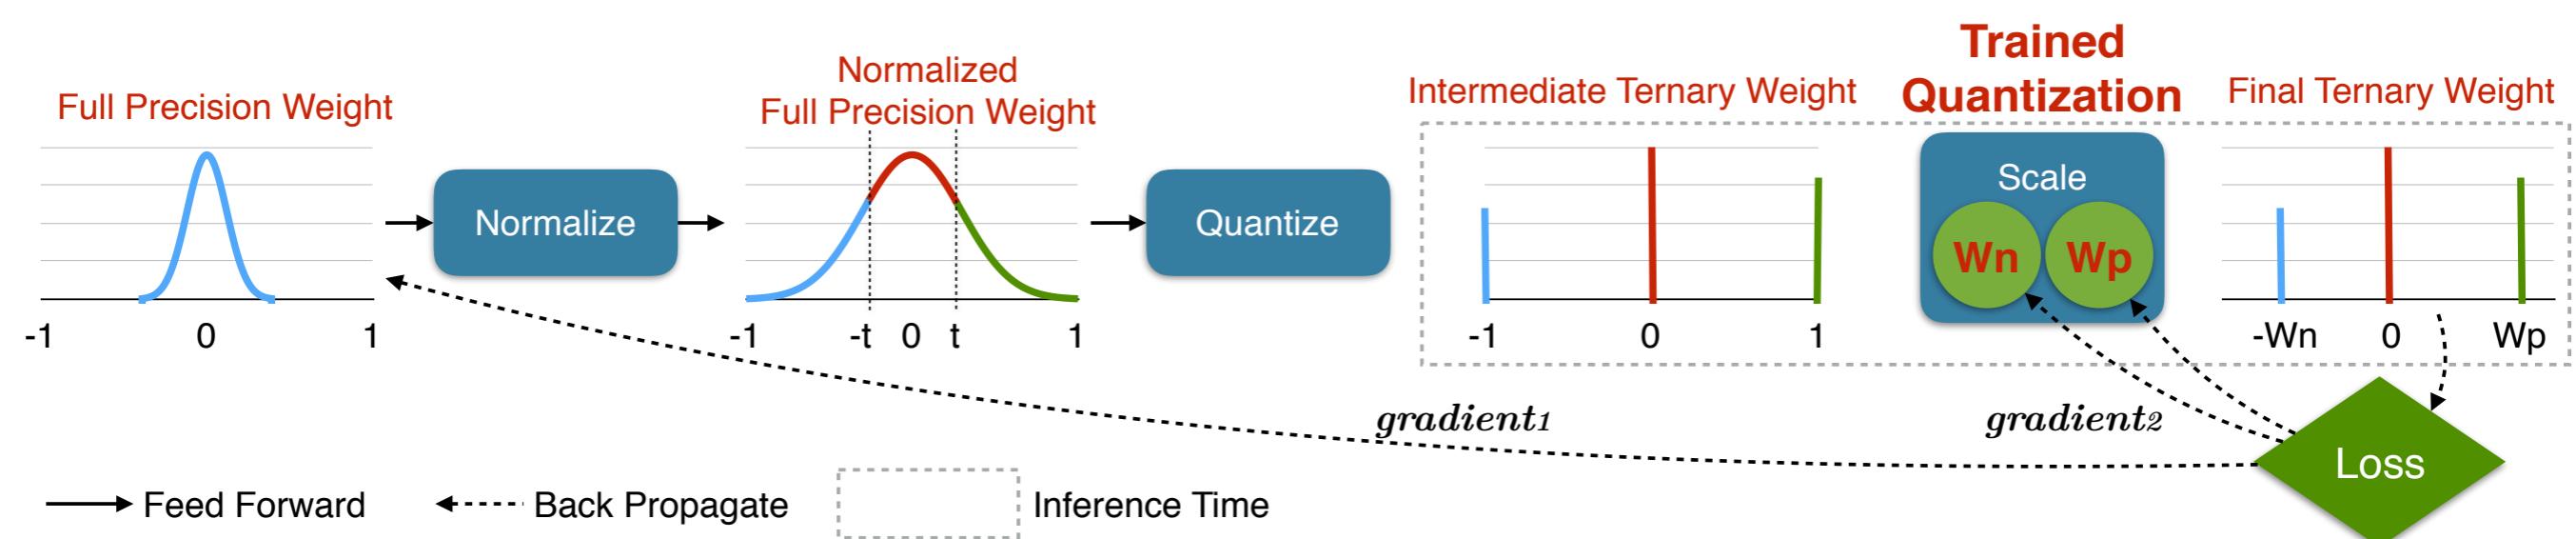
# Part 1: Algorithms for Efficient Inference

- 1. Pruning
- 2. Weight Sharing
- 3. Quantization
- 4. Low Rank Approximation
- 5. Binary / Ternary Net
- 6. Winograd Transformation

# Binary / Ternary Net: Motivation



# Trained Ternary Quantization



Zhu, Han, Mao, Dally. Trained Ternary Quantization, ICLR'17

# Weight Evolution during Training

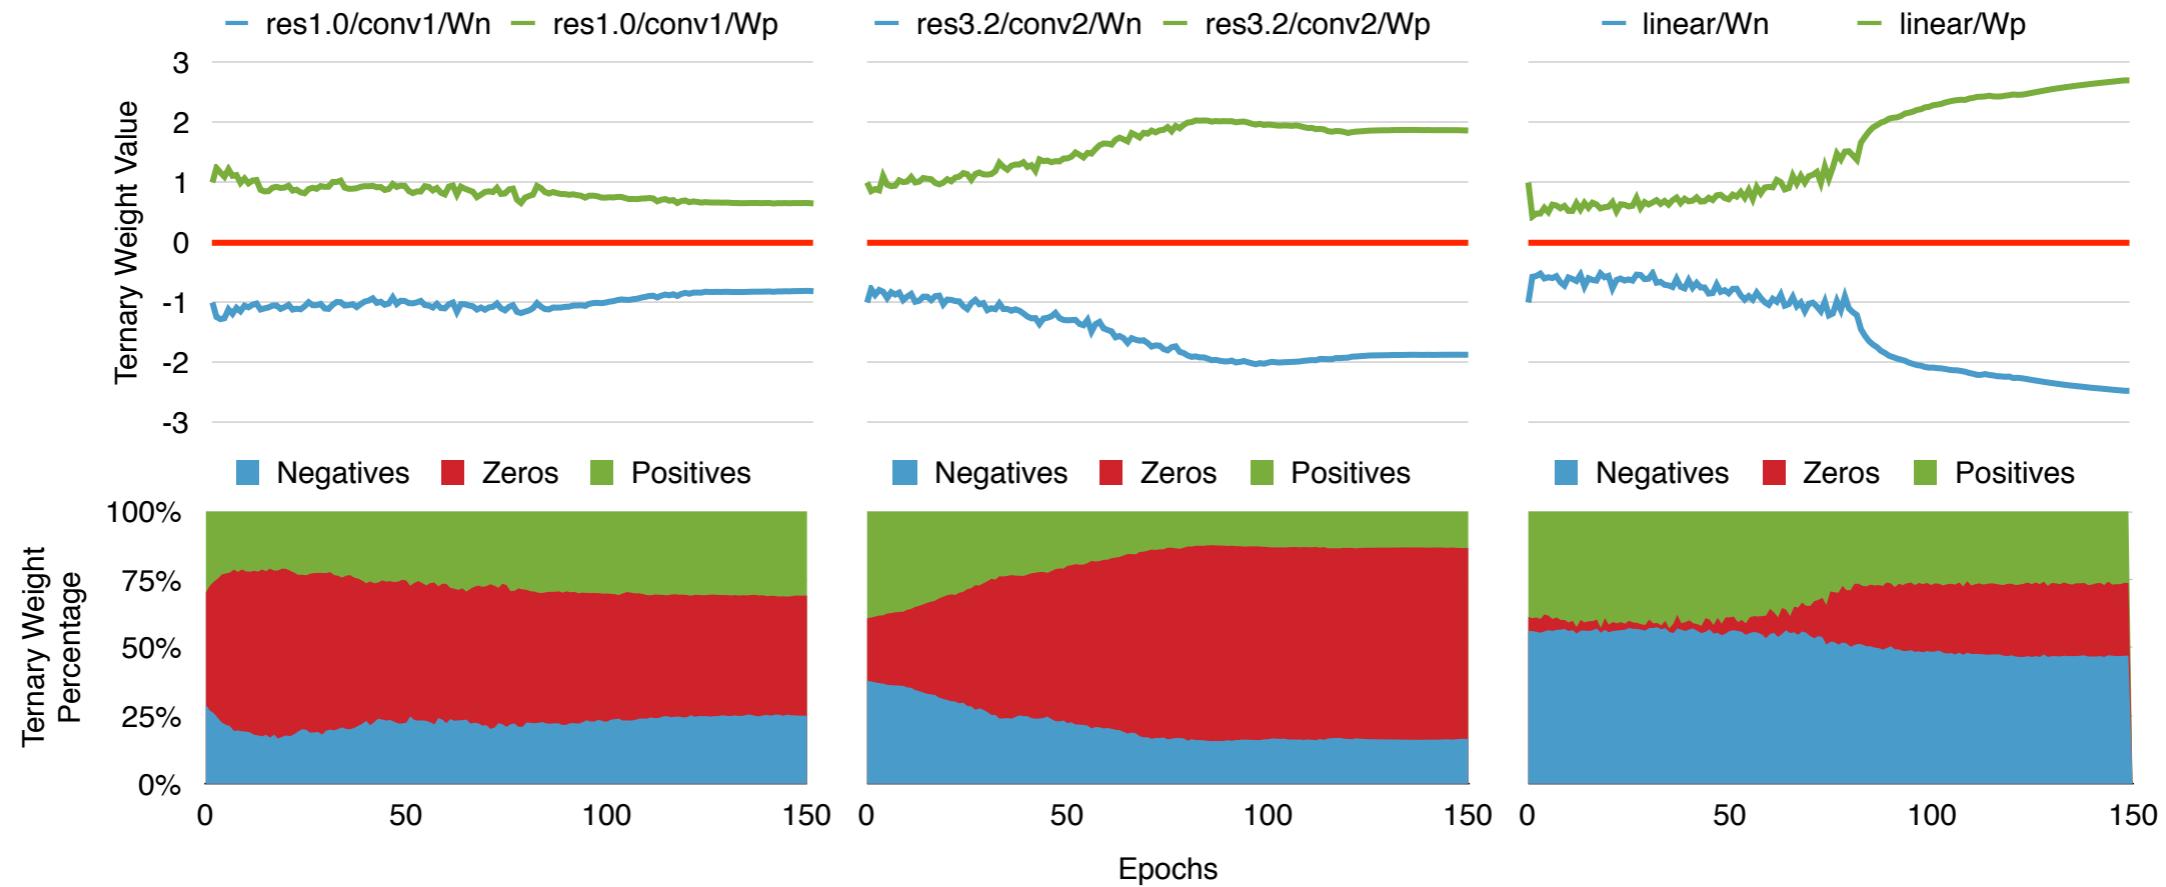
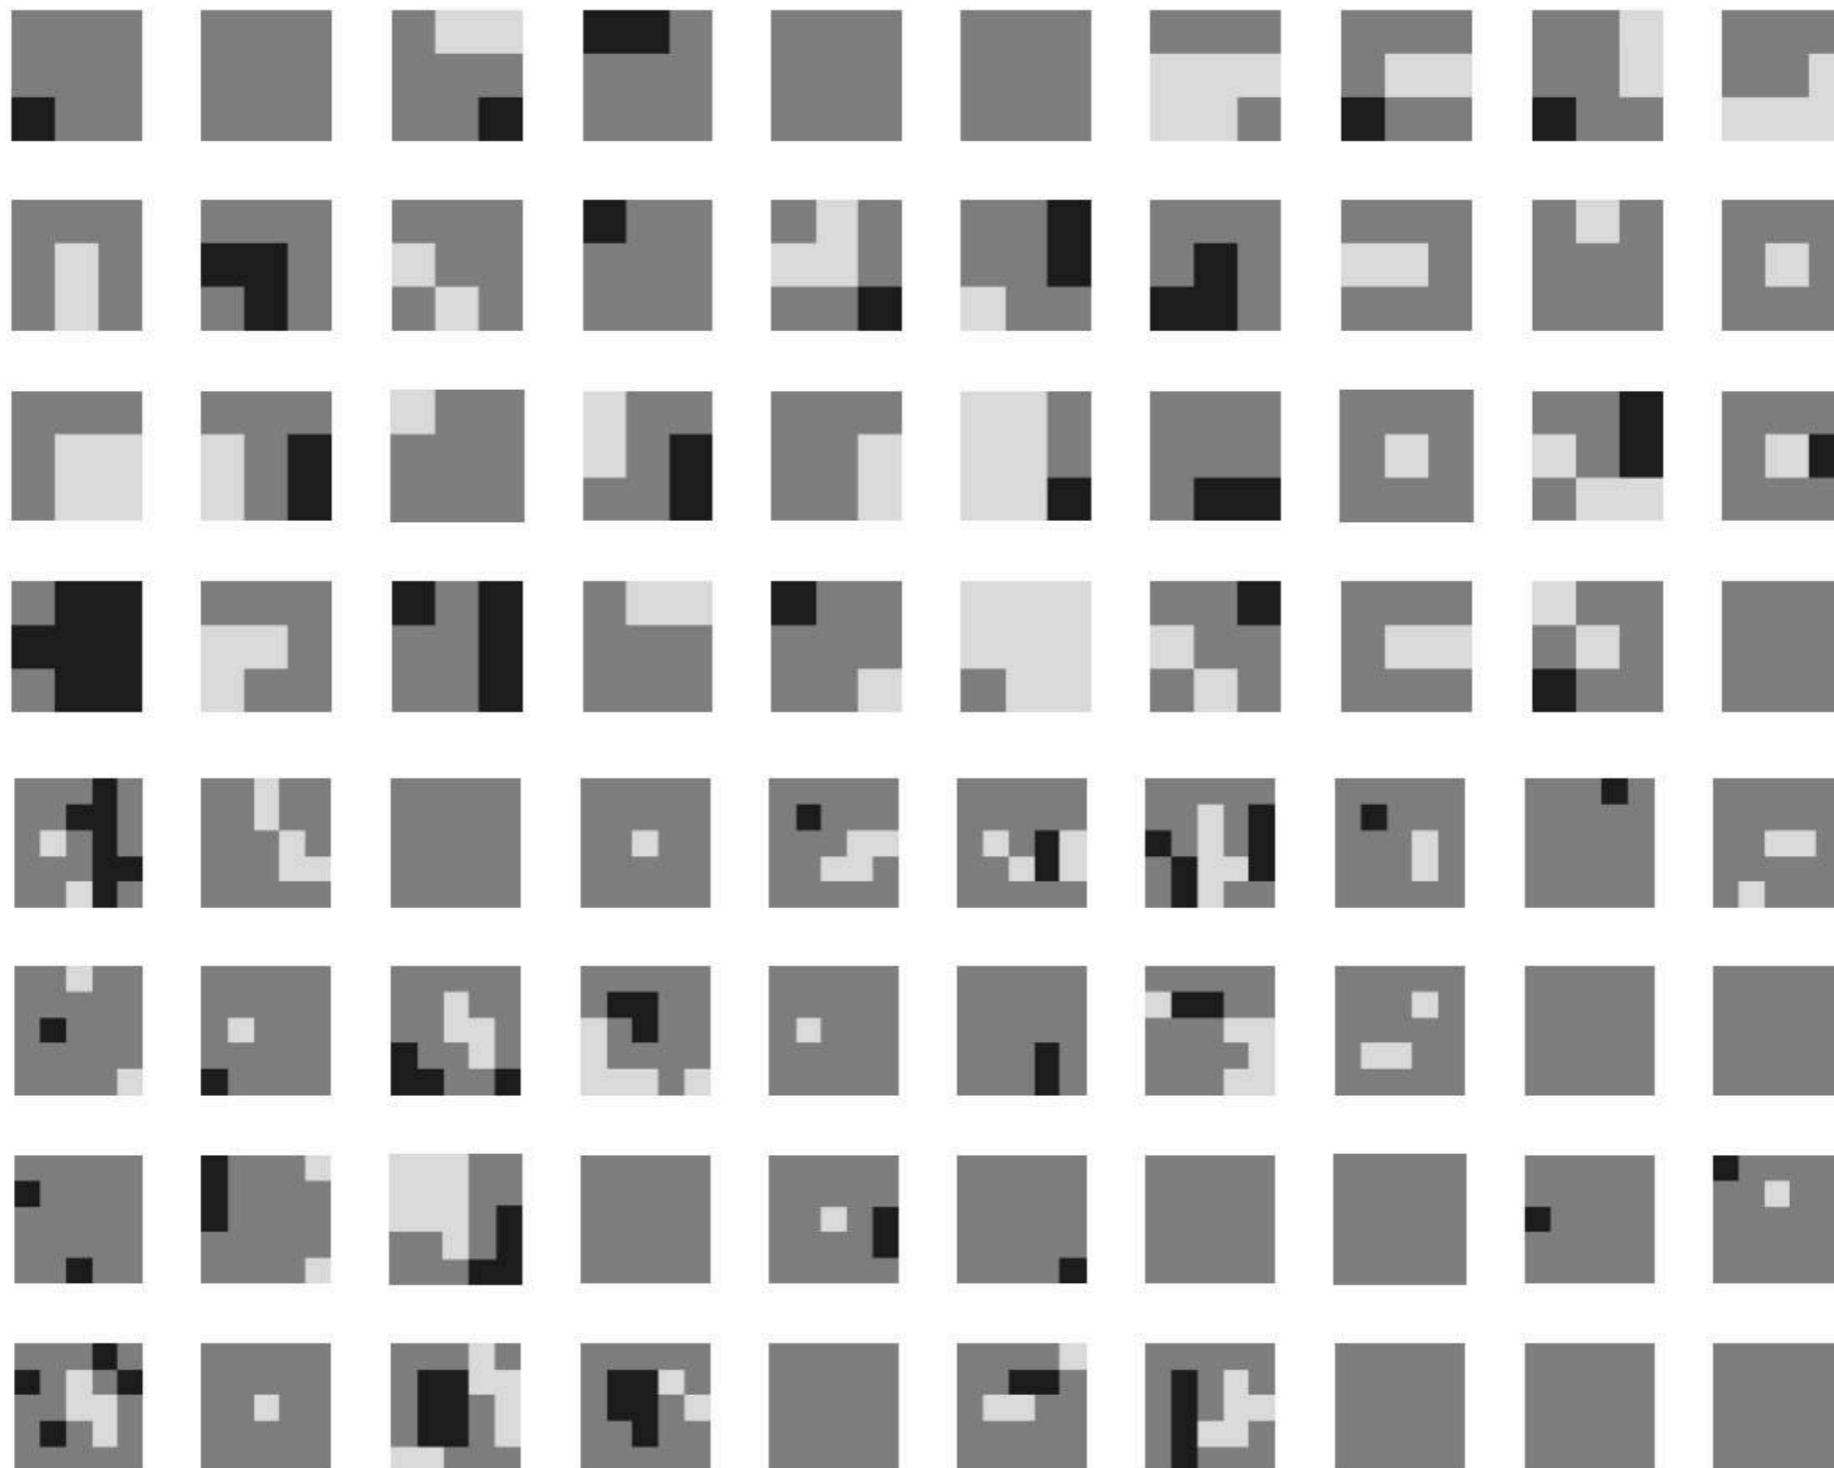


Figure 2: Ternary weights value (above) and distribution (below) with iterations for different layers of ResNet-20 on CIFAR-10.

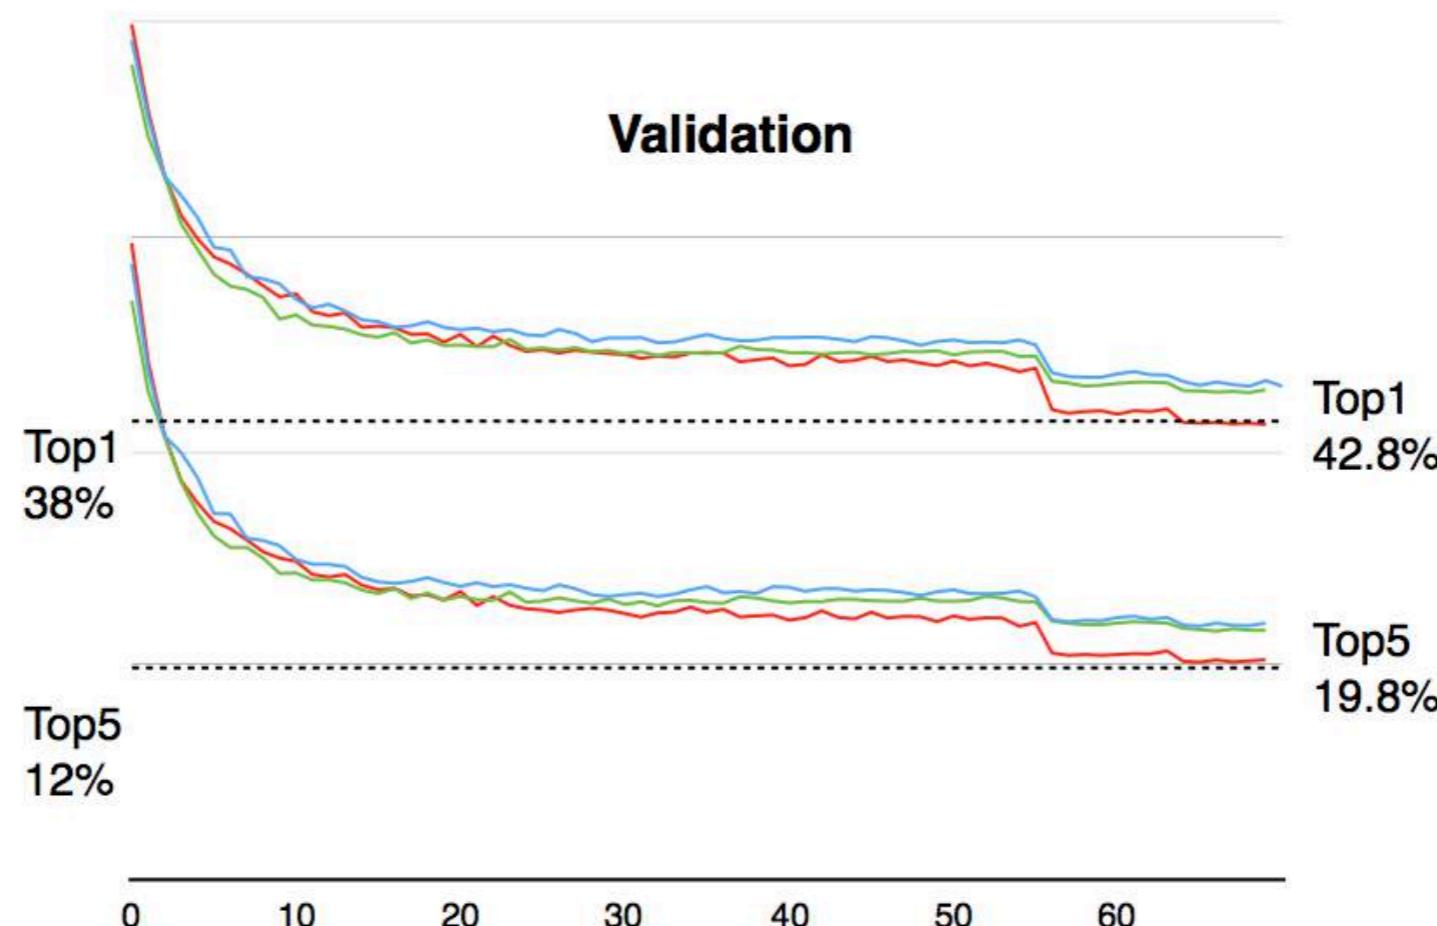
# Visualization of the TTQ Kernels



Zhu, Han, Mao, Dally. Trained Ternary Quantization, ICLR'17

# Error Rate on ImageNet

— DoReFa-Net    — TWN    — Ours    - - - Full precision (with Dropout)



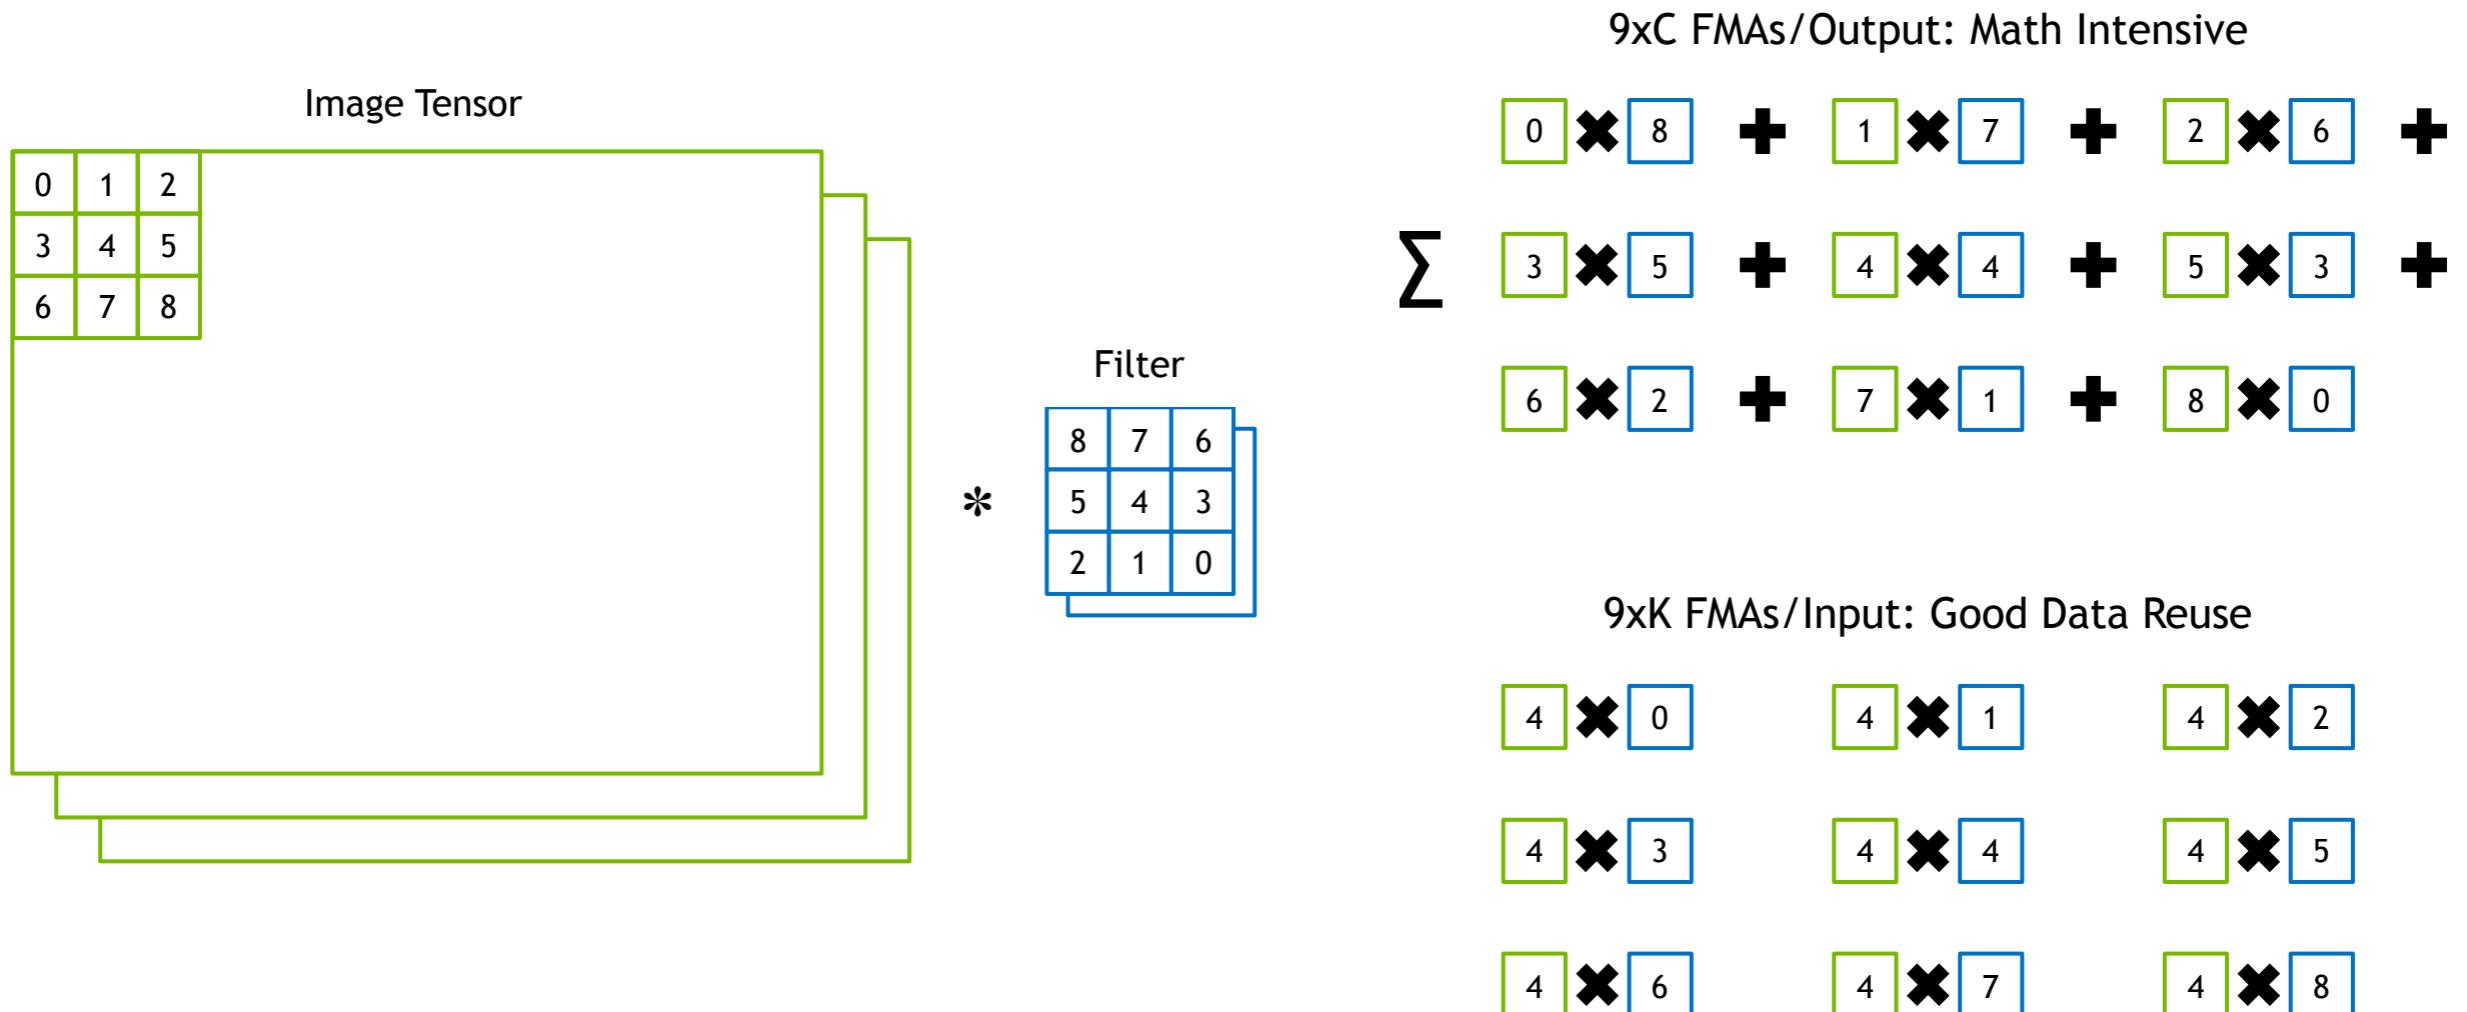
Zhu, Han, Mao, Dally. Trained Ternary Quantization, ICLR'17

# Part 1: Algorithms for Efficient Inference

- 1. Pruning
- 2. Weight Sharing
- 3. Quantization
- 4. Low Rank Approximation
- 5. Binary / Ternary Net
- 6. Winograd Transformation

# 3x3 DIRECT Convolutions

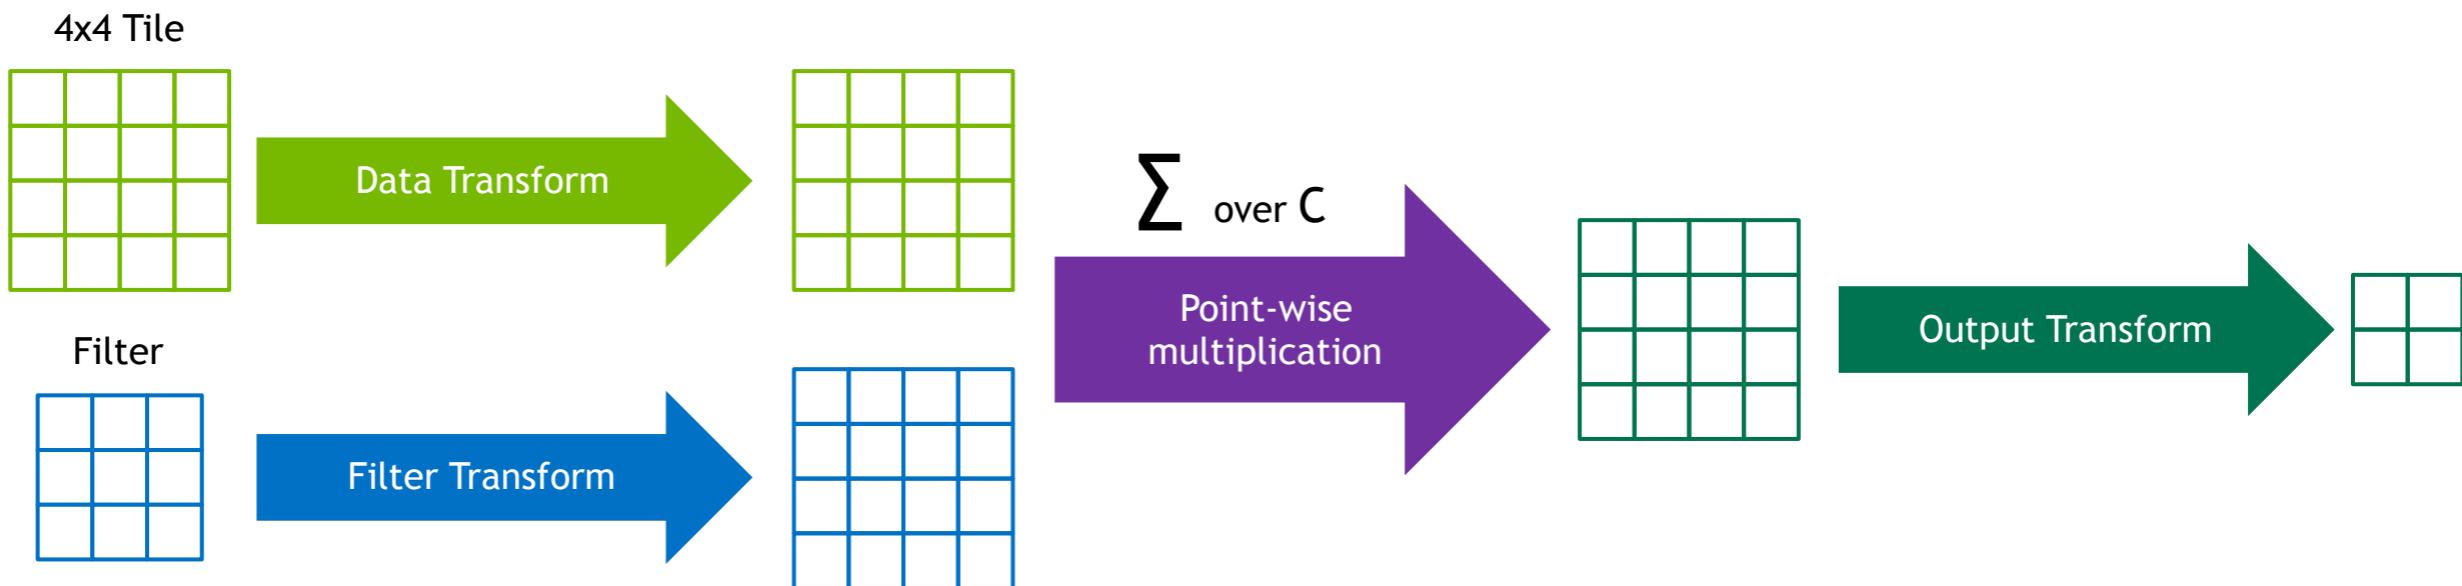
Compute Bound



Direct convolution: we need  $9 \times C \times 4 = 36 \times C$  FMAs for 4 outputs

# 3x3 WINOGRAD Convolutions

Transform Data to Reduce Math Intensity



Direct convolution: we need  $9 \times C \times 4 = 36 \times C$  FMAs for 4 outputs

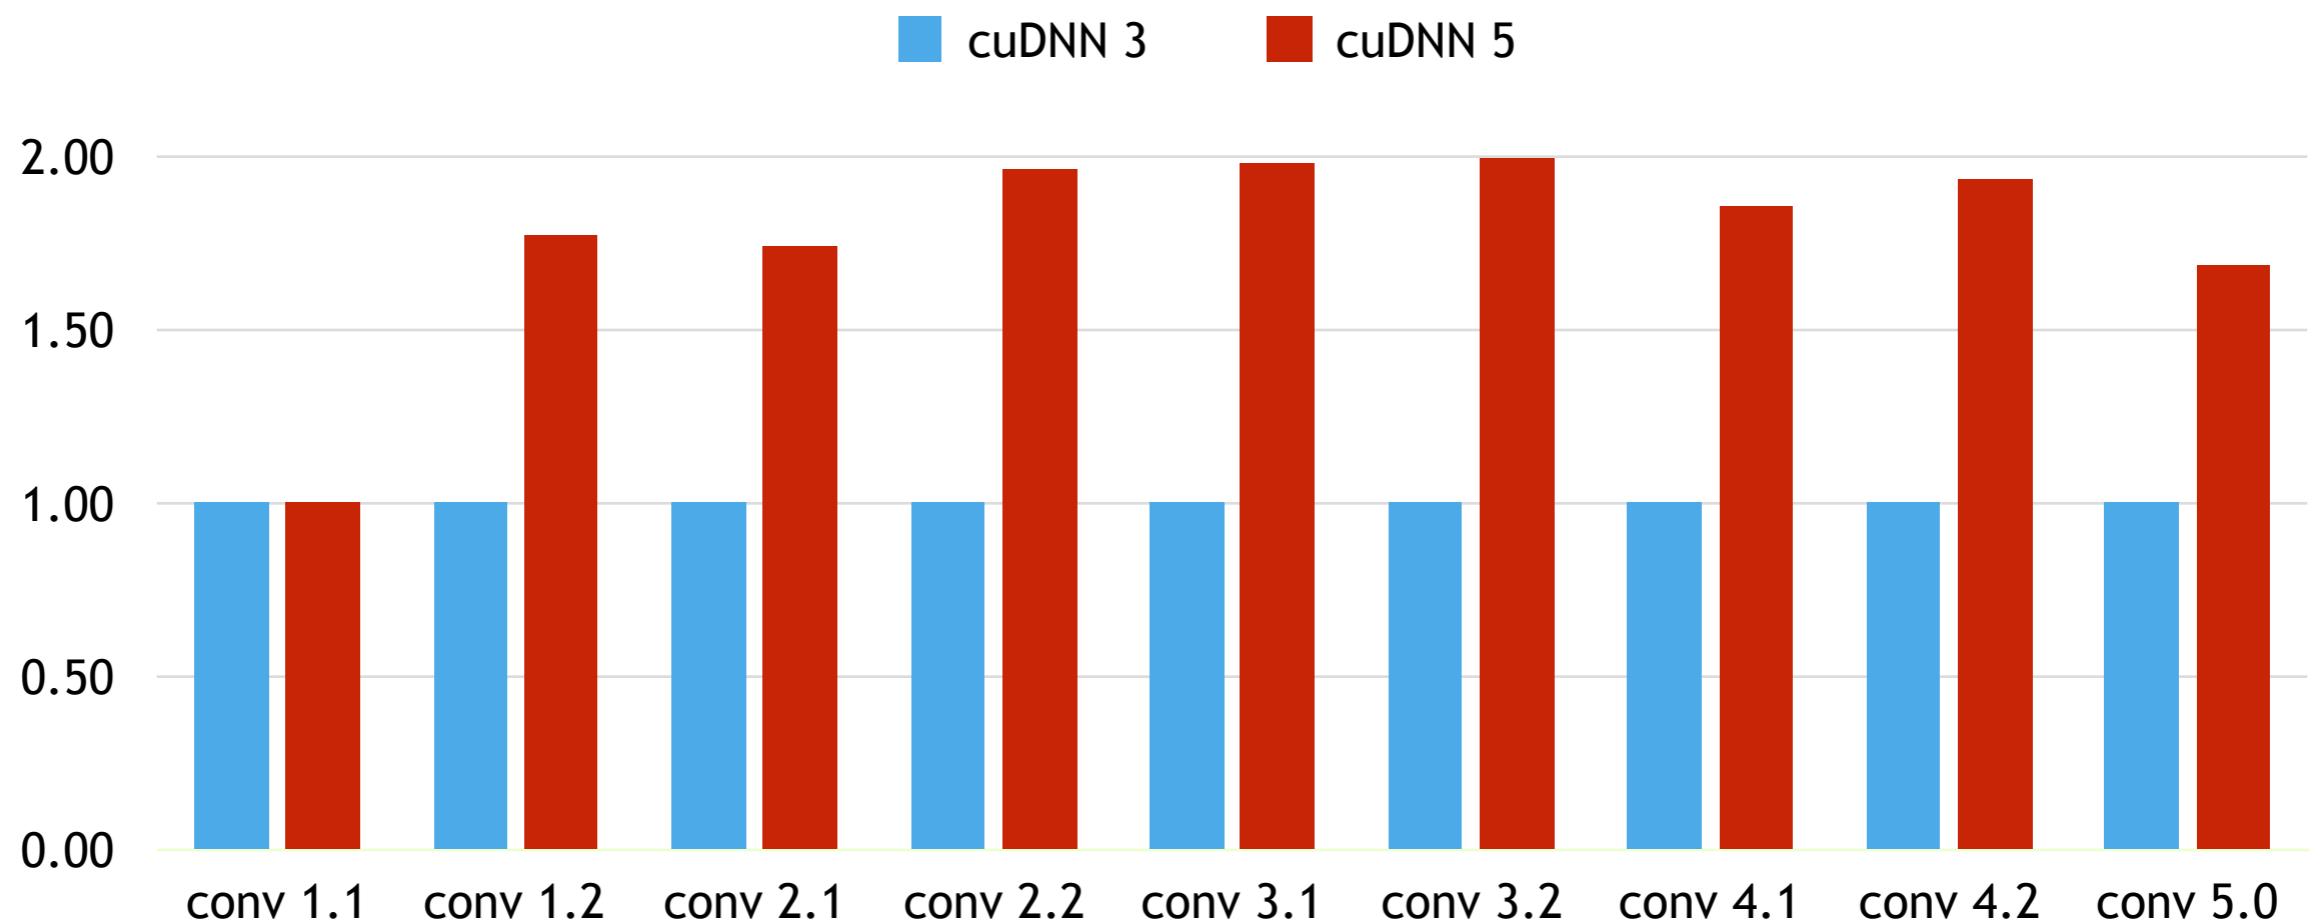
Winograd convolution: we need  $16 \times C$  FMAs for 4 outputs: **2.25x** fewer FMAs

See A. Lavin & S. Gray, "Fast Algorithms for Convolutional Neural Networks

Julien Demouth, Convolution OPTIMIZATION: Winograd, NVIDIA

# Speedup of Winograd Convolution

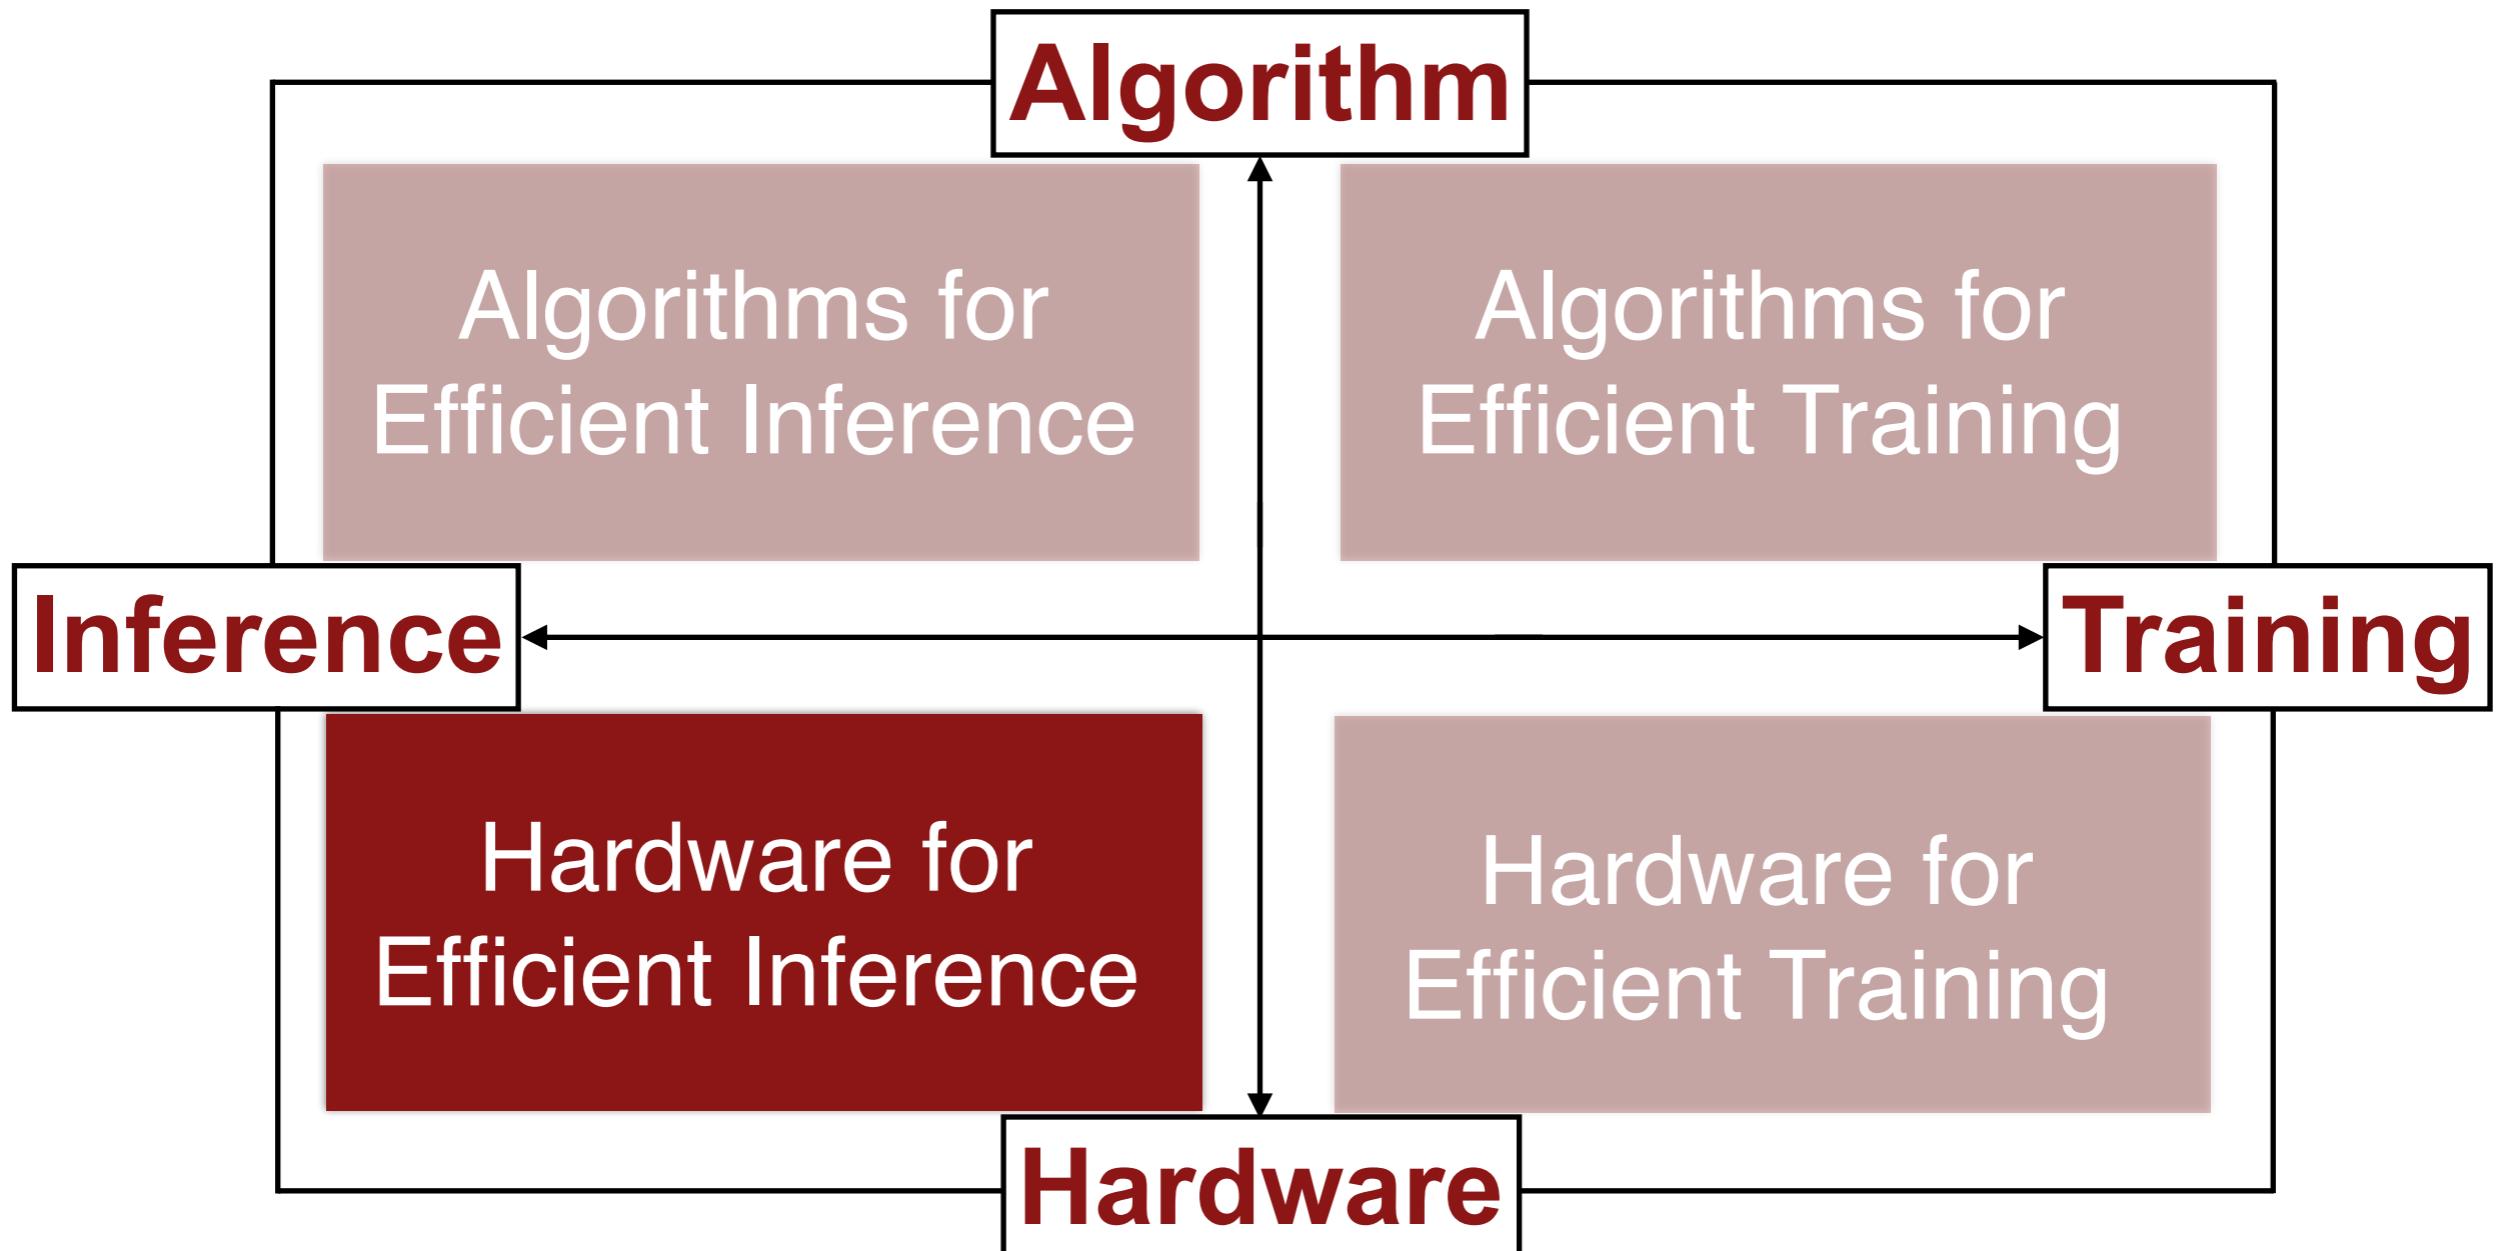
VGG16, Batch Size 1 - Relative Performance



Measured on Maxwell TITAN X

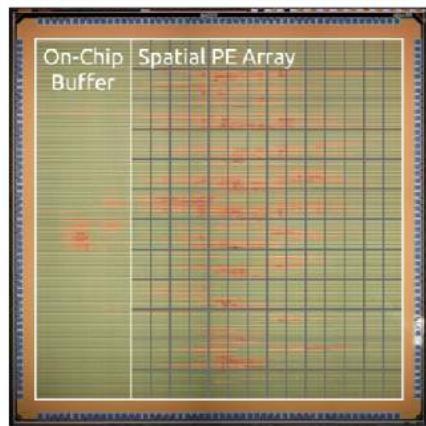
Julien Demouth, Convolution OPTIMIZATION: Winograd, NVIDIA

# Agenda

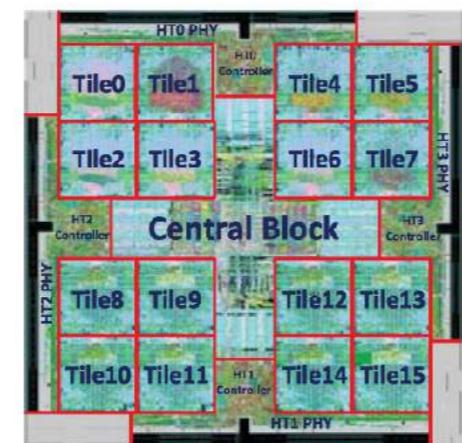


# Hardware for Efficient Inference

a common goal: minimize memory access



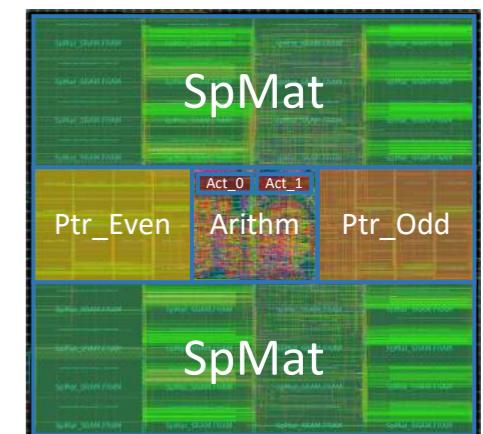
Eyeriss  
MIT  
RS Dataflow



DaDianao  
CAS  
eDRAM



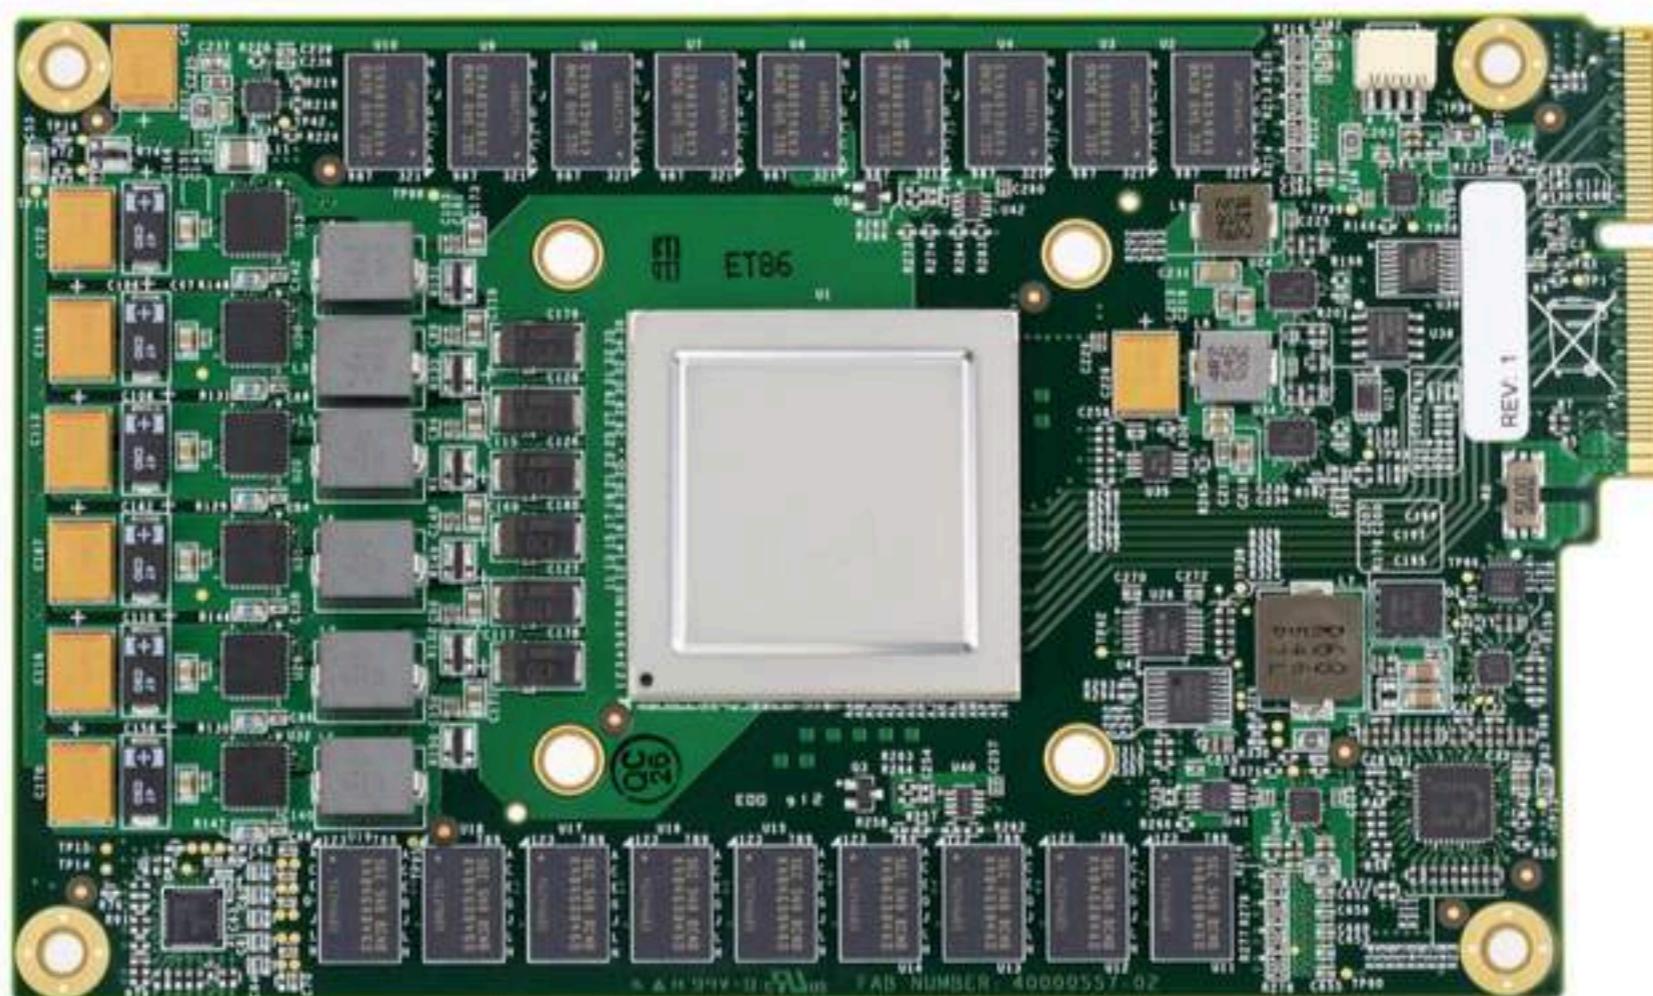
TPU  
Google  
8-bit Integer



EIE  
Stanford  
Compression/  
Sparsity

“This unit is designed for dense matrices. Sparse architectural support was omitted for time-to-deploy reasons. Sparsity will have high priority in future designs”

# Google TPU



TPU Card to replace a disk

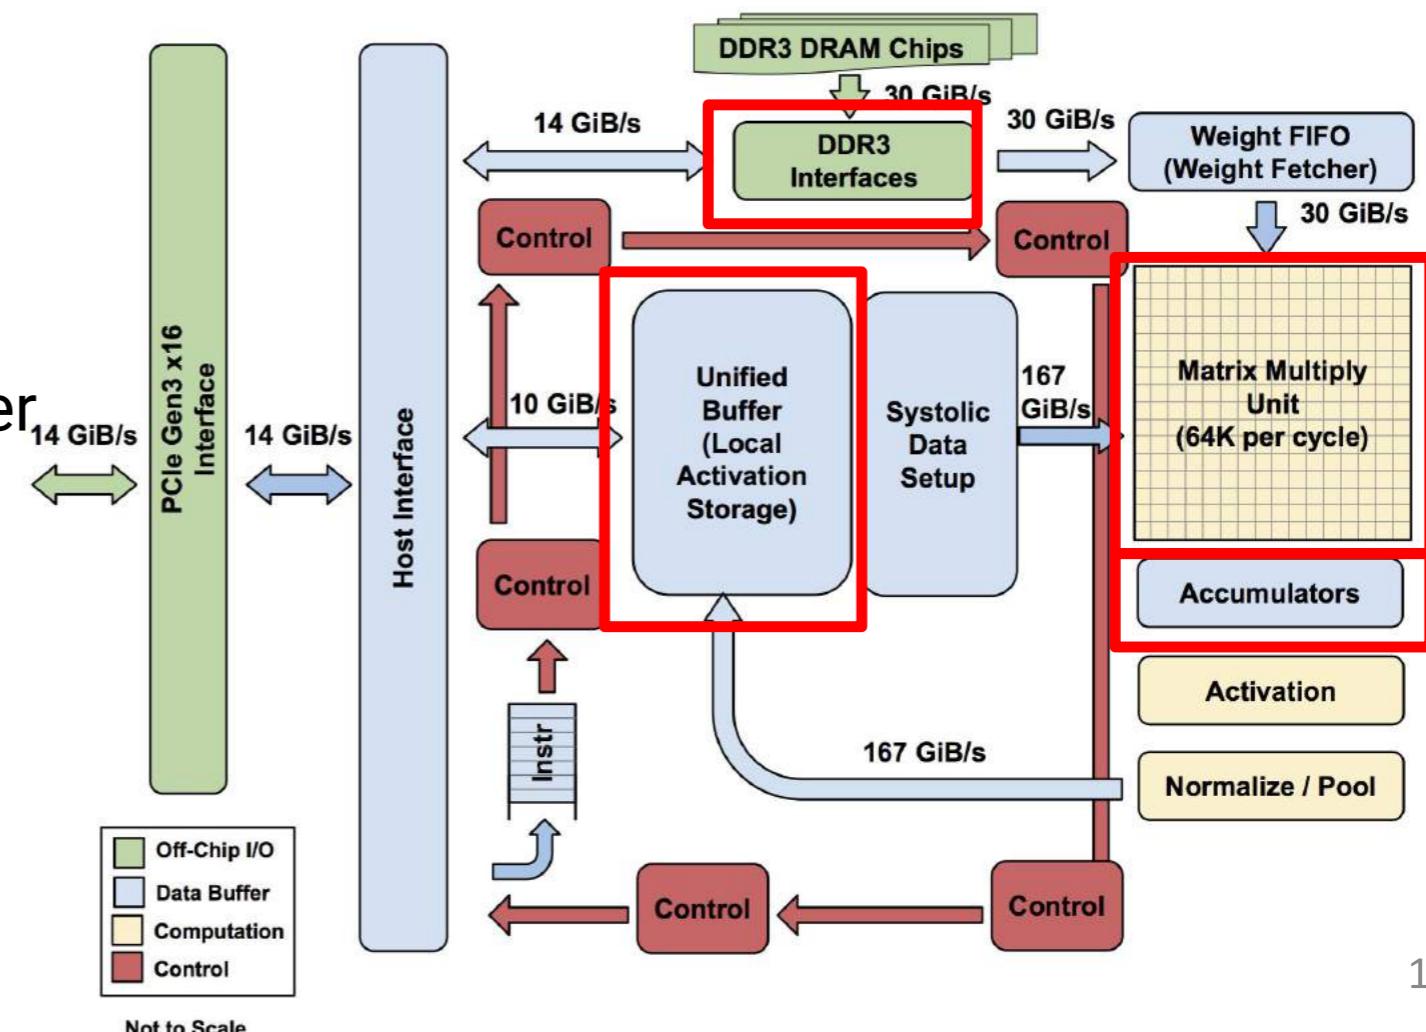
Up to 4 cards / server

David Patterson and the Google TPU Team, In-Data Center Performance Analysis of a Tensor Processing Unit

# Google TPU

- The Matrix Unit: 65,536 (256x256) 8-bit multiply-accumulate units
- 700 MHz clock rate
- Peak: 92T operations/second
  - $65,536 * 2 * 700M$
- >25X as many MACs vs GPU
- >100X as many MACs vs CPU
- 4 MiB of on-chip Accumulator memory
- 24 MiB of on-chip Unified Buffer (activation memory)
- 3.5X as much on-chip memory vs GPU
- Two 2133MHz DDR3 DRAM channels
- 8 GiB of off-chip weight DRAM memory

## TPU: High-level Chip Architecture



15

David Patterson and the Google TPU Team, In-Data Center Performance Analysis of a Tensor Processing Unit

# Google TPU

Processor	<i>mm<sup>2</sup></i>	Clock MHz	TDP Watts	Idle Watts	Memory GB/sec	Peak TOPS/chip	
						8b int.	32b FP
CPU: Haswell (18 core)	662	2300	145	41	51	2.6	1.3
GPU: Nvidia K80 (2 / card)	561	560	150	25	160	--	2.8
TPU	<331*	700	75	28	34	91.8	--

\*TPU is less than half die size of the Intel Haswell processor

K80 and TPU in 28 nm process; Haswell fabbed in Intel 22 nm process

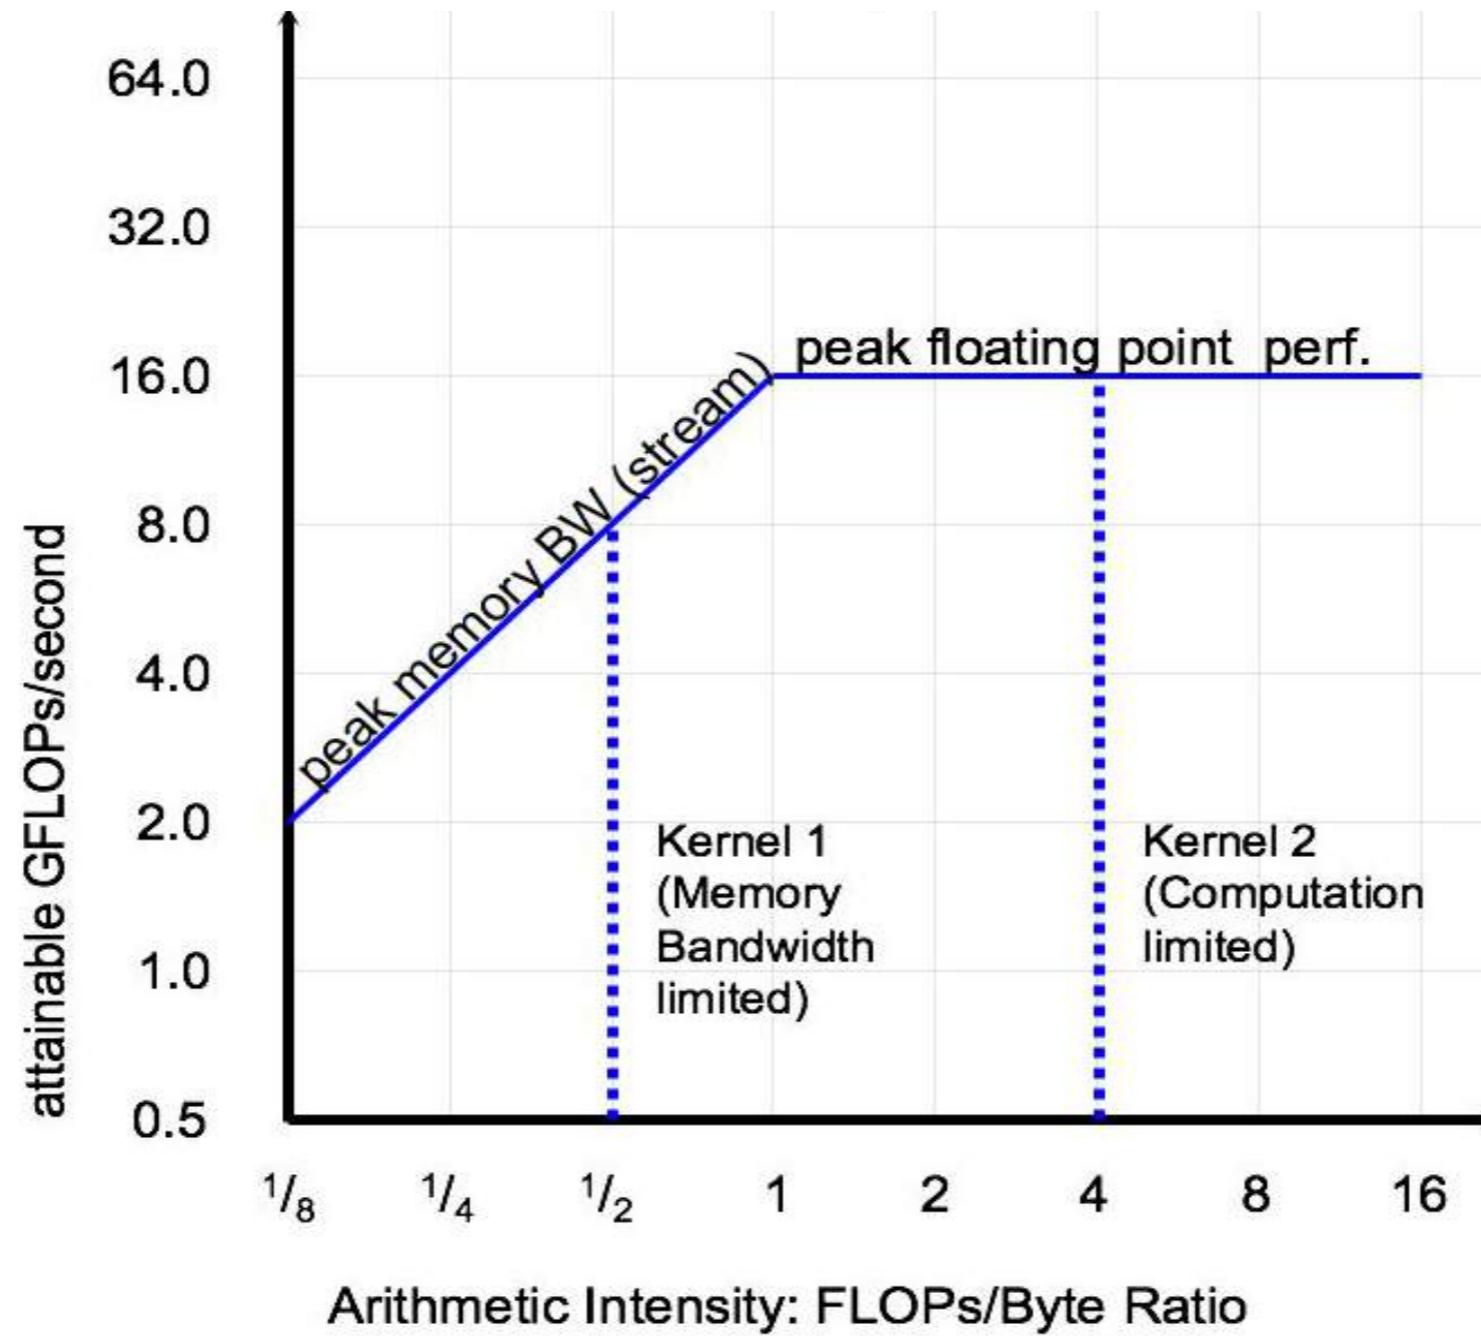
These chips and platforms chosen for comparison because widely deployed in Google data centers

# Inference Datacenter Workload

Name	LOC	Layers					Nonlinear function	Weights	TPU Ops / Weight Byte	TPU Batch Size	% Deployed
		FC	Conv	Vector	Pool	Total					
MLP0	0.1k	5				5	ReLU	20M	200	200	61%
MLP1	1k	4				4	ReLU	5M	168	168	
LSTM0	1k	24		34		58	sigmoid, tanh	52M	64	64	29%
LSTM1	1.5k	37		19		56	sigmoid, tanh	34M	96	96	
CNN0	1k		16			16	ReLU	8M	2888	8	5%
CNN1	1k	4	72		13	89	ReLU	100M	1750	32	

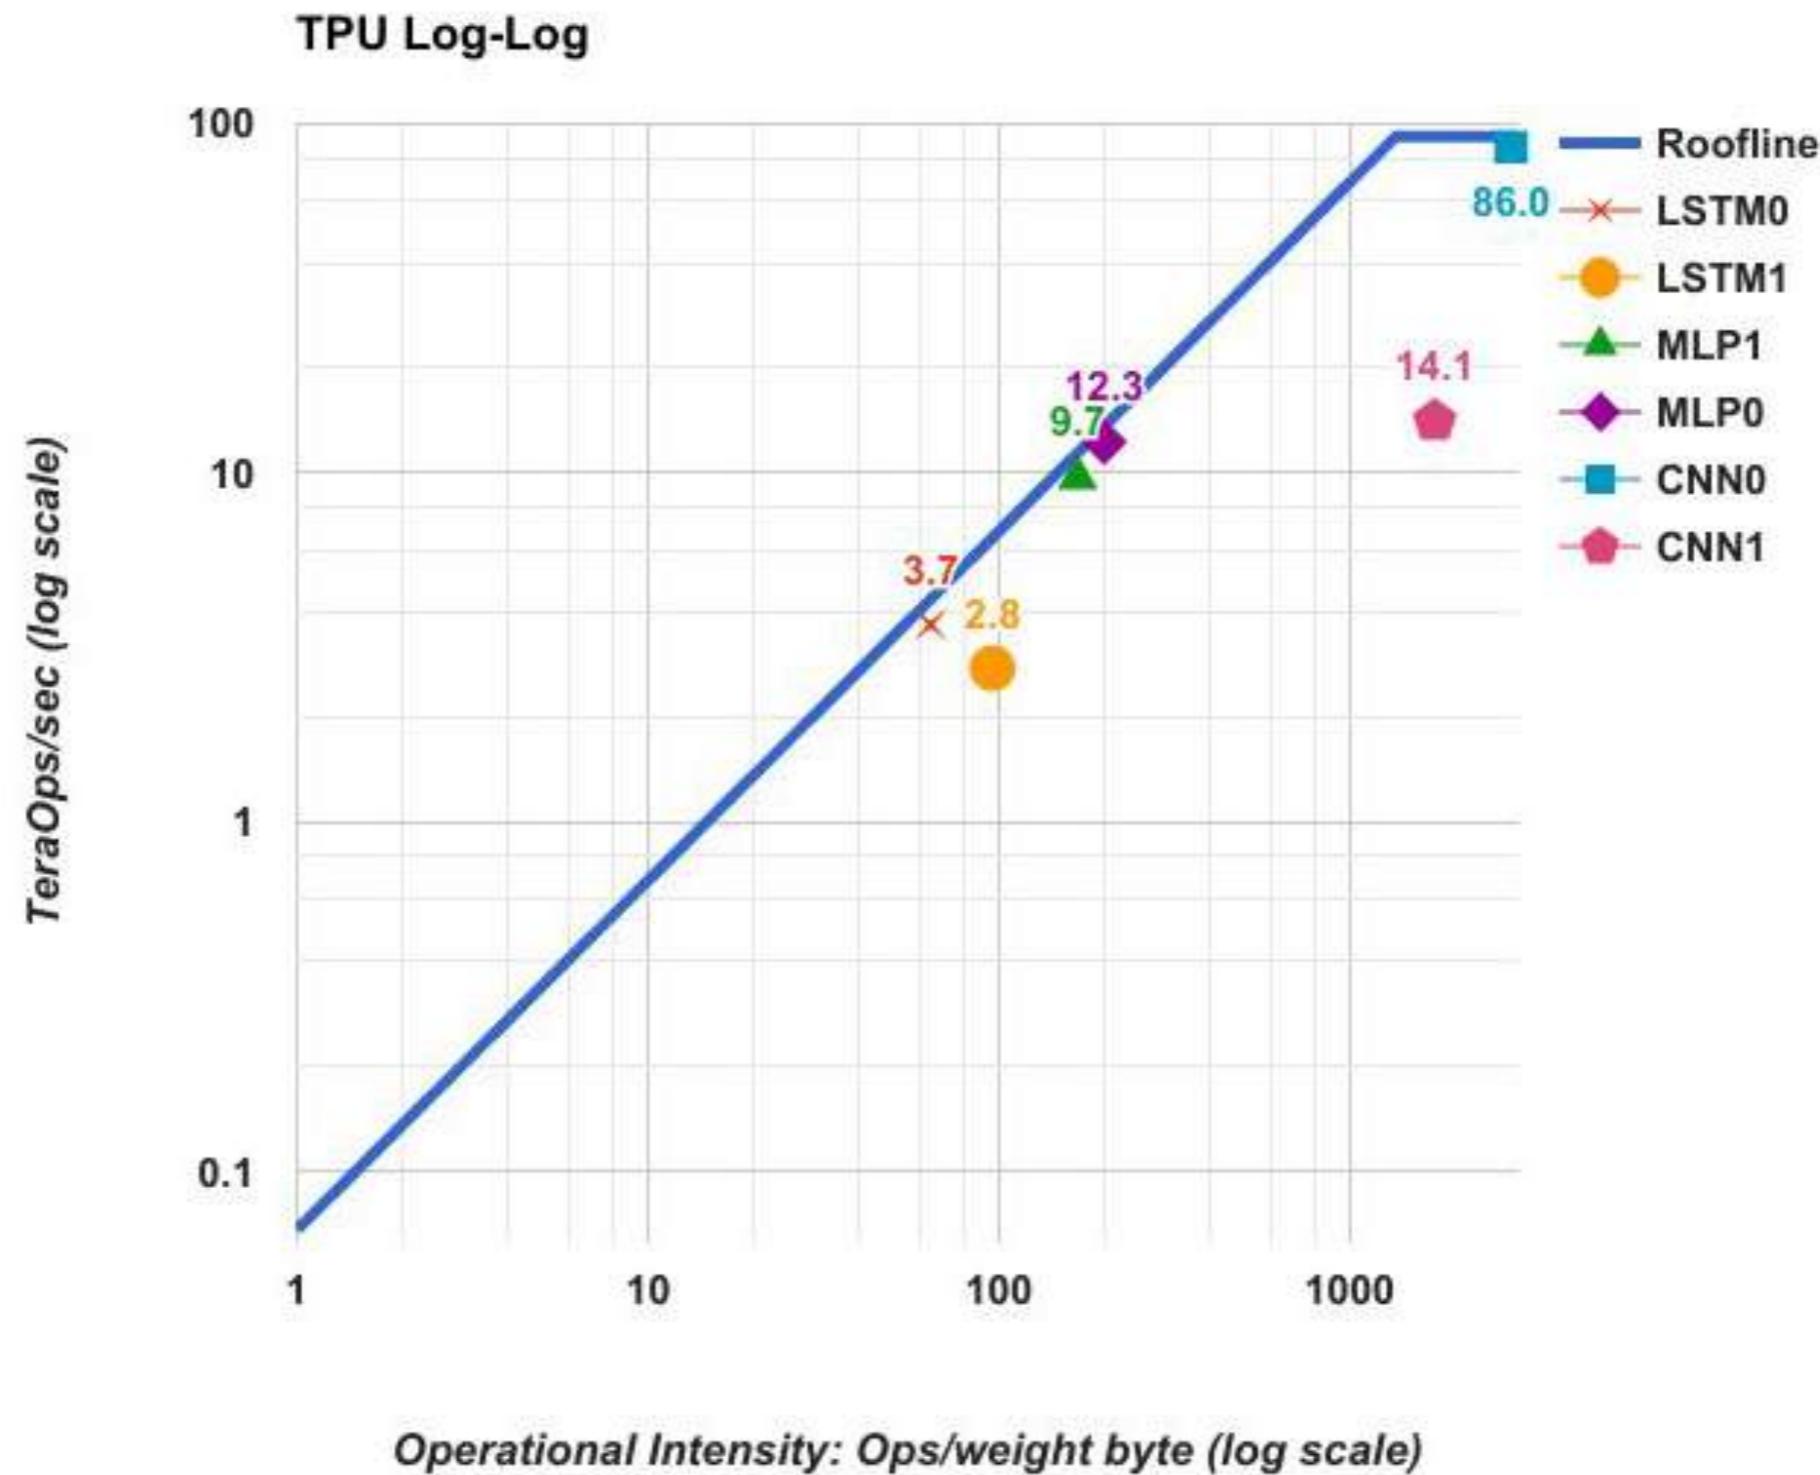
# Roofline Model: Identify Performance Bottleneck

$$\text{GFLOP/s} = \text{Min}(\text{Peak GFLOP/s}, \text{Peak GB/s} \times \text{AI})$$



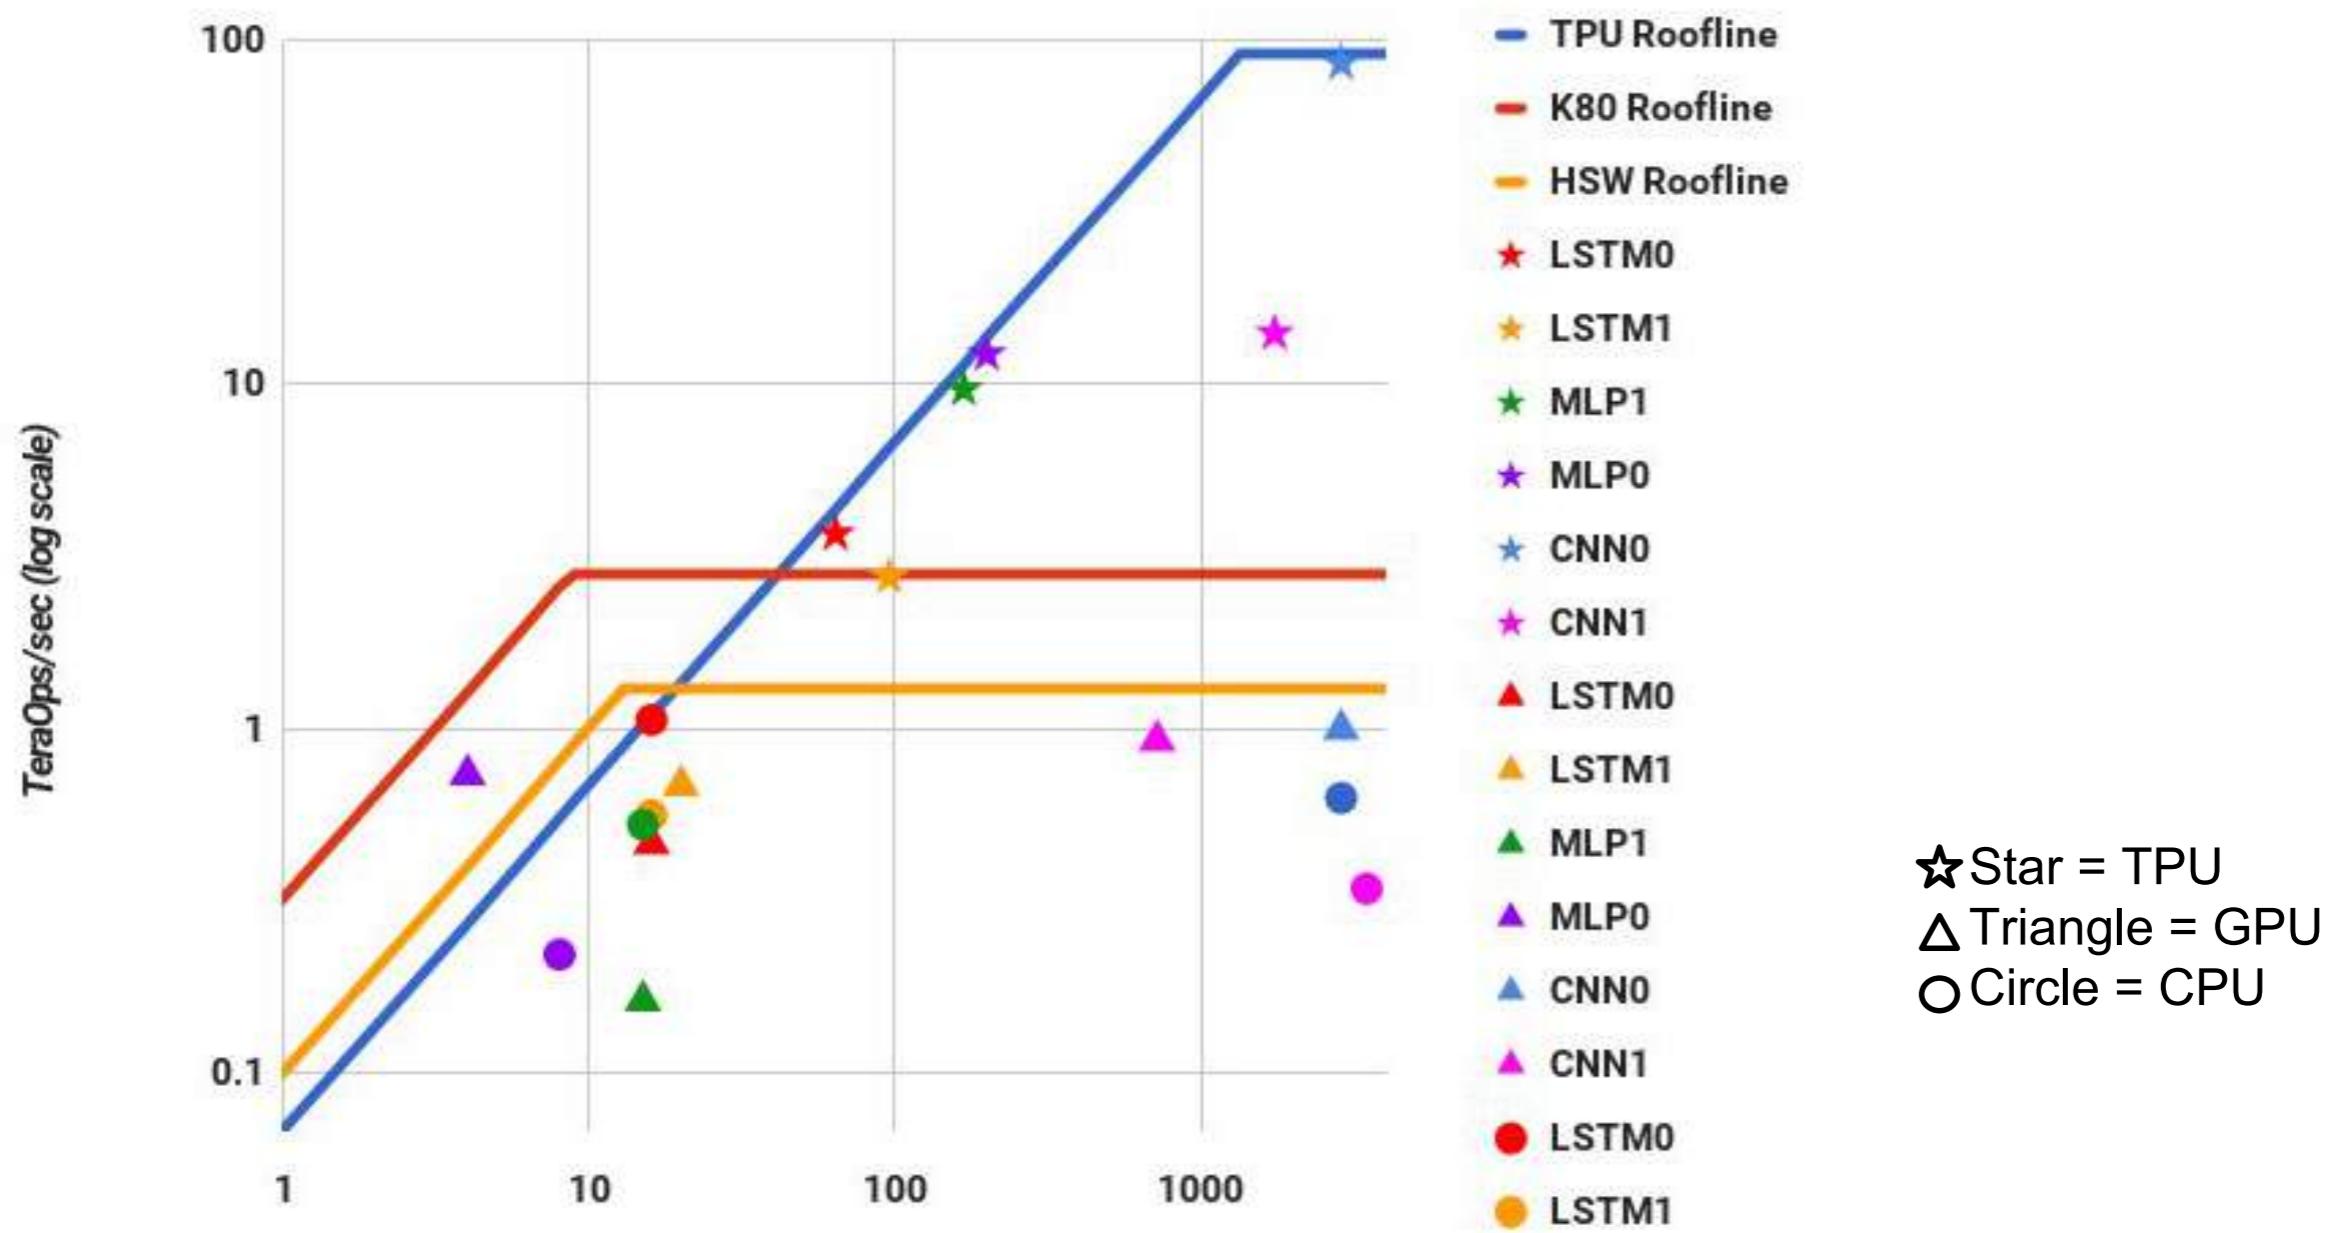
David Patterson and the Google TPU Team, In-Data Center Performance Analysis of a Tensor Processing Unit

# TPU Roofline



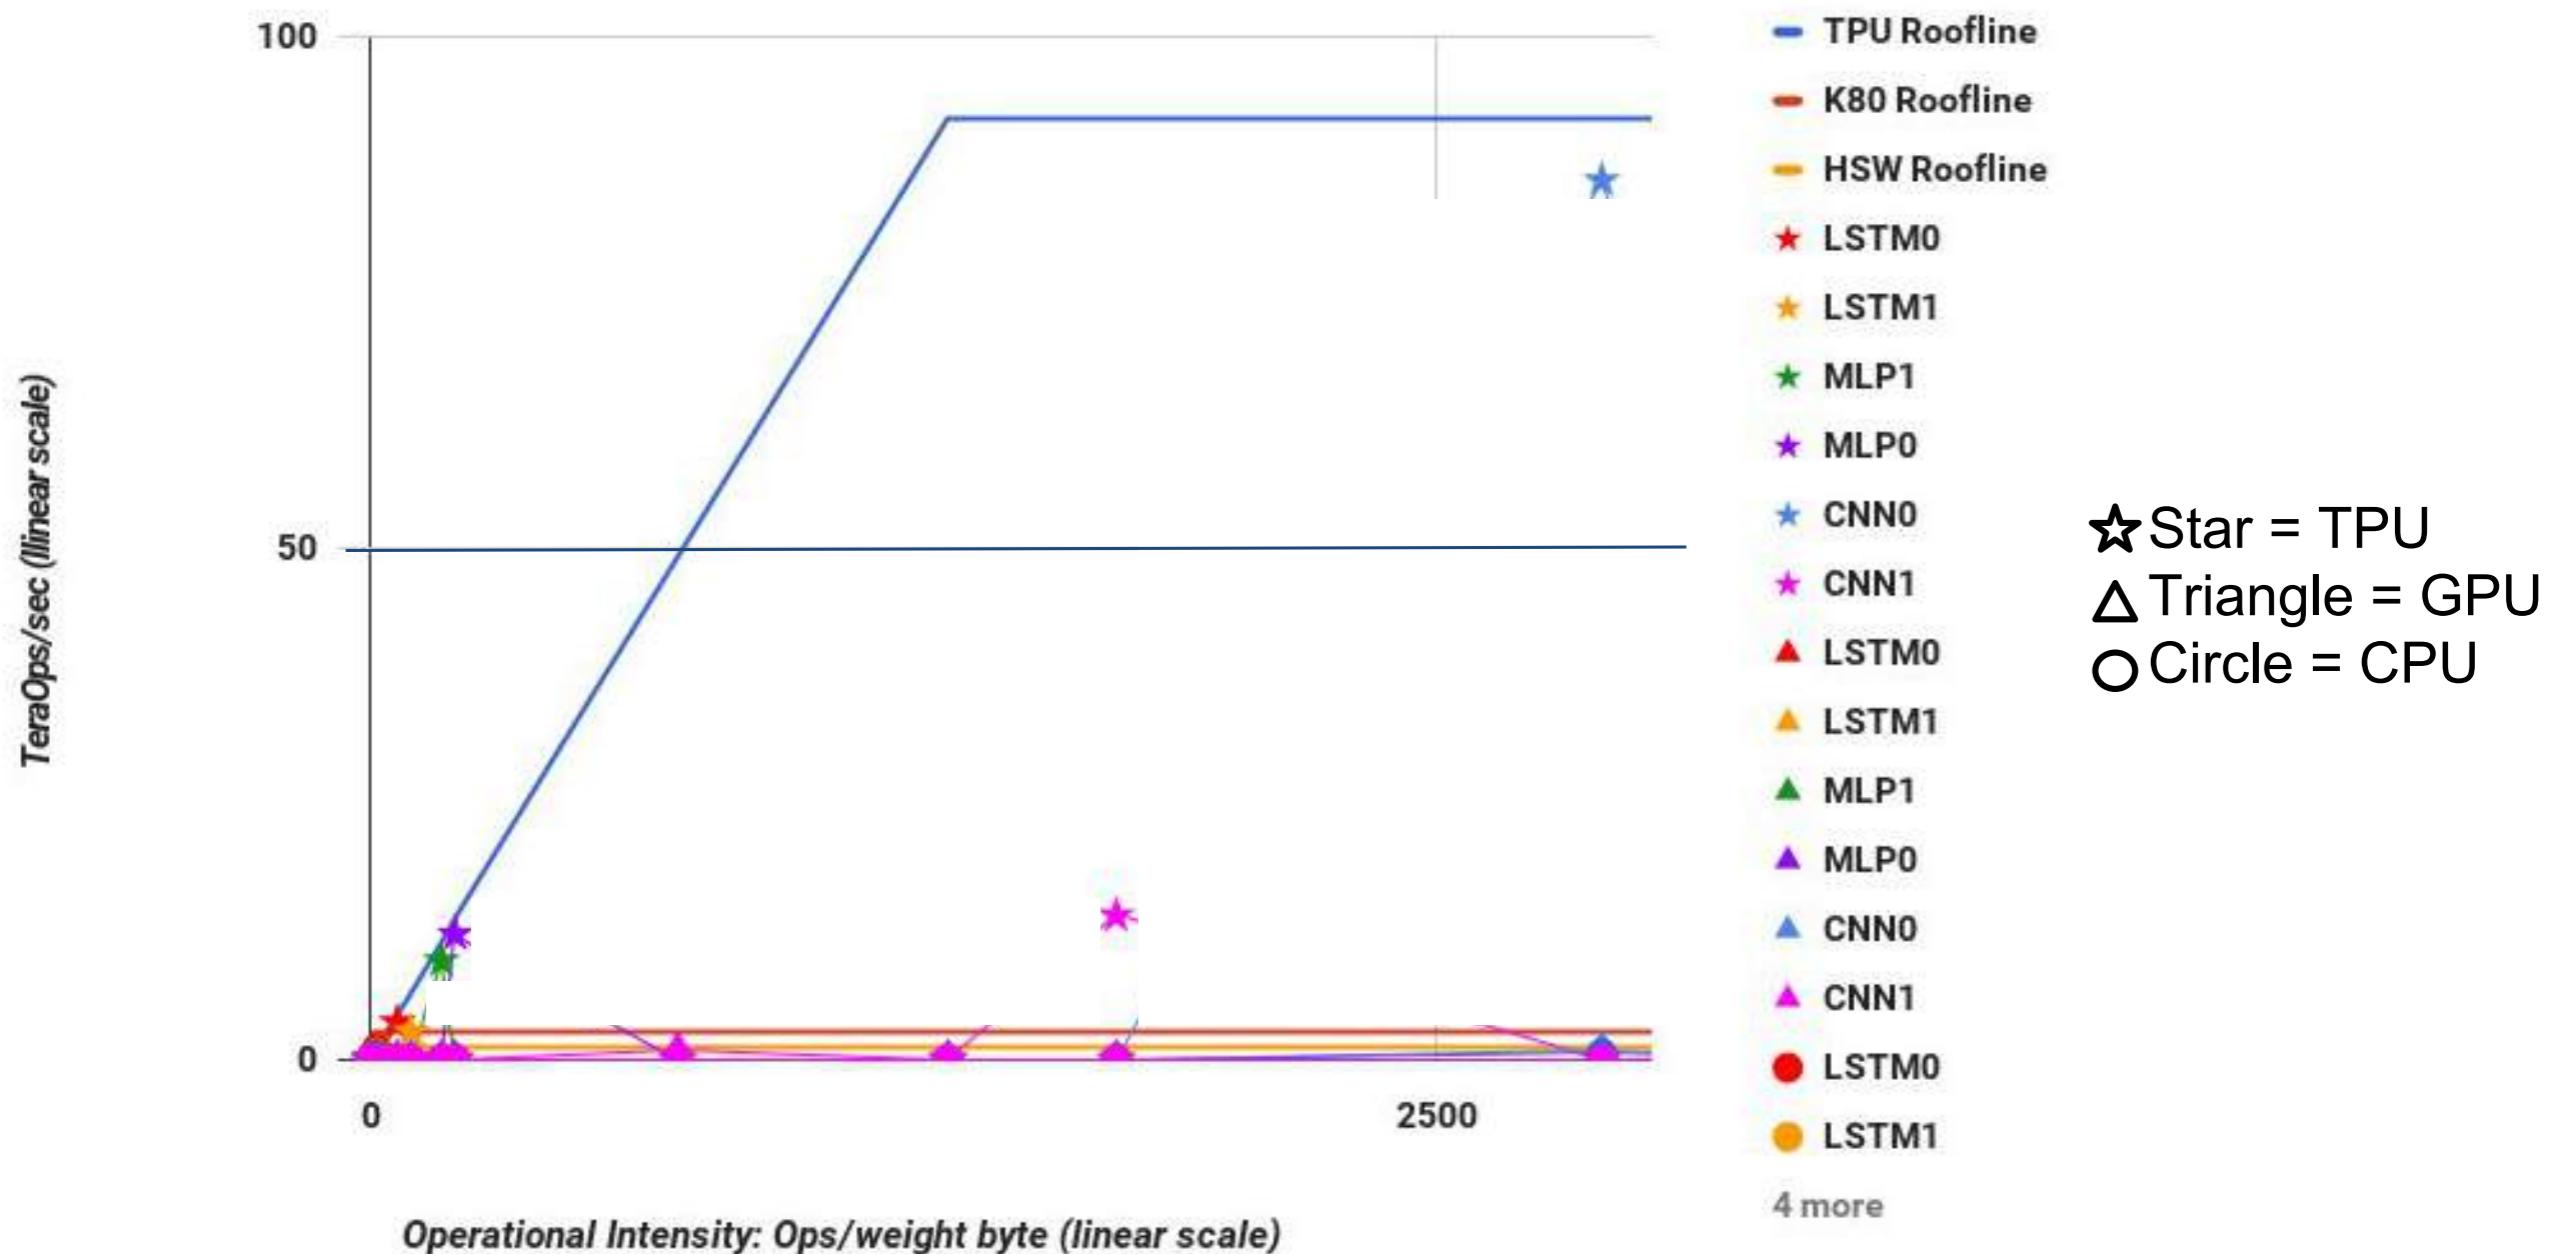
David Patterson and the Google TPU Team, In-Data Center Performance Analysis of a Tensor Processing Unit

# Log Rooflines for CPU, GPU, TPU



David Patterson and the Google TPU Team, In-Data Center Performance Analysis of a Tensor Processing Unit

# Linear Rooflines for CPU, GPU, TPU



David Patterson and the Google TPU Team, In-Data Center Performance Analysis of a Tensor Processing Unit

## Why so far below Rooflines?

Low latency requirement => Can't batch more => low ops/byte

## How to Solve this?

less memory footprint => need compress the model

## Challenge:

Hardware that can infer on compressed model

# EIE: the First DNN Accelerator for Sparse, Compressed Model

# EIE: the First DNN Accelerator for Sparse, Compressed Model

$$0 * A = 0$$

**Sparse Weight**

90% *static* sparsity

$$W * 0 = 0$$

**Sparse Activation**

70% *dynamic* sparsity

~~$$2.09, 1.92 \Rightarrow 2$$~~

**Weight Sharing**

4-bit weights



10x less computation



3x less computation



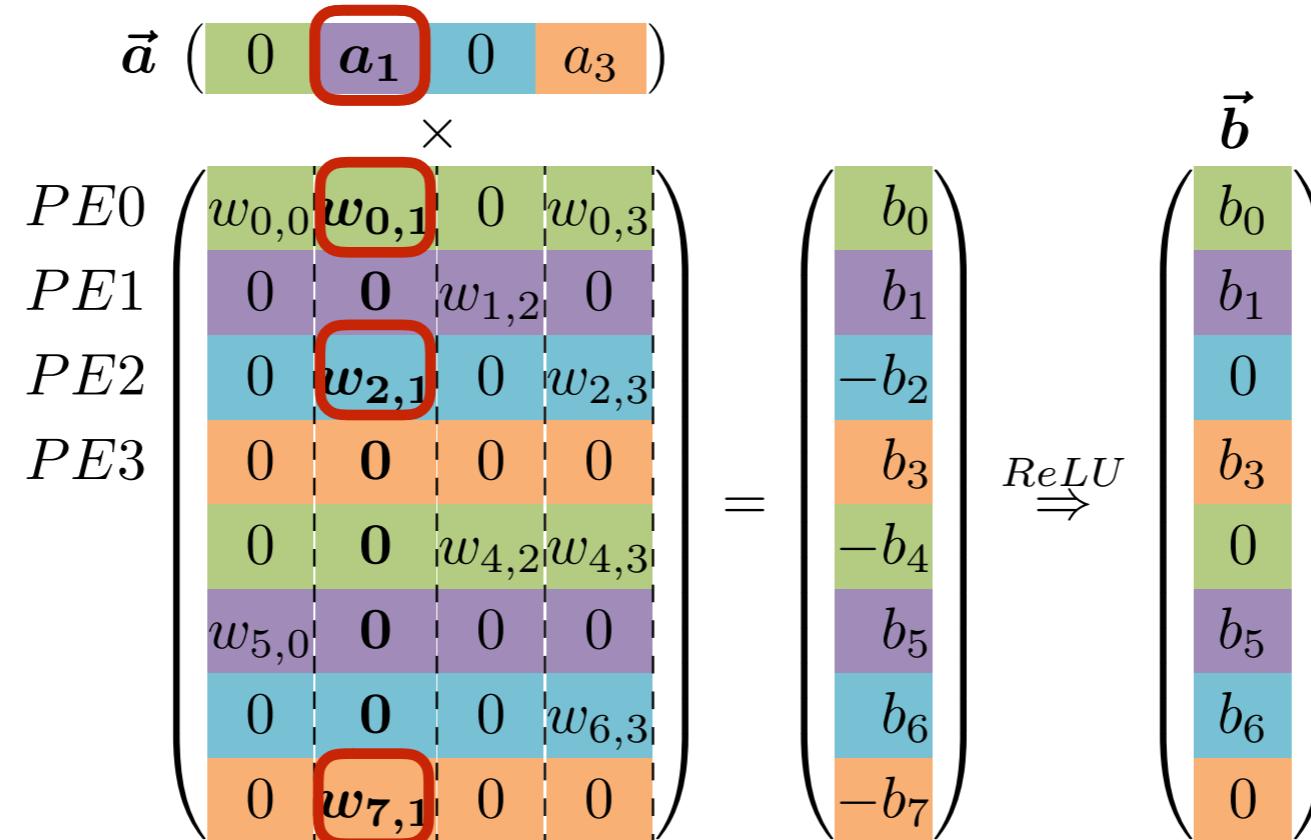
5x less memory footprint



8x less memory footprint

# EIE: Reduce Memory Access by Compression

logically



physically

Virtual Weight	$W_{0,0}$	$W_{0,1}$	$W_{4,2}$	$W_{0,3}$	$W_{4,3}$
Relative Index	0	1	2	0	0
Column Pointer	0	1	2	3	

# Dataflow

$$\begin{array}{c}
 \vec{a} \quad (0 \quad a_1 \quad 0 \quad a_3) \\
 \times \\
 PE0 \left( \begin{array}{|c|c|c|c|} \hline w_{0,0} & w_{0,1} & 0 & w_{0,3} \\ \hline 0 & 0 & w_{1,2} & 0 \\ \hline 0 & w_{2,1} & 0 & w_{2,3} \\ \hline 0 & 0 & 0 & 0 \\ \hline 0 & 0 & w_{4,2} & w_{4,3} \\ \hline w_{5,0} & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & w_{6,3} \\ \hline 0 & w_{7,1} & 0 & 0 \\ \hline \end{array} \right) = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \xrightarrow{\text{ReLU}} \vec{b} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

$$\begin{array}{c}
 \vec{a} \left( \begin{array}{cccc} 0 & \boxed{a_1} & 0 & a_3 \end{array} \right) \\
 \times \\
 PE0 \left( \begin{array}{cc|cc} w_{0,0} & w_{0,1} & 0 & w_{0,3} \\ \hline 0 & 0 & w_{1,2} & 0 \end{array} \right) \\
 PE1 \left( \begin{array}{cc|cc} 0 & 0 & w_{1,2} & 0 \\ \hline 0 & w_{2,1} & 0 & w_{2,3} \end{array} \right) \\
 PE2 \left( \begin{array}{cc|cc} 0 & 0 & 0 & 0 \\ \hline 0 & 0 & w_{4,2} & w_{4,3} \end{array} \right) \\
 PE3 \left( \begin{array}{cc|cc} w_{5,0} & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & w_{6,3} \\ 0 & w_{7,1} & 0 & 0 \end{array} \right) \\
 = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \xrightarrow{\text{ReLU}} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

$$\begin{array}{c}
 \vec{a} \left( \begin{array}{cccc} 0 & \boxed{a_1} & 0 & a_3 \end{array} \right) \\
 \times \\
 PE0 \left( \begin{array}{cccc} w_{0,0} & \boxed{w_{0,1}} & 0 & w_{0,3} \\ 0 & 0 & w_{1,2} & 0 \\ 0 & \boxed{w_{2,1}} & 0 & w_{2,3} \\ 0 & 0 & 0 & 0 \\ 0 & 0 & w_{4,2} & w_{4,3} \\ w_{5,0} & 0 & 0 & 0 \\ 0 & 0 & 0 & w_{6,3} \\ 0 & \boxed{w_{7,1}} & 0 & 0 \end{array} \right) = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \xrightarrow{\text{ReLU}} \vec{b} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

$$\begin{array}{c}
 \vec{a} \left( \begin{array}{cccc} 0 & a_1 & 0 & a_3 \end{array} \right) \\
 \times \\
 PE0 \left( \begin{array}{cccc} w_{0,0} & w_{0,1} & 0 & w_{0,3} \\ 0 & 0 & w_{1,2} & 0 \\ 0 & w_{2,1} & 0 & w_{2,3} \\ 0 & 0 & 0 & 0 \\ 0 & 0 & w_{4,2} & w_{4,3} \\ w_{5,0} & 0 & 0 & 0 \\ 0 & 0 & 0 & w_{6,3} \\ 0 & w_{7,1} & 0 & 0 \end{array} \right) = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \xrightarrow{\text{ReLU}} \vec{b} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

$$\begin{array}{c}
 \vec{a} \left( \begin{array}{cccc} 0 & a_1 & 0 & a_3 \end{array} \right) \\
 \times \\
 PE0 \left( \begin{array}{cc|cc} w_{0,0} & w_{0,1} & 0 & w_{0,3} \\ 0 & 0 & w_{1,2} & 0 \\ \hline 0 & w_{2,1} & 0 & w_{2,3} \\ 0 & 0 & 0 & 0 \\ 0 & 0 & w_{4,2} & w_{4,3} \\ w_{5,0} & 0 & 0 & 0 \\ 0 & 0 & 0 & w_{6,3} \\ 0 & w_{7,1} & 0 & 0 \end{array} \right) = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \xrightarrow{\text{ReLU}} \vec{b} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

$$\begin{array}{c}
 \vec{a} \left( \begin{array}{cccc} 0 & a_1 & 0 & \boxed{a_3} \end{array} \right) \\
 \times \\
 PE0 \left( \begin{array}{cccc} w_{0,0} & \boxed{w_{0,1}} & 0 & w_{0,3} \\ 0 & 0 & w_{1,2} & 0 \\ 0 & \boxed{w_{2,1}} & 0 & w_{2,3} \\ 0 & 0 & 0 & 0 \\ 0 & 0 & w_{4,2} & w_{4,3} \\ w_{5,0} & 0 & 0 & 0 \\ 0 & 0 & 0 & w_{6,3} \\ 0 & \boxed{w_{7,1}} & 0 & 0 \end{array} \right) \\
 = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \\
 \xrightarrow{\text{ReLU}} \vec{b} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

$$\begin{array}{c}
 \vec{a} \left( \begin{array}{cccc} 0 & a_1 & 0 & a_3 \end{array} \right) \\
 \times \\
 PE0 \left( \begin{array}{cc|cc} w_{0,0} & w_{0,1} & 0 & w_{0,3} \\ \hline 0 & 0 & w_{1,2} & 0 \\ 0 & w_{2,1} & 0 & w_{2,3} \\ \hline 0 & 0 & 0 & 0 \\ 0 & 0 & w_{4,2} & w_{4,3} \\ \hline w_{5,0} & 0 & 0 & 0 \\ 0 & 0 & 0 & w_{6,3} \\ \hline 0 & w_{7,1} & 0 & 0 \end{array} \right) \\
 = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \\
 \xrightarrow{\text{ReLU}} \vec{b} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

$$\vec{a} \begin{pmatrix} 0 & a_1 & 0 & a_3 \end{pmatrix} \times \\
 PE0 \begin{pmatrix} w_{0,0} & w_{0,1} & 0 & w_{0,3} \\ 0 & 0 & w_{1,2} & 0 \\ 0 & w_{2,1} & 0 & w_{2,3} \\ 0 & 0 & 0 & 0 \\ 0 & 0 & w_{4,2} & w_{4,3} \\ w_{5,0} & 0 & 0 & 0 \\ 0 & 0 & 0 & w_{6,3} \\ 0 & w_{7,1} & 0 & 0 \end{pmatrix} = \begin{pmatrix} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{pmatrix} \xrightarrow{\text{ReLU}} \vec{b} \begin{pmatrix} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{pmatrix}$$

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

# Dataflow

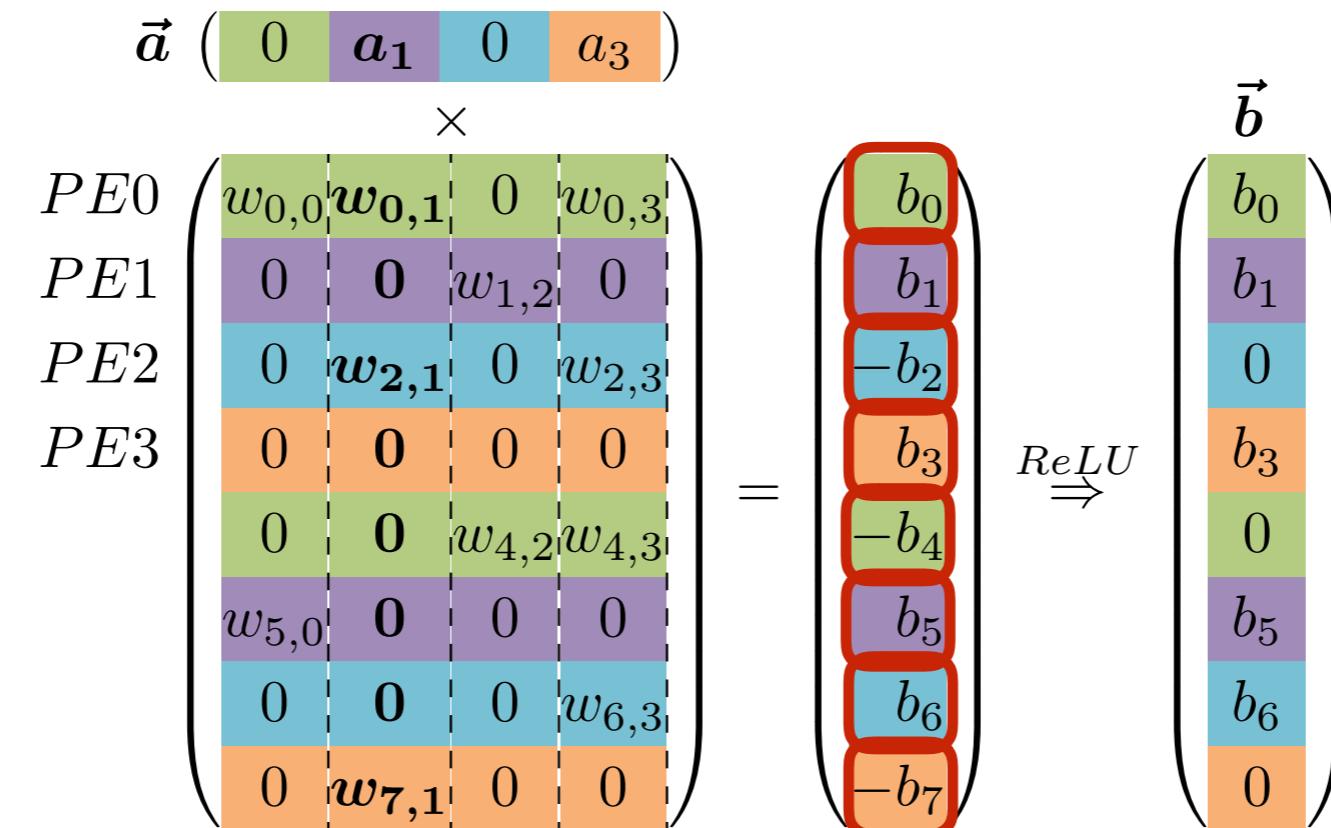
$$\begin{array}{c}
 \vec{a} \left( \begin{array}{cccc} 0 & a_1 & 0 & a_3 \end{array} \right) \\
 \times \\
 PE0 \left( \begin{array}{cc|cc} w_{0,0} & w_{0,1} & 0 & w_{0,3} \\ \hline 0 & 0 & w_{1,2} & 0 \end{array} \right) \\
 PE1 \left( \begin{array}{cc|cc} 0 & 0 & w_{1,2} & 0 \\ \hline 0 & w_{2,1} & 0 & w_{2,3} \end{array} \right) \\
 PE2 \left( \begin{array}{cc|cc} 0 & 0 & 0 & 0 \\ \hline 0 & 0 & w_{4,2} & w_{4,3} \end{array} \right) \\
 PE3 \left( \begin{array}{cc|cc} w_{5,0} & 0 & 0 & 0 \\ \hline 0 & 0 & 0 & w_{6,3} \\ 0 & w_{7,1} & 0 & 0 \end{array} \right) \\
 = \left( \begin{array}{c} b_0 \\ b_1 \\ -b_2 \\ b_3 \\ -b_4 \\ b_5 \\ b_6 \\ -b_7 \end{array} \right) \\
 \xrightarrow{\text{ReLU}} \vec{b} \left( \begin{array}{c} b_0 \\ b_1 \\ 0 \\ b_3 \\ 0 \\ b_5 \\ b_6 \\ 0 \end{array} \right)
 \end{array}$$

The diagram illustrates a sparse matrix-vector multiplication. The input vector  $\vec{a}$  has elements  $0, a_1, 0, a_3$ . The weight matrix is partitioned into 8 parallel processing elements (PE0 to PE3). The multiplication results in intermediate values  $b_0, b_1, -b_2, b_3, -b_4, b_5, b_6, -b_7$ , which are then passed through a ReLU activation function to produce the final output vector  $\vec{b}$  with elements  $b_0, b_1, 0, b_3, 0, b_5, b_6, 0$ . Specific weights and results are highlighted with red boxes.

rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

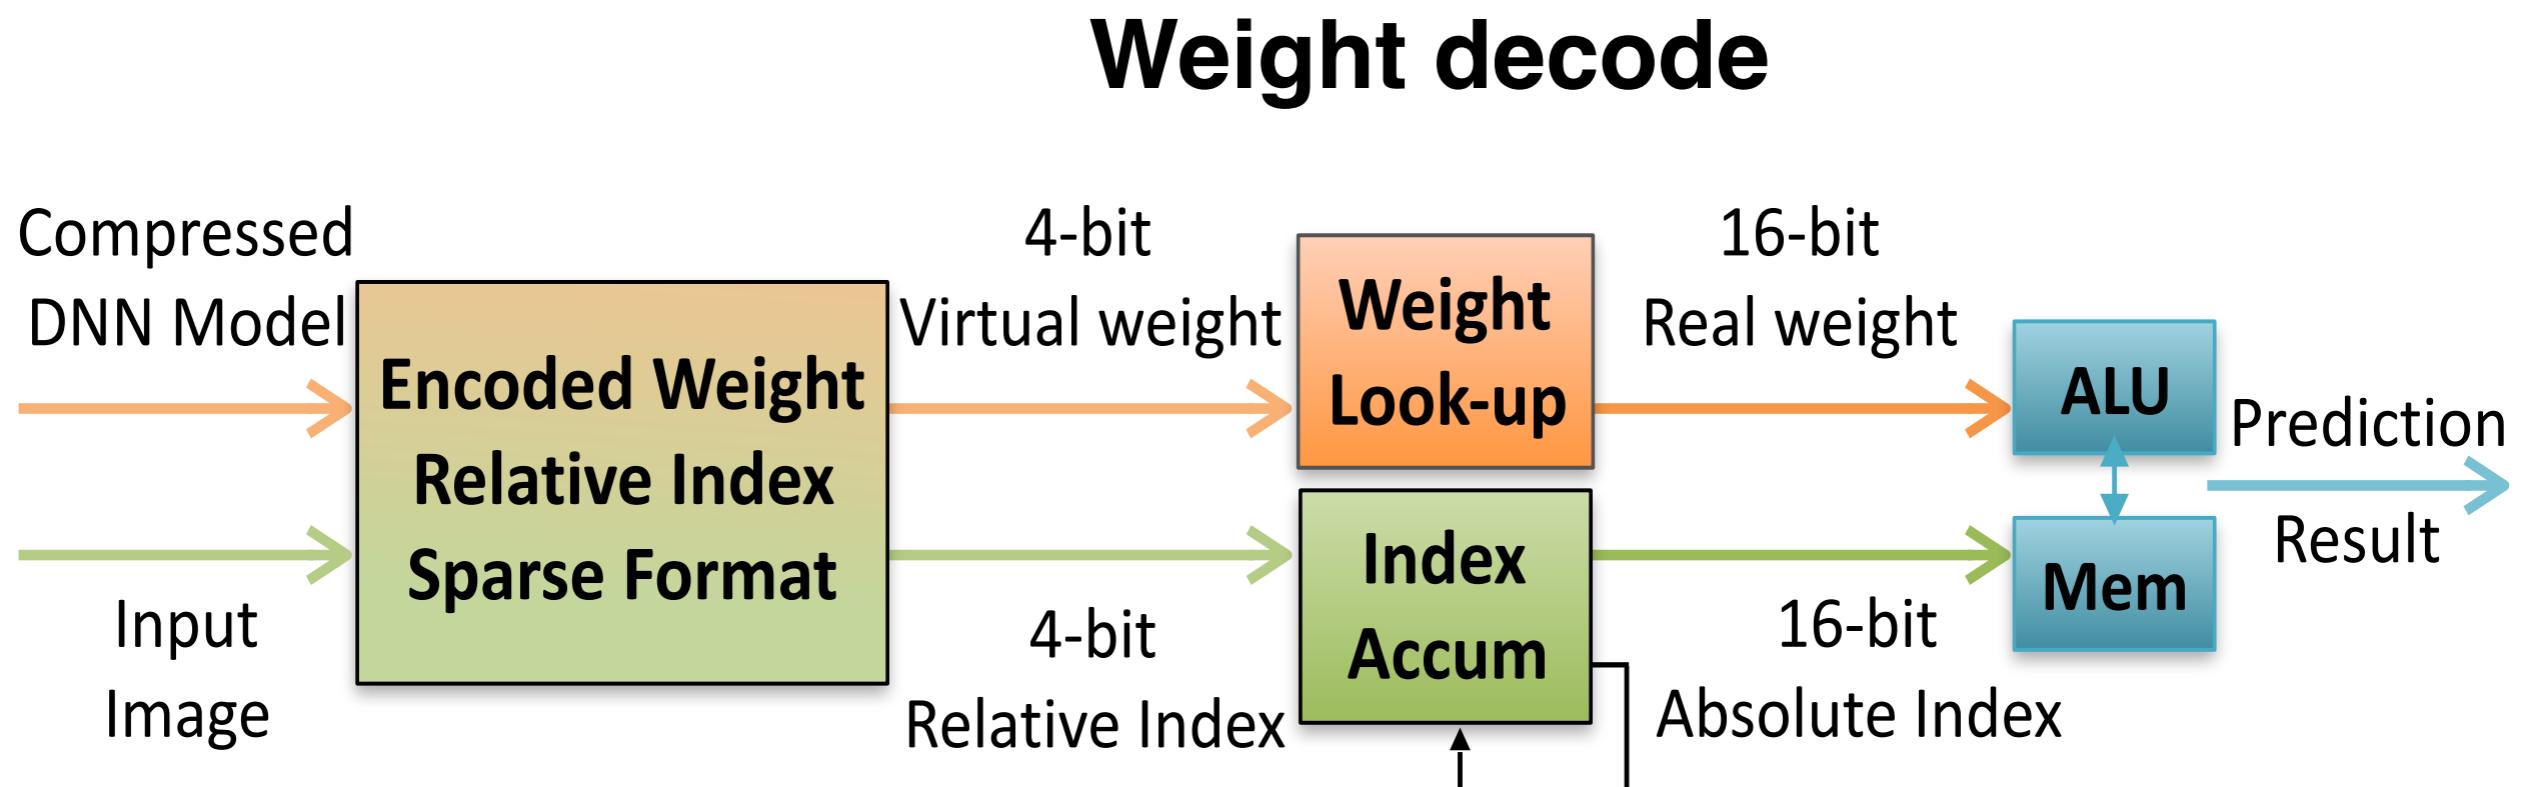
# Dataflow



rule of thumb:

$$0 * A = 0 \quad W * 0 = 0$$

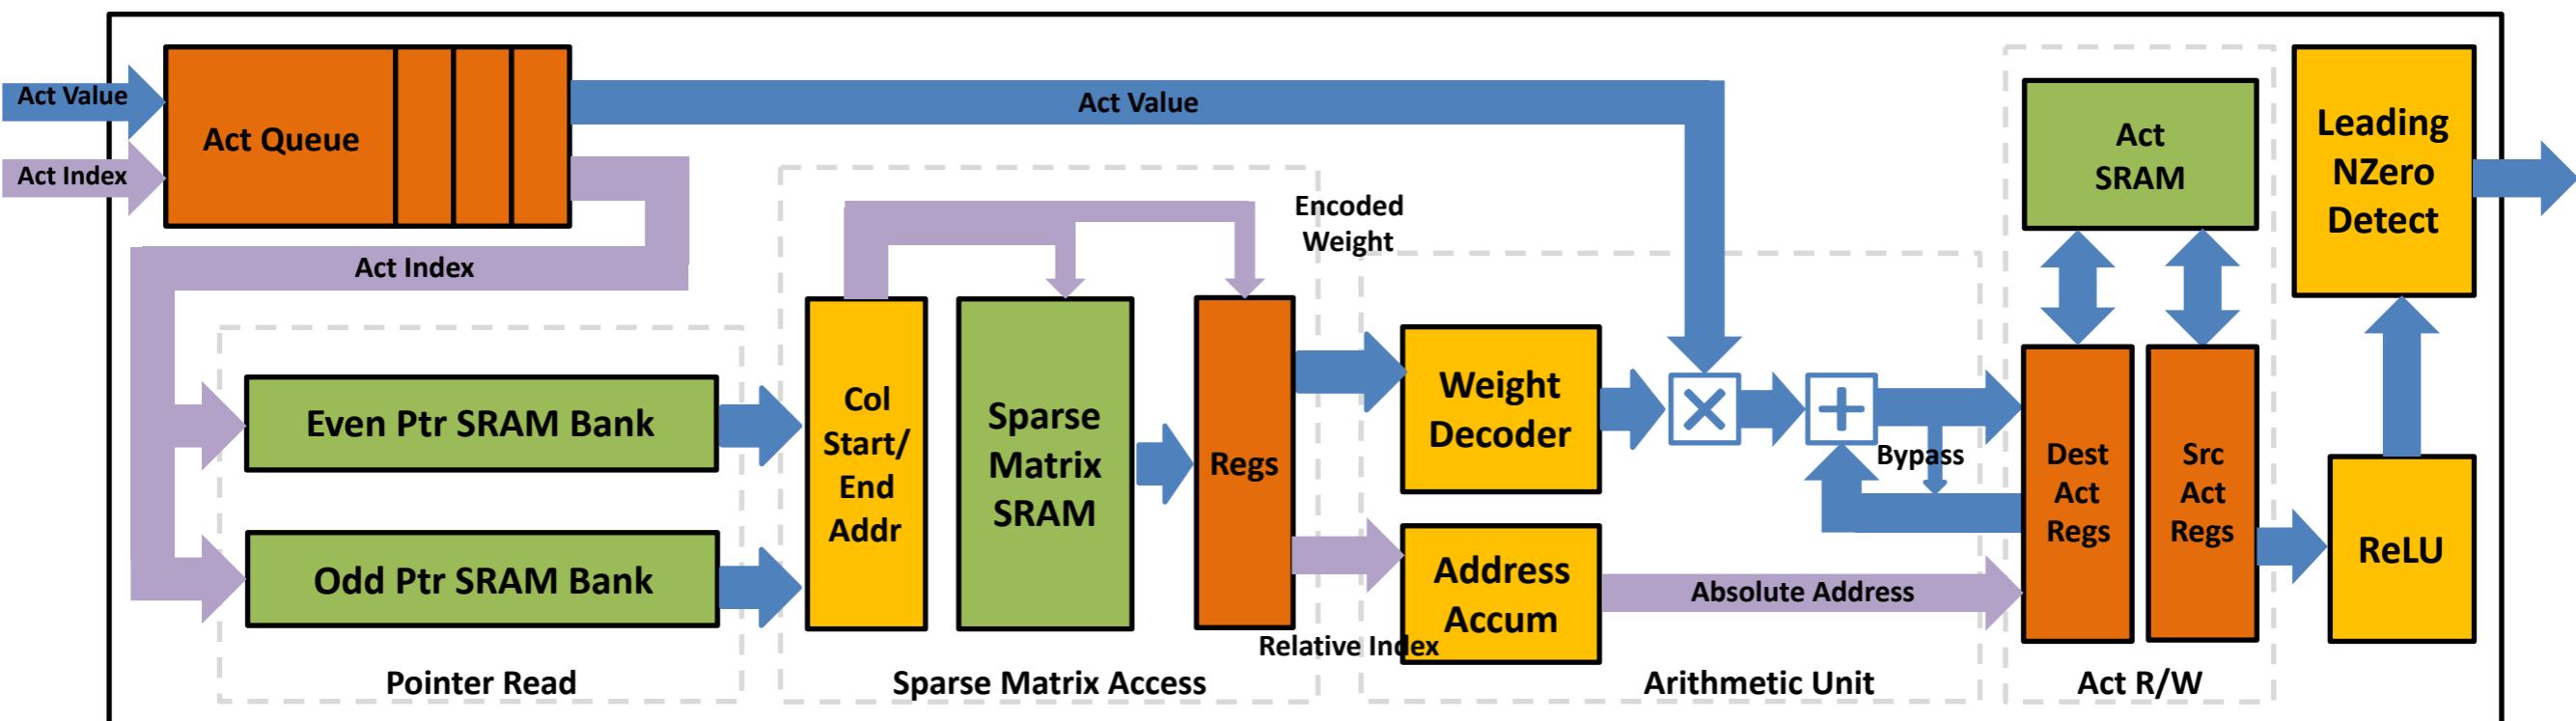
# EIE Architecture



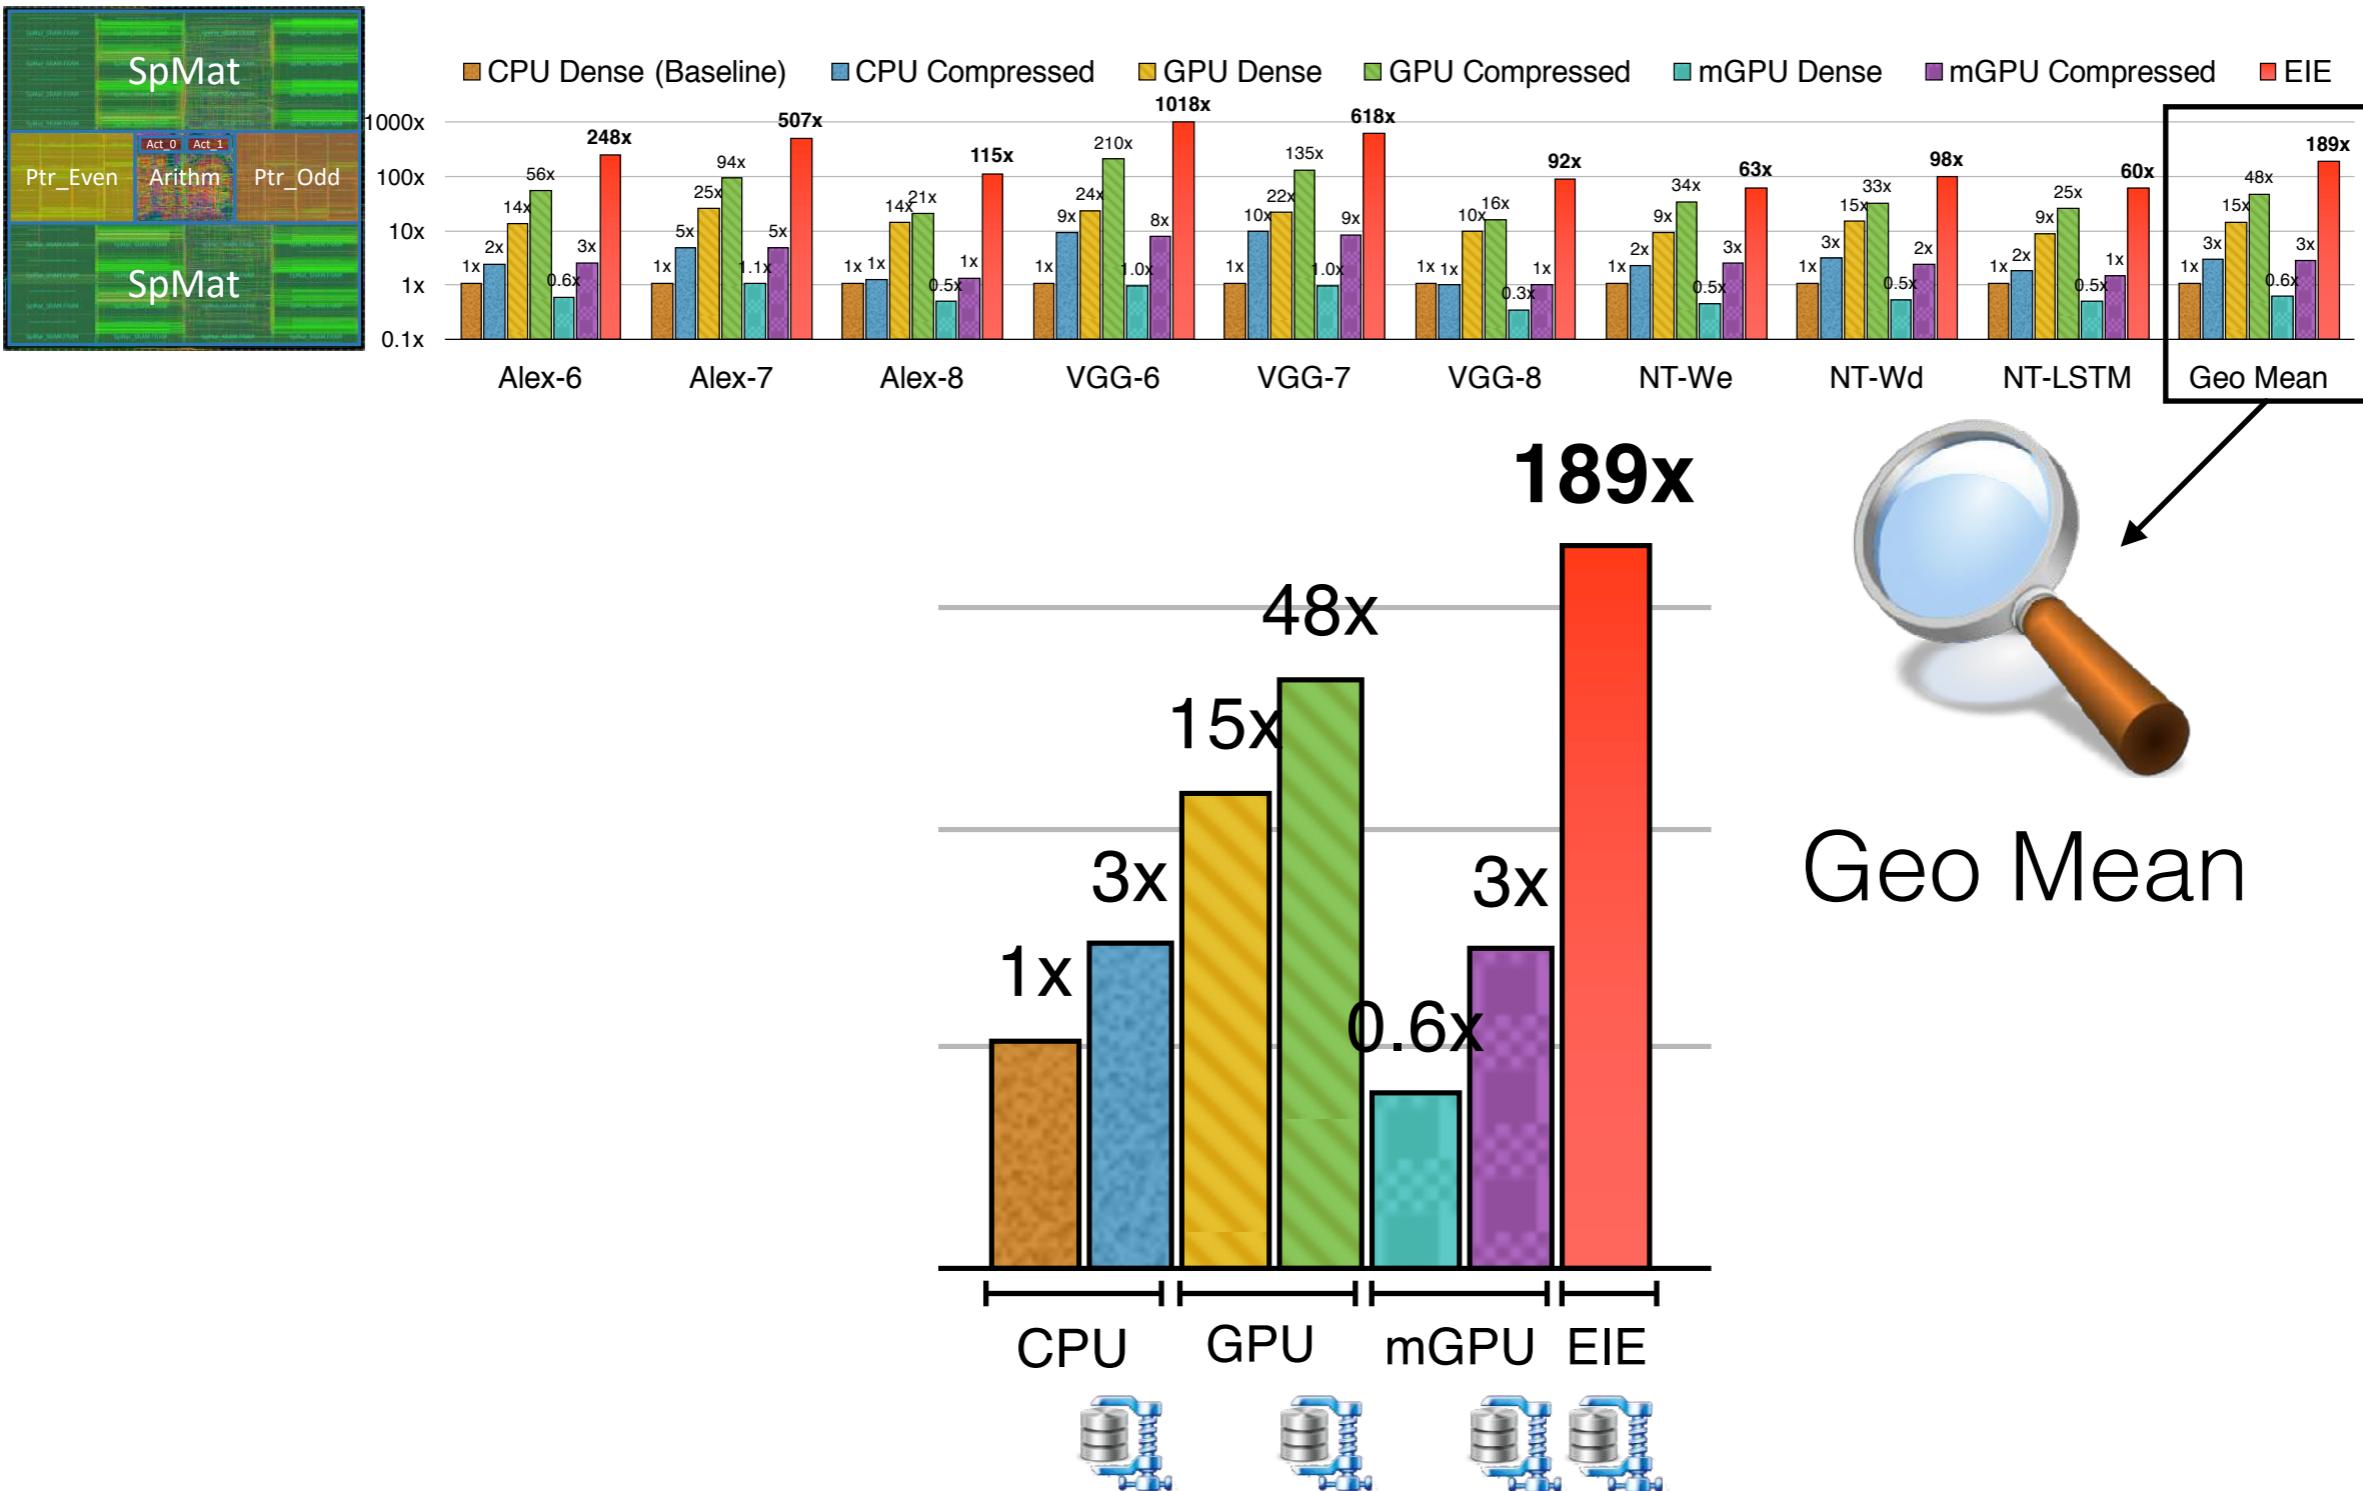
## Address Accumulate

rule of thumb:  $0 * A = 0$        $W * 0 = 0$        ~~$2.09, 1.92 \Rightarrow 2$~~

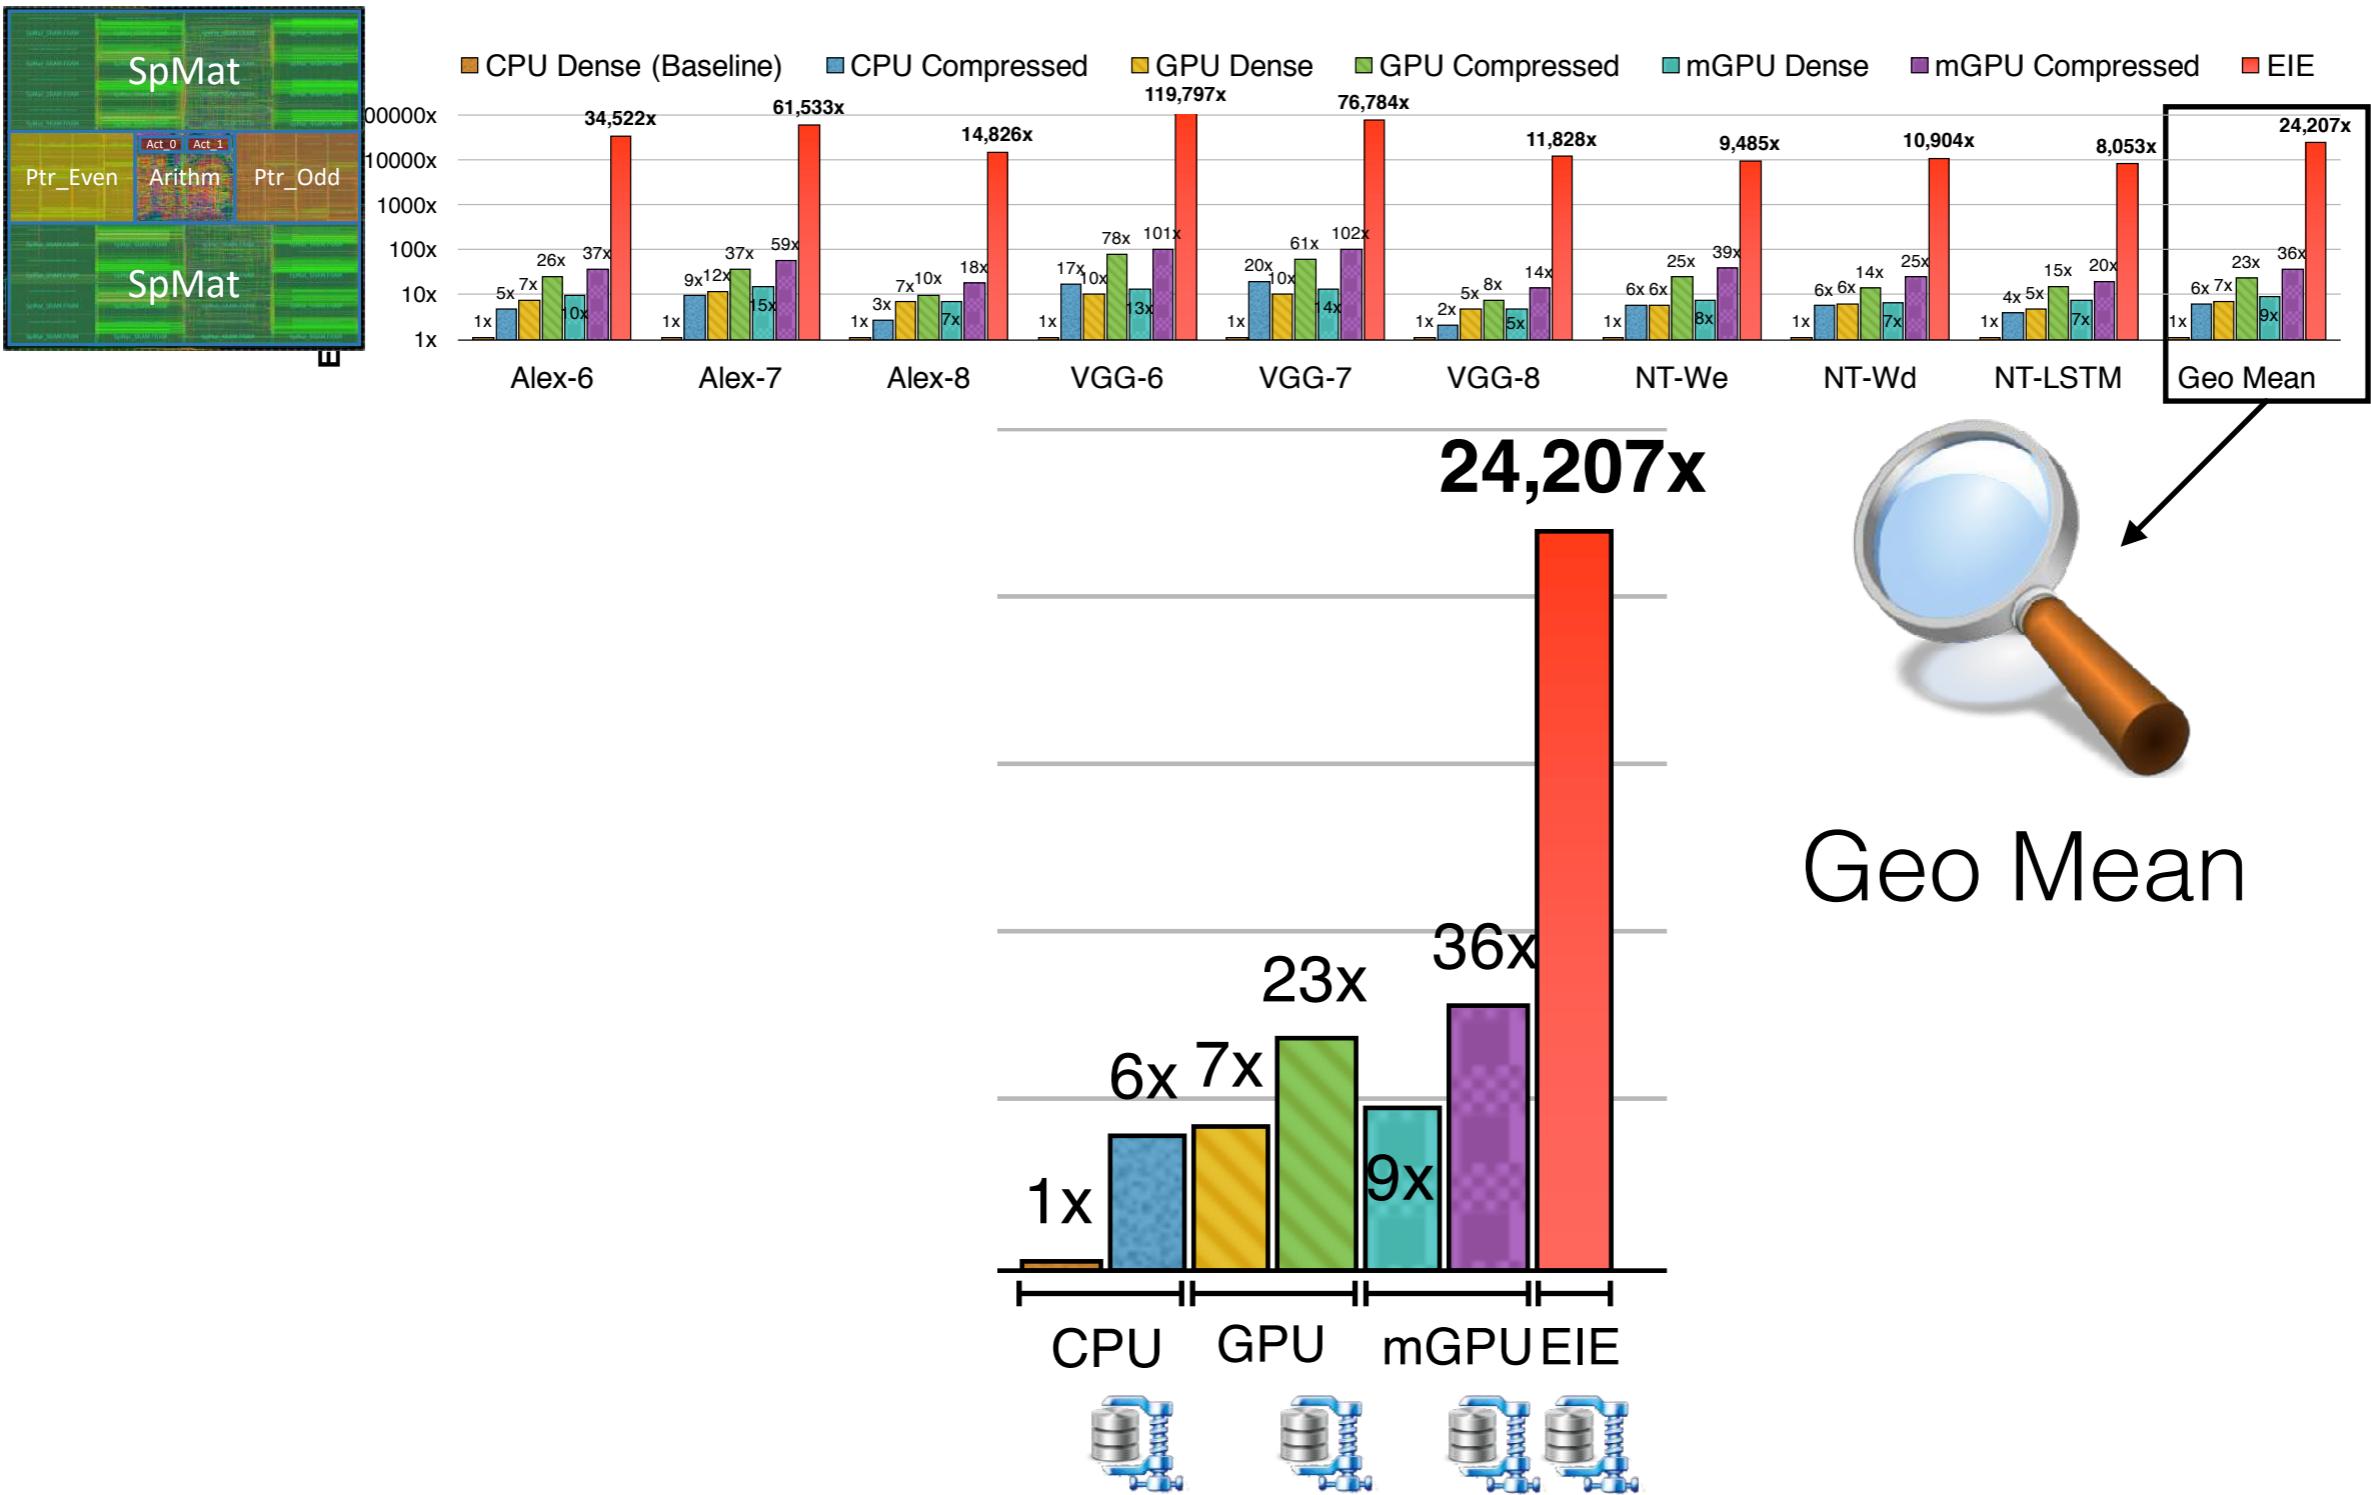
# Micro Architecture for each PE



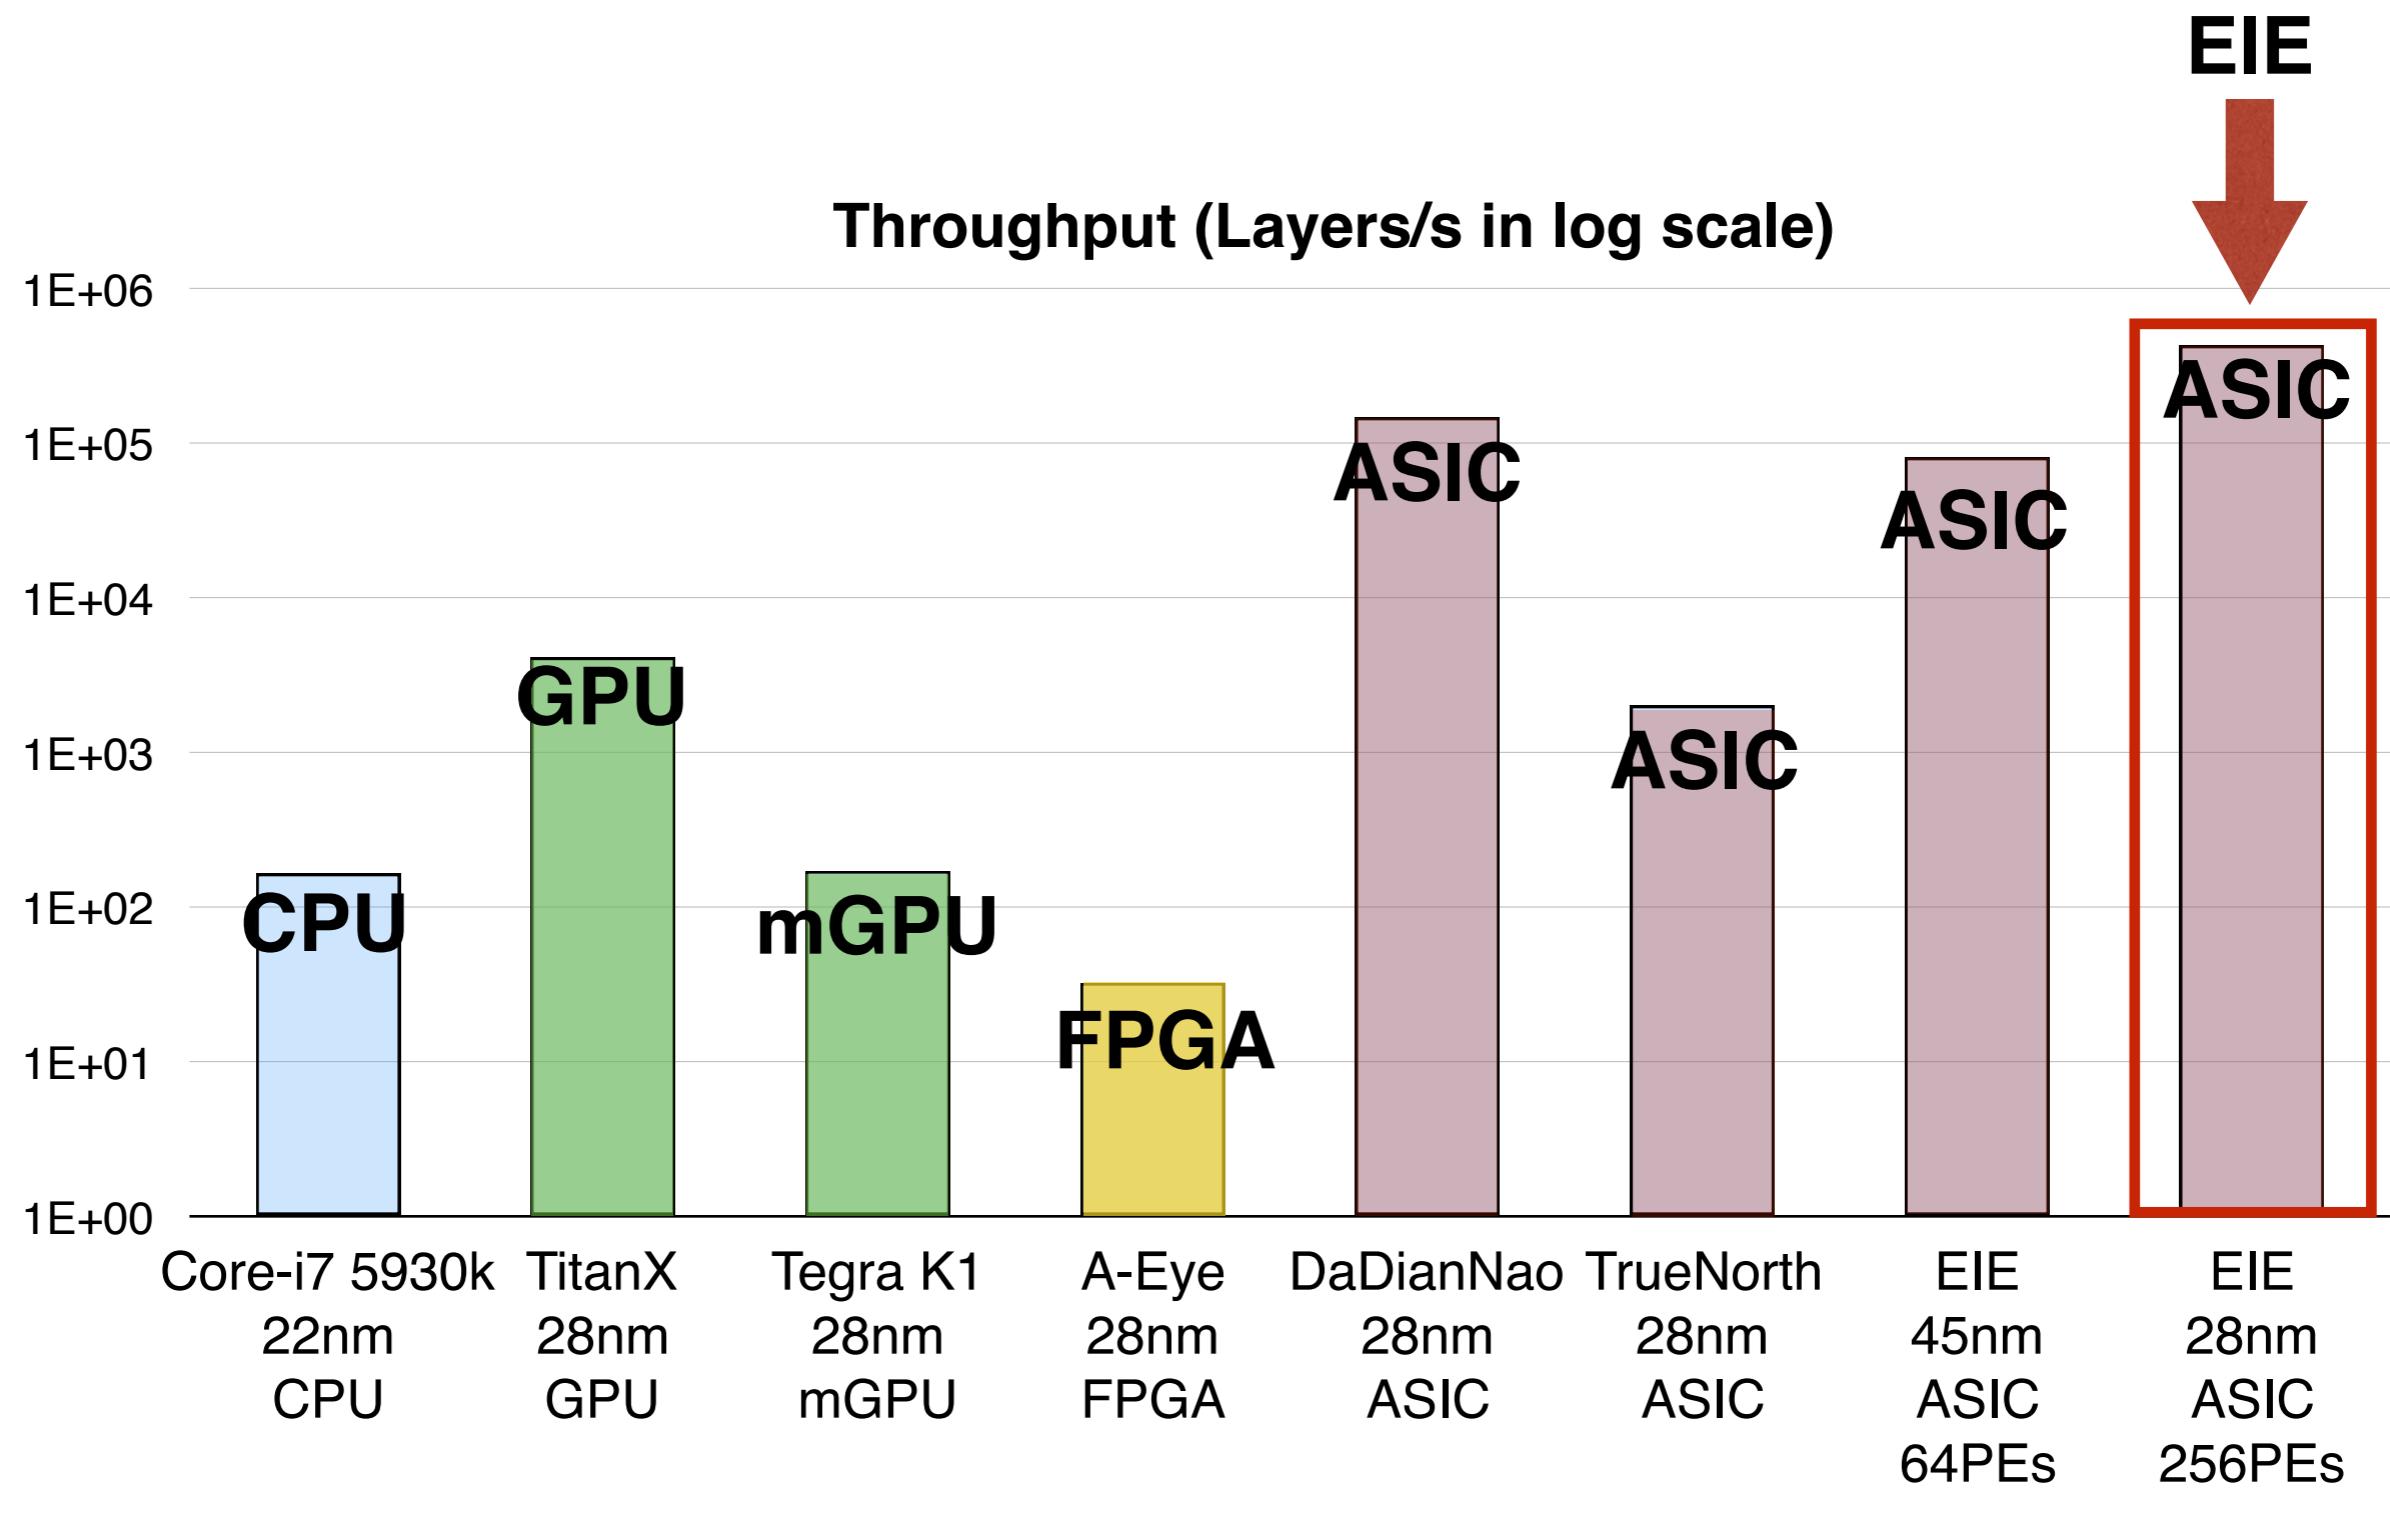
# Speedup on EIE



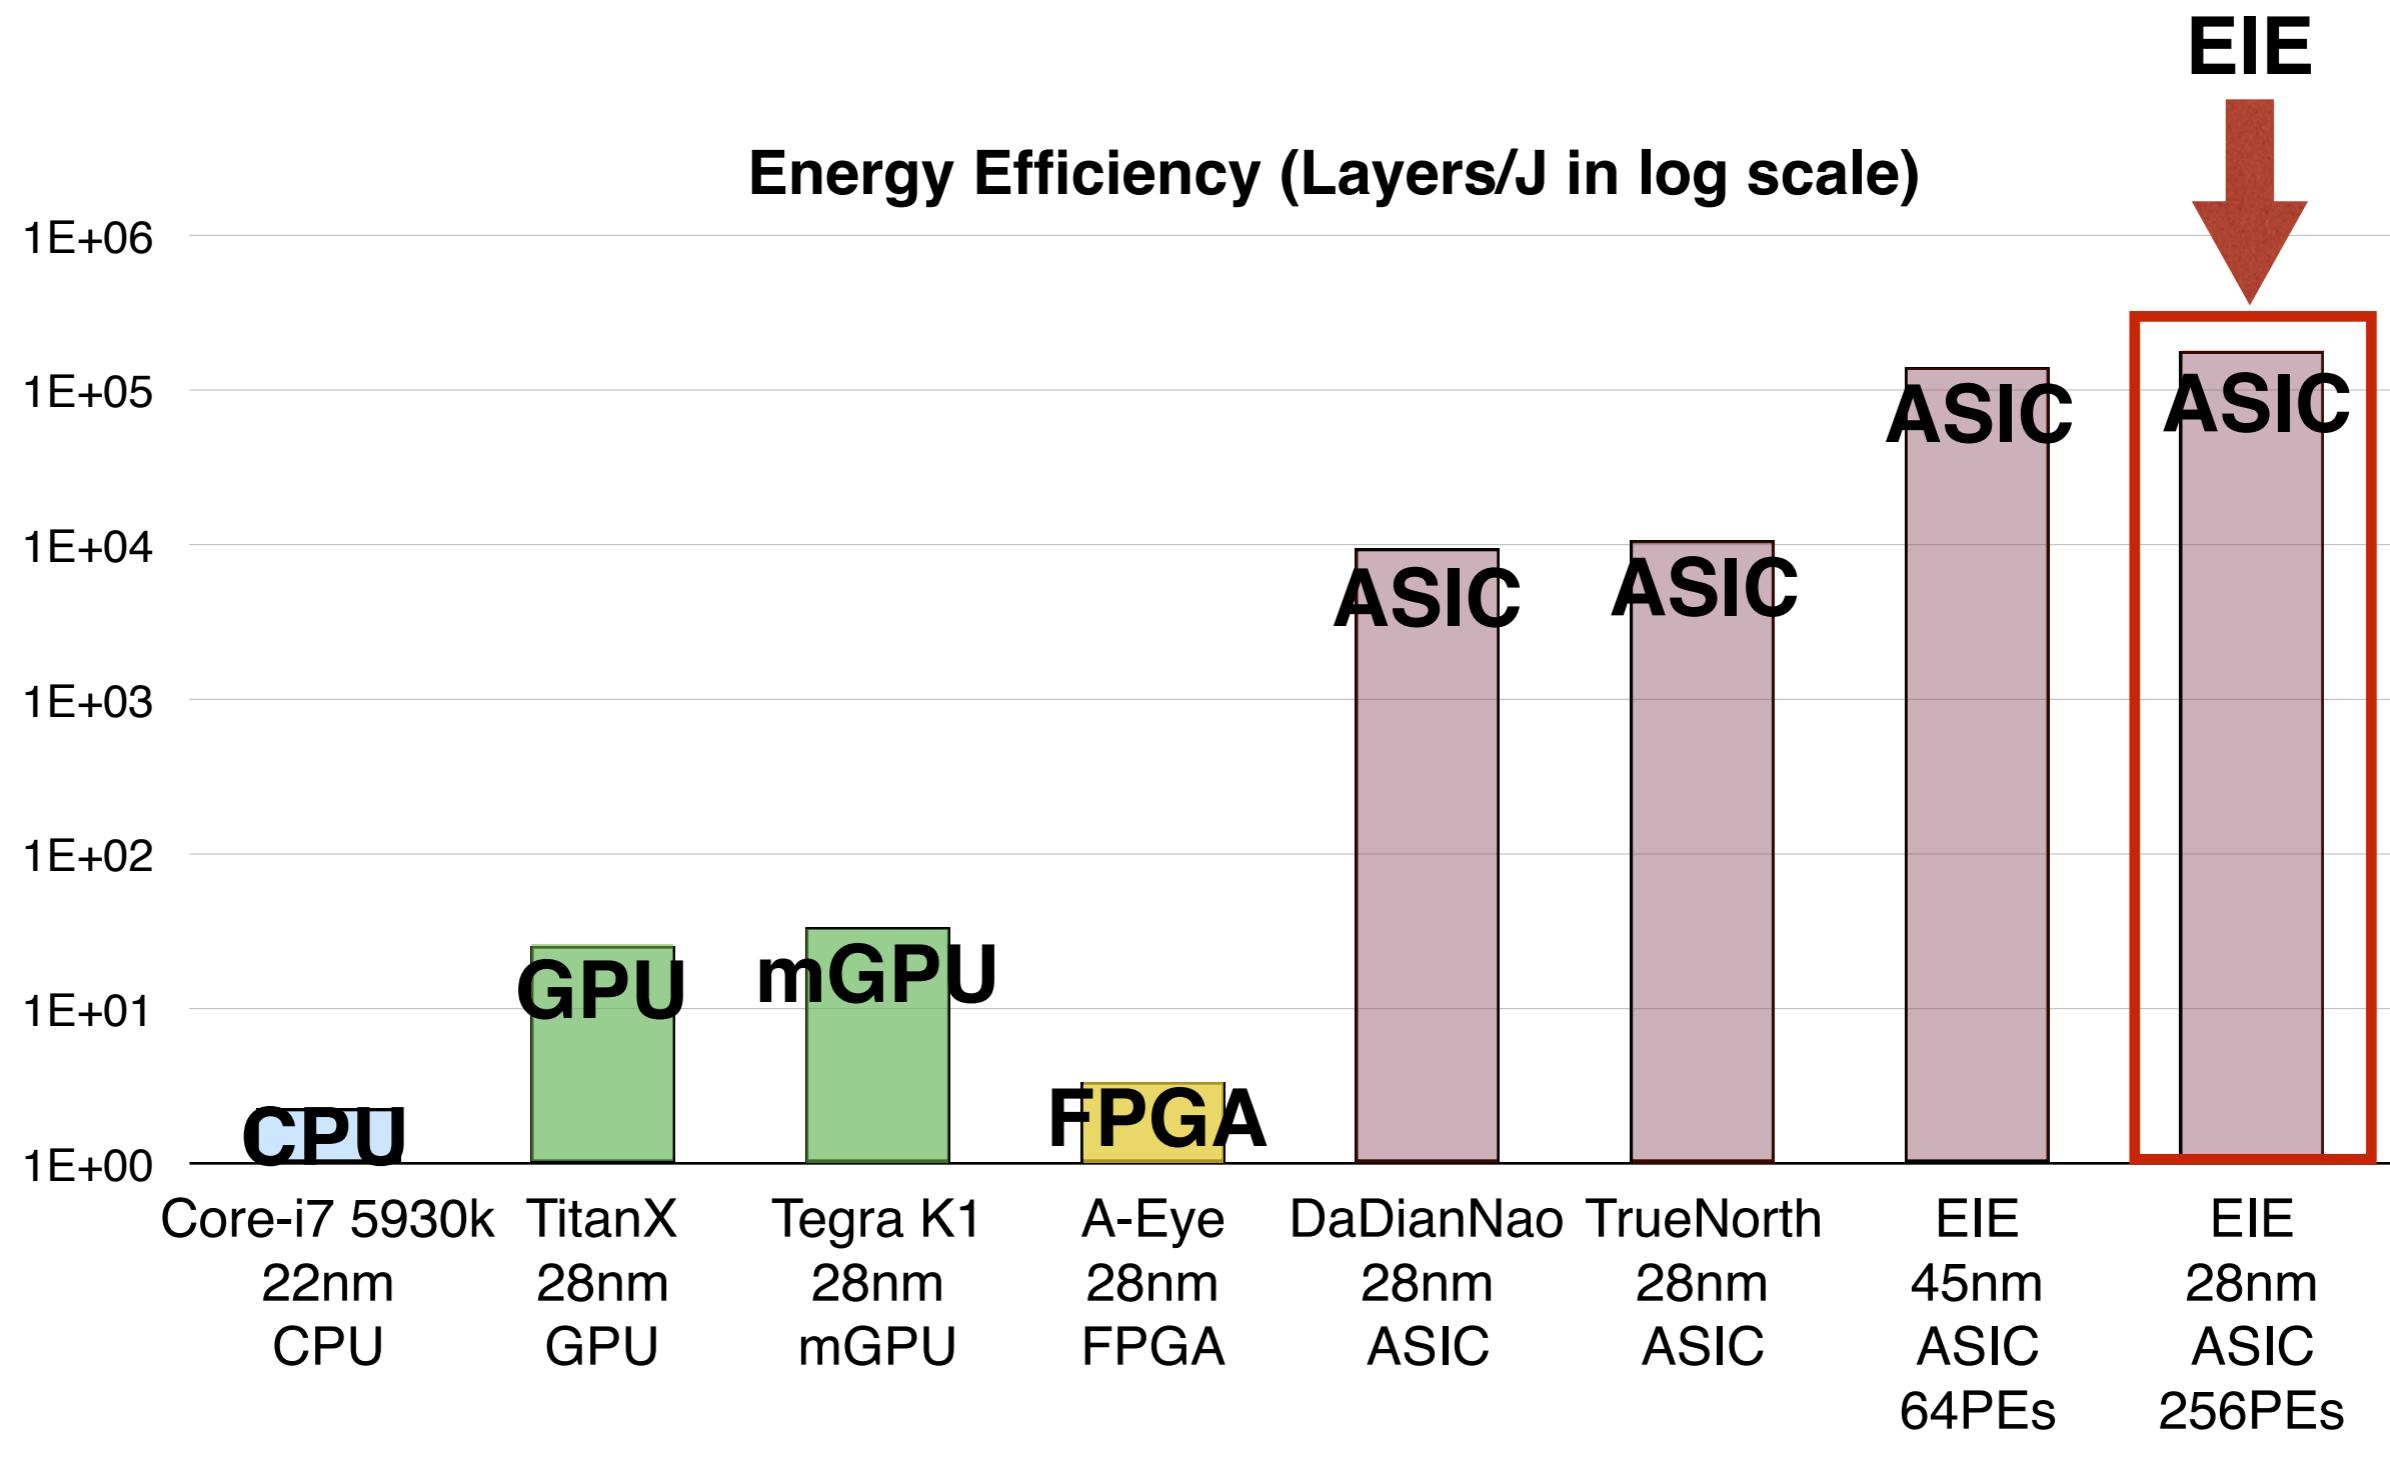
# Energy Efficiency on EIE



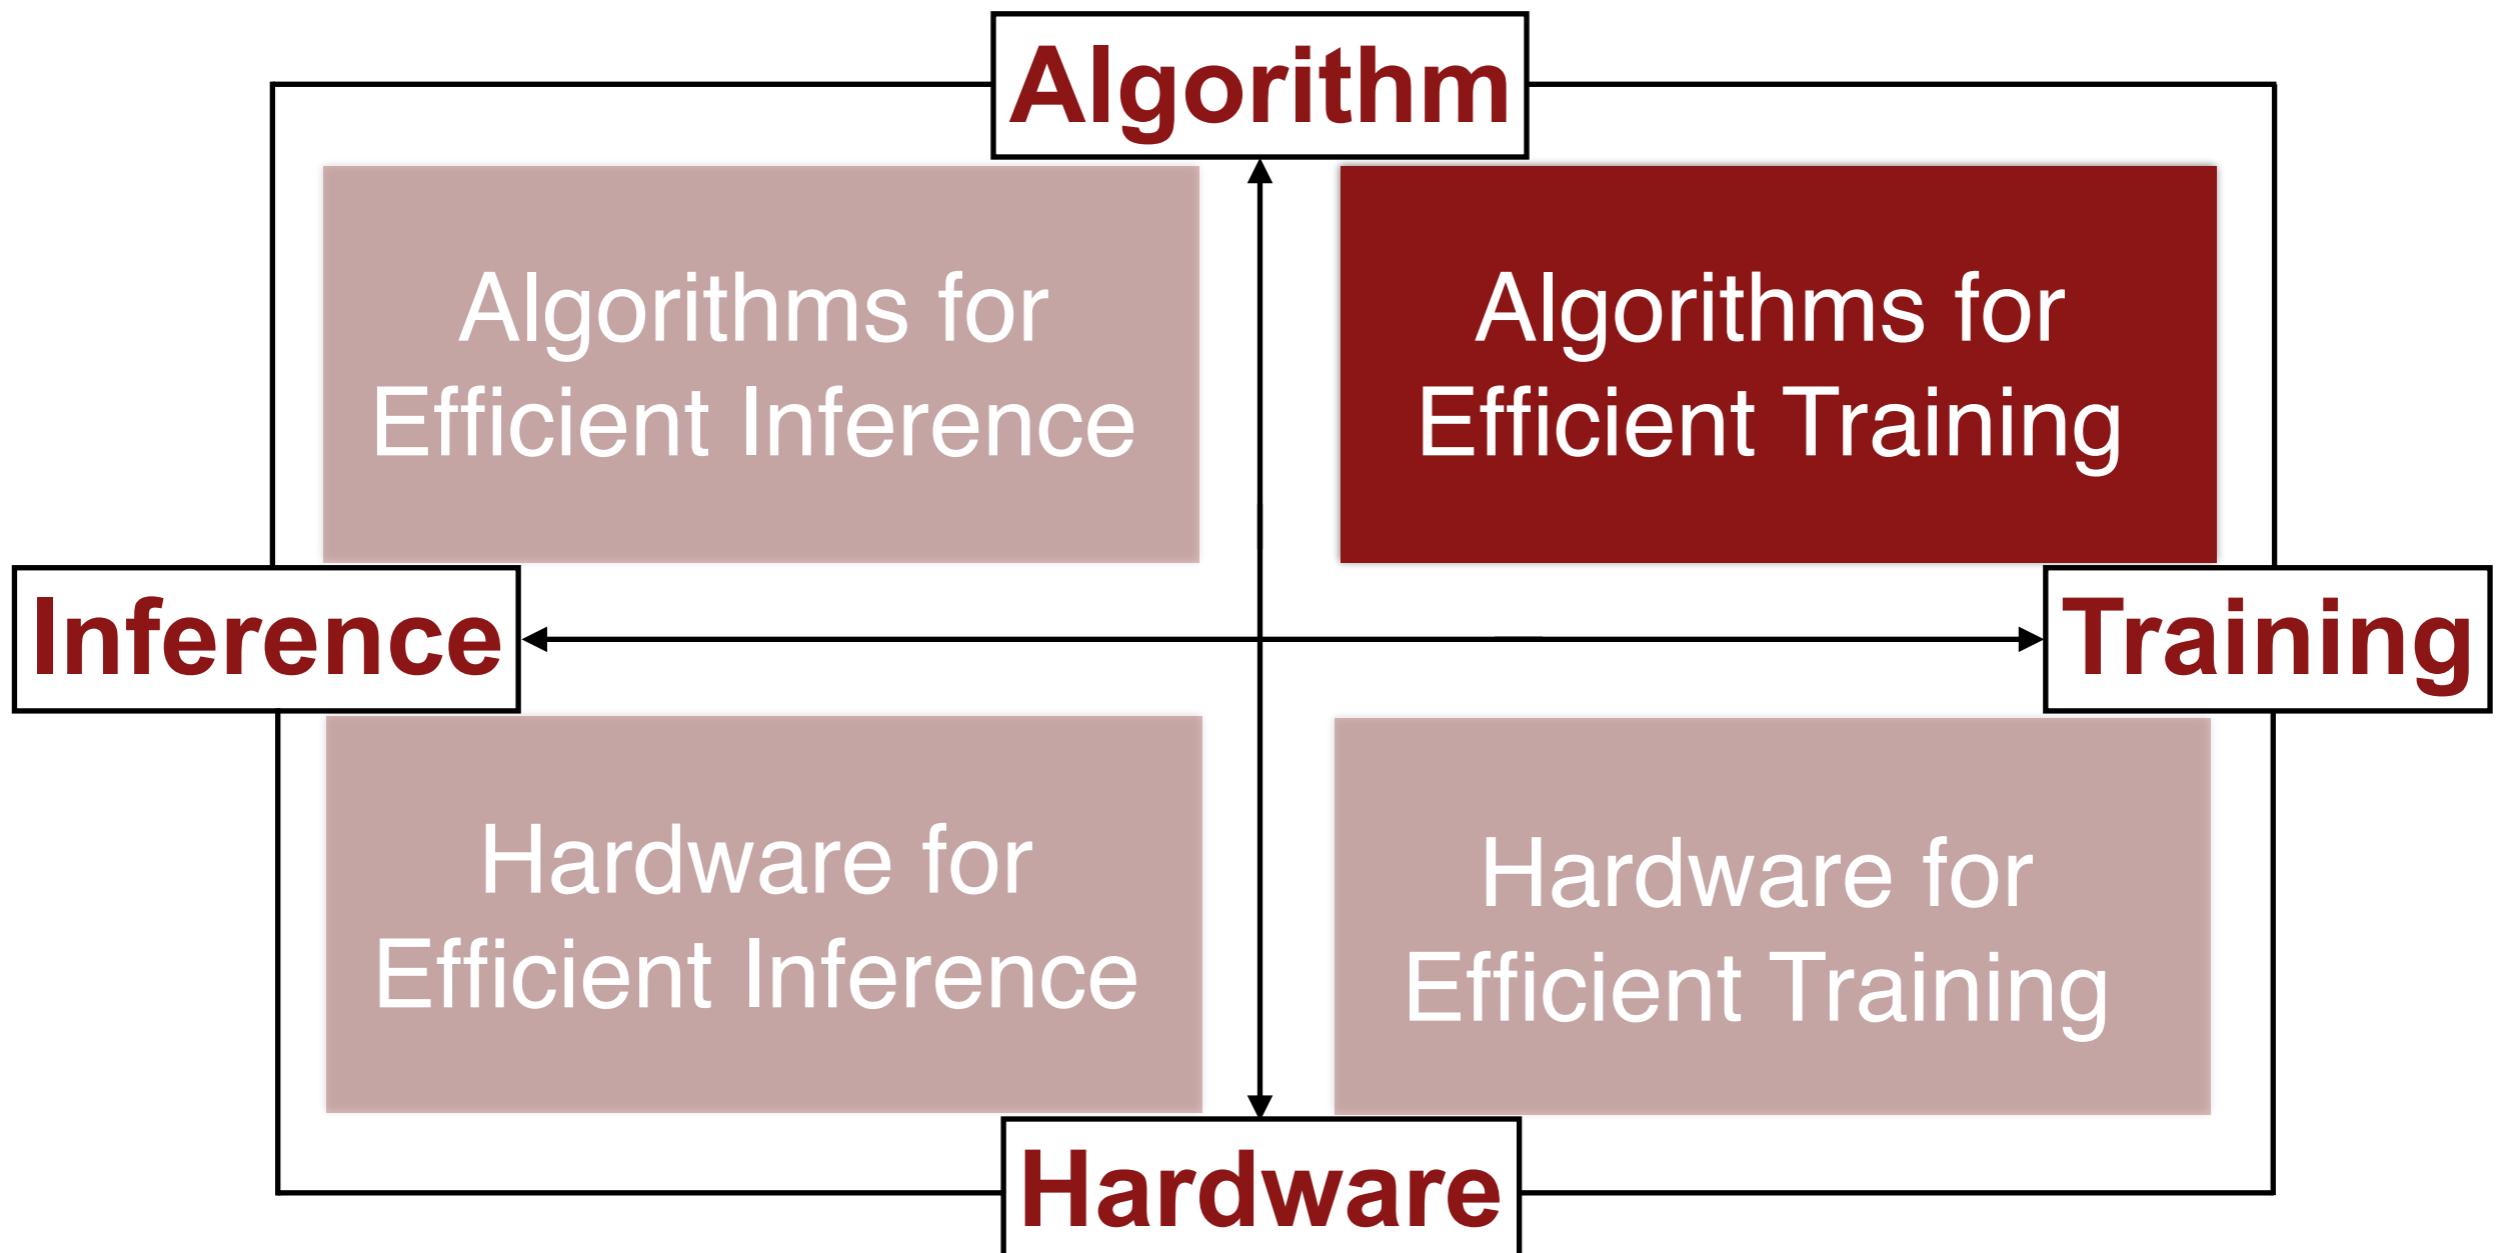
# Comparison: Throughput



# Comparison: Energy Efficiency



# Agenda



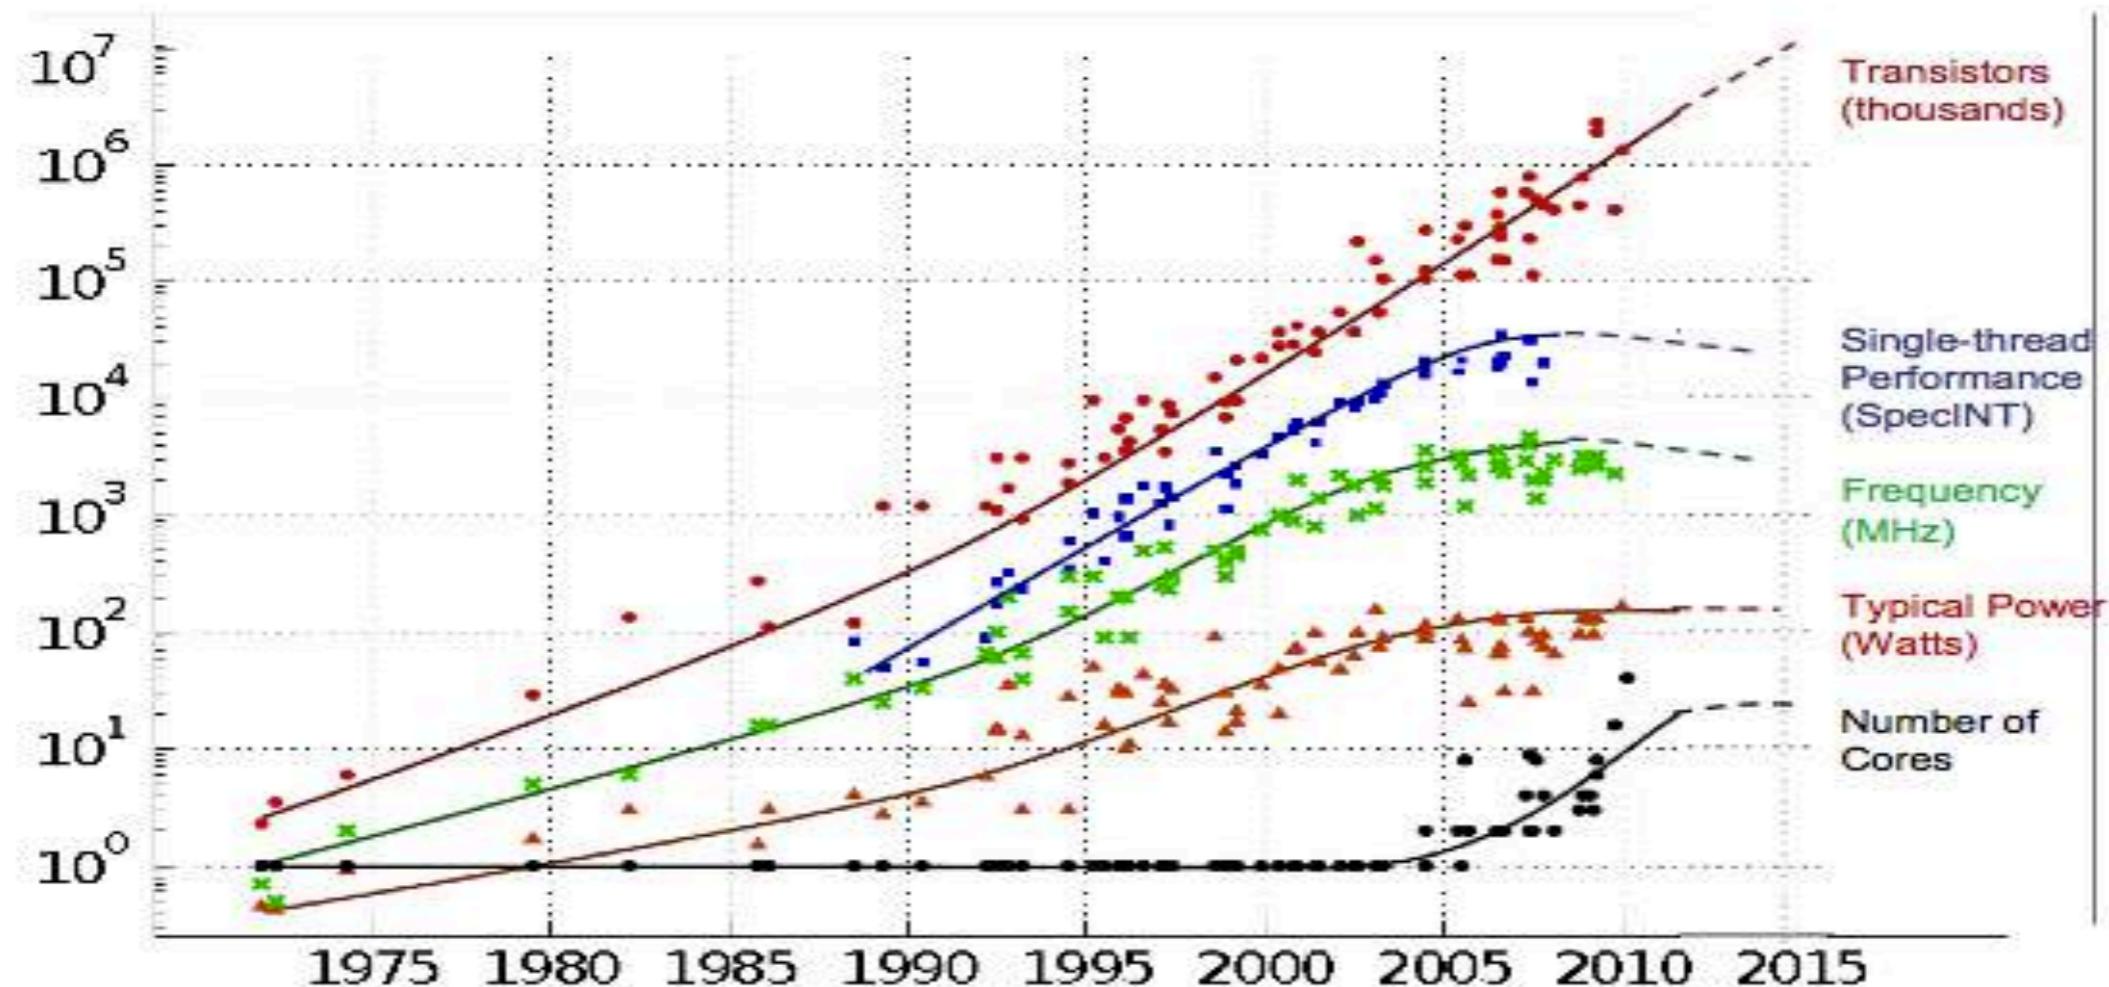
# Part 3: Efficient Training — Algorithms

- 1. Parallelization
- 2. Mixed Precision with FP16 and FP32
- 3. Model Distillation
- 4. DSD: Dense-Sparse-Dense Training

# Part 3: Efficient Training — Algorithms

- 1. Parallelization
- 2. Mixed Precision with FP16 and FP32
- 3. Model Distillation
- 4. DSD: Dense-Sparse-Dense Training

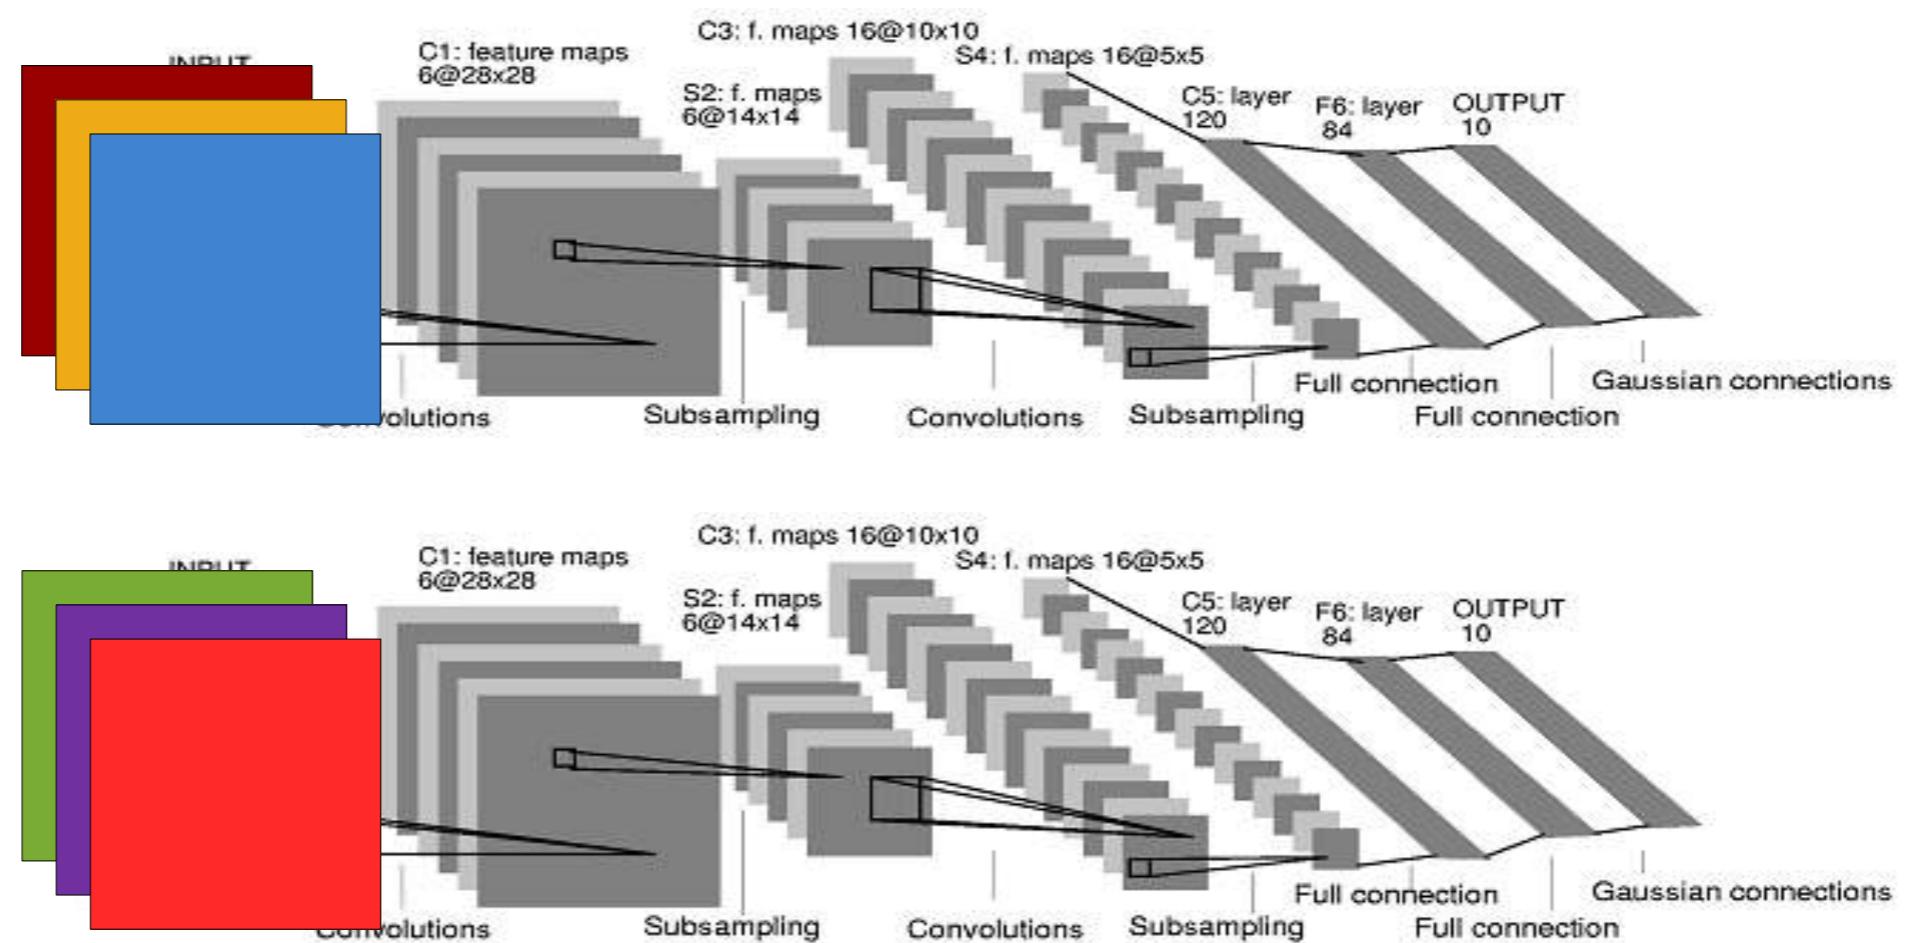
# Moore's law made CPUs 300x faster than in 1990 But its over...



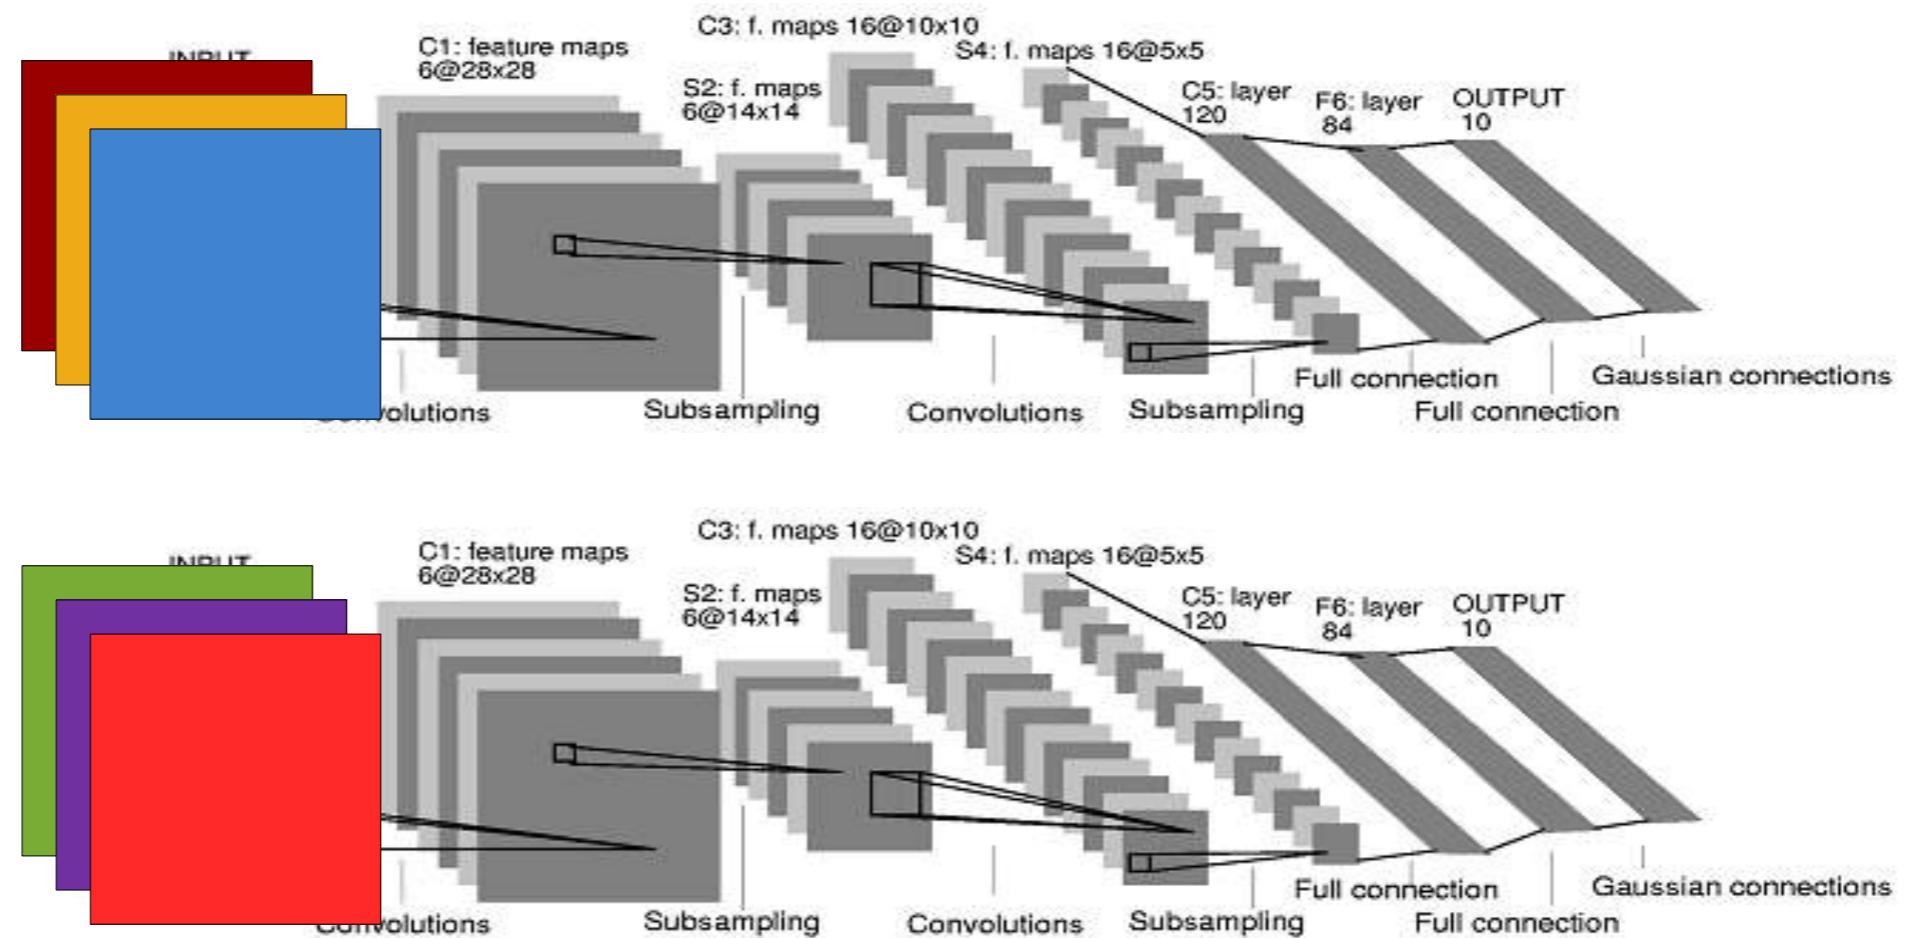
Original data collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond and C. Batten  
Dotted line extrapolations by C. Moore

C Moore, Data Processing in ExaScale-Class Computer Systems, Salishan, April 2011

# Data Parallel – Run multiple inputs in parallel



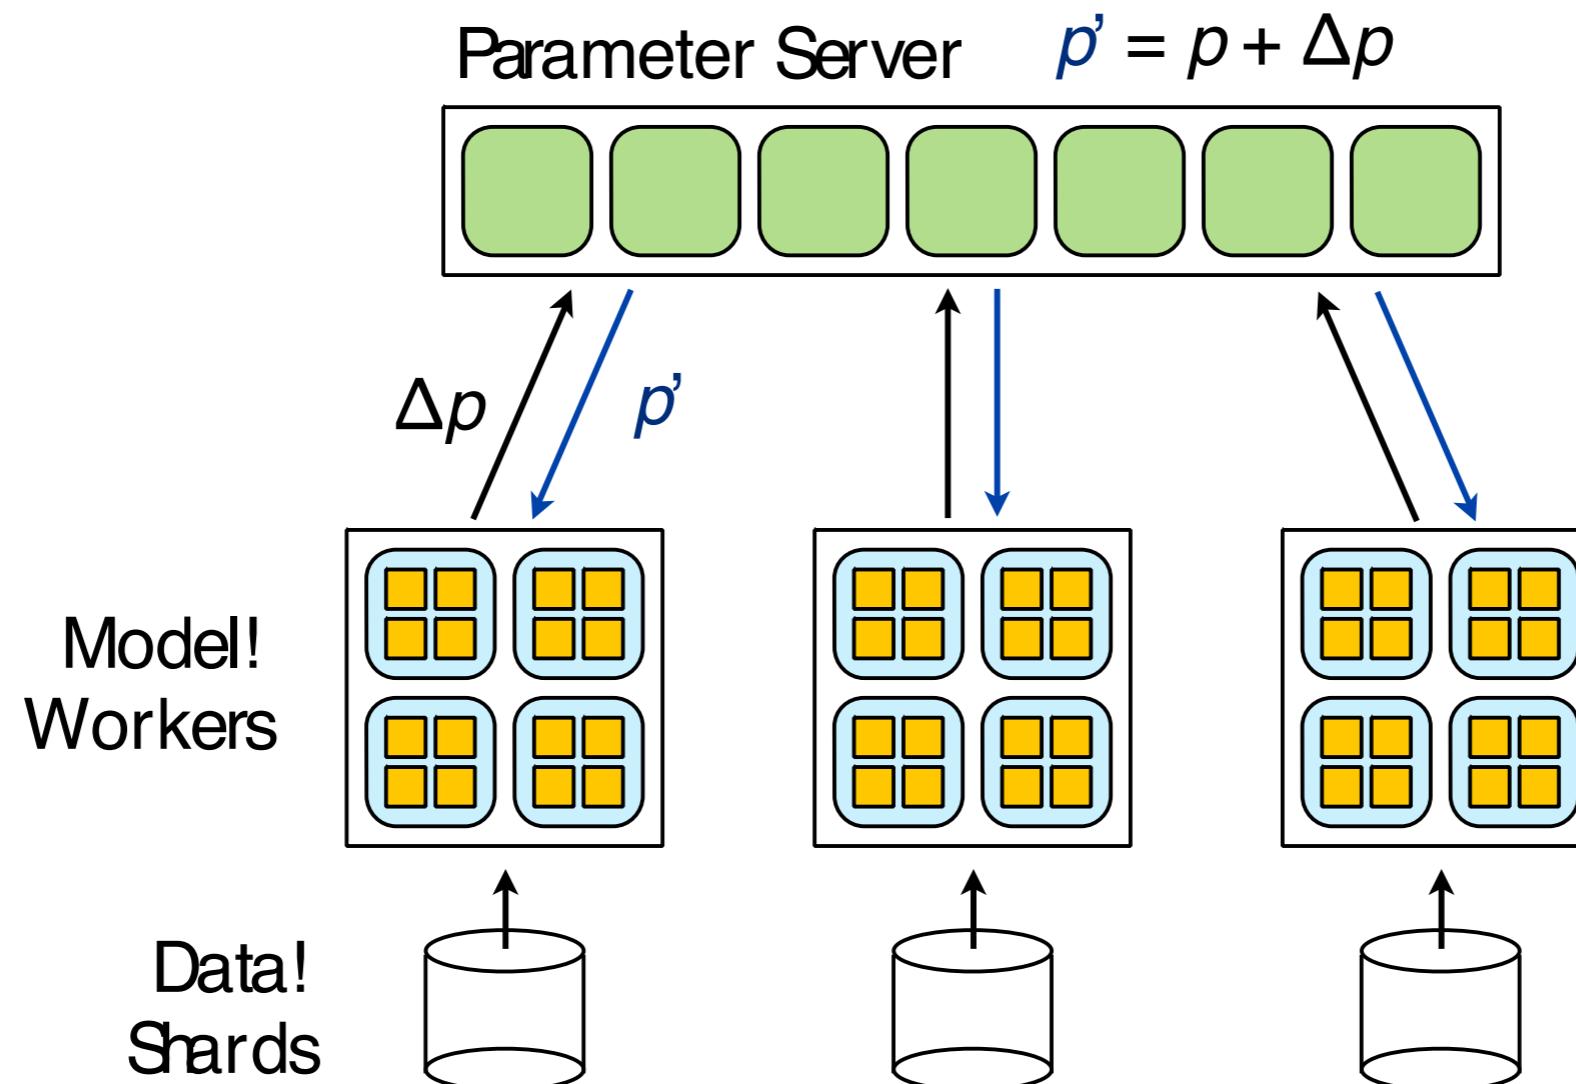
# Data Parallel – Run multiple inputs in parallel



- Doesn't affect latency for one input
- Requires P-fold larger batch size
- For training requires coordinated weight update

Dally, High Performance Hardware for Machine Learning, NIPS'2015

# Parameter Update



Large Scale Distributed Deep Networks, Jeff Dean et al., 2013

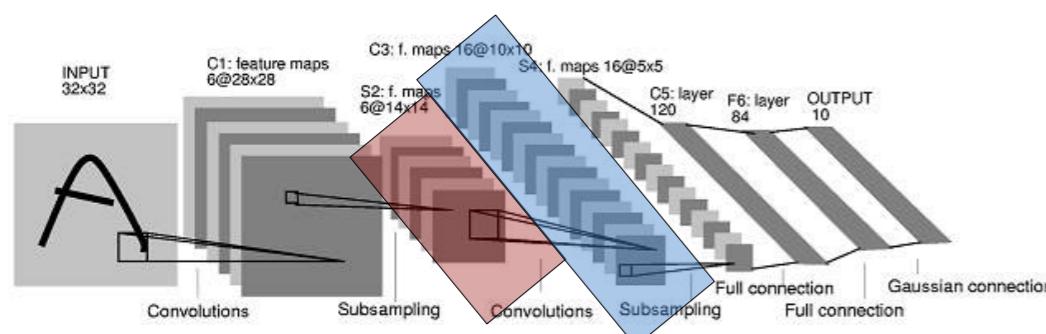
# **Model Parallel**

## **Split up the Model – i.e. the network**

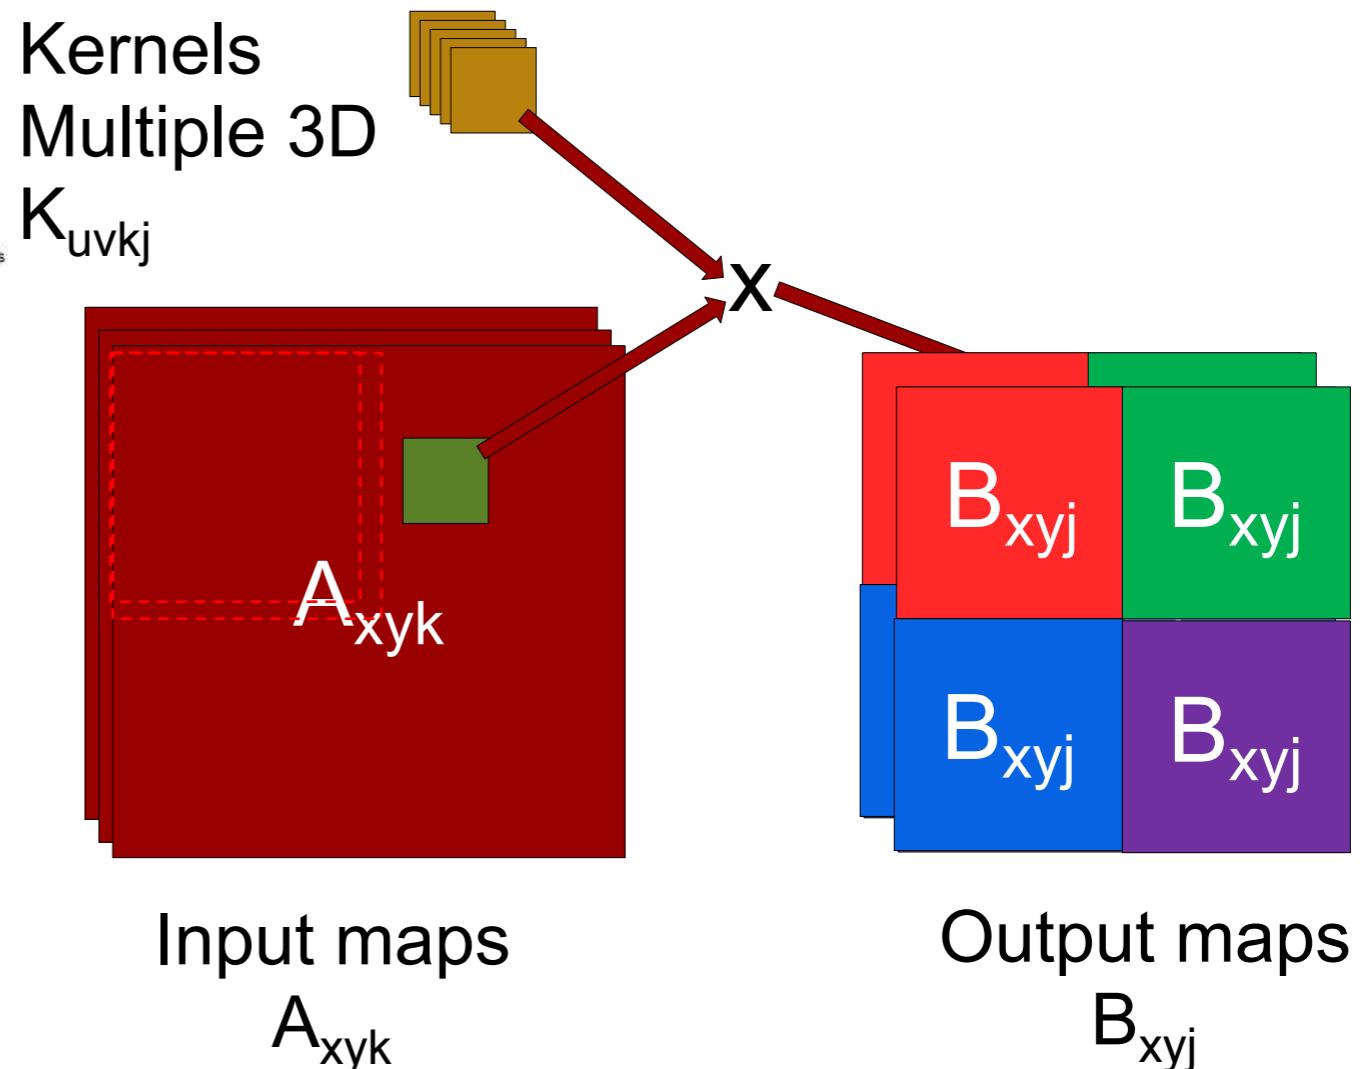
Dally, High Performance Hardware for Machine Learning, NIPS'2015

Stanford University

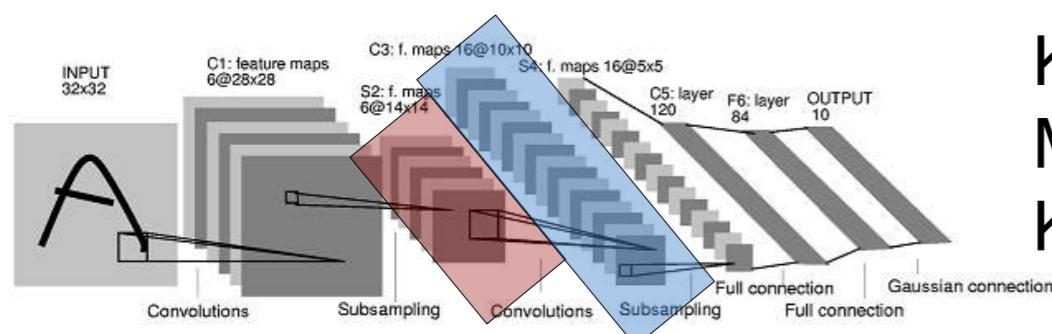
# Model-Parallel Convolution – by output region (x,y)



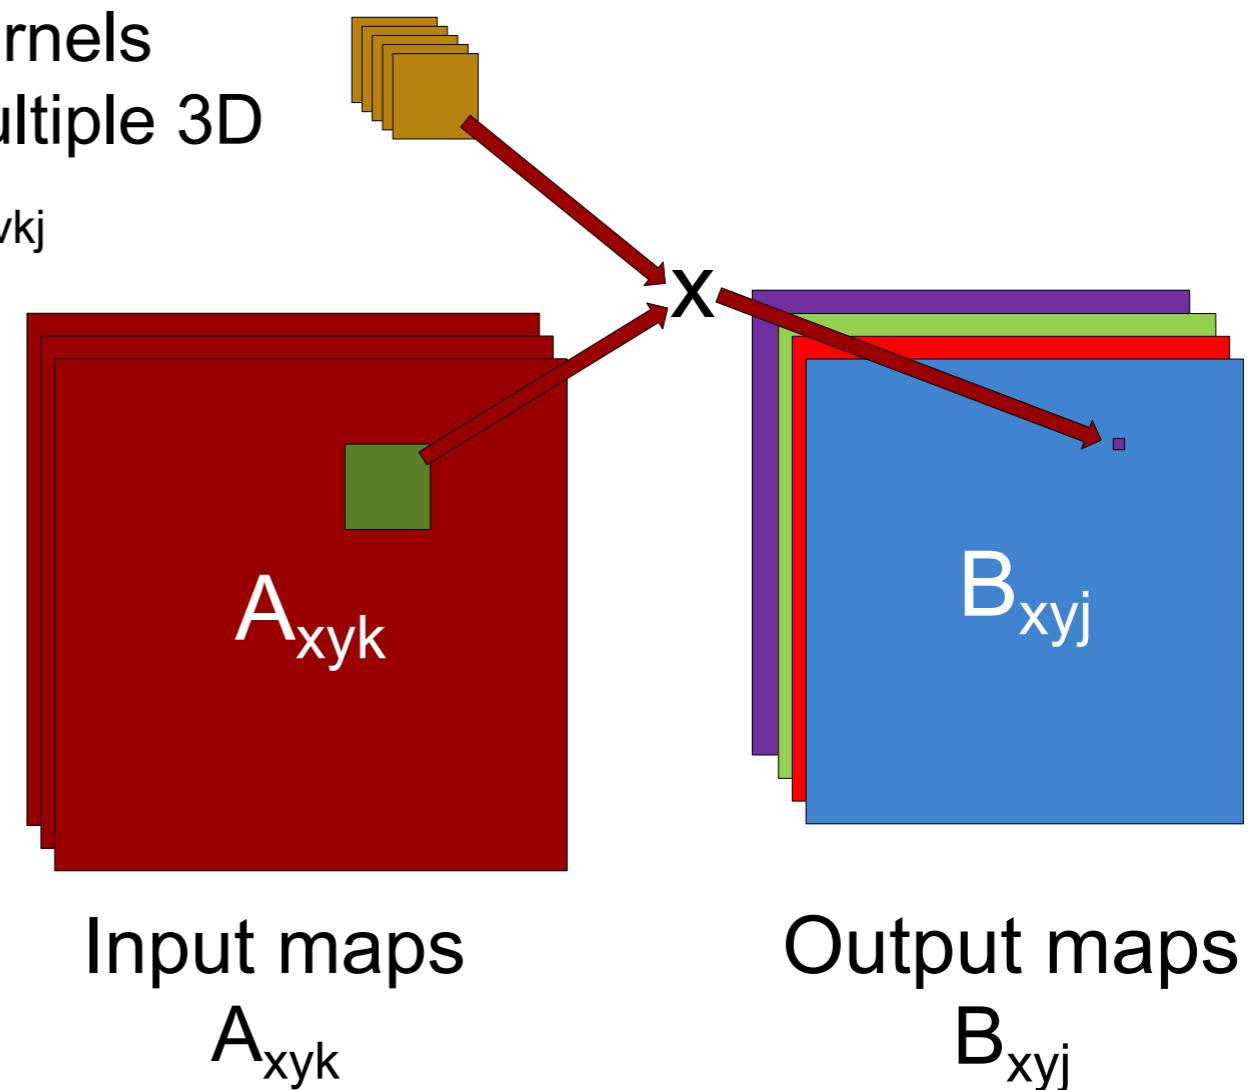
6D Loop  
For all output map  $j$   
For each input map  $k$   
For each pixel  $x, y$   
For each kernel element  $u, v$   
 $B_{xyj} += A_{(x-u)(y-v)k} \times K_{uvkj}$



# Model-Parallel Convolution – By output map j (filter)

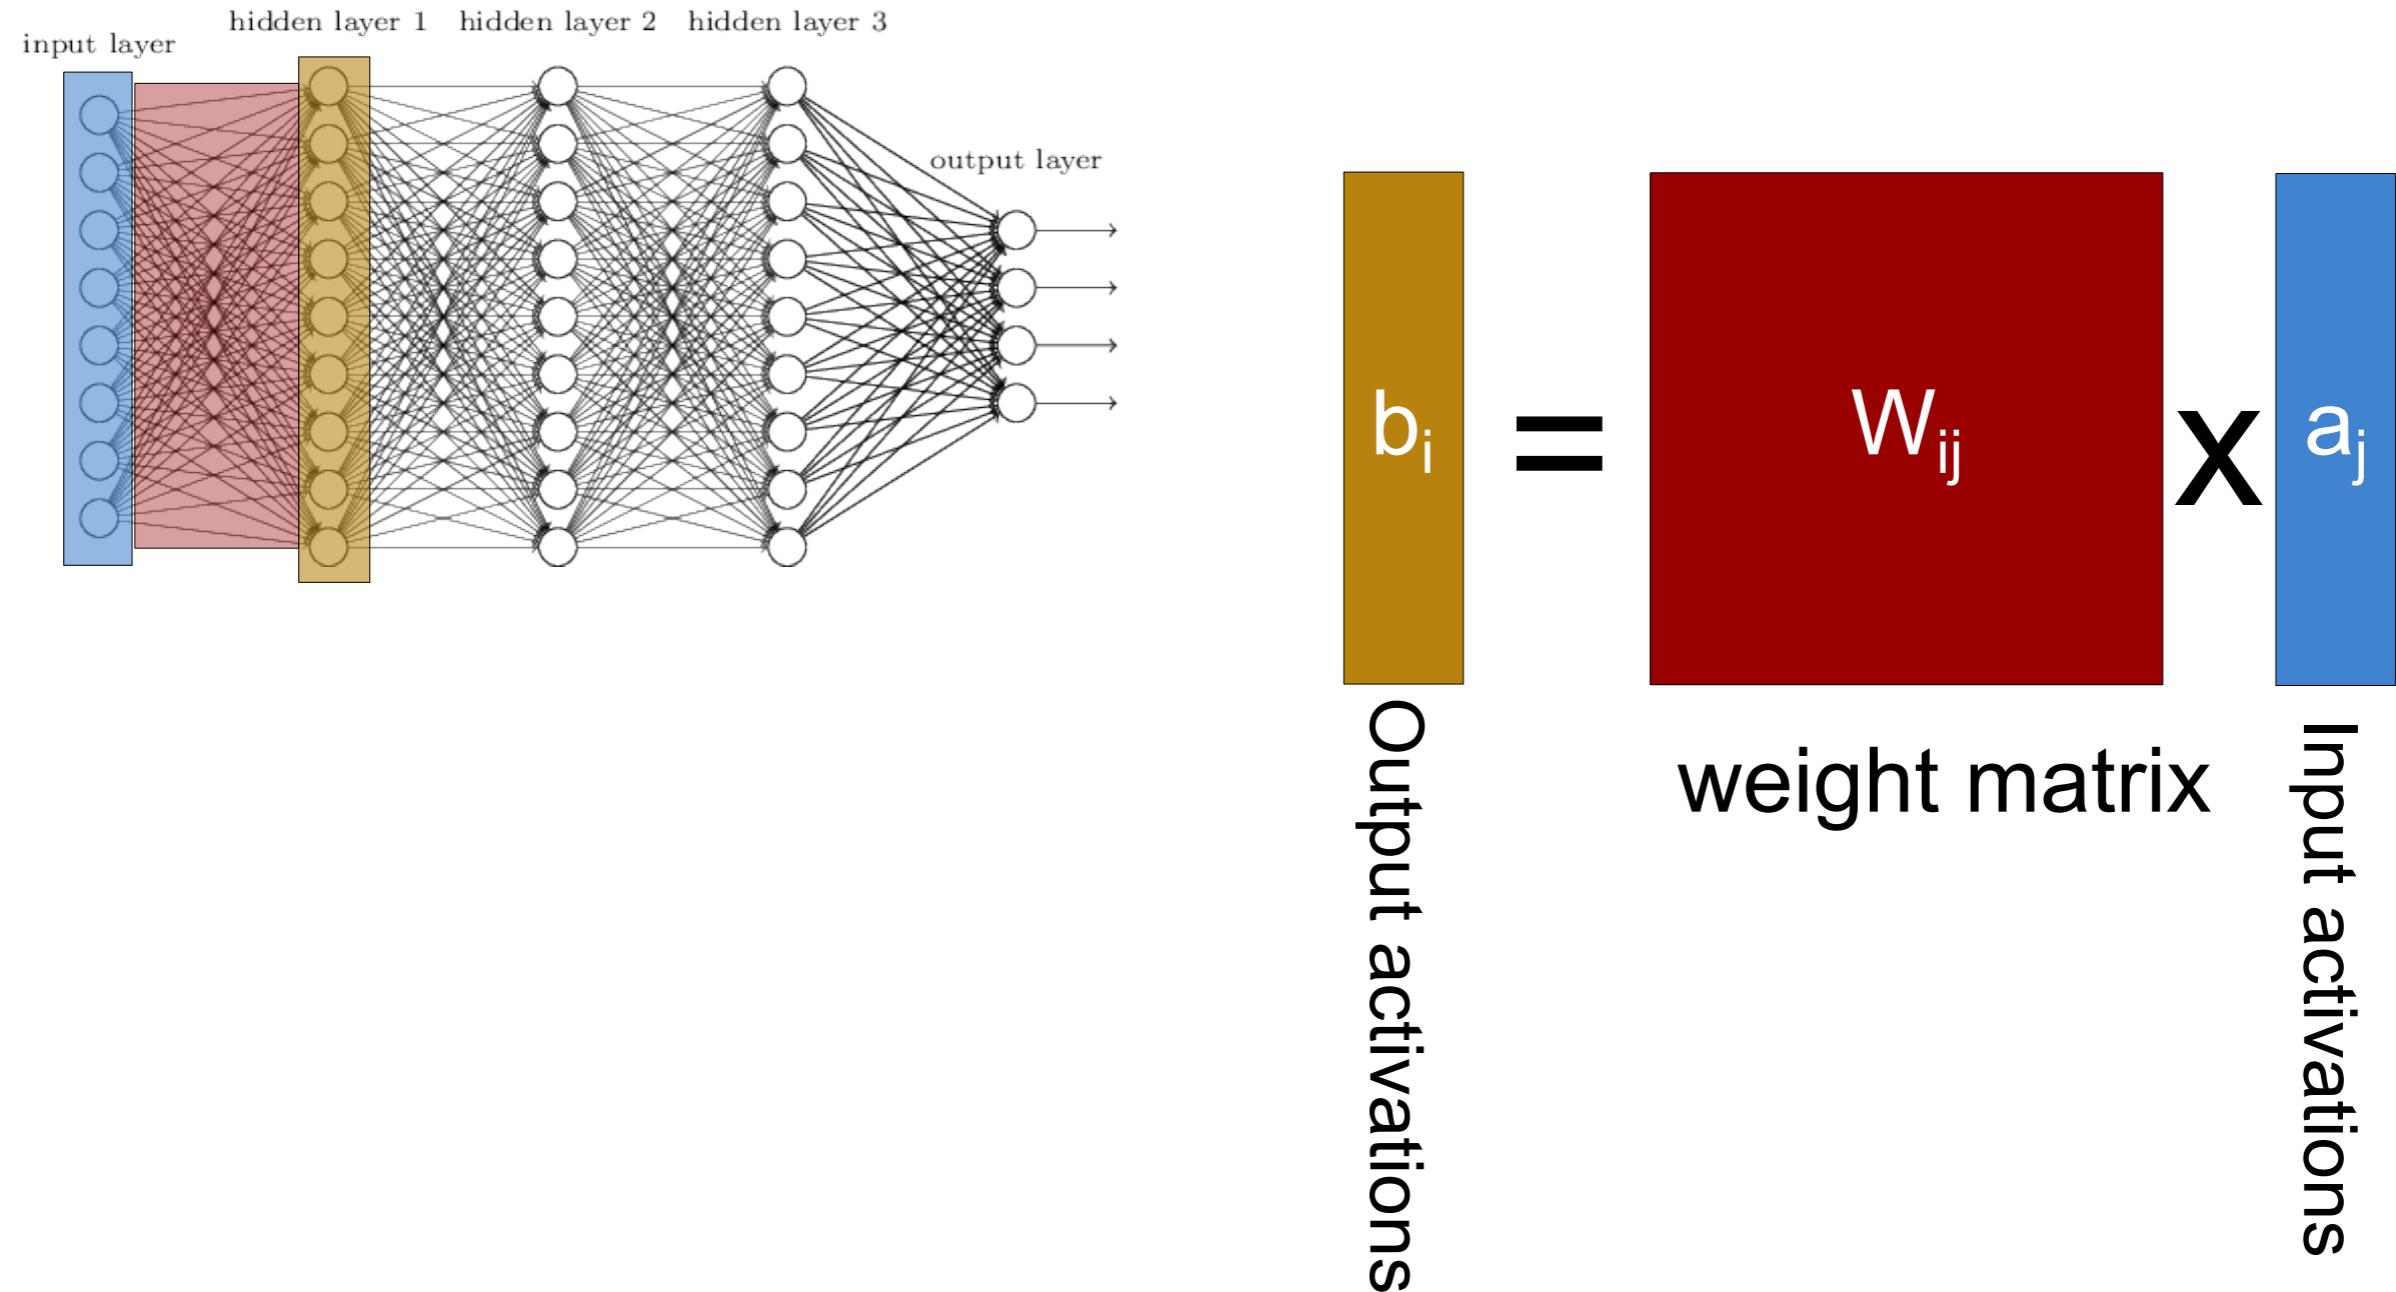


Kernels  
Multiple 3D  
 $K_{uvkj}$



6D Loop  
**Forall** output map  $j$   
For each input map  $k$   
For each pixel  $x,y$   
For each kernel element  $u,v$   
$$B_{xyj} += A_{(x-u)(y-v)k} \times K_{uvkj}$$

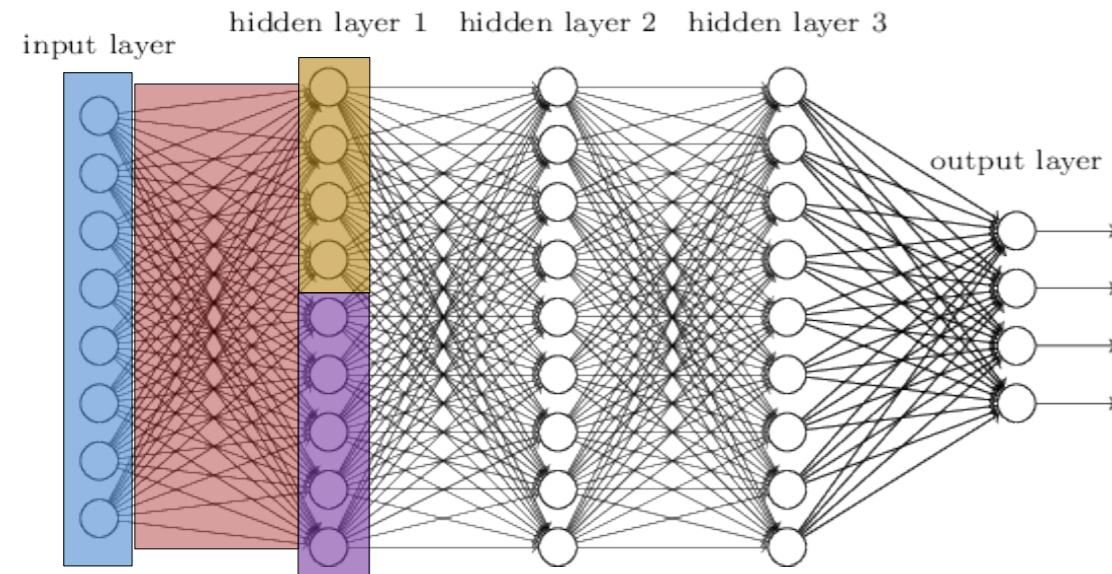
# Model Parallel Fully-Connected Layer ( $M \times V$ )



Dally, High Performance Hardware for Machine Learning, NIPS'2015

Stanford University

# Model Parallel Fully-Connected Layer ( $M \times V$ )



$$b_i = W_{ij} \text{ weight matrix } X a_j$$

Output activations

Input activations

# **Hyper-Parameter Parallel**

**Try many alternative networks in parallel**

Dally, High Performance Hardware for Machine Learning, NIPS'2015

Stanford University

# Summary of Parallelism

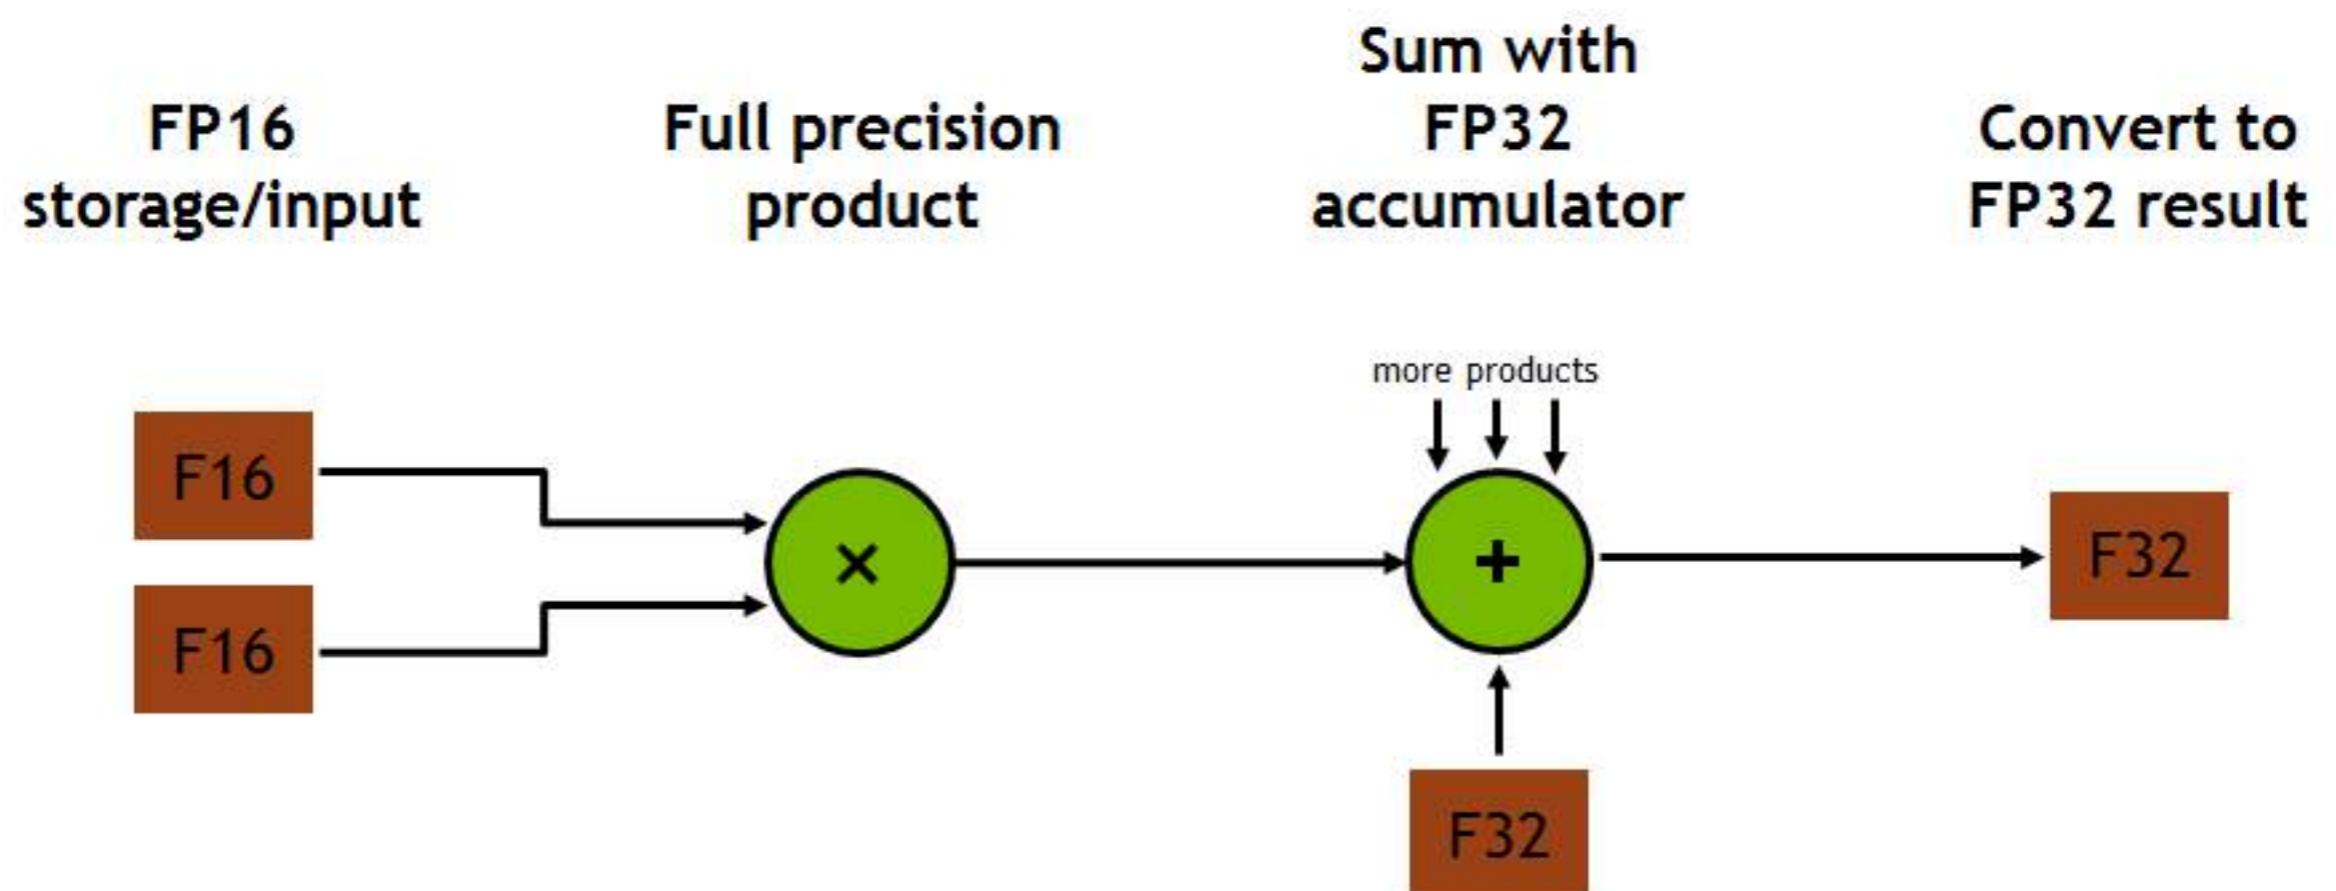
- Lots of parallelism in DNNs
  - 16M independent multiplies in one FC layer
  - Limited by overhead to exploit a fraction of this
- Data parallel
  - Run multiple training examples in parallel
  - Limited by batch size
- Model parallel
  - Split model over multiple processors
  - By layer
  - Conv layers by map region
  - Fully connected layers by output activation
- Easy to get 16-64 GPUs training one model in parallel

Dally, High Performance Hardware for Machine Learning, NIPS'2015

# Part 3: Efficient Training — Algorithms

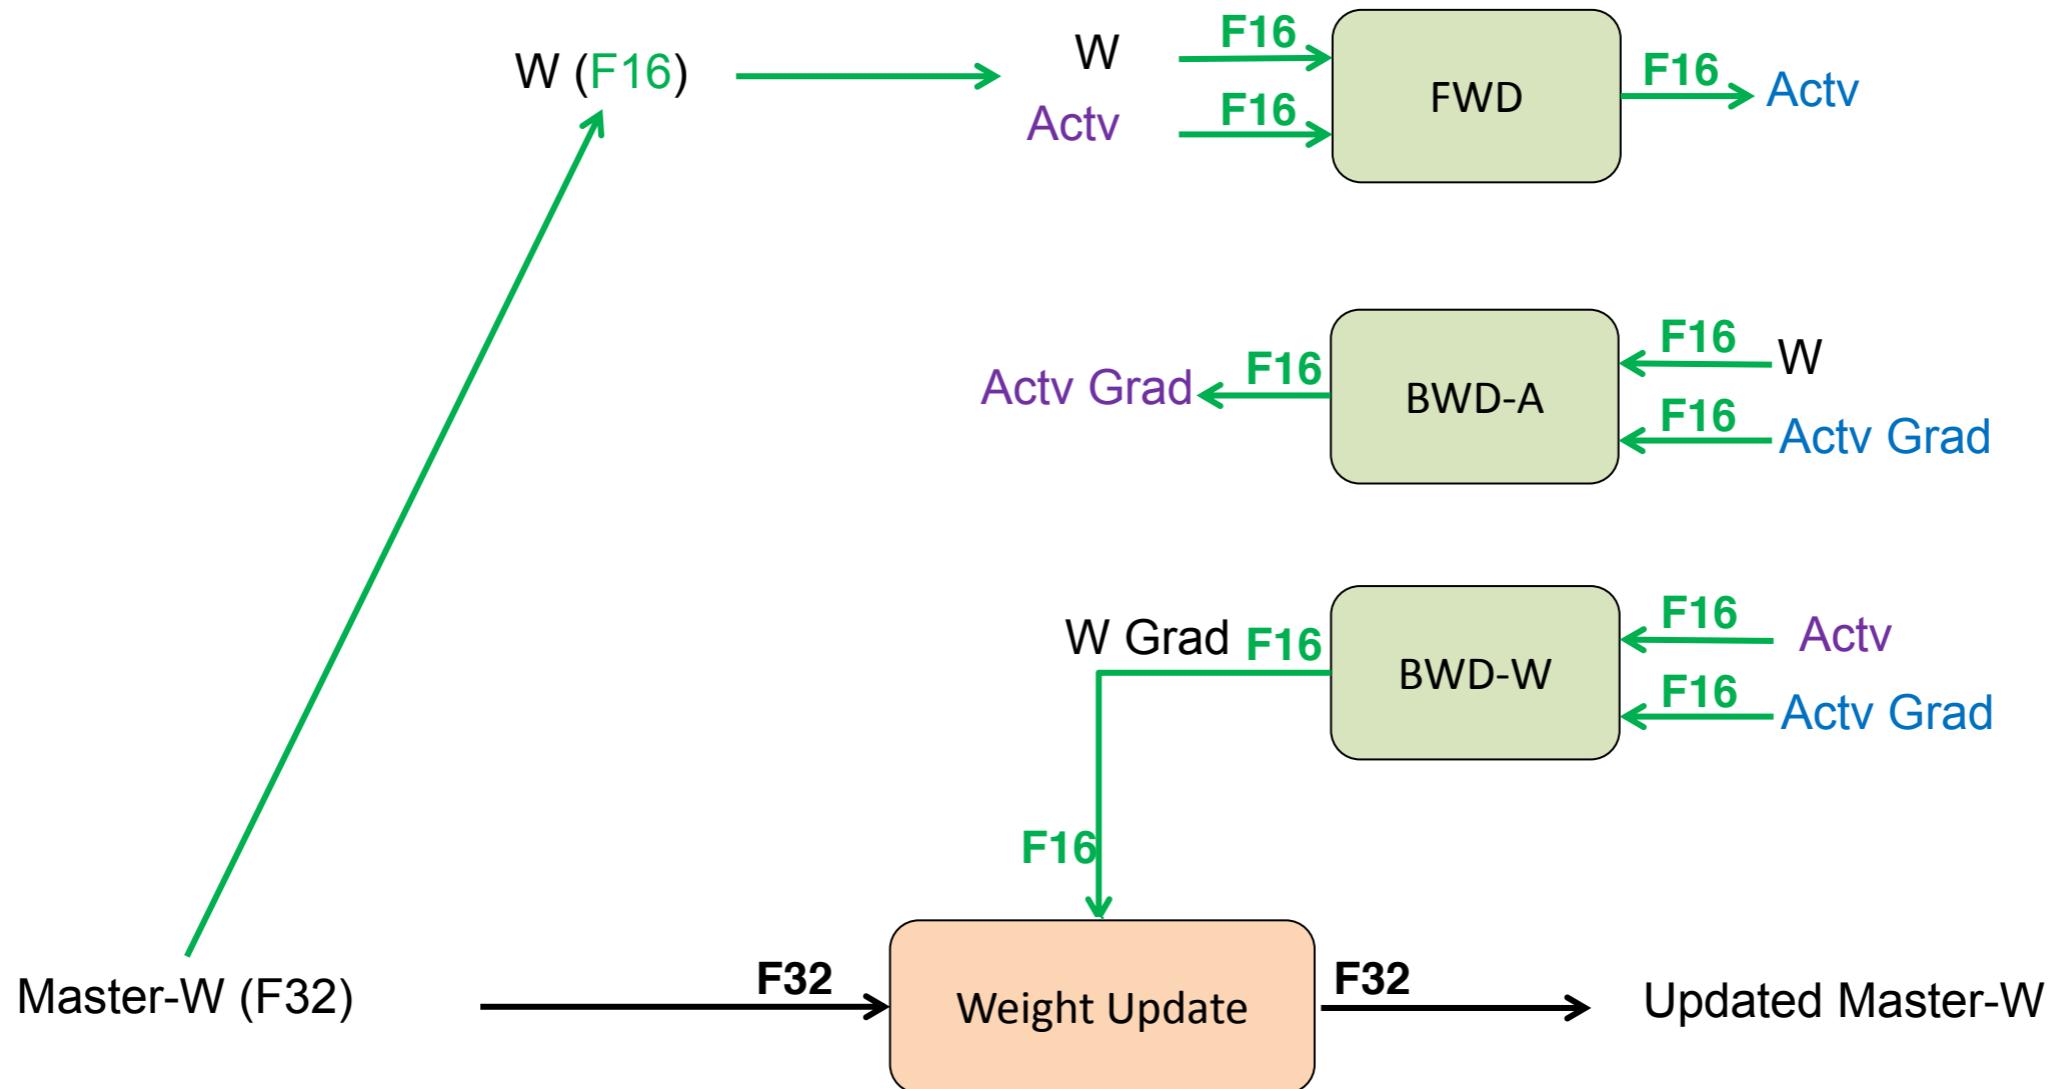
- 1. Parallelization
- 2. Mixed Precision with FP16 and FP32
- 3. Model Distillation
- 4. DSD: Dense-Sparse-Dense Training

# Mixed Precision



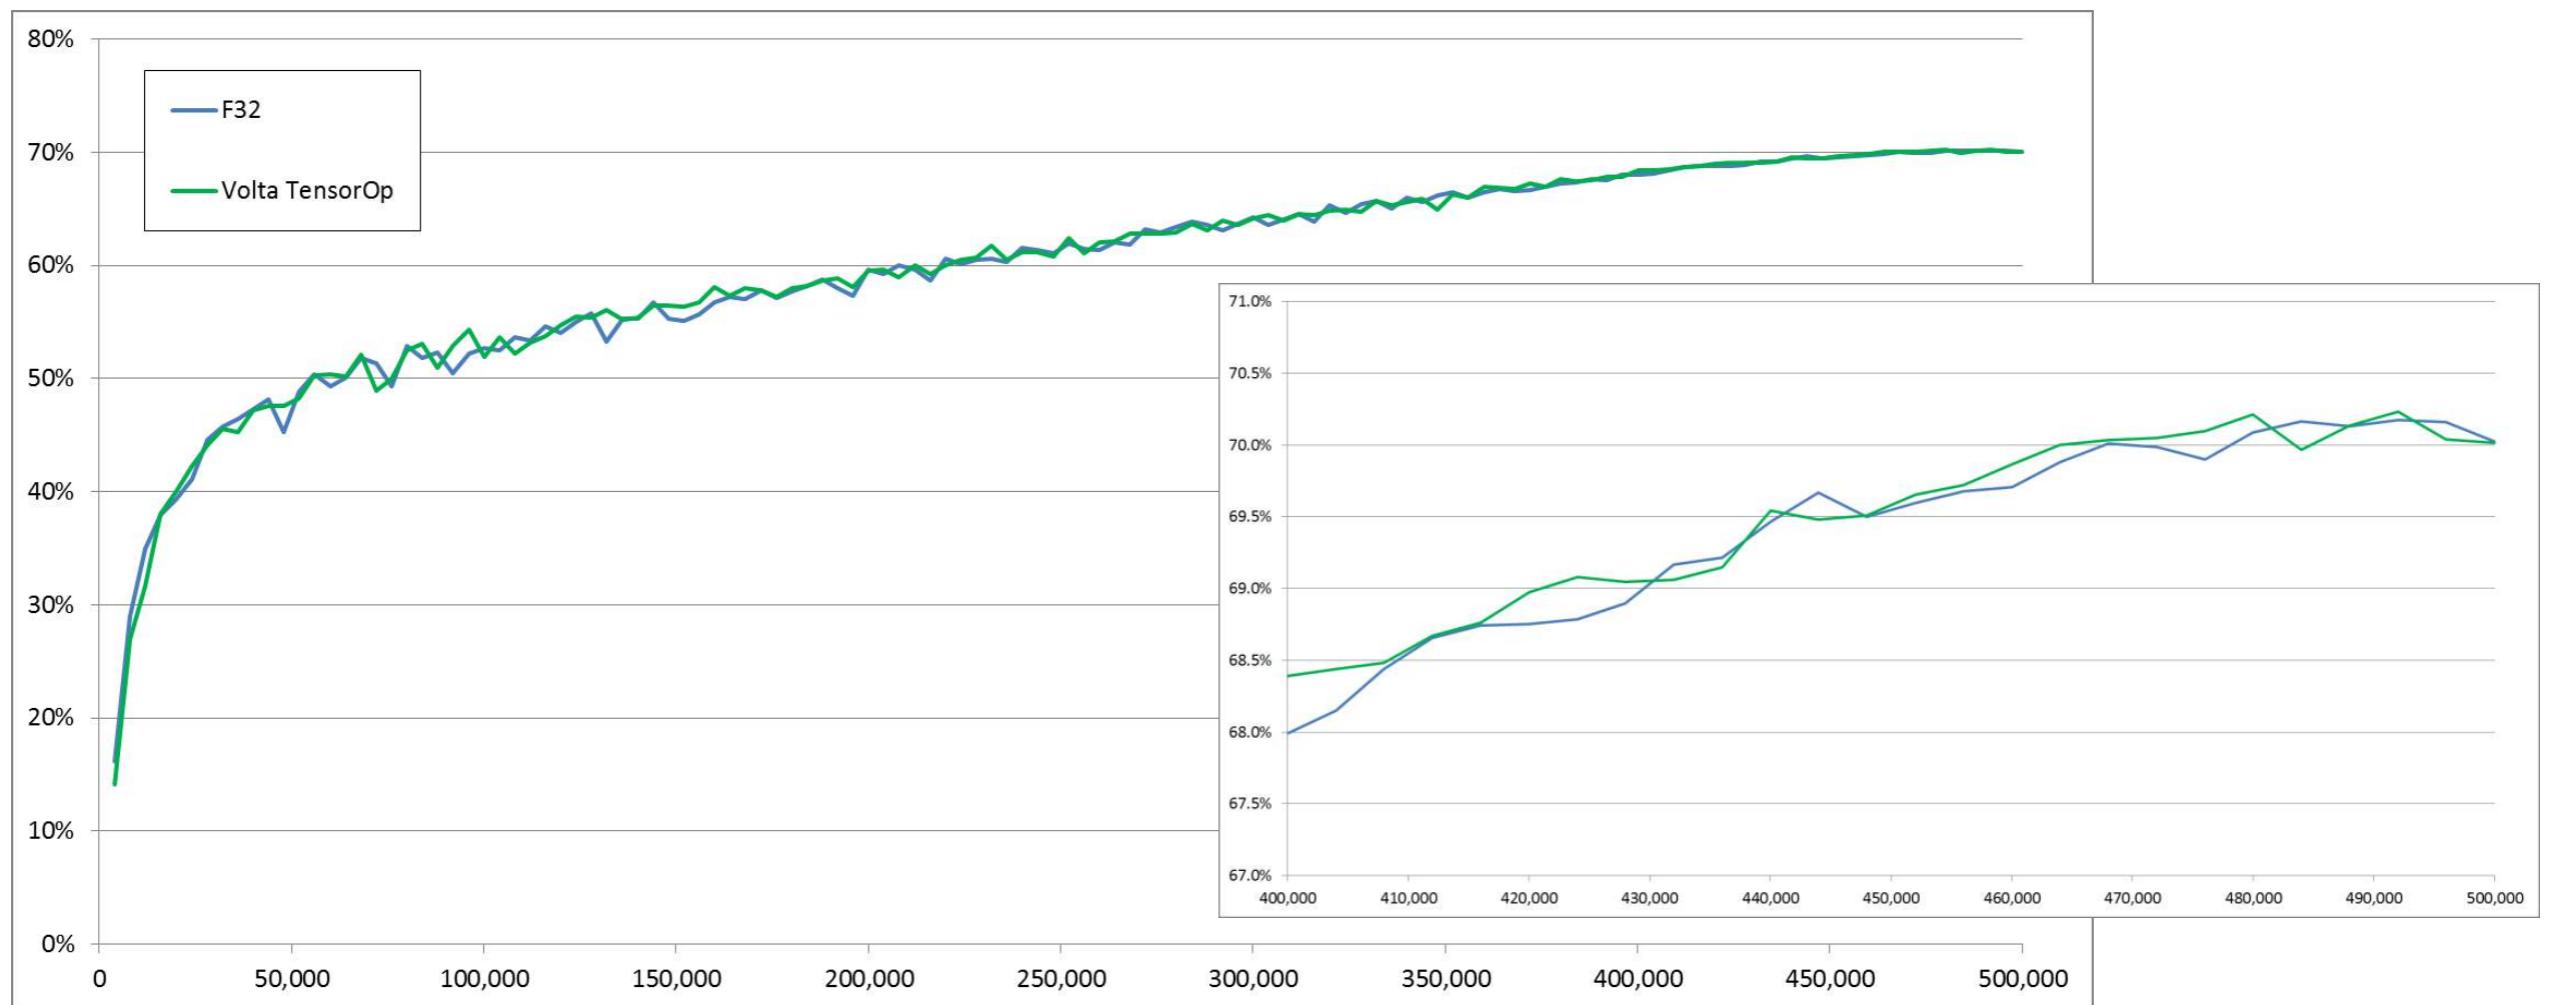
<https://devblogs.nvidia.com/parallelforall/cuda-9-features-revealed/>

# Mixed Precision Training



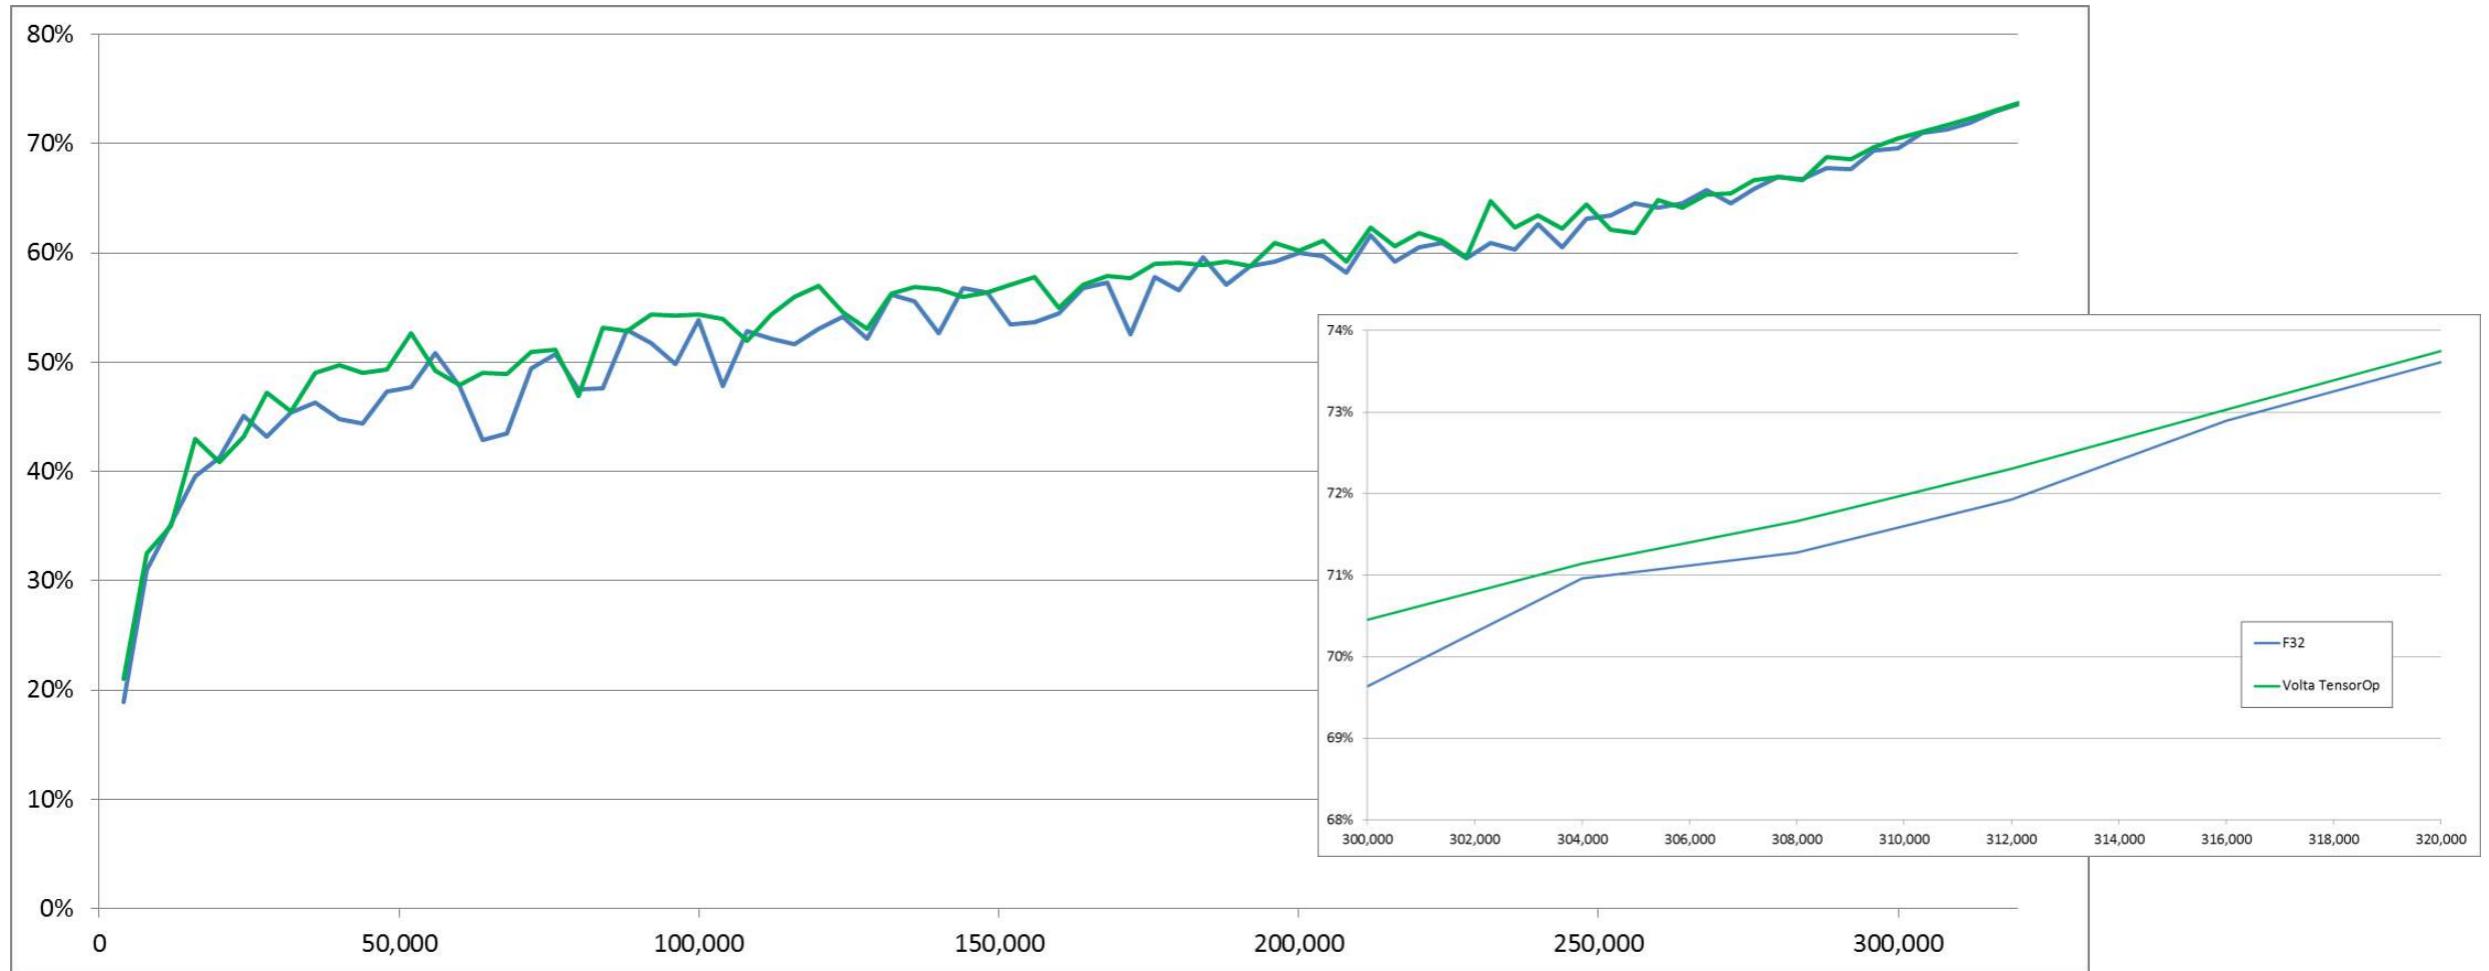
Boris Ginsburg, Sergei Nikolaev, Paulius Micikevicius, “Training with mixed precision”, NVIDIA GTC 2017

# Inception V1



Boris Ginsburg, Sergei Nikolaev, Paulius Micikevicius, “Training with mixed precision”, NVIDIA GTC 2017

# ResNet



Boris Ginsburg, Sergei Nikolaev, Paulius Micikevicius, “Training with mixed precision”, NVIDIA GTC 2017

## AlexNet

Mode	Top1 accuracy, %	Top5 accuracy, %
Fp32	58.62	81.25
Mixed precision training	58.12	80.71

## Inception V3

Mode	Top1 accuracy, %	Top5 accuracy, %
Fp32	71.75	90.52
Mixed precision training	71.17	90.10

## ResNet-50

Mode	Top1 accuracy, %	Top5 accuracy, %
Fp32	73.85	91.44
Mixed precision training	73.6	91.11

Boris Ginsburg, Sergei Nikolaev, Paulius Micikevicius, “Training with mixed precision”, NVIDIA GTC 2017

# Part 3: Efficient Training Algorithm

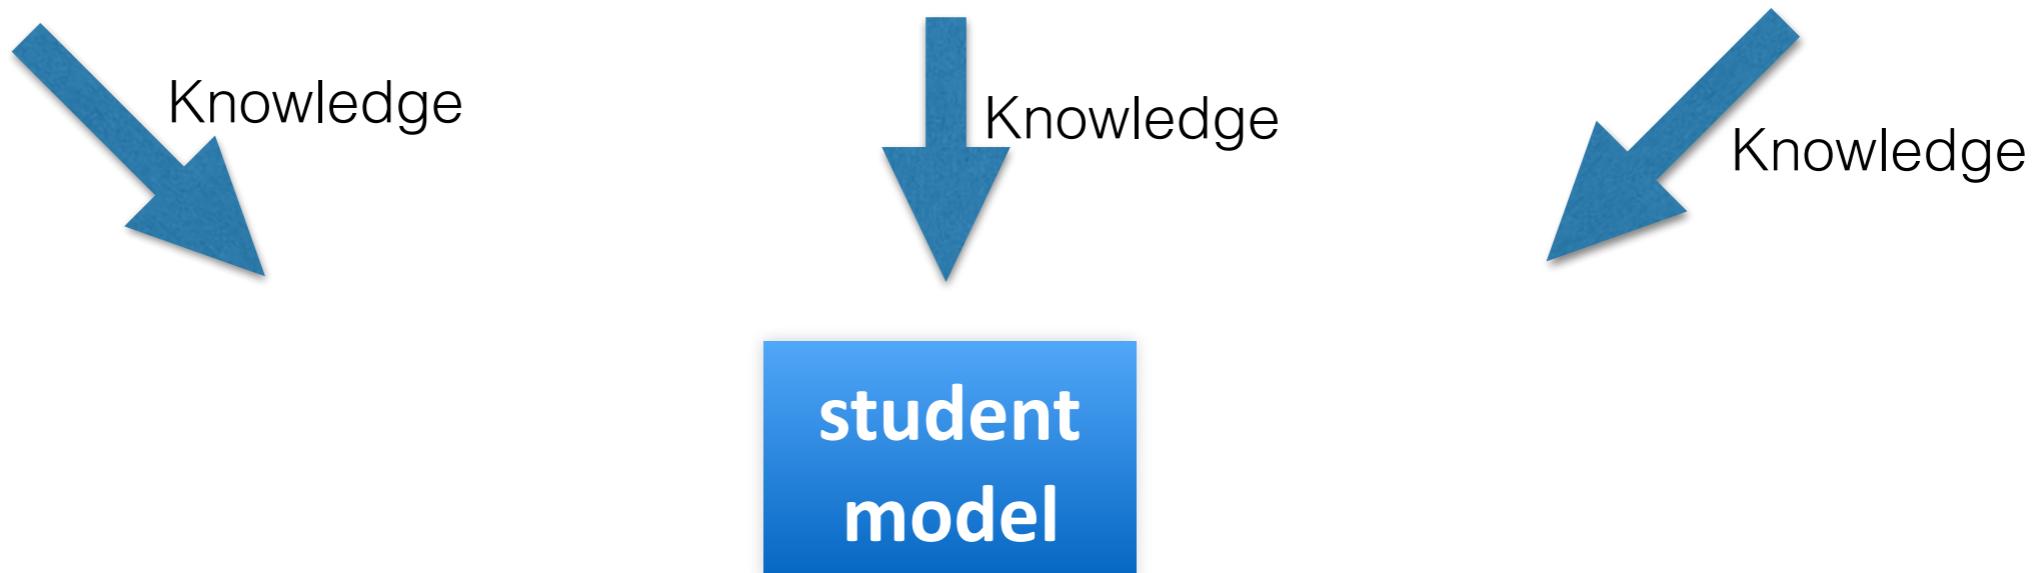
- 1. Parallelization
- 2. Mixed Precision with FP16 and FP32
- 3. Model Distillation
- 4. DSD: Dense-Sparse-Dense Training

# Model Distillation

Teacher model 1  
(Googlenet)

Teacher model 2  
(Vggnet)

Teacher model 3  
(Resnet)



student model has much smaller model size

# Softened outputs reveal the dark knowledge

cow	dog	cat	car
0	1	0	0

original hard  
targets

cow	dog	cat	car
$10^{-6}$	.9	.1	$10^{-9}$

output of  
geometric  
ensemble

cow	dog	cat	car
.05	.3	.2	.005

softened output  
of ensemble

# Softened outputs reveal the dark knowledge

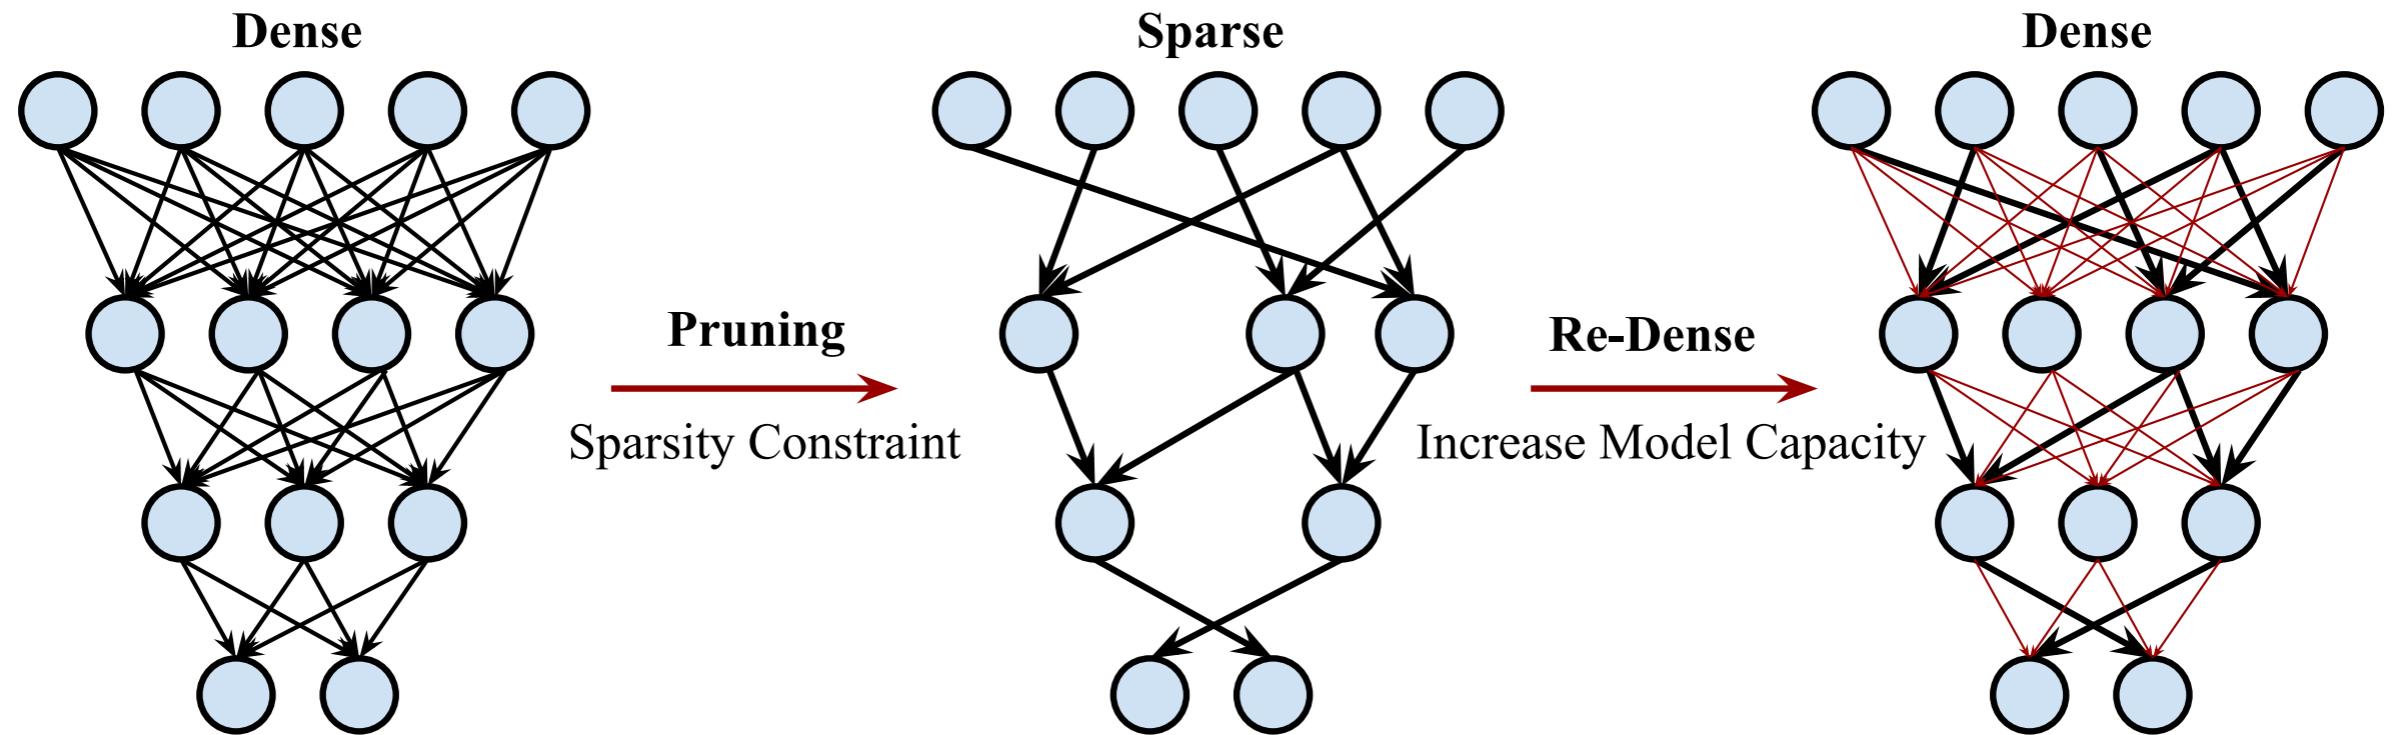
$$p_i = \frac{\exp\left(\frac{z_i}{T}\right)}{\sum_j \exp\left(\frac{z_j}{T}\right)}$$

- Method: Divide score by a “temperature” to get a much softer distribution
- Result: Start with a trained model that classifies **58.9%** of the test frames correctly. The new model converges to **57.0%** correct even when it is only trained on **3%** of the data

# Part 3: Efficient Training Algorithm

- 1. Parallelization
- 2. Mixed Precision with FP16 and FP32
- 3. Model Distillation
- 4. DSD: Dense-Sparse-Dense Training

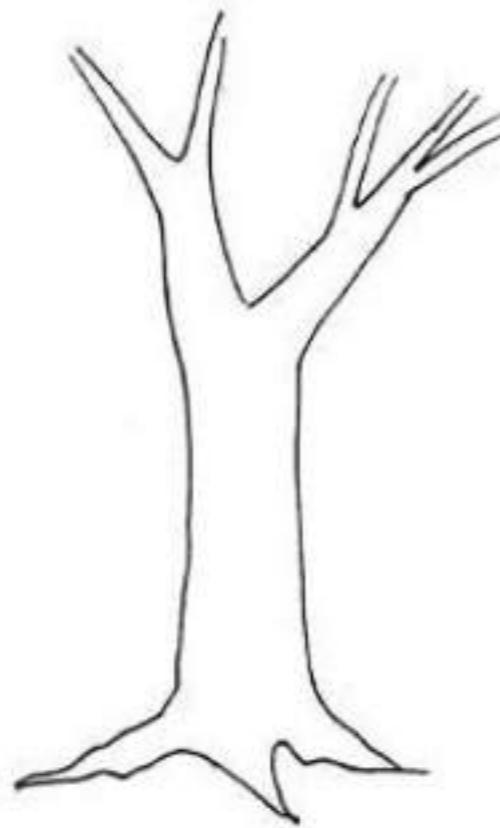
# DSD: Dense Sparse Dense Training



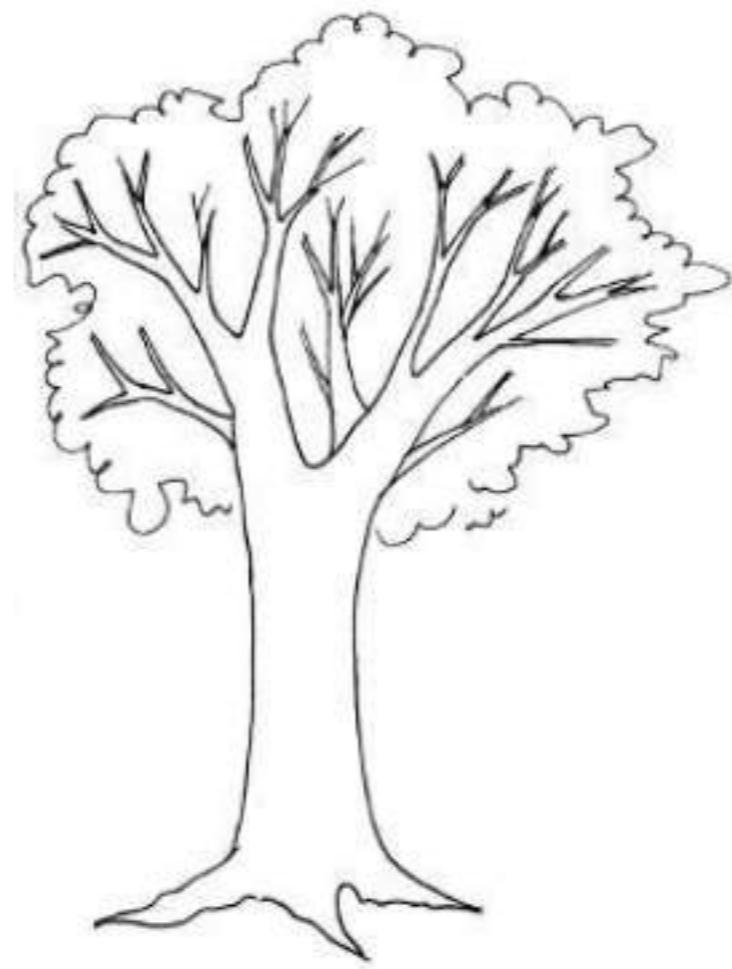
DSD produces same model architecture but can find better optimization solution, arrives at better local minima, and achieves higher prediction accuracy across a wide range of deep neural networks on CNNs / RNNs / LSTMs.

Han et al. "DSD: Dense-Sparse-Dense Training for Deep Neural Networks", ICLR 2017

# DSD: Intuition



learn the trunk first



then learn the leaves

Han et al. "DSD: Dense-Sparse-Dense Training for Deep Neural Networks", ICLR 2017

Stanford University

# DSD is General Purpose: Vision, Speech, Natural Language

Network	Domain	Dataset	Type	Baseline	DSD	Abs. Imp.	Rel. Imp.
GoogleNet	Vision	ImageNet	CNN	31.1%	→ <b>30.0%</b>	1.1%	3.6%
VGG-16	Vision	ImageNet	CNN	31.5%	→ <b>27.2%</b>	4.3%	13.7%
ResNet-18	Vision	ImageNet	CNN	30.4%	→ <b>29.3%</b>	1.1%	3.7%
ResNet-50	Vision	ImageNet	CNN	24.0%	→ <b>23.2%</b>	0.9%	3.5%

Open Sourced DSD Model Zoo: <https://songhan.github.io/DSD>

The baseline results of AlexNet, VGG16, GoogleNet, SqueezeNet are from [Caffe Model Zoo](#). ResNet18, ResNet50 are from [fb.resnet.torch](#).

# DSD is General Purpose: Vision, Speech, Natural Language

Network	Domain	Dataset	Type	Baseline	DSD	Abs. Imp.	Rel. Imp.
GoogleNet	Vision	ImageNet	CNN	31.1%	→ 30.0%	1.1%	3.6%
VGG-16	Vision	ImageNet	CNN	31.5%	→ 27.2%	4.3%	13.7%
ResNet-18	Vision	ImageNet	CNN	30.4%	→ 29.3%	1.1%	3.7%
ResNet-50	Vision	ImageNet	CNN	24.0%	→ 23.2%	0.9%	3.5%
NeuralTalk	Caption	Flickr-8K	LSTM	16.8	→ 18.5	1.7	10.1%

Open Sourced DSD Model Zoo: <https://songhan.github.io/DSD>

The baseline results of AlexNet, VGG16, GoogleNet, SqueezeNet are from [Caffe Model Zoo](#). ResNet18, ResNet50 are from [fb.resnet.torch](#).

# DSD is General Purpose: Vision, Speech, Natural Language

Network	Domain	Dataset	Type	Baseline	DSD	Abs. Imp.	Rel. Imp.
GoogleNet	Vision	ImageNet	CNN	31.1%	→ <b>30.0%</b>	1.1%	3.6%
VGG-16	Vision	ImageNet	CNN	31.5%	→ <b>27.2%</b>	4.3%	13.7%
ResNet-18	Vision	ImageNet	CNN	30.4%	→ <b>29.3%</b>	1.1%	3.7%
ResNet-50	Vision	ImageNet	CNN	24.0%	→ <b>23.2%</b>	0.9%	3.5%
NeuralTalk	Caption	Flickr-8K	LSTM	16.8	→ <b>18.5</b>	1.7	10.1%
DeepSpeech	Speech	WSJ'93	RNN	33.6%	→ <b>31.6%</b>	2.0%	5.8%
DeepSpeech-2	Speech	WSJ'93	RNN	14.5%	→ <b>13.4%</b>	1.1%	7.4%

Open Sourced DSD Model Zoo: <https://songhan.github.io/DSD>

The baseline results of AlexNet, VGG16, GoogleNet, SqueezeNet are from [Caffe Model Zoo](#). ResNet18, ResNet50 are from [fb.resnet.torch](#).

# *DSD Model Zoo*

DSD model zoo. Better accuracy models from DSD training on Imagenet with same model architecture.

[View DSD Model Zoo on GitHub](#)

 Download

 Download

## DSD Model Zoo

This repo contains pre-trained models by Dense-Sparse-Dense(DSD) training on Imagenet.

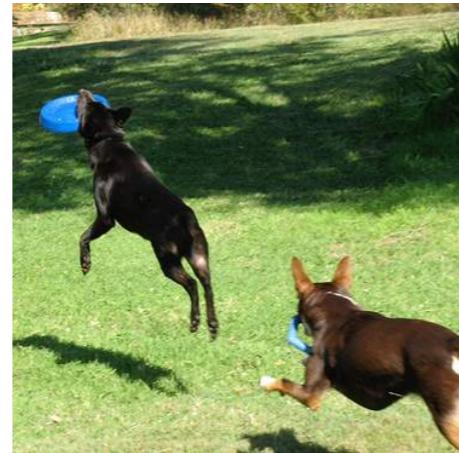
Compared to conventional training method, dense→sparse→dense (DSD) training yielded higher accuracy with same model architecture.

Sparsity is a powerful form of regularization. Our intuition is that, once the network arrives at a local minimum given the sparsity constraint, relaxing the constraint gives the network more freedom to escape the saddle point and arrive at a higher-accuracy local minimum.

### Download:

<https://songhan.github.io/DSD>

# DSD on Caption Generation



✗ **Baseline**: a boy in a red shirt is climbing a rock wall.

✗ **Sparse**: a young girl is jumping off a tree.

✓ **DSD**: a young girl in a pink shirt is swinging on a swing.

○ **Baseline**: a basketball player in a red uniform is playing with a ball.

○ **Sparse**: a basketball player in a blue uniform is jumping over the goal.

✓ **DSD**: a basketball player in a white uniform is trying to make a shot.

✓ **Baseline**: two dogs are playing together in a field.

✓ **Sparse**: two dogs are playing in a field.

✓ **DSD**: two dogs are playing in the grass.

✗ **Baseline**: a man and a woman are sitting on a bench.

○ **Sparse**: a man is sitting on a bench with his hands in the air.

○ **DSD**: a man is sitting on a bench with his arms folded.

✗ **Baseline**: a person in a red jacket is riding a bike through the woods.

○ **Sparse**: a car drives through a mud puddle.

✓ **DSD**: a car drives through a forest.

Baseline model: Andrej Karpathy, [Neural Talk model zoo](#).

Han et al. "DSD: Dense-Sparse-Dense Training for Deep Neural Networks", ICLR 2017

# DSD on Caption Generation



✗ **Baseline:** a boy is swimming in a pool.  
○ **Sparse:** a small black dog is jumping into a pool.  
✓ **DSD:** a black and white dog is swimming in a pool.

✗ **Baseline:** a group of people are standing in front of a building.  
○ **Sparse:** a group of people are standing in front of a building.  
✓ **DSD:** a group of people are walking in a park.

✗ **Baseline:** two girls in bathing suits are playing in the water.  
✓ **Sparse:** two children are playing in the sand.  
✓ **DSD:** two children are playing in the sand.

○ **Baseline:** a man in a red shirt and jeans is riding a bicycle down a street.  
○ **Sparse:** a man in a red shirt and a woman in a wheelchair.  
✓ **DSD:** a man and a woman are riding on a street.



✗ **Baseline:** a group of people sit on a bench in front of a building.  
○ **Sparse:** a group of people are standing in front of a building.  
✓ **DSD:** a group of people are standing in a fountain.

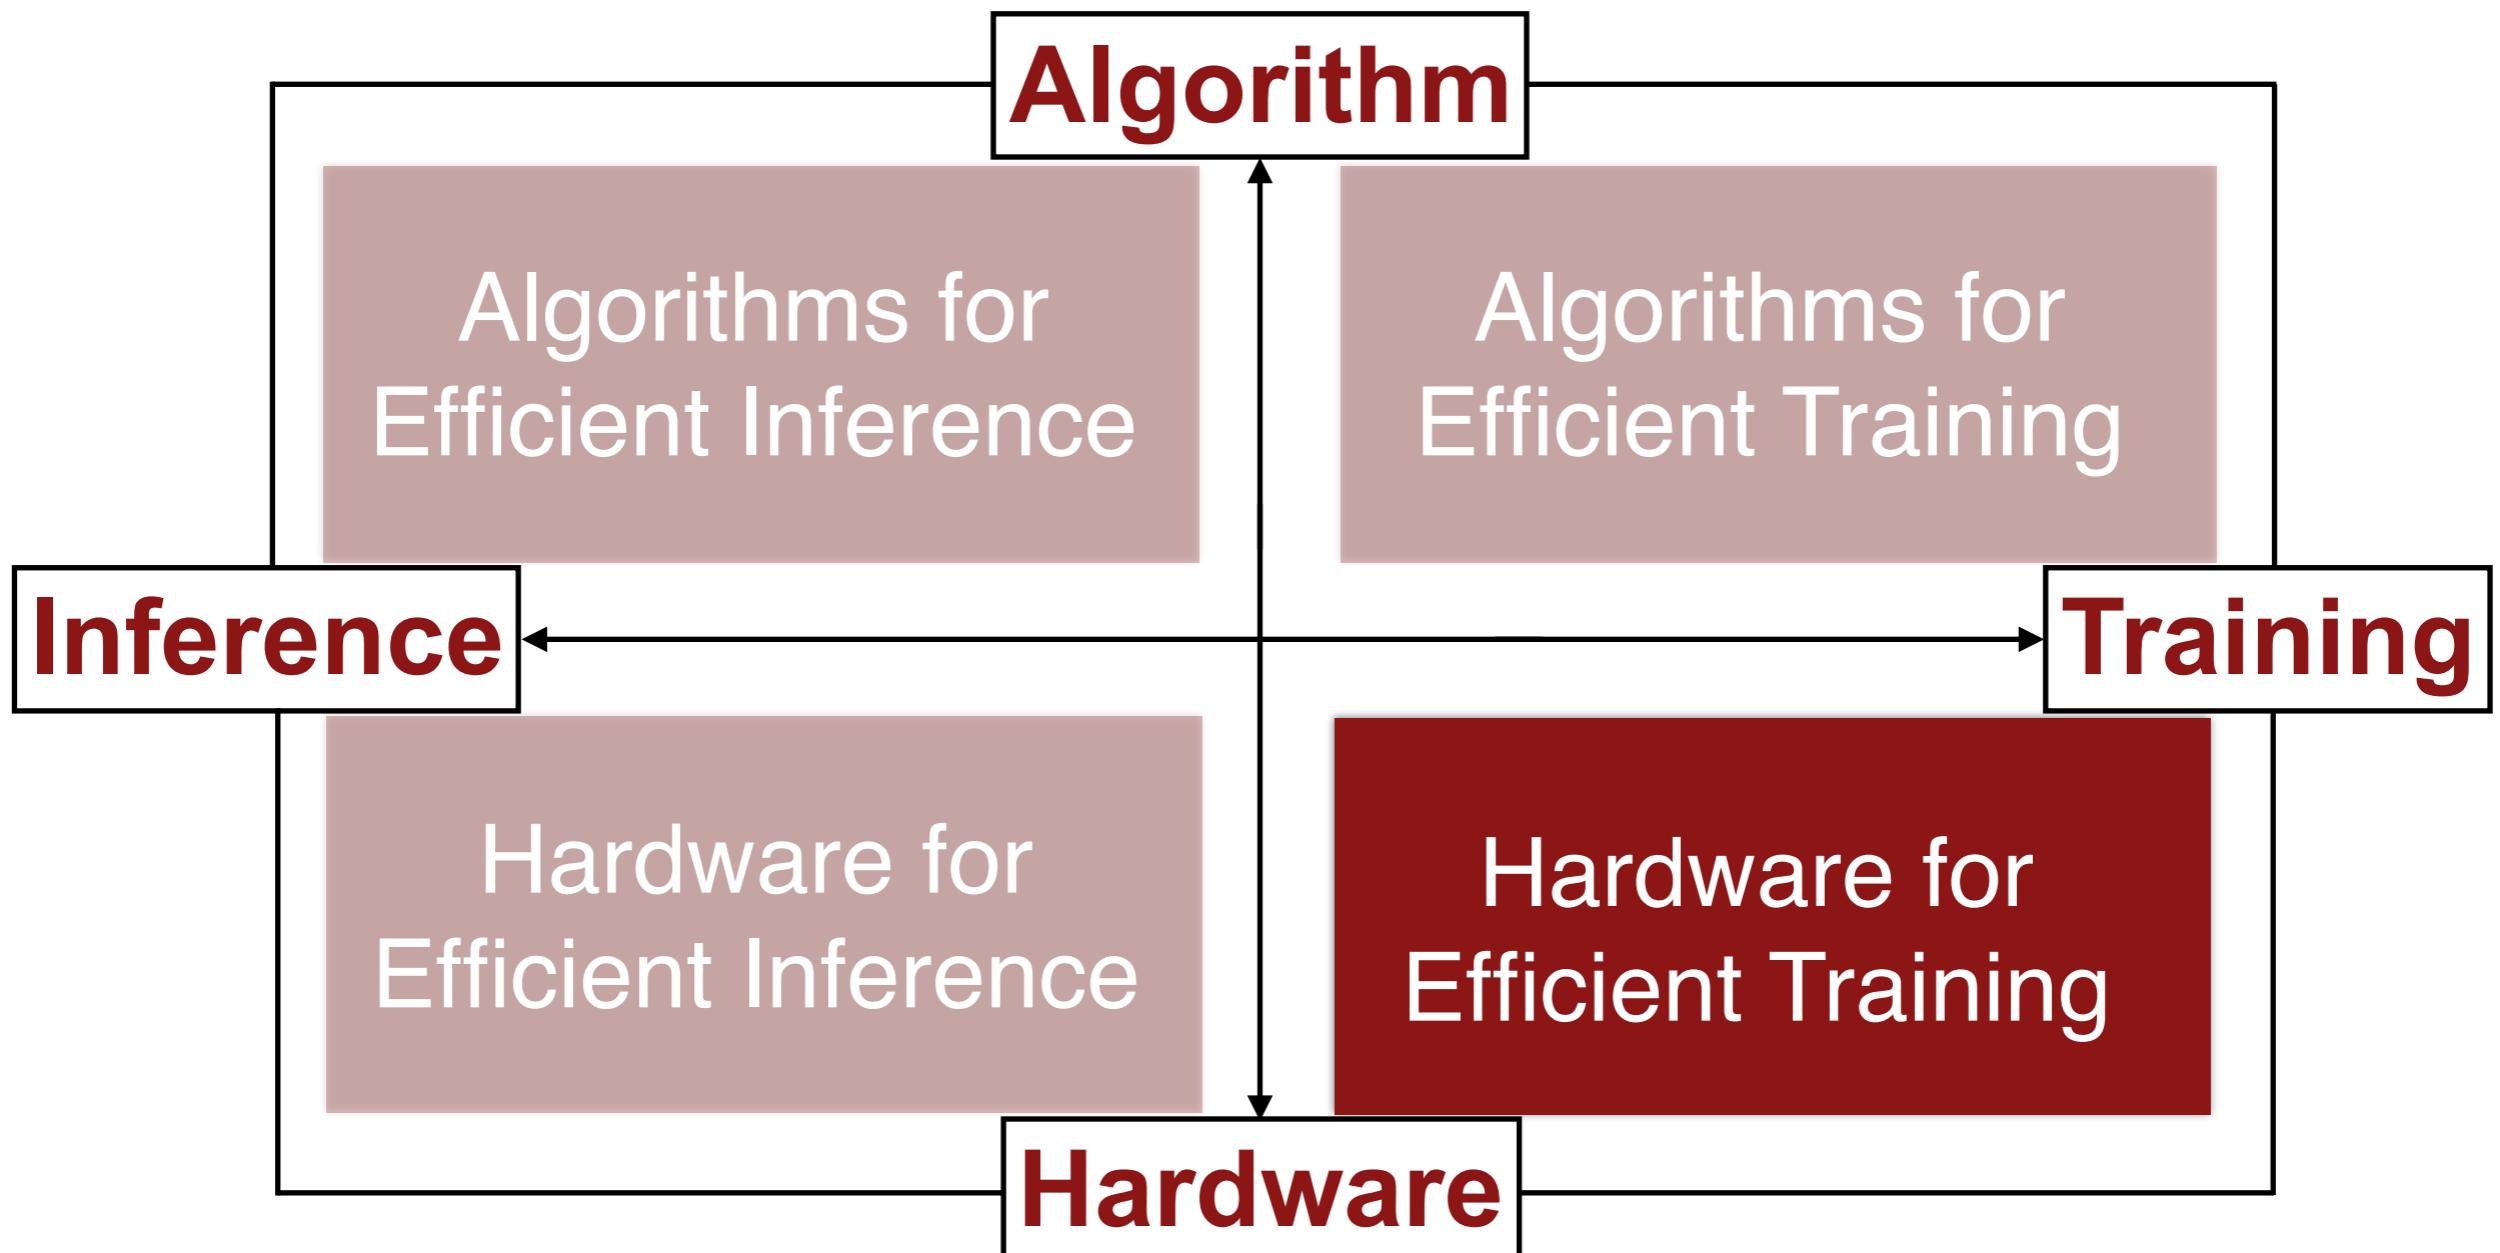
✗ **Baseline:** a man in a black jacket and a black jacket is smiling.  
○ **Sparse:** a man and a woman are standing in front of a mountain.  
✓ **DSD:** a man in a black jacket is standing next to a man in a black shirt.

○ **Baseline:** a group of football players in red uniforms.  
○ **Sparse:** a group of football players in a field.  
✓ **DSD:** a group of football players in red and white uniforms.

○ **Baseline:** a dog runs through the grass.  
○ **Sparse:** a dog runs through the grass.  
✓ **DSD:** a white and brown dog is running through the grass.

Baseline model: Andrej Karpathy, [Neural Talk model zoo](#).

# Agenda



# CPUs for Training

## Intel Knights Landing (2016)



- 7 TFLOPS FP32
- 16GB MCDRAM – 400 GB/s
- 245W TDP
- 29 GFLOPS/W (FP32)
- 14nm process

**Knights Mill:** next gen Xeon Phi “optimized for deep learning”

Intel announced the addition of new vector instructions for deep learning (AVX512-4VNNIW and AVX512-4FMAPS), October 2016

Slide Source: Sze et al Survey of DNN Hardware, MICRO'16 Tutorial.  
Image Source: Intel, Data Source: Next Platform

# GPUs for Training

## Nvidia PASCAL GP100 (2016)

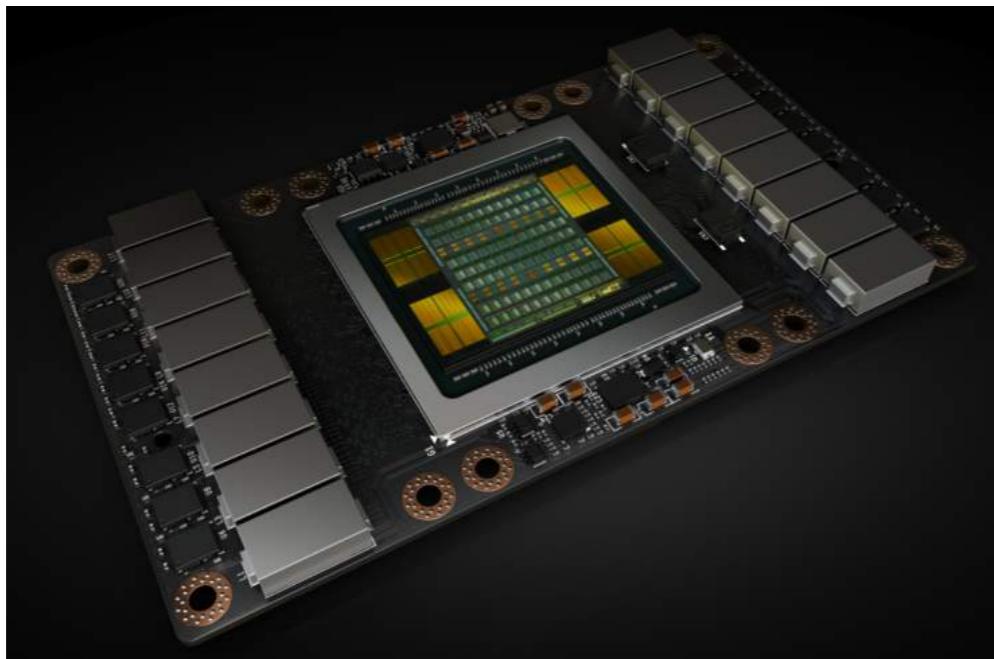


- 10/20 TFLOPS FP32/FP16
- 16GB HBM – 750 GB/s
- 300W TDP
- 67 GFLOPS/W (FP16)
- 16nm process
- 160GB/s NV Link

Slide Source: Sze et al Survey of DNN Hardware, MICRO'16 Tutorial.  
Data Source: NVIDIA

# GPUs for Training

## Nvidia Volta GV100 (2017)



- 15 FP32 TFLOPS
- 120 Tensor TFLOPS
- 16GB HBM2 @ 900GB/s
- 300W TDP
- 12nm process
- 21B Transistors
- die size: 815 mm<sup>2</sup>
- 300GB/s NVLink

Data Source: NVIDIA

# What's new in Volta: Tensor Core

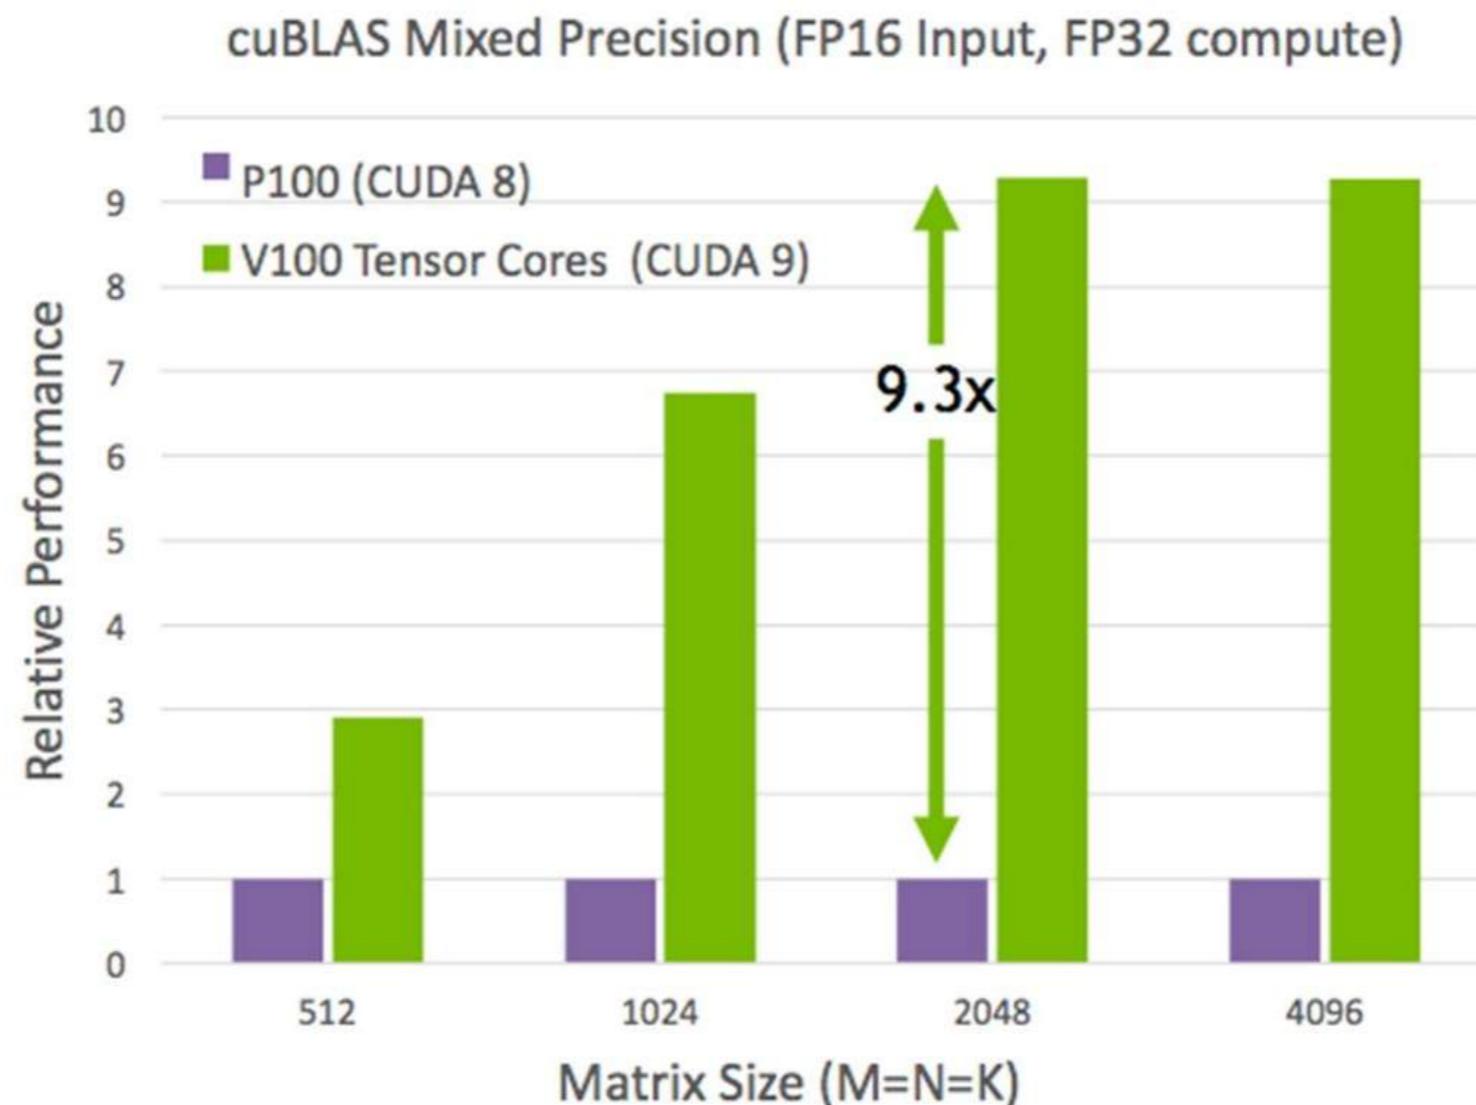
$$D = \left( \begin{array}{cccc} A_{0,0} & A_{0,1} & A_{0,2} & A_{0,3} \\ A_{1,0} & A_{1,1} & A_{1,2} & A_{1,3} \\ A_{2,0} & A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,0} & A_{3,1} & A_{3,2} & A_{3,3} \end{array} \right) \left( \begin{array}{cccc} B_{0,0} & B_{0,1} & B_{0,2} & B_{0,3} \\ B_{1,0} & B_{1,1} & B_{1,2} & B_{1,3} \\ B_{2,0} & B_{2,1} & B_{2,2} & B_{2,3} \\ B_{3,0} & B_{3,1} & B_{3,2} & B_{3,3} \end{array} \right) + \left( \begin{array}{cccc} C_{0,0} & C_{0,1} & C_{0,2} & C_{0,3} \\ C_{1,0} & C_{1,1} & C_{1,2} & C_{1,3} \\ C_{2,0} & C_{2,1} & C_{2,2} & C_{2,3} \\ C_{3,0} & C_{3,1} & C_{3,2} & C_{3,3} \end{array} \right)$$

FP16 or FP32                  FP16                  FP16 or FP32

a new instruction that performs 4x4x4 FMA mixed-precision operations per clock  
12X increase in throughput for the Volta V100 compared to the Pascal P100

<https://devblogs.nvidia.com/parallelforall/cuda-9-features-revealed/>

# Pascal v.s. Volta

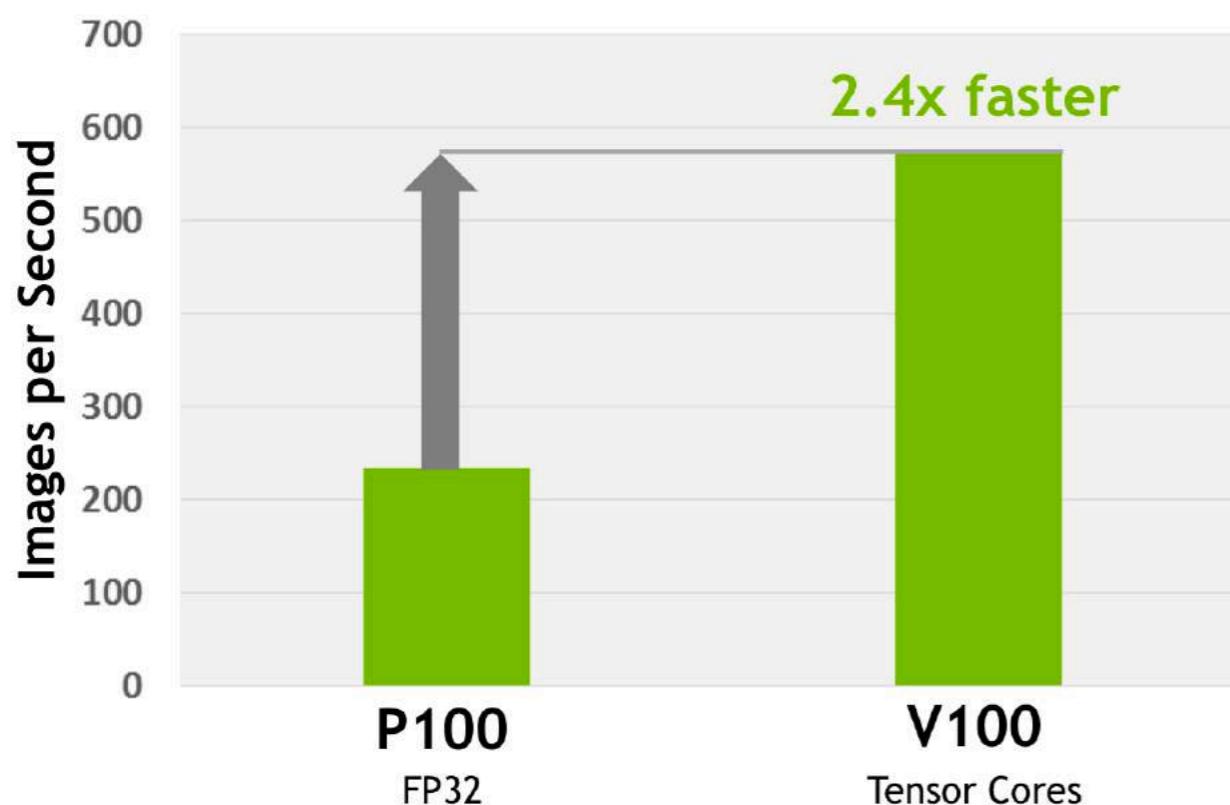


Tesla V100 Tensor Cores and CUDA 9 deliver up to 9x higher performance for GEMM operations.

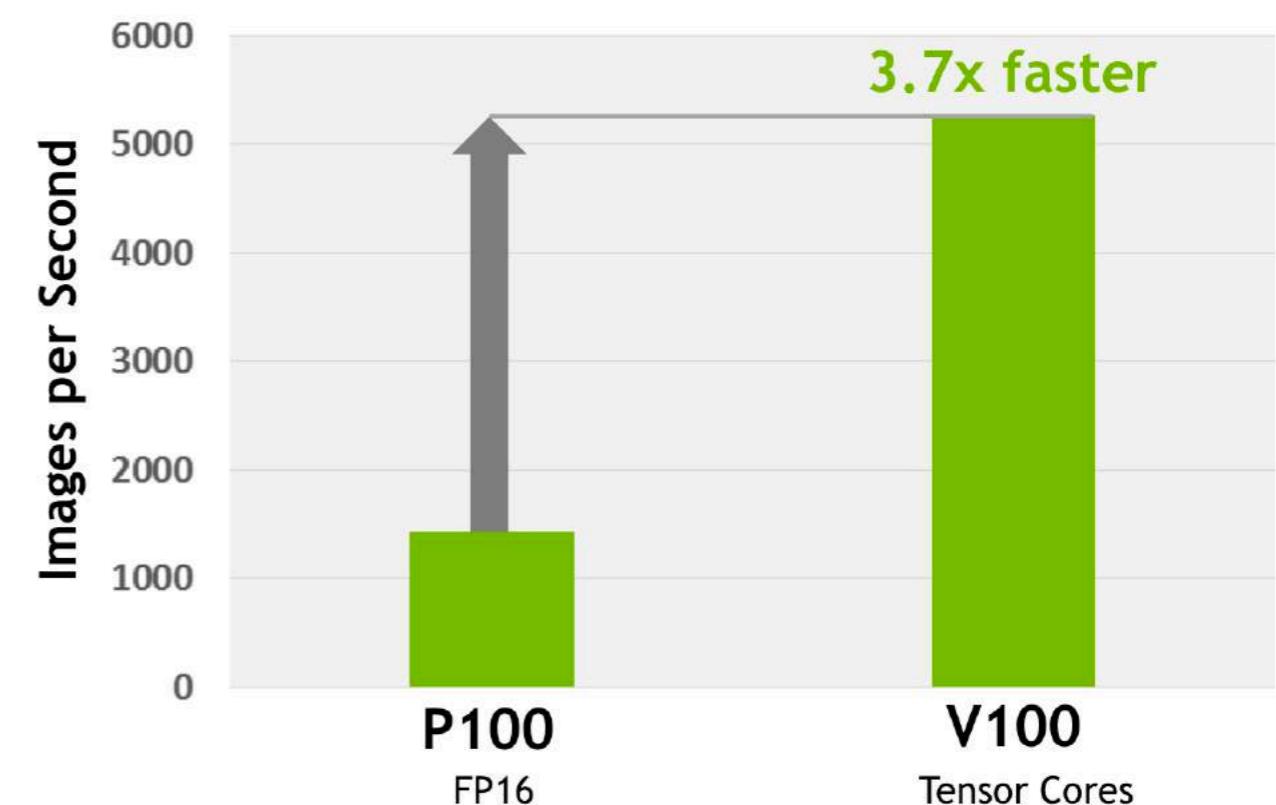
<https://devblogs.nvidia.com/parallelforall/cuda-9-features-revealed/>

# Pascal v.s. Volta

## ResNet-50 Training



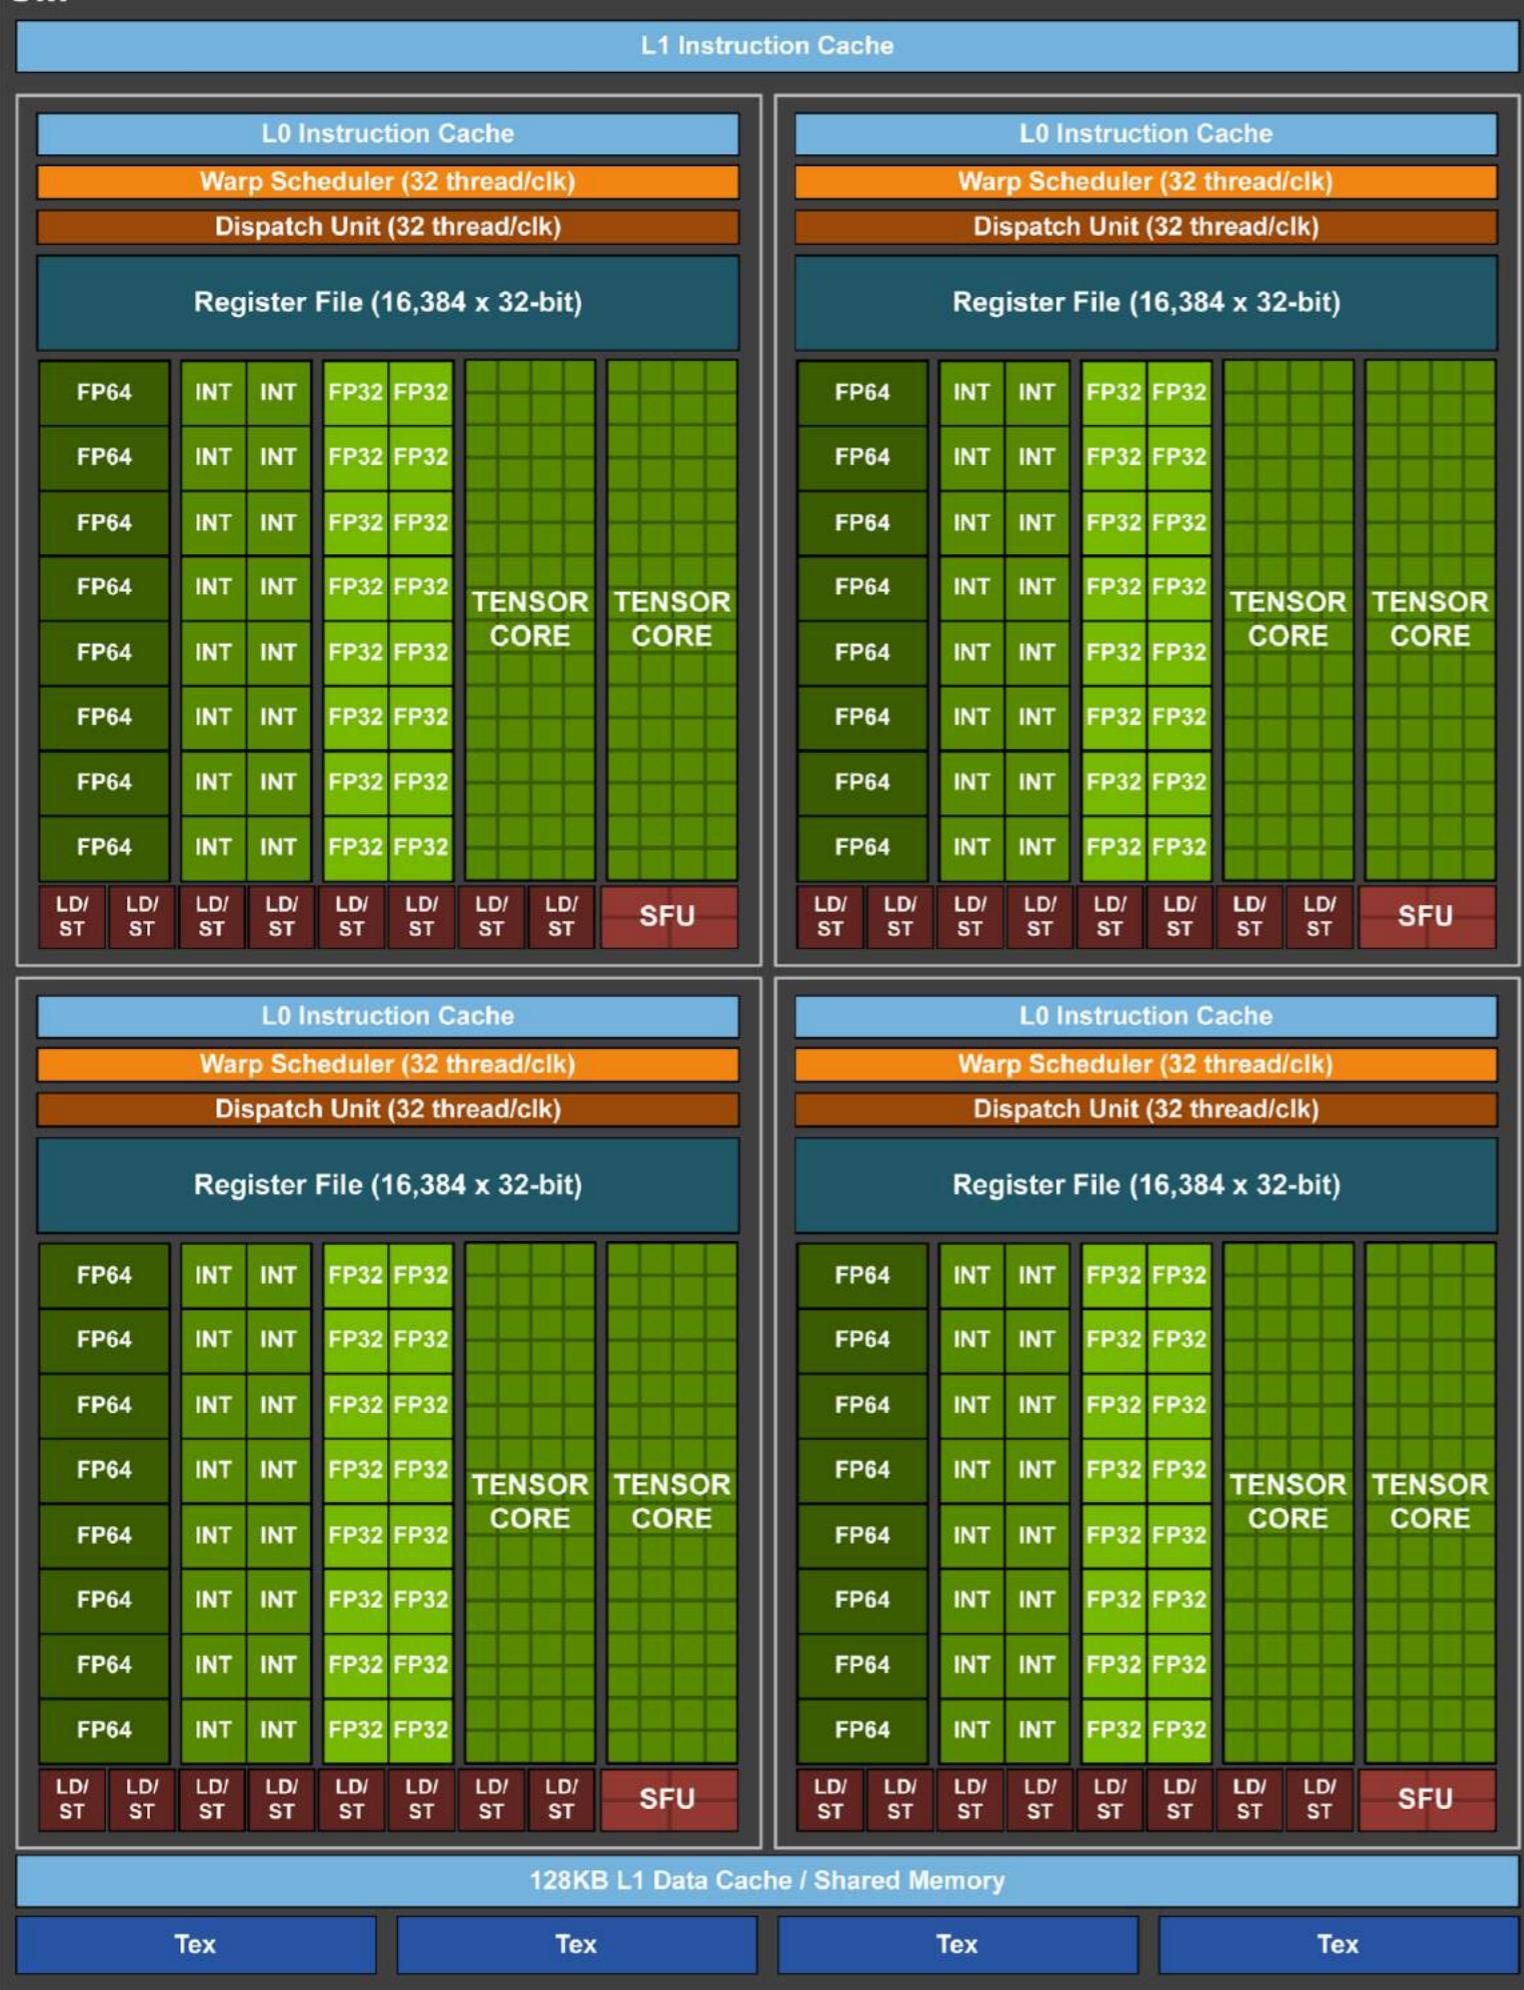
## ResNet-50 Inference TensorRT - 7ms Latency



Left: Tesla V100 trains the ResNet-50 deep neural network 2.4x faster than Tesla P100.

Right: Given a target latency per image of 7ms, Tesla V100 is able to perform inference using the ResNet-50 deep neural network 3.7x faster than Tesla P100.

<https://devblogs.nvidia.com/parallelforall/cuda-9-features-revealed/>



The GV100 SM is partitioned into four processing blocks, each with:

- 8 FP64 Cores
- 16 FP32 Cores
- 16 INT32 Cores
- two of the new mixed-precision Tensor Cores for deep learning
- a new L0 instruction cache
- one warp scheduler
- one dispatch unit
- a 64 KB Register File.

<https://devblogs.nvidia.com/parallelforall/cuda-9-features-revealed/>

Tesla Product	Tesla K40	Tesla M40	Tesla P100	Tesla V100
GPU	GK110 (Kepler)	GM200 (Maxwell)	GP100 (Pascal)	GV100 (Volta)
GPU Boost Clock	810/875 MHz	1114 MHz	1480 MHz	1455 MHz
Peak FP32 TFLOP/s*	5.04	6.8	10.6	15
Peak Tensor Core TFLOP/s*	-	-	-	120
Memory Interface	384-bit GDDR5	384-bit GDDR5	4096-bit HBM2	4096-bit HBM2
Memory Size	Up to 12 GB	Up to 24 GB	16 GB	16 GB
TDP	235 Watts	250 Watts	300 Watts	300 Watts
Transistors	7.1 billion	8 billion	15.3 billion	21.1 billion
GPU Die Size	551 mm <sup>2</sup>	601 mm <sup>2</sup>	610 mm <sup>2</sup>	815 mm <sup>2</sup>
Manufacturing Process	28 nm	28 nm	16 nm FinFET+	12 nm FFN

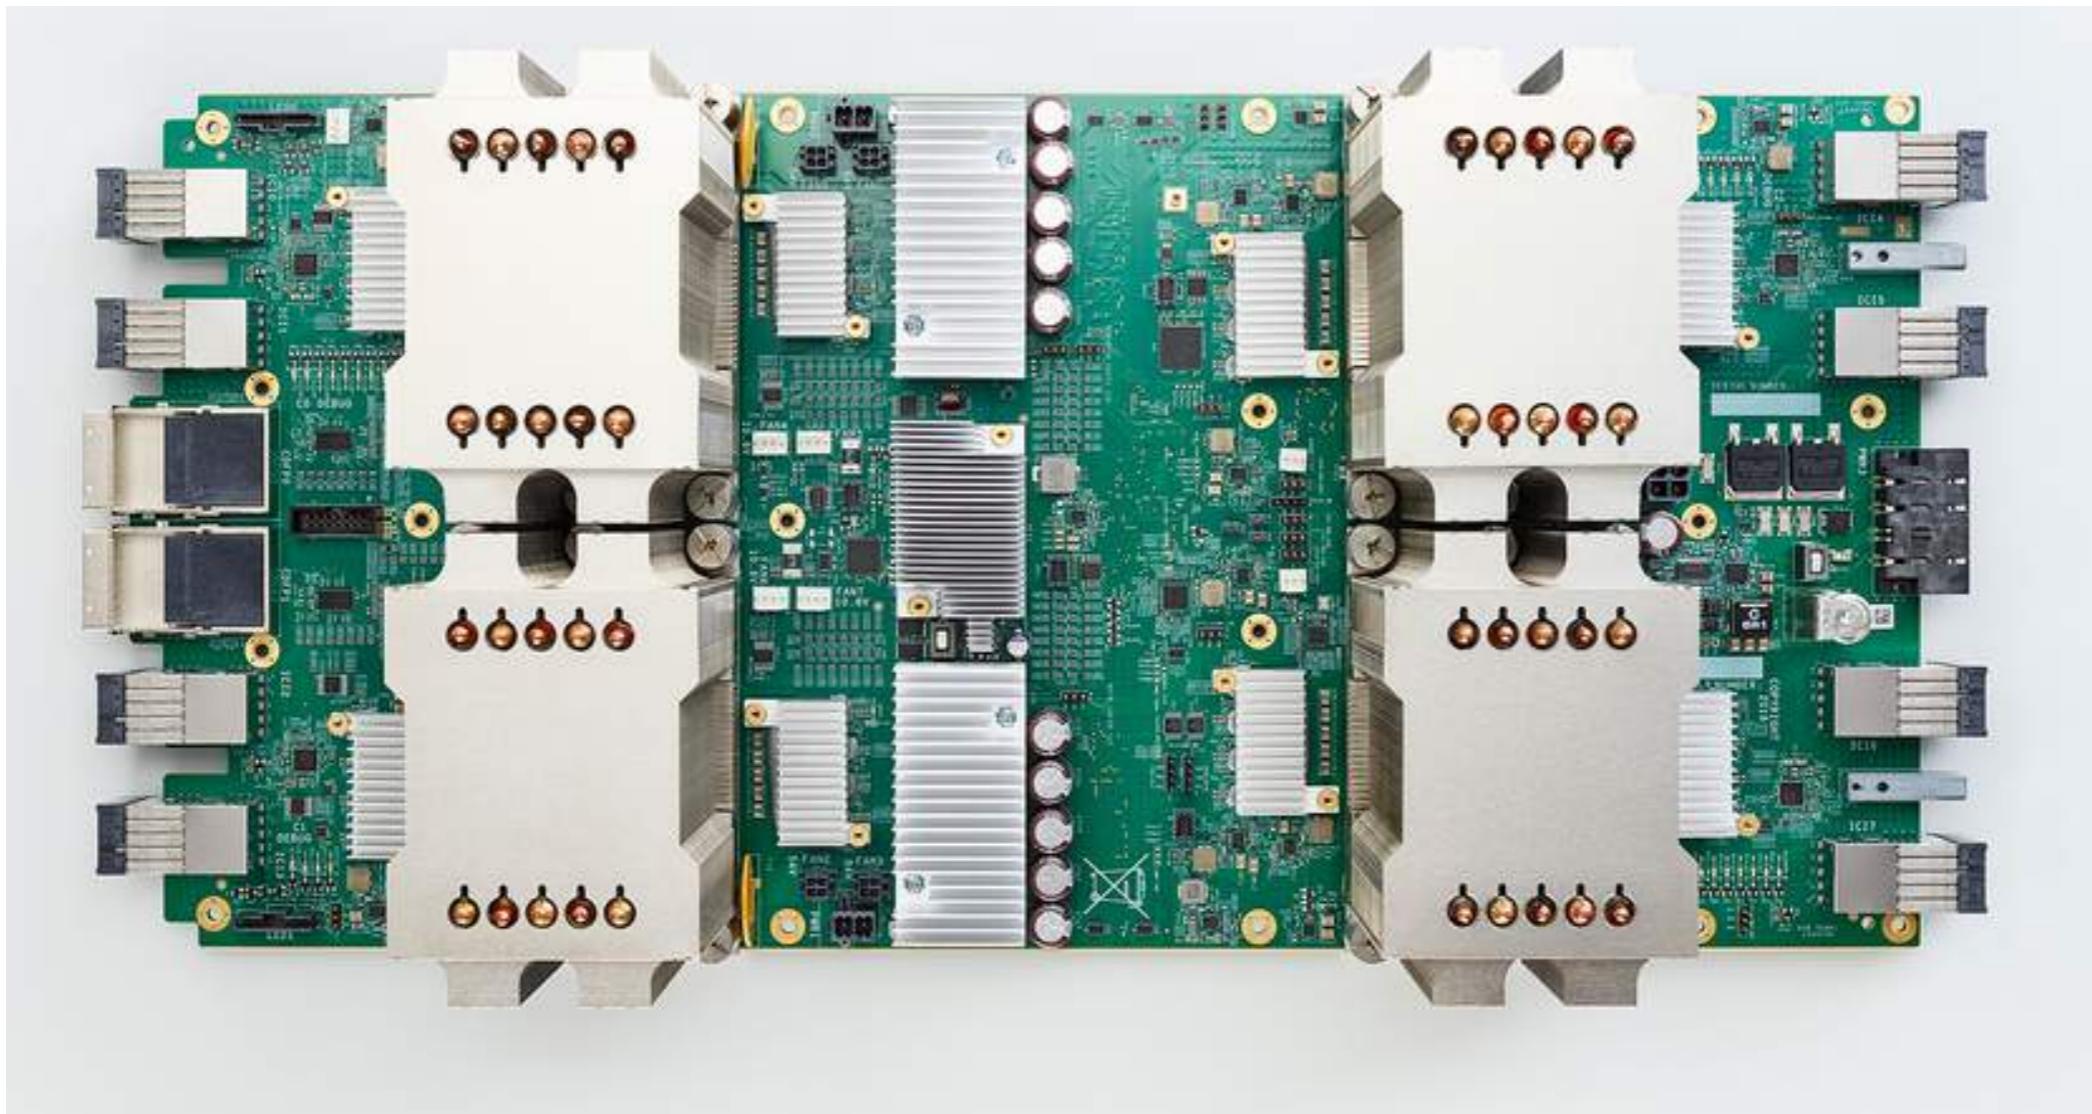
<https://devblogs.nvidia.com/parallelforall/cuda-9-features-revealed/>

# GPU / TPU

	K80 2012	TPU 2015	P40 2016
<b>Inferences/Sec &lt;10ms latency</b>	$\frac{1}{13}X$	1X	2X
<b>Training TOPS</b>	6 FP32	NA	12 FP32
<b>Inference TOPS</b>	6 FP32	90 INT8	48 INT8
<b>On-chip Memory</b>	16 MB	24 MB	11 MB
<b>Power</b>	300W	75W	250W
<b>Bandwidth</b>	320 GB/S	34 GB/S	350 GB/S

<https://blogs.nvidia.com/blog/2017/04/10/ai-drives-rise-accelerated-computing-datacenter/>

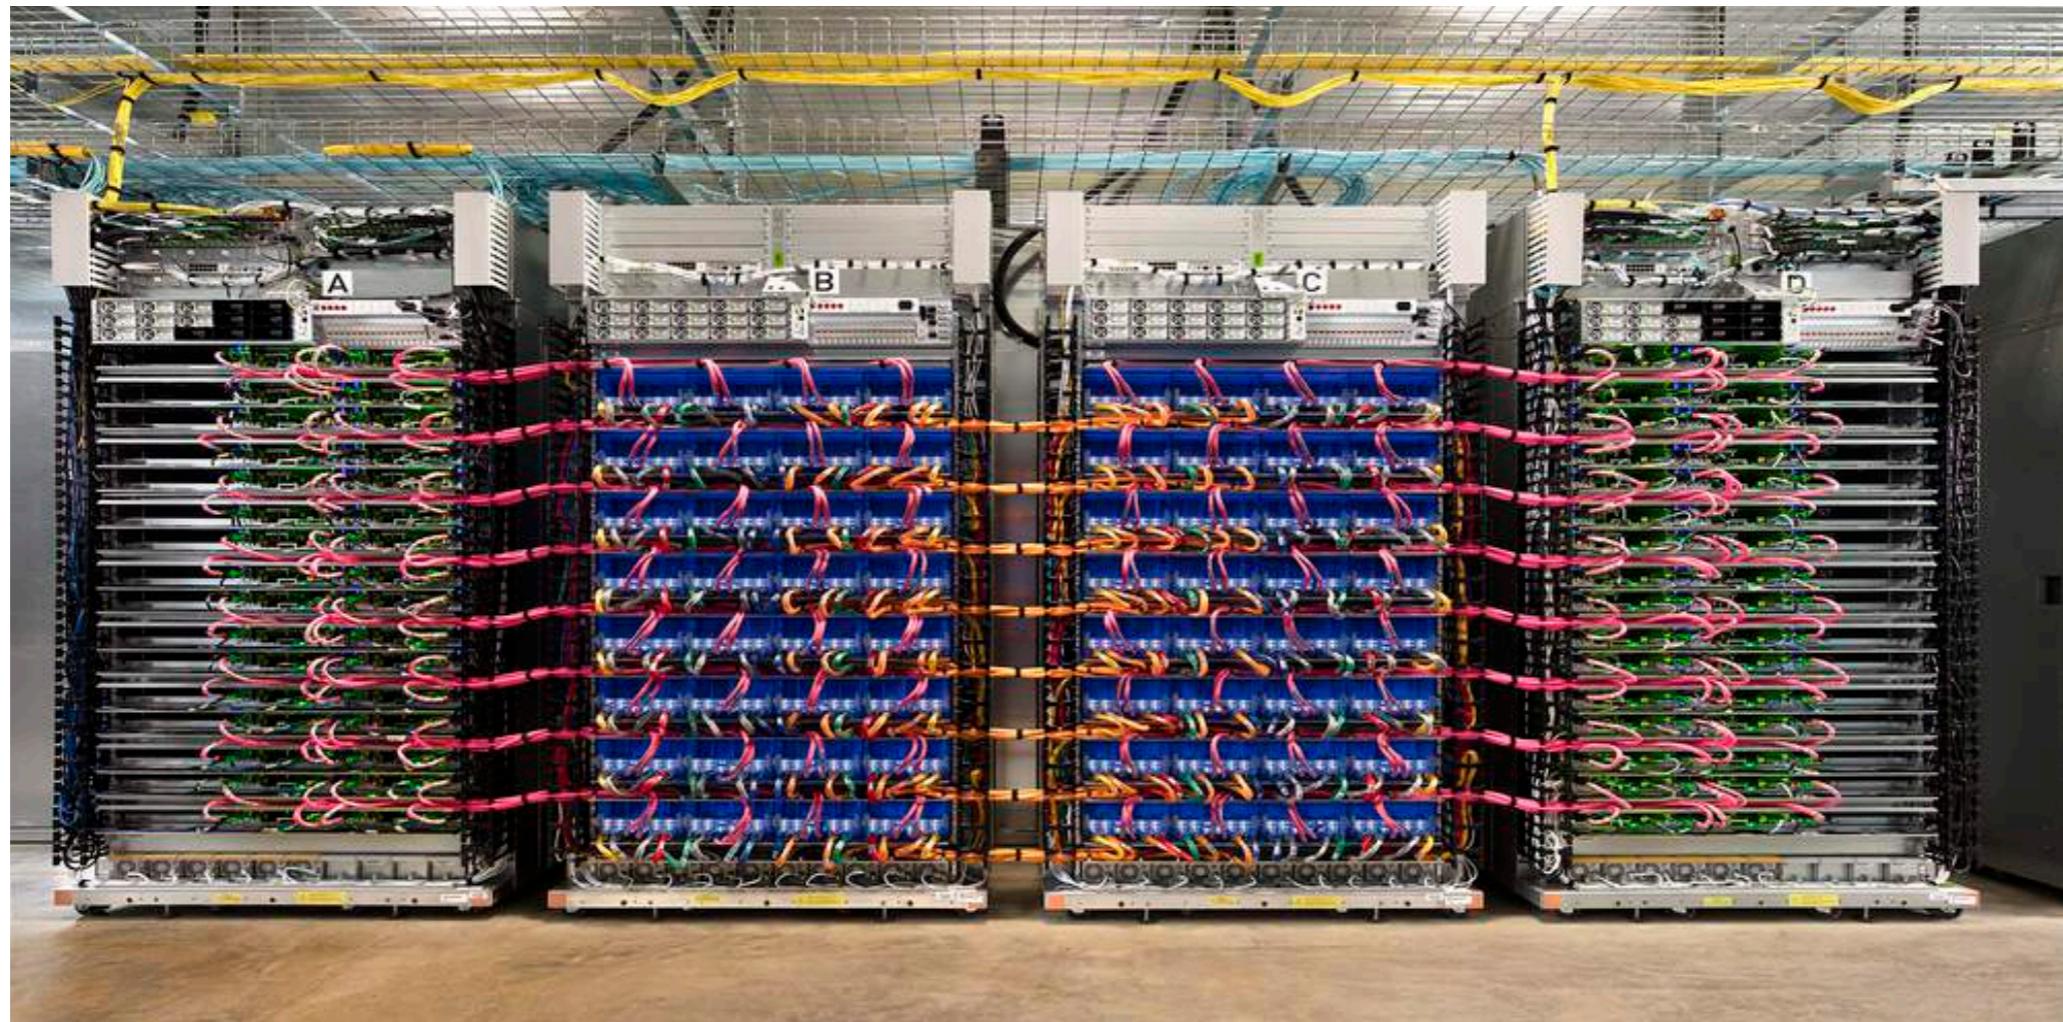
# Google Cloud TPU



Cloud TPU delivers up to 180 teraflops to train and run machine learning models.

source: Google Blog

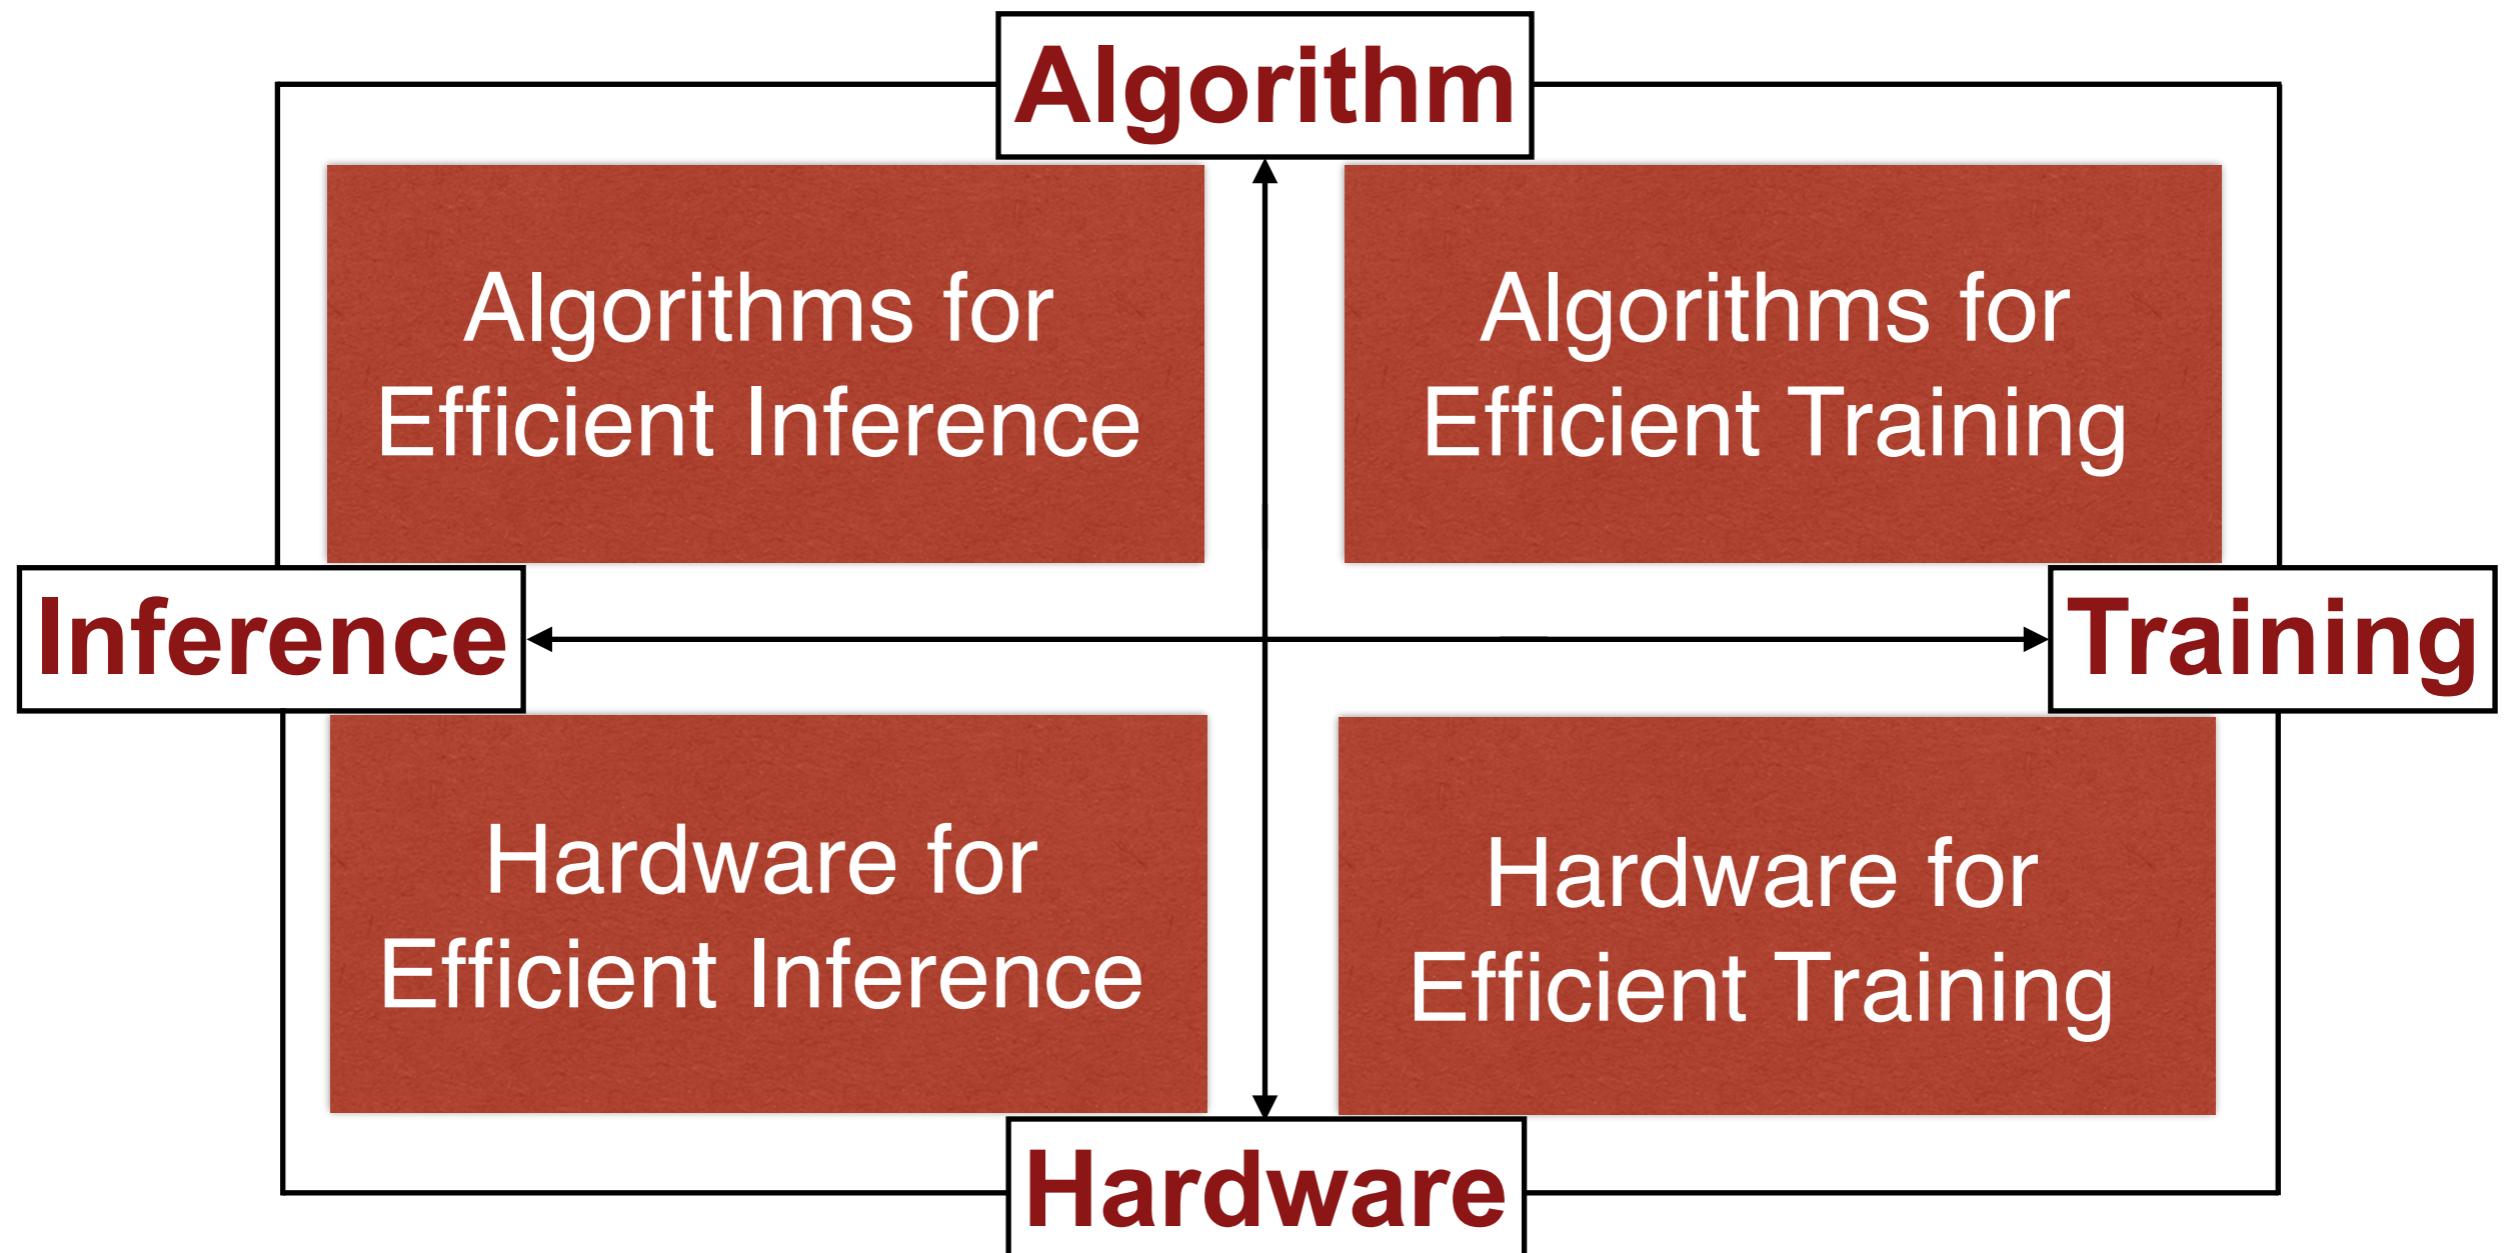
# Google Cloud TPU



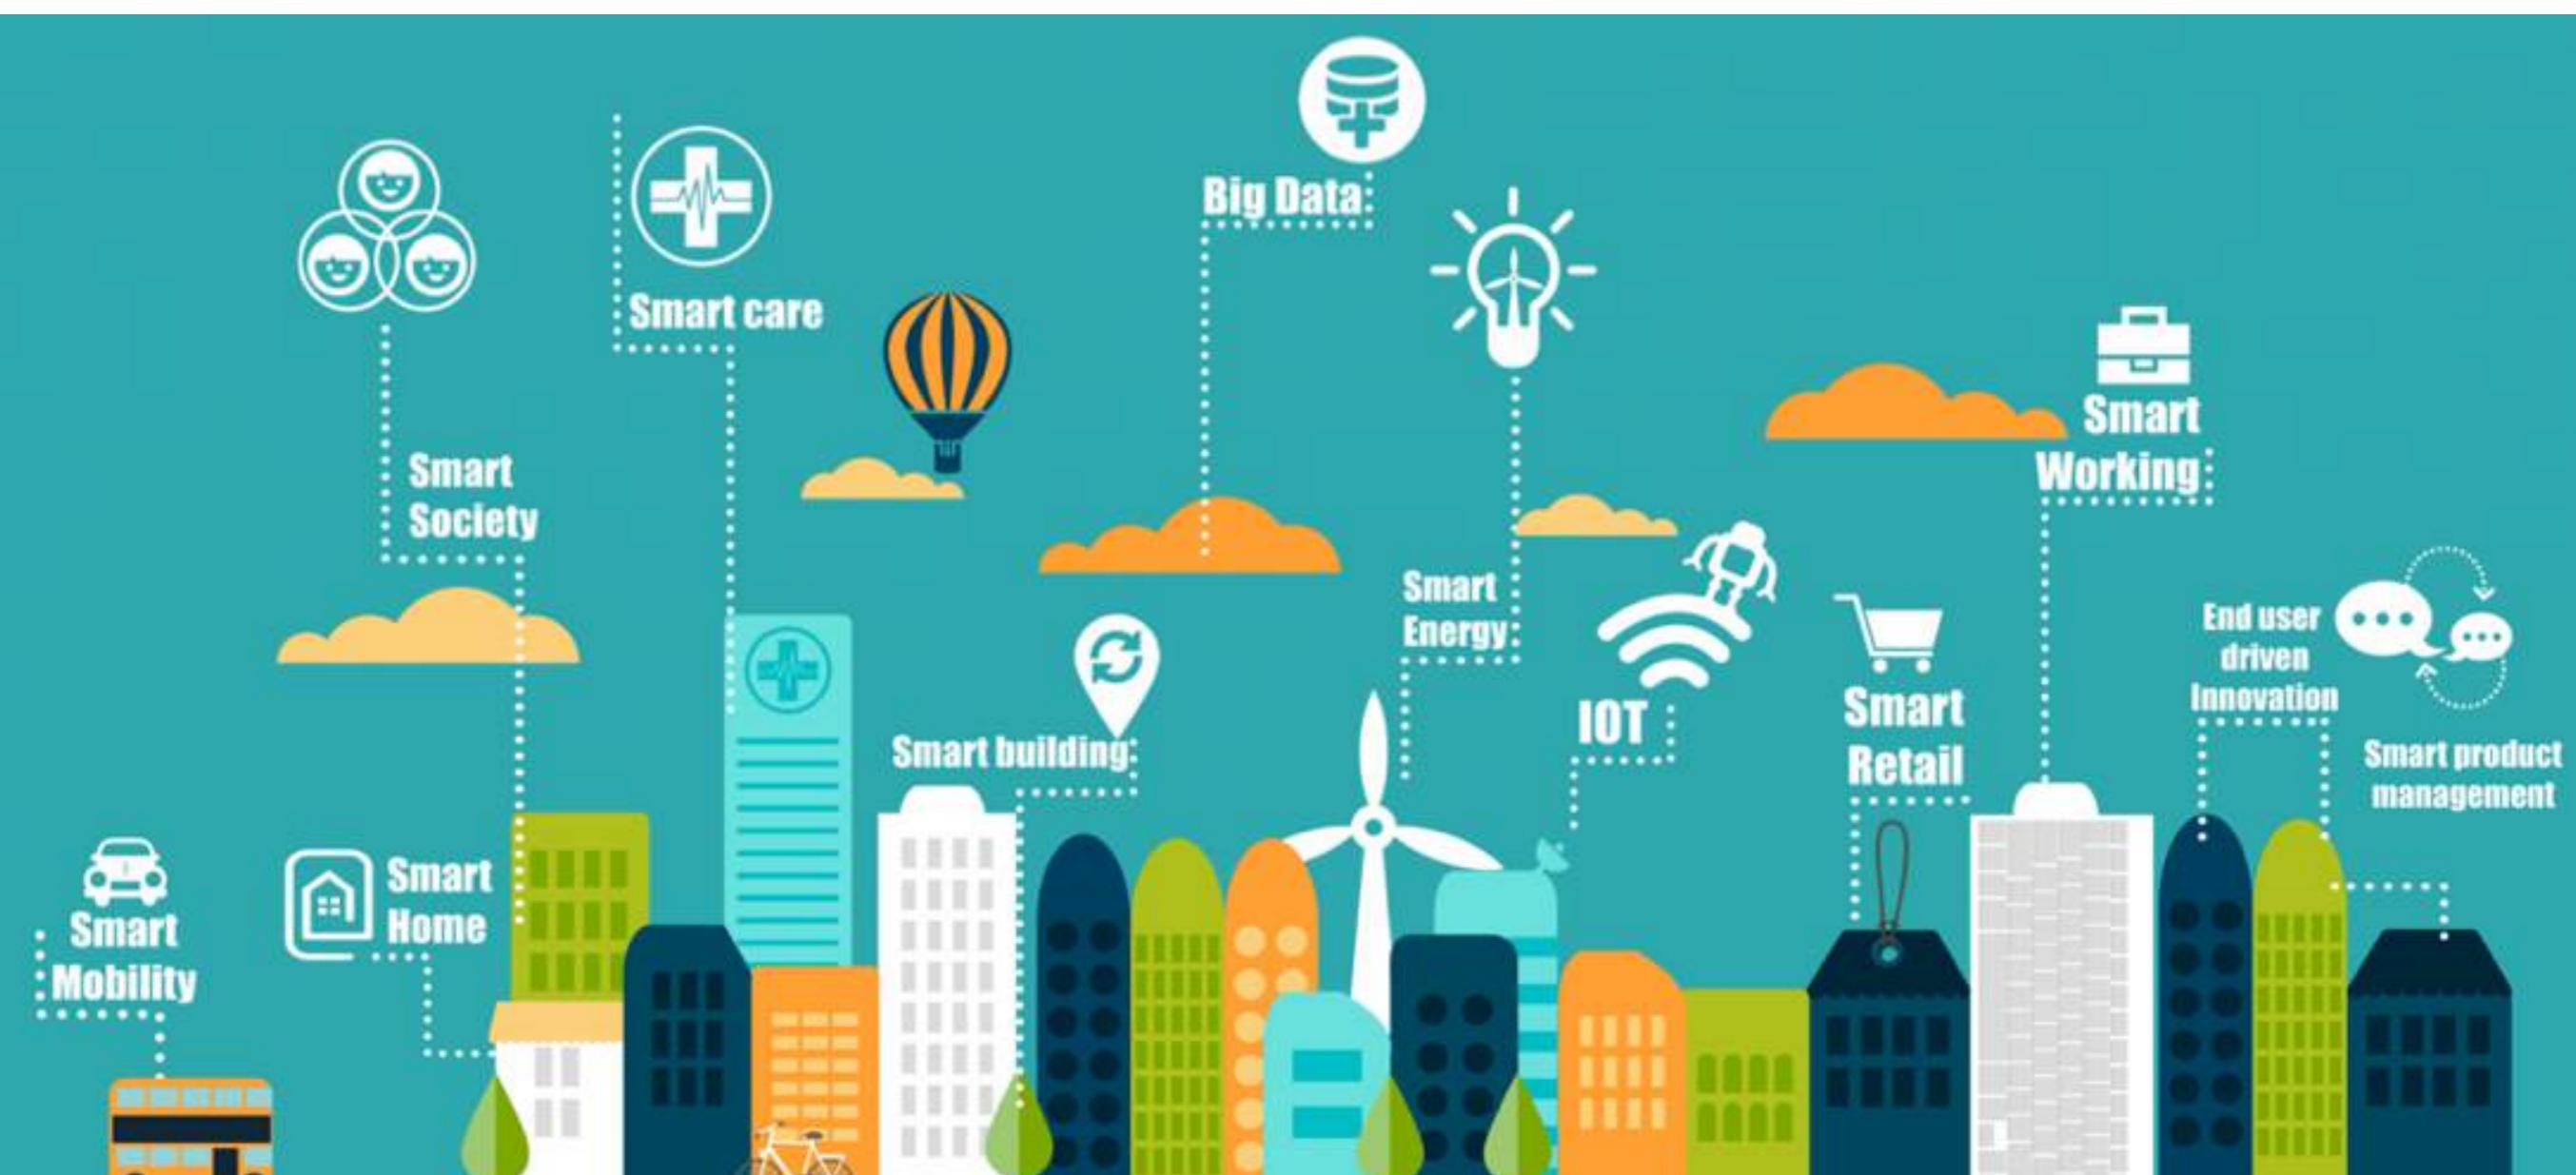
A “TPU pod” built with 64 second-generation TPUs delivers up to 11.5 petaflops of machine learning acceleration.

“One of our new large-scale translation models used to take a **full day** to train on 32 of the best commercially-available GPUs—now it trains to the same accuracy in **an afternoon** using just **one eighth** of a TPU pod.” — Google Blog

# Wrap-Up



# Future



Smart

Low Latency

Privacy

Mobility

Energy-Efficient

# Outlook: the Focus for Computation



PC Era



Mobile-First Era



AI-First Era

Computing



Mobile  
Computing



Brain-Inspired  
Cognitive  
Computing

Sundar Pichai, Google IO, 2016

# Thank you!

[stanford.edu/~songhan](http://stanford.edu/~songhan)

